

# XuanTie-C910-C920-UserManual

Nov 06, 2023

**Copyright 2021 T-Head Semiconductor Co., Ltd.**

Licensed under the Apache License, Version 2.0 (the "License"); you may not use this file except in compliance with the License. You may obtain a copy of the License at

<http://www.apache.org/licenses/LICENSE-2.0>

Unless required by applicable law or agreed to in writing, software distributed under the License is distributed on an "AS IS" BASIS, WITHOUT WARRANTIES OR CONDITIONS OF ANY KIND, either express or implied. See the License for the specific language governing permissions and limitations under the License.

# History

Version	Description	Date
01	C910-C920 first version	2021.10.19
02	change pmpaddr table	2021.10.25
03	C910-C920/C920 second version	2023.9.1

<b>1</b>	<b>Overview</b>	<b>1</b>
1.1	Introduction . . . . .	1
1.2	Features . . . . .	1
1.2.1	Architectural features of C910/C920MP . . . . .	1
1.2.2	Features of the C910/C920 core . . . . .	2
1.3	Configurations . . . . .	3
1.4	XuanTie extended architecture . . . . .	3
1.5	Version compatibility . . . . .	3
1.6	Naming conventions . . . . .	4
1.6.1	Terms . . . . .	4
<b>2</b>	<b>C910/C920MP Overview</b>	<b>5</b>
2.1	Structure . . . . .	5
2.2	In-core subsystems . . . . .	5
2.2.1	IFU . . . . .	5
2.2.2	IDU . . . . .	5
2.2.3	Execution units . . . . .	7
2.2.4	LSU . . . . .	7
2.2.5	RTU . . . . .	7
2.2.6	MMU . . . . .	7
2.2.7	PMP . . . . .	8
2.3	Multi-core subsystems . . . . .	8
2.3.1	CIU . . . . .	8
2.3.2	L2 cache . . . . .	8
2.3.3	Master device interface unit . . . . .	8
2.3.4	PLIC . . . . .	9
2.3.5	Timer . . . . .	9

2.4	Interface overview . . . . .	9
<b>3</b>	<b>Instruction Sets</b>	<b>10</b>
3.1	RV base instruction sets . . . . .	10
3.1.1	Integer instruction set (RV64I) . . . . .	10
3.1.2	Multiply/Divide instruction set (RV64M) . . . . .	13
3.1.3	Atomic instruction set (RV64A) . . . . .	14
3.1.4	Single-precision floating-point instruction set (RV64F) . . . . .	15
3.1.5	Compressed instruction set (RV64C) . . . . .	17
3.2	XuanTie extended instruction sets . . . . .	19
3.2.1	Arithmetic operation instructions . . . . .	19
3.2.2	Bit operation instructions . . . . .	20
3.2.3	Memory access instructions . . . . .	21
3.2.4	Cache instructions . . . . .	24
3.2.5	Multi-core synchronization instructions . . . . .	25
3.2.6	Half-precision floating-point instructions . . . . .	26
<b>4</b>	<b>CPU Modes and Registers</b>	<b>29</b>
4.1	CPU modes . . . . .	29
4.2	Register view . . . . .	30
4.3	General-purpose registers . . . . .	30
4.4	Floating-point registers . . . . .	31
4.4.1	Transmit data between floating-point and general-purpose registers . . . . .	32
4.4.2	Maintain consistency of register precision . . . . .	32
4.5	System control registers . . . . .	32
4.5.1	Standard control registers . . . . .	32
4.5.2	Extended control registers . . . . .	35
4.6	Data formats . . . . .	37
4.6.1	Integer data format . . . . .	38
4.6.2	Floating-point data format . . . . .	38
4.7	Big-endian and little-endian . . . . .	38
<b>5</b>	<b>Exceptions and Interrupts</b>	<b>41</b>
5.1	Overview . . . . .	41
5.2	Exceptions . . . . .	43
5.2.1	Exception handling . . . . .	43
5.2.2	Return from exceptions . . . . .	44
5.2.3	Imprecise exceptions . . . . .	45
5.3	Interrupts . . . . .	45
5.3.1	Interrupt priorities . . . . .	45
5.3.2	Interrupt responses . . . . .	45
5.3.3	Return from interrupts . . . . .	46

<b>6</b>	<b>Memory Model</b>	<b>47</b>
6.1	Overview . . . . .	47
6.1.1	Memory attributes . . . . .	47
6.1.2	Memory ordering model . . . . .	48
6.2	MMU . . . . .	49
6.2.1	Overview . . . . .	49
6.2.2	TLB . . . . .	49
6.2.3	Address translation process . . . . .	50
6.2.4	System control registers . . . . .	51
6.2.4.1	Supervisor address translation and protection (satp) register . . . . .	52
6.2.4.2	smcir register . . . . .	52
6.2.4.3	smir register . . . . .	53
6.2.4.4	smeh register . . . . .	54
6.2.4.5	smel register . . . . .	55
6.3	PMP . . . . .	57
6.3.1	Overview . . . . .	57
6.3.2	PMP control registers . . . . .	58
6.3.2.1	Physical memory protection configuration (pmpcfg) register . . . . .	58
6.3.2.2	Physical memory protection address (pmpaddr) register . . . . .	60
6.4	Memory access order . . . . .	60
<b>7</b>	<b>Memory Subsystem</b>	<b>62</b>
7.1	Memory Subsystem Overview . . . . .	62
7.2	L1 I-Cache . . . . .	62
7.2.1	Overview . . . . .	62
7.2.2	Way prediction . . . . .	63
7.2.3	Loop acceleration buffer . . . . .	63
7.2.4	Branch history table . . . . .	63
7.2.5	Branch and jump target predictor . . . . .	63
7.2.6	Indirect branch predictor . . . . .	64
7.2.7	Return address predictor . . . . .	64
7.2.8	Fast jump target predictor . . . . .	65
7.3	L1 D-Cache . . . . .	65
7.3.1	Overview . . . . .	65
7.3.2	Cache coherence . . . . .	66
7.3.3	Exclusive access . . . . .	66
7.4	L2 Cache . . . . .	67
7.4.1	Overview . . . . .	67
7.4.2	Cache coherence . . . . .	67
7.4.3	Structure . . . . .	68
7.4.4	RAM latency . . . . .	68
7.5	Accelerated memory access . . . . .	70

7.5.1	Instruction prefetch of the L1 I-Cache . . . . .	70
7.5.2	Multi-channel data prefetch of the L1 D-Cache . . . . .	70
7.5.3	L1 adaptive write allocation mechanism . . . . .	71
7.5.4	L2 prefetch mechanism . . . . .	71
7.6	L1/L2 cache operation instructions and registers . . . . .	72
7.6.1	Extended registers of the L1 cache . . . . .	72
7.6.2	Extended registers of the L2 cache . . . . .	72
7.6.3	L1/L2 cache operation instructions . . . . .	73
<b>8</b>	<b>Vector Computations(not supportC910 )</b>	<b>74</b>
8.1	Vector programming model . . . . .	74
8.2	Vector control registers . . . . .	74
8.3	Vector exceptions . . . . .	75
<b>9</b>	<b>Security Design</b>	<b>77</b>
9.1	Security Requirements . . . . .	77
9.2	Processor Security Model . . . . .	77
9.3	System Security Architecture . . . . .	79
9.3.1	Secure memory management . . . . .	79
9.3.2	Secure interrupts . . . . .	82
9.3.3	Secure Access Control . . . . .	84
9.3.4	Secure Debugging . . . . .	86
<b>10</b>	<b>Interrupt Controllers</b>	<b>87</b>
10.1	Core local interrupt (CLINT) controller . . . . .	87
10.1.1	CLINT register address mapping . . . . .	87
10.1.2	Software interrupts . . . . .	90
10.1.3	Timer interrupts . . . . .	91
10.2	Platform-level interrupt controller (PLIC) . . . . .	92
10.2.1	Interrupt arbitration . . . . .	92
10.2.2	Interrupt request and response . . . . .	93
10.2.3	Interrupt completion . . . . .	93
10.2.4	PLIC register address mapping . . . . .	94
10.2.5	PLIC_PRIO register . . . . .	96
10.2.6	PLIC_IP register . . . . .	96
10.2.7	PLIC_IE register . . . . .	97
10.2.8	PLIC_CTRL register . . . . .	97
10.2.9	PLIC_TH register . . . . .	98
10.2.10	PLIC_CLAIM register . . . . .	98
10.3	Multi-core interrupts . . . . .	99
10.3.1	Multiple cores respond to external interrupts in parallel . . . . .	99
10.3.2	Send software interrupts across cores . . . . .	99

<b>11 Bus Interface</b>	<b>100</b>
11.1 AXI master device interface . . . . .	100
11.1.1 Features of the AXI master device interface . . . . .	100
11.1.2 Outstanding capability of the AXI master device interface . . . . .	100
11.1.3 Supported transmission types . . . . .	102
11.1.4 Supported response types . . . . .	102
11.1.5 CPU behavior in different bus responses . . . . .	102
11.1.6 Signals supported by the AXI master device interface . . . . .	103
<b>12 Debug</b>	<b>107</b>
12.1 Features of the debug unit . . . . .	107
12.2 Connection between the debug unit and CPU cores . . . . .	108
12.3 Debug interface signals . . . . .	109
<b>13 Power Management</b>	<b>112</b>
13.1 Power domain . . . . .	112
13.2 Overview of low-power modes . . . . .	112
13.3 Core WFI process . . . . .	113
13.4 Individual-core power-off process . . . . .	113
13.5 Cluster power-off process (hardware clearing of the L2 cache) . . . . .	114
13.6 Cluster power-off process (software clearing of the L2 cache) . . . . .	115
13.7 Simplified scenario: overall cluster power-off process (hardware clearing of the L2 cache) . . . . .	116
13.8 Simplified scenario: overall cluster power-off process (software clearing of the L2 cache) . . . . .	117
13.9 Low power consumption related programming models and interface signals . . . . .	118
13.9.1 Programming models . . . . .	118
13.9.2 Interface signals . . . . .	119
<b>14 Performance Monitoring Unit</b>	<b>120</b>
14.1 PMU overview . . . . .	120
14.2 PMU programming model . . . . .	120
14.2.1 PMU functions . . . . .	120
14.2.2 PMU event overflow interrupt . . . . .	121
14.3 PMU related control registers . . . . .	121
14.3.1 M-mode counter access enable register (mcounteren) . . . . .	121
14.3.2 S-mode counter access enable register (scounteren) . . . . .	122
14.3.3 M-mode count inhibit register (mcountinhibit) . . . . .	123
14.3.4 S-mode write enable register (mcounterwen) . . . . .	124
14.3.5 Performance monitoring event select register (mhpmevent3-31) . . . . .	124
14.3.6 Event counters . . . . .	126
<b>15 Program Examples</b>	<b>128</b>
15.1 Optimal CPU performance configuration . . . . .	128
15.2 MMU setting example . . . . .	129



15.3	PMP setting example . . . . .	133
15.4	Cache examples . . . . .	134
15.4.1	Cache enabling example . . . . .	134
15.4.2	Example of synchronization between the instruction and data caches . . . . .	135
15.4.3	Example of synchronization between the TLB and the data cache . . . . .	135
15.5	Synchronization primitive examples . . . . .	136
15.6	PLIC setting example . . . . .	136
15.7	PMU setting example . . . . .	137
<b>16</b>	<b>Appendix A Standard Instructions</b>	<b>139</b>
16.1	Appendix A-1 I instructions . . . . .	139
16.1.1	ADD: a signed add instruction . . . . .	139
16.1.2	ADDI: a signed add immediate instruction . . . . .	140
16.1.3	ADDIW: a signed add immediate instruction that operates on the lower 32 bits . . .	140
16.1.4	ADDW: a signed add instruction that operates on the lower 32 bits . . . . .	141
16.1.5	AND: a bitwise AND instruction . . . . .	141
16.1.6	ANDI: an immediate bitwise AND instruction . . . . .	142
16.1.7	AUIPC: an instruction that adds the immediate in the upper bits to the PC . . . .	142
16.1.8	BEQ: a branch-if-equal instruction . . . . .	142
16.1.9	BGE: a signed branch-if-greater-than-or-equal instruction . . . . .	143
16.1.10	BGEU: an unsigned branch-if-greater-than-or-equal instruction . . . . .	144
16.1.11	BLT: a signed branch-if-less-than instruction . . . . .	145
16.1.12	BLTU: an unsigned branch-if-less-than instruction . . . . .	145
16.1.13	BNE: a branch-if-not-equal instruction . . . . .	146
16.1.14	CSRRC: a move instruction that clears control registers . . . . .	147
16.1.15	CSRRCI: a move instruction that clears immediates in control registers . . . . .	147
16.1.16	CSRRS: a move instruction for setting control registers . . . . .	148
16.1.17	CSRRSI: a move instruction for setting immediates in control registers . . . . .	148
16.1.18	CSRRW: a move instruction that reads/writes control registers . . . . .	149
16.1.19	CSRRWI: a move instruction that reads/writes immediates in control registers . . .	150
16.1.20	EBREAK: a breakpoint instruction . . . . .	150
16.1.21	ECALL: an environment call instruction . . . . .	151
16.1.22	FENCE: a memory synchronization instruction . . . . .	151
16.1.23	FENCE.I: an instruction stream synchronization instruction . . . . .	152
16.1.24	JAL: an instruction for directly jumping to a subroutine . . . . .	152
16.1.25	JALR: an instruction for jumping to a subroutine by using an address in a register .	153
16.1.26	LB: a sign-extended byte load instruction . . . . .	153
16.1.27	LBU: an unsign-extended byte load instruction . . . . .	154
16.1.28	LD: a doubleword load instruction . . . . .	154
16.1.29	LH: a sign-extended halfword load instruction . . . . .	155
16.1.30	LHU: an unsign-extended halfword load instruction . . . . .	155
16.1.31	LUI: an instruction for loading the immediate in the upper bits . . . . .	156

16.1.32	LW: a sign-extended word load instruction . . . . .	156
16.1.33	LWU: an unsign-extended word load instruction . . . . .	157
16.1.34	MRET: an instruction for returning from exceptions in M-mode . . . . .	157
16.1.35	OR: a bitwise OR instruction . . . . .	158
16.1.36	ORI: an immediate bitwise OR instruction . . . . .	158
16.1.37	SB: a byte store instruction . . . . .	158
16.1.38	SD: a doubleword store instruction . . . . .	159
16.1.39	SFENCE.VMA: a virtual memory synchronization instruction . . . . .	159
16.1.40	SH: a halfword store instruction . . . . .	160
16.1.41	SLL: a logical left shift instruction . . . . .	161
16.1.42	SLLI: an immediate logical left shift instruction . . . . .	161
16.1.43	SLLIW: an immediate logical left shift instruction that operates on the lower 32 bits	162
16.1.44	SLLW: a logical left shift instruction that operates on the lower 32 bits . . . . .	162
16.1.45	SLT: a signed set-if-less-than instruction . . . . .	163
16.1.46	SLTI: a signed set-if-less-than-immediate instruction . . . . .	163
16.1.47	SLTIU: an unsigned set-if-less-than-immediate instruction . . . . .	164
16.1.48	SLTU: an unsigned set-if-less-than instruction . . . . .	164
16.1.49	SRA: an arithmetic right shift instruction . . . . .	165
16.1.50	SRAI: an immediate arithmetic right shift instruction . . . . .	165
16.1.51	SLLIW: an immediate arithmetic right shift instruction that operates on the lower 32 bits . . . . .	166
16.1.52	SRAW: an arithmetic right shift instruction that operates on the lower 32 bits . . . .	166
16.1.53	SRET: an instruction for returning from exceptions in S-mode . . . . .	167
16.1.54	SRL: a logical right shift instruction . . . . .	167
16.1.55	SRLI: an immediate logical right shift instruction . . . . .	167
16.1.56	SRLIW: an immediate logical right shift instruction that operates on the lower 32 bits	168
16.1.57	SRLW: a logical right shift instruction that operates on the lower 32 bits . . . . .	168
16.1.58	SUB: a signed subtract instruction . . . . .	169
16.1.59	SUBW: a signed subtract instruction that operates on the lower 32 bits . . . . .	169
16.1.60	SW: a word store instruction . . . . .	170
16.1.61	WFI: an instruction for entering the low power mode . . . . .	170
16.1.62	XOR: a bitwise XOR instruction . . . . .	171
16.1.63	XORI: an immediate bitwise XOR instruction . . . . .	171
16.2	Appendix A-2 M instructions . . . . .	171
16.2.1	DIV: a signed divide instruction . . . . .	172
16.2.2	DIVU: an unsigned divide instruction . . . . .	172
16.2.3	DIVUW: an unsigned divide instruction that operates on the lower 32 bits . . . . .	173
16.2.4	DIVW: a signed divide instruction that operates on the lower 32 bits . . . . .	173
16.2.5	MUL: a signed multiply instruction . . . . .	174
16.2.6	MULH: a signed multiply instruction that extracts the upper bits . . . . .	174
16.2.7	MULHSU: a signed-unsigned multiply instruction that extracts the upper bits . . . .	174
16.2.8	MULHU: an unsigned multiply instruction that extracts the upper bits . . . . .	175

16.2.9	MULW: a signed multiply instruction that operates on the lower 32 bits . . . . .	175
16.2.10	REM: a signed remainder instruction . . . . .	176
16.2.11	REMU: an unsigned remainder instruction . . . . .	176
16.2.12	REMUW: an unsigned remainder instruction that operates on the lower 32 bits . . .	177
16.2.13	REMW: a signed remainder instruction that operates on the lower 32 bits . . . . .	177
16.3	Appendix A-3 A instructions . . . . .	178
16.3.1	AMOADD.D: an atomic add instruction . . . . .	178
16.3.2	AMOADD.W: an atomic add instruction that operates on the lower 32 bits . . . . .	179
16.3.3	AMOAND.D: an atomic bitwise AND instruction . . . . .	180
16.3.4	AMOAND.W: an atomic bitwise AND instruction that operates on the lower 32 bits	181
16.3.5	AMOMAX.D: an atomic signed MAX instruction . . . . .	182
16.3.6	AMOMAX.W: an atomic signed MAX instruction that operates on the lower 32 bits	183
16.3.7	MOMAXU.DA: an atomic unsigned MAX instruction . . . . .	184
16.3.8	AMOMAXU.W: an atomic unsigned MAX instruction that operates on the lower 32 bits. . . . .	185
16.3.9	AMOMIN.D: an atomic signed MIN instruction . . . . .	186
16.3.10	AMOMIN.W: an atomic signed MIN instruction that operates on the lower 32 bits .	187
16.3.11	AMOMINU.D: an atomic unsigned MIN instruction . . . . .	188
16.3.12	AMOMINU.W: an atomic unsigned MIN instruction that operates on the lower 32 bits	189
16.3.13	AMOOR.D: an atomic bitwise OR instruction. . . . .	190
16.3.14	AMOOR.W: an atomic bitwise OR instruction that operates on the lower 32 bits . .	190
16.3.15	AMOSWAP.D: an atomic swap instruction . . . . .	191
16.3.16	AMOSWAP.W: an atomic swap instruction that operates on the lower 32 bits . . . .	192
16.3.17	AMOXOR.D: an atomic bitwise XOR instruction . . . . .	193
16.3.18	AMOXOR.W: an atomic bitwise XOR instruction that operates on the lower 32 bits	194
16.3.19	LR.D: a doubleword load-reserved instruction . . . . .	195
16.3.20	LR.W: a word load-reserved instruction . . . . .	196
16.3.21	SC.D: a doubleword store-conditional instruction . . . . .	197
16.3.22	SC.W: a word store-conditional instruction . . . . .	198
16.4	Appendix A-4 F instructions . . . . .	199
16.4.1	FADD.S: a single-precision floating-point add instruction . . . . .	199
16.4.2	FCLASS.S: a single-precision floating-point classify instruction . . . . .	200
16.4.3	FCVT.L.S: an instruction that converts a single-precision floating-point number into a signed long integer . . . . .	201
16.4.4	FCVT.LU.S: an instruction that converts a single-precision floating-point number into an unsigned long integer . . . . .	202
16.4.5	FCVT.S.L: an instruction that converts a signed long integer into a single-precision floating-point number . . . . .	203
16.4.6	FCVT.S.LU: an instruction that converts an unsigned long integer into a single- precision floating-point number . . . . .	204
16.4.7	FCVT.S.W: an instruction that converts a signed integer into a single-precision floating-point number . . . . .	205

16.4.8	FCVT.S.WU: an instruction that converts an unsigned integer into a single-precision floating-point number . . . . .	206
16.4.9	FCVT.W.S: an instruction that converts a single-precision floating-point number into a signed integer . . . . .	207
16.4.10	FCVT.WU.S: an instruction that converts a single-precision floating-point number into an unsigned integer . . . . .	208
16.4.11	FDIV.S: a single-precision floating-point divide instruction . . . . .	209
16.4.12	FEQ.S: a single-precision floating-point compare equal instruction . . . . .	210
16.4.13	FLE.S: a single-precision floating-point compare less than or equal to instruction . . . . .	210
16.4.14	FLT.S: a single-precision floating-point compare less than instruction . . . . .	211
16.4.15	FLW: a single-precision floating-point load instruction . . . . .	212
16.4.16	FMADD.S: a single-precision floating-point multiply-add instruction . . . . .	212
16.4.17	FMAX.S: a single-precision floating-point MAX instruction . . . . .	213
16.4.18	FMIN.S: a single-precision floating-point MIN instruction . . . . .	214
16.4.19	FMSUB.S: a single-precision floating-point multiply-subtract instruction . . . . .	214
16.4.20	FMUL.S: a single-precision floating-point multiply instruction . . . . .	215
16.4.21	FMV.W.X: a single-precision floating-point write move instruction . . . . .	216
16.4.22	FMV.X.H: a single-precision floating-point read move instruction . . . . .	217
16.4.23	FNMADD.S: a single-precision floating-point negate-(multiply-add) instruction . . . . .	217
16.4.24	FNMSUB.S: a single-precision floating-point negate-(multiply-subtract) instruction . . . . .	218
16.4.25	FSGNJ.S: a single-precision floating-point sign-injection instruction . . . . .	219
16.4.26	FSGNJN.S: a single-precision floating-point negate sign-injection instruction . . . . .	220
16.4.27	FSGNJX.S: a single-precision floating-point XOR sign-injection instruction . . . . .	220
16.4.28	FSQRT.S: a single-precision floating-point square-root instruction . . . . .	221
16.4.29	FSUB.S: a single-precision floating-point subtract instruction . . . . .	222
16.4.30	FSW: a single-precision floating-point store instruction . . . . .	223
16.5	Appendix A-5 D instructions . . . . .	223
16.5.1	FADD.D: a double-precision floating-point add instruction . . . . .	223
16.5.2	FCLASS.D: a double-precision floating-point classify instruction . . . . .	224
16.5.3	FCVT.D.L: an instruction that converts a signed long integer into a double-precision floating-point number . . . . .	225
16.5.4	FCVT.D.LU: an instruction that converts an unsigned long integer into a double-precision floating-point number . . . . .	226
16.5.5	FCVT.D.S: an instruction that converts a single-precision floating-point number into a double-precision floating-point number . . . . .	227
16.5.6	FCVT.D.W: an instruction that converts a signed integer into a double-precision floating-point number . . . . .	228
16.5.7	FCVT.D.WU: an instruction that converts an unsigned integer into a double-precision floating-point number . . . . .	228
16.5.8	FCVT.L.D: an instruction that converts a double-precision floating-point number into a signed long integer . . . . .	229

16.5.9	FCVT.LU.D: an instruction that converts a double-precision floating-point number into an unsigned long integer . . . . .	230
16.5.10	FCVT.S.D: an instruction that converts a double-precision floating-point number into a single-precision floating-point number . . . . .	231
16.5.11	FCVT.W.D: an instruction that converts a double-precision floating-point number into a signed integer . . . . .	232
16.5.12	FCVT.WU.D: an instruction that converts a double-precision floating-point number into an unsigned integer . . . . .	233
16.5.13	FDIV.D: a double-precision floating-point divide instruction . . . . .	234
16.5.14	FEQ.D: a double-precision floating-point compare equal instruction . . . . .	235
16.5.15	FLD: a double-precision floating-point load instruction . . . . .	235
16.5.16	FLE.D: a double-precision floating-point compare less than or equal to instruction . . . . .	236
16.5.17	FLT.D: a double-precision floating-point compare less than instruction . . . . .	236
16.5.18	FMADD.D: a double-precision floating-point multiply-add instruction . . . . .	237
16.5.19	FMAX.D: a double-precision floating-point MAX instruction . . . . .	238
16.5.20	FMIN.D: a double-precision floating-point MIN instruction . . . . .	239
16.5.21	FMSUB.D: a double-precision floating-point multiply-subtract instruction . . . . .	239
16.5.22	FMUL.D: a double-precision floating-point multiply instruction . . . . .	240
16.5.23	FMV.D.X: a double-precision floating-point write move instruction . . . . .	241
16.5.24	FMV.X.D: a double-precision floating-point read move instruction . . . . .	242
16.5.25	FNMADD.D: a double-precision floating-point negate-(multiply-add) instruction . . . . .	242
16.5.26	FNMSUB.D: a double-precision floating-point negate-(multiply-subtract) instruction . . . . .	243
16.5.27	FSD: a double-precision floating-point store instruction . . . . .	244
16.5.28	FSGNJ.D: a double-precision floating-point sign-injection instruction . . . . .	245
16.5.29	FSGNJN.D: a double-precision floating-point negate sign-injection instruction . . . . .	245
16.5.30	FSGNJX.D: a double-precision floating-point XOR sign-injection instruction . . . . .	246
16.5.31	FSQRT.D: a double-precision floating-point square-root instruction . . . . .	246
16.5.32	FSUB.D: a double-precision floating-point subtract instruction . . . . .	247
16.6	Appendix A-6 C Instructions . . . . .	248
16.6.1	C.ADD: a signed add instruction . . . . .	248
16.6.2	C.ADDI: a signed add immediate instruction . . . . .	249
16.6.3	C.ADDIW: an add immediate instruction that operates on the lower 32 bits . . . . .	249
16.6.4	C.ADDI4SPN: an instruction that adds an immediate scaled by 4 to the stack pointer . . . . .	250
16.6.5	C.ADDI16SP: an instruction that adds an immediate scaled by 16 to the stack pointer . . . . .	251
16.6.6	C.ADDW: a signed add instruction that operates on the lower 32 bits . . . . .	251
16.6.7	C.AND: a bitwise AND instruction . . . . .	252
16.6.8	C.ANDI: an immediate bitwise AND instruction . . . . .	253
16.6.9	C.BEQZ: a branch-if-equal-to-zero instruction . . . . .	254
16.6.10	C.BNEZ: a branch-if-not-equal-to-zero instruction . . . . .	255
16.6.11	C.EBREAK: a break instruction . . . . .	256
16.6.12	C.FLD: a floating-point load doubleword instruction . . . . .	256
16.6.13	C.FLDSP: a floating-point doubleword load stack instruction . . . . .	257

16.6.14	C.FSD: a floating-point store doubleword instruction . . . . .	258
16.6.15	C.FSDSP: a floating-point store doubleword stack pointer instruction . . . . .	259
16.6.16	C.J: a unconditional jump instruction . . . . .	259
16.6.17	C.JALR: a jump and link register instruction . . . . .	260
16.6.18	C.JR: a jump register instruction . . . . .	260
16.6.19	C.LD: a load doubleword instruction . . . . .	261
16.6.20	C.LDSP: a load doubleword instruction . . . . .	262
16.6.21	C.LI: a load immediate instruction . . . . .	262
16.6.22	C.LUI: a load upper immediate instruction . . . . .	263
16.6.23	C.LW: a load word instruction . . . . .	264
16.6.24	C.LWSP: a load word stack pointer instruction . . . . .	264
16.6.25	C.MV: an instruction that copies the value in rs to rd . . . . .	265
16.6.26	C.NOP: a no-operation instruction . . . . .	266
16.6.27	C.OR: a bitwise OR instruction . . . . .	266
16.6.28	C.SD: a store doubleword instruction . . . . .	267
16.6.29	C.SDSP: a store doubleword stack pointer instruction . . . . .	268
16.6.30	C.SLLI: an immediate logical left shift instruction . . . . .	268
16.6.31	C.SRAI: a right shift arithmetic immediate instruction . . . . .	269
16.6.32	C.SRLI: an immediate right shift instruction . . . . .	270
16.6.33	C.SW: a store word instruction . . . . .	270
16.6.34	C.SWSP: a store word stack pointer instruction . . . . .	271
16.6.35	C.SUB: a signed subtract instruction . . . . .	272
16.6.36	C.SUBW: a signed subtract instruction that operates on the lower 32 bits . . . . .	273
16.6.37	C.XOR: a bitwise XOR instruction . . . . .	273
16.7	Appendix A-8 Pseudo instructions . . . . .	274

## **17 Appendix B T-Head Extended Instructions 278**

17.1	Appendix B-1 Cache instructions . . . . .	278
17.1.1	DCACHE.CALL: an instruction that clears all dirty page table entries in the D-Cache	278
17.1.2	DCACHE.CIALL: an instruction that clears all dirty page table entries in the D-Cache and invalidates the D-Cache . . . . .	279
17.1.3	DCACHE.CIPA: clears dirty page table entries that match the specified physical addresses from the D-Cache and invalidates the the D-Cache . . . . .	280
17.1.4	DCACHE.CISW: an instruction that clears dirty page table entries in the D-Cache based on the specified way and set and invalidates the D-Cache . . . . .	280
17.1.5	DCACHE.CIVA: an instruction that clears dirty page table entries that match the specified virtual addresses in the D-Cache and invalidates the D-Cache . . . . .	281
17.1.6	DCACHE.CPA: an instruction that clears dirty page table entries that match the specified physical addresses from the D-Cache . . . . .	282
17.1.7	DCACHE.CPAL1: an instruction that clears dirty page table entries that match the specified physical addresses from the L1 D-Cache . . . . .	282

17.1.8	DCACHE.CVA: an instruction that clears dirty page table entries that match the specified virtual addresses in the D-Cache . . . . .	283
17.1.9	DCACHE.CVAL1: an instruction that clears dirty page table entries that match the specified virtual addresses in the L1 D-Cache . . . . .	284
17.1.10	DCACHE.IPA: an instruction that invalidates page table entries that match the specified physical addresses in the D-Cache . . . . .	284
17.1.11	DCACHE.ISW: an instruction that invalidates page table entries in the D-Cache based on the specified way and set and invalidates the D-Cache . . . . .	285
17.1.12	DCACHE.IVA: an instruction that invalidates the D-Cache based on the specified virtual address . . . . .	286
17.1.13	DCACHE.IALL: an instruction that invalidates all page table entries in the D-Cache.	286
17.1.14	ICACHE.IALL: an instruction that invalidates all page table entries in the I-Cache .	287
17.1.15	ICACHE.IALLS: an instruction that invalidates all page table entries in the I-Cache through broadcasting . . . . .	287
17.1.16	ICACHE.IPA: an instruction that invalidates page table entries that match the specified physical addresses in the I-Cache . . . . .	288
17.1.17	ICACHE.IVA: an instruction that invalidates page table entries that match the specified virtual addresses in the I-Cache . . . . .	289
17.1.18	L2CACHE.CALL: an instruction that clears all dirty page table entries in the L2 Cache	289
17.1.19	L2CACHE.CIALL: an instruction that clears all dirty page table entries in the L2 Cache and invalidates the L2 Cache . . . . .	290
17.1.20	L2CACHE.IALL: an instruction that invalidates the L2 Cache . . . . .	290
17.1.21	DCACHE.CSW: an instruction that clears dirty page table entries in the D-Cache based on the specified set and way . . . . .	291
17.2	Appendix B-2 Multi-core synchronization instructions . . . . .	292
17.2.1	SFENCE.VMAS: a broadcast instruction that synchronizes the virtual memory address	292
17.2.2	SYNC: an instruction that performs the synchronization operation . . . . .	293
17.2.3	SYNC.I: an instruction that synchronizes the clearing operation. . . . .	293
17.2.4	SYNC.IS: a broadcast instruction that synchronizes the clearing operation . . . . .	294
17.2.5	SYNC.S: a broadcast instruction that performs a synchronization operation . . . . .	294
17.3	Appendix B-3 Arithmetic operation instructions . . . . .	295
17.3.1	ADDSSL: an add register instruction that shifts registers . . . . .	295
17.3.2	MULA: a multiply-add instruction . . . . .	295
17.3.3	MULAH: a multiply-add instruction that operates on the lower 16 bits . . . . .	296
17.3.4	MULAW: a multiply-add instruction that operates on the lower 32 bits . . . . .	296
17.3.5	MULS: a multiply-subtract instruction . . . . .	297
17.3.6	MULSH: a multiply-subtract instruction that operates on the lower 16 bits . . . . .	297
17.3.7	MULSW: a multiply-subtract instruction that operates on the lower 32 bits . . . . .	298
17.3.8	MVEQZ: an instruction that sends a message when the register is 0 . . . . .	298
17.3.9	MVNEZ: an instruction that sends a message when the register is not 0 . . . . .	299
17.3.10	SRRI: an instruction that implements a cyclic right shift operation on a linked list .	299



17.3.11	SRRIW: an instruction that implements a cyclic right shift operation on a linked list of low 32 bits of registers. . . . .	300
17.4	Appendix B-4 Bitwise operation instructions . . . . .	300
17.4.1	EXT: a signed extension instruction that extracts consecutive bits of a register . . .	300
17.4.2	EXTU: a zero extension instruction that extracts consecutive bits of a register . . .	301
17.4.3	FF0: an instruction that finds the first bit with the value of 0 in a register . . . . .	301
17.4.4	FF1: an instruction that finds the bit with the value of 1 . . . . .	302
17.4.5	REV: an instruction that reverses the byte order in a word stored in the register . .	302
17.4.6	RE VW: an instruction that reverses the byte order in a low 32-bit word . . . . .	303
17.4.7	TST: an instruction that tests bits with the value of 0 . . . . .	303
17.4.8	TSTNBZ: an instruction that tests bytes with the value of 0 . . . . .	304
17.5	Appendix B-5 Storage instructions . . . . .	305
17.5.1	FLRD: a load doubleword instruction that shifts floating-point registers . . . . .	305
17.5.2	FLRW: a load word instruction that shifts floating-point registers . . . . .	305
17.5.3	FLURD: a load doubleword instruction that shifts low 32 bits of floating-point registers	306
17.5.4	FLURW: a load word instruction that shifts low 32 bits of floating-point registers . .	306
17.5.5	FSRD: a store doubleword instruction that shifts floating-point registers . . . . .	307
17.5.6	FSRW: a store word instruction that shifts floating-point registers. . . . .	307
17.5.7	FSURD: a store doubleword instruction that shifts low 32 bits of floating-point registers	308
17.5.8	FSURW: a store word instruction that shifts low 32 bits of floating-point registers . .	308
17.5.9	LBIA: a base-address auto-increment instruction that extends signed bits and loads bytes . . . . .	309
17.5.10	LBIB: a load byte instruction that auto-increments the base address and extends signed bits . . . . .	310
17.5.11	LBUIA: a base-address auto-increment instruction that extends zero bits and loads bytes . . . . .	310
17.5.12	LBUIB: a load byte instruction that auto-increments the base address and extends zero bits . . . . .	311
17.5.13	LDD: an instruction that loads double registers . . . . .	311
17.5.14	LDIA: a base-address auto-increment instruction that loads doublewords and extends signed bits . . . . .	312
17.5.15	LDIB: a load doubleword instruction that auto-increments the base address and extends the signed bits . . . . .	312
17.5.16	LHIA: a base-address auto-increment instruction that loads halfwords and extends signed bits . . . . .	313
17.5.17	LHIB: a load halfword instruction that auto-increments the base address and extends signed bits . . . . .	314
17.5.18	LHUIA: a base-address auto-increment instruction that extends zero bits and loads halfwords . . . . .	314
17.5.19	LHUIB: a load halfword instruction that auto-increments the base address and extends zero bits . . . . .	315
17.5.20	LRB: a load byte instruction that shifts registers and extends signed bits . . . . .	315



17.5.21	LRBU: a load byte instruction that shifts registers and extends zero bits . . . . .	316
17.5.22	LRD: a load doubleword instruction that shifts registers . . . . .	316
17.5.23	LRH: a load halfword instruction that shifts registers and extends signed bits . . . .	316
17.5.24	LRHU: a load halfword instruction that shifts registers and extends zero bits . . . .	317
17.5.25	LRW: a load word instruction that shifts registers and extends signed bits . . . . .	317
17.5.26	LRWU: a load word instruction that shifts registers and extends zero bits . . . . .	318
17.5.27	LURB: a load byte instruction that shifts low 32 bits of registers and extends signed bits . . . . .	318
17.5.28	LURBU: a load byte instruction that shifts low 32 bits of registers and extends zero bits . . . . .	319
17.5.29	LURD: a load doubleword instruction that shifts low 32 bits of registers . . . . .	319
17.5.30	LURH: a load halfword instruction that shifts low 32 bits of registers and extends signed bits . . . . .	320
17.5.31	LURHU: a load halfword instruction that shifts low 32 bits of registers and extends zero bits . . . . .	320
17.5.32	LURW: a load word instruction that shifts low 32 bits of registers and extends signed bits . . . . .	321
17.5.33	LURWU: a load word instruction that shifts 32 bits of registers and extends zero bits	321
17.5.34	LWD: a load word instruction that loads double registers and extends signed bits . .	322
17.5.35	LWIA: a base-address auto-increment instruction that extends signed bits and loads words . . . . .	323
17.5.36	LWIB: a load word instruction that auto-increments the base address and extends signed bits . . . . .	323
17.5.37	LWUD: a load word instruction that loads double registers and extends zero bits . .	324
17.5.38	LWUIA: a base-address auto-increment instruction that extends zero bits and loads words . . . . .	324
17.5.39	LWUIB: a load word instruction that auto-increments the base address and extends zero bits . . . . .	325
17.5.40	SBIA: a base-address auto-increment instruction that stores bytes . . . . .	325
17.5.41	SBIB: a store byte instruction that auto-increments the base address . . . . .	326
17.5.42	SDD: an instruction that stores double registers . . . . .	326
17.5.43	SDIA: a base-address auto-increment instruction that stores doublewords . . . . .	327
17.5.44	SDIB: a store doubleword instruction that auto-increments the base address . . . . .	327
17.5.45	SHIA: a base-address auto-increment instruction that stores halfwords . . . . .	328
17.5.46	SHIB: a store halfword instruction that auto-increments the base address . . . . .	328
17.5.47	SRB: a store byte instruction that shifts registers . . . . .	329
17.5.48	SRD: a store doubleword instruction that shifts registers . . . . .	329
17.5.49	SRH: a store halfword instruction that shifts registers . . . . .	329
17.5.50	SRW: a store word instruction that shifts registers . . . . .	330
17.5.51	SURB: a store byte instruction that shifts low 32 bits of registers . . . . .	330
17.5.52	SURD: a store doubleword instruction that shifts low 32 bits of registers . . . . .	331
17.5.53	SURH: a store halfword instruction that shifts low 32 bits of registers . . . . .	331

17.5.54	SURW: a store word instruction that shifts low 32 bits of registers . . . . .	332
17.5.55	SWIA: a base-address auto-increment instruction that stores words . . . . .	332
17.5.56	SWIB: a store word instruction that auto-increments the base address . . . . .	333
17.5.57	SWD: an instruction that stores the low 32 bits of double registers . . . . .	333
17.6	Appendix B-6 Half-precision floating-point instructions . . . . .	334
17.6.1	FADD.H: a half-precision floating-point add instruction . . . . .	334
17.6.2	FCLASS.H: a half-precision floating-point classification instruction . . . . .	335
17.6.3	FCVT.H.L: an instruction that converts a signed long integer into a half-precision floating-point number . . . . .	336
17.6.4	FCVT.H.LU: an instruction that converts an unsigned long integer into a half-precision floating-point number . . . . .	337
17.6.5	FCVT.H.S: an instruction that converts single precision floating-point data to half-precision floating-point data . . . . .	338
17.6.6	FCVT.H.W: an instruction that converts a signed integer into a half-precision floating-point number . . . . .	339
17.6.7	FCVT.H.WU: an instruction that converts an unsigned integer into a half-precision floating-point number . . . . .	340
17.6.8	FCVT.L.H: an instruction that converts a half-precision floating-point number to a signed long integer . . . . .	341
17.6.9	FCVT.LU.H: an instruction that converts a half-precision floating-point number to an unsigned long integer . . . . .	342
17.6.10	FCVT.S.H: an instruction that converts half-precision floating-point data to single precision floating-point data . . . . .	343
17.6.11	FCVT.W.H: an instruction that converts a half-precision floating-point number to a signed integer . . . . .	343
17.6.12	FCVT.WU.H: an instruction that converts a half-precision floating-point number to an unsigned integer . . . . .	344
17.6.13	FDIV.H: a half-precision floating-point division instruction . . . . .	345
17.6.14	FEQ.H: an equal instruction that compares two half-precision numbers . . . . .	346
17.6.15	FLE.H: a less than or equal to instruction that compares two half-precision floating-point numbers . . . . .	347
17.6.16	FLH: an instruction that loads half-precision floating-point data . . . . .	347
17.6.17	FLT.H: a less than instruction that compares two half-precision floating-point numbers	348
17.6.18	FMADD.H: a half-precision floating-point multiply-add instruction . . . . .	349
17.6.19	FMAX.H: a half-precision floating-point maximum instruction . . . . .	350
17.6.20	FMIN.H: a half-precision floating-point minimum instruction . . . . .	350
17.6.21	FMSUB.H: a half-precision floating-point multiply-subtract instruction . . . . .	351
17.6.22	FMUL.H: a half-precision floating-point multiply instruction . . . . .	352
17.6.23	FMV.H.X: a half-precision floating-point write transmit instruction . . . . .	353
17.6.24	FMV.X.H: a transmission instruction that reads half-precision floating-point registers	353
17.6.25	FNMADD.H: a half-precision floating-point negate-(multiply-add) instruction . . . .	354
17.6.26	FNMSUB.H: a half-precision floating-point negate-(multiply-subtract) instruction . .	355

17.6.27	FSGNJ.H: a half-precision floating-point sign-injection instruction . . . . .	356
17.6.28	FSGNJN.H: a half-precision floating-point sign-injection negate instruction . . . . .	356
17.6.29	FSGNJX.H: a half-precision floating-point sign-injection XOR instruction . . . . .	357
17.6.30	FSH: an instruction that stores half-precision floating point numbers . . . . .	357
17.6.31	FSQRT.H: an instruction that calculates the square root of the half-precision floating-point number . . . . .	358
17.6.32	FSUB.H: a half-precision floating-point subtract instruction . . . . .	359
<b>18</b>	<b>Appendix C Control Registers</b>	<b>361</b>
18.1	Appendix C-1 M-mode control registers . . . . .	361
18.1.1	M-mode information register group . . . . .	361
18.1.1.1	Machine vendor ID register (mvendorid) . . . . .	361
18.1.1.2	Machine architecture ID register (marchid) . . . . .	361
18.1.1.3	Machine implementation ID register (mimpid) . . . . .	362
18.1.1.4	Machine hart ID register (mhartid) . . . . .	362
18.1.2	M-mode exception configuration register group . . . . .	362
18.1.2.1	Machine status register (mstatus) . . . . .	362
18.1.2.2	M-mode instruction set architecture register (misa) . . . . .	365
18.1.2.3	M-mode exception downgrade control register (medeleg) . . . . .	365
18.1.2.4	M-mode interrupt downgrade control register (mideleg) . . . . .	366
18.1.2.5	M-mode interrupt-enable register (mie) . . . . .	366
18.1.2.6	M-mode trap vector base address register (mtvec) . . . . .	367
18.1.2.7	M-mode counter access enable register (mcounteren) . . . . .	368
18.1.3	M-mode exception handling register group . . . . .	368
18.1.3.1	M-mode scratch register (mscratch) . . . . .	368
18.1.3.2	M-mode exception program counter register (mepc) . . . . .	368
18.1.3.3	M-mode cause register (mcause) . . . . .	368
18.1.3.4	M-mode interrupt-pending register (mip) . . . . .	369
18.1.4	M-mode memory protection registers . . . . .	370
18.1.4.1	Physical memory protection configuration register (pmpcfg) . . . . .	370
18.1.4.2	Physical memory address register (pmpaddr) . . . . .	371
18.1.5	M-mode counter registers . . . . .	371
18.1.5.1	M-mode cycle counter (mcycle) . . . . .	371
18.1.5.2	M-mode instructions-retired counter (minstret) . . . . .	371
18.1.5.3	M-mode event counter (mhpmcountern) . . . . .	371
18.1.6	M-mode counter configuration registers . . . . .	371
18.1.6.1	M-mode event selector (mhpmeventn) . . . . .	372
18.1.7	M-mode CPU control and status extension registers . . . . .	372
18.1.7.1	M-mode extension status register (mxstatus) . . . . .	372
18.1.7.2	M-mode hardware configuration register (mhcr) . . . . .	374
18.1.7.3	M-mode hardware operation register (mcor) . . . . .	375
18.1.7.4	M-mode L2 Cache control register (mccr2) . . . . .	377

18.1.7.5	M-mode implicit operation register (mhint)	378
18.1.7.6	M-mode reset vector base address register (mrvbr)	380
18.1.7.7	S-mode counter write enable register (mcounterwen)	381
18.1.7.8	M-mode event interrupt enable register (mcounterinten)	381
18.1.7.9	M-mode event overflow mark register (mcounteren)	382
18.1.8	M-mode cache access extension registers	382
18.1.8.1	M-mode cache instruction register (mcins)	382
18.1.8.2	M-mode cache access index register (mcindex)	383
18.1.8.3	M-mode cache data register (mcdata0/1)	384
18.1.9	M-mode CPU model registers	384
18.1.9.1	M-mode CPU model register (mcpuid)	384
18.1.9.2	On-chip bus base address register (mapbaddr)	385
18.1.10	Multi-core extension registers	385
18.1.10.1	Snoop listening enable register (msmpr)	385
18.2	Appendix C-2 S-mode control registers	385
18.2.1	S-mode exception configuration registers	385
18.2.1.1	S-mode status register (sstatus)	385
18.2.1.2	S-mode interrupt-enable register (sie)	386
18.2.1.3	S-mode trap vector base address register (stvec)	386
18.2.1.4	S-mode counter access enable register (scounteren)	387
18.2.2	S-mode exception handling registers	387
18.2.2.1	S-mode scratch register (sscratch)	387
18.2.2.2	S-mode exception program counter register (sepc)	387
18.2.2.3	S-mode cause register (scause)	387
18.2.2.4	S-mode interrupt-pending register (sip)	388
18.2.3	S-mode address translation registers	388
18.2.3.1	S-mode address translation register (satp)	388
18.2.4	S-mode CPU control and status extension registers	388
18.2.4.1	S-mode extension status register (sxstatus)	388
18.2.4.2	S-mode hardware control register (shcr)	389
18.2.4.3	S-mode event overflow interrupt enable register (scounterinten)	389
18.2.4.4	S-mode event overflow mark register (scounterof)	389
18.2.4.5	S-mode cycle counter (scycle)	389
18.2.4.6	S-mode instructions-retired counter (sinstret)	389
18.2.4.7	S-mode event counter (shpmcountern)	390
18.2.5	S-mode MMU extension register	390
18.2.5.1	S-mode MMU control register (smcir)	390
18.2.5.2	S-mode MMU control register (smir)	390
18.2.5.3	S-mode MMU control register (smeh)	390
18.2.5.4	S-mode MMU control register (smel)	390
18.3	Appendix C-3 U-mode control registers	390
18.3.1	U-mode floating-point control registers	391

18.3.1.1	Floating-point accrued exceptions register (fflags)	391
18.3.1.2	Floating-point dynamic rounding mode register (frm)	391
18.3.1.3	Floating-point control and status register (fcsr)	391
18.3.2	U-mode counter/timer registers	392
18.3.2.1	User cycle register (cycle)	392
18.3.2.2	U-mode timer register (time)	392
18.3.2.3	User instructions-retired counter (instret)	393
18.3.2.4	User event counter (hpmcountern)	393
18.3.3	U-mode floating-point extension control registers	393
18.3.3.1	U-mode floating-point extension control register (fxcr)	393

## 1.1 Introduction

C910/C920MP is a high-performance 64-bit multi-core CPU built on the RISC-V architecture. It is oriented to edge computing that requires high performance. For example, it can be applied to edge servers, edge computing cards, advanced machine vision, advanced video surveillance, self-driving, mobile smart terminals, and 5G base stations. C910/C920MP adopts a homogeneous multi-core architecture and supports dual cores. Each C910/C920 core runs on a microsystem architecture developed by Alibaba Cloud and is optimized for high performance. High-performance technologies are introduced, such as a superscalar architecture with a 3-to-8 line decoder and multi-channel data prefetch. In addition, the C910/C920 core performs real-time detection and shuts down internal idle function modules to reduce dynamic power consumption of the CPU.

## 1.2 Features

### 1.2.1 Architectural features of C910/C920MP

- Homogeneous multi-core architecture and configuration of 1 to 4 cores supported;
- Power-off for each core or the entire cluster supported;
- One AXI4 master interface and a 128-bit bus supported;
- Two levels of caches provided: L1 cache running on the Harvard architecture and L2 shared cache;
- L1 instruction/data cache size: 64 KB, with a cache line size of 64 bytes;

- The Modified, Exclusive, Shared, Invalid (MESI) cache coherence protocol supported for L1 cache, and the Modified, Owned, Exclusive, Shared, Invalid (MOESI) cache coherence protocol supported for L2 cache;
- L2 cache: a 16-way set-associative cache sized 1 MB and with a cache line size of 64 bytes;
- Core local interrupt (CLINT) controller and platform-level interrupt controller (PLIC) supported;
- Timers supported;
- Custom multi-core debug frameworks with RISC-V-compatible interfaces supported.

### 1.2.2 Features of the C910/C920 core

- RISC-V 64GC instruction architecture;
- Little-endian mode supported;
- 9-stage to 12-stage pipelining architecture;
- Superscalar architecture with a 3-to-8 line decoder, fully transparent to software;
- In-order fetch, out-of-order issue, out-of-order completion, and in-order retirement;
- Two-level translation lookaside buffer (TLB) memory management units for virtual/physical address translation and memory management;
- I-Cache/D-Cache size: 64 KB, with a cache line size of 64 bytes;
- Instruction prefetch, and automatic detection and dynamic startup of hardware;
- Low-power access technology with I-Cache way prediction;
- Low-power execution technology with short-loop buffer;
- 64 KB two-level multi-way parallel branch predictor;
- Branch target buffer with 1024/2048 entries;
- 12-layer hardware return address stack supported;
- Indirect branch predictor with 256 entries;
- Non-blocking issue and speculative execution;
- Renaming technology based on physical registers;
- 0-latency move instructions supported;
- Dual issue and full out-of-order execution for load/store instructions;
- Concurrent bus access for up to 8 read requests and 8 write requests;
- Write combining supported;
- Strided 8-channel hardware prefetch supported;

- Half-precision, single-precision floating-point units supported.

## 1.3 Configurations

C910/C920MP adopts the following configurations:

- Number of cores: 2
- L1 I-Cache: 64KB
- L1 D-Cache: 64KB
- L2 Cache: 1MB
- Master Interface: AXI4
- Number of external interrupts: 144

## 1.4 XuanTie extended architecture

C910/C920 is compatible with XuanTie C-series extended architecture 1.0, which provides extensions in the following aspects:

- Operation instructions: C910/C920 improves operation capabilities with integer, floating-point, and load/store instructions, well supplementing the RISC-V base instruction sets.
- Cache operations: C910/C920 enables you to easily maintain caches to improve cache efficiency.
- Memory model: C910/C920 manages address attributes efficiently to improve memory access efficiency.
- Control registers: C910/C920 extends the features of control registers based on the standard RISC-V architecture.
- Multi-core synchronization instructions: C910/C920 adopts multi-core synchronization instructions to keep multi-core consistency.
- PLIC: C910/C920 is integrated with a built-in PLIC.

## 1.5 Version compatibility

C910/C920 is compatible with the following RISC-V standard versions:

- *The RISC-V Instruction Set Manual, Volume I: RISC-V User-Level ISA, Version 2.2.*
- *The RISC-V Instruction Set Manual, Volume II: Privileged Architecture, Version 1.10.*
- The mcountinhibit register in the *RISC-V Instruction Set Manual, Volume II: Privileged Architecture, Version 20190125-Public-Review-draft* is added to the performance monitoring unit (PMU).



## 1.6 Naming conventions

### 1.6.1 Terms

- **Logic 1** :The level value corresponding to the Boolean logic value TRUE.
- **Logic 0** :The level value corresponding to the Boolean logic value FALSE.
- **Set** :The action of setting one or more bits to the level value corresponding to logic 1.
- **Clear** :The action of setting one or more bits to the level value corresponding to logic 0.
- **Reserved bit** :A bit reserved for feature extension. The value of a reserved bit is 0 unless otherwise specified.
- **Signal** :An electrical value used to transfer information based on its state or state transition.
- **Pin** :An external electrical and physical connection. Multiple signals can connect to one pin.
- **Enable** :The action of switching a discrete signal to a valid state:
  - Switch a valid low-level signal from a high level to a low level.
  - Switch a valid high-level signal from a low level to a high level.
- **Disable** :The action of switching the state of an enabled signal:
  - Switch a valid low-level signal from a low level to a high level.
  - Switch a valid high-level signal from a high level to a low level.
- **LSB** :The least significant bit. **MSB**: The most significant bit.
- **Signal, bit field, and control bit**: Expressed based on a general rule.
- **Identifier followed by a value range**: Indicates a group of signals from the most significant bit to the least significant bit.

For example, `addr[4:0]` indicates a group of address buses, where `addr[4]` indicates the most significant bit, and `addr[0]` indicates the least significant bit.

- **Single identifier**: Indicates a single signal.

For example, `pad_cpu_rst_b` indicates a single signal.

In some cases, an identifier followed by a number is used to express a specific meaning. For example, `addr15` indicates the 16th bit of a group of buses.

## 2.1 Structure

The structure of C910/C920MP is shown in [Fig. 2.1](#) .

## 2.2 In-core subsystems

C910/C920 consists of the following in-core subsystems: instruction fetch unit (IFU), instruction decoding unit (IDU), integer unit (IU), floating-point unit (FPU), load/store unit (LSU), retirement unit (RTU), memory management unit (MMU), and physical memory protection (PMP) unit.

### 2.2.1 IFU

The IFU can fetch and parallel process up to eight instructions at a time. It improves access efficiency with a variety of technologies, for example, I-Cache way prediction, instruction registers, loop acceleration buffers, and direct/indirect branch prediction. The IFU features low power consumption, high branch prediction accuracy, and high prefetch efficiency.

### 2.2.2 IDU

The IDU can decode up to three instructions and detect data correlation at a time. The IDU detects data correlation between instructions by using the physical register renaming technology, and sends the

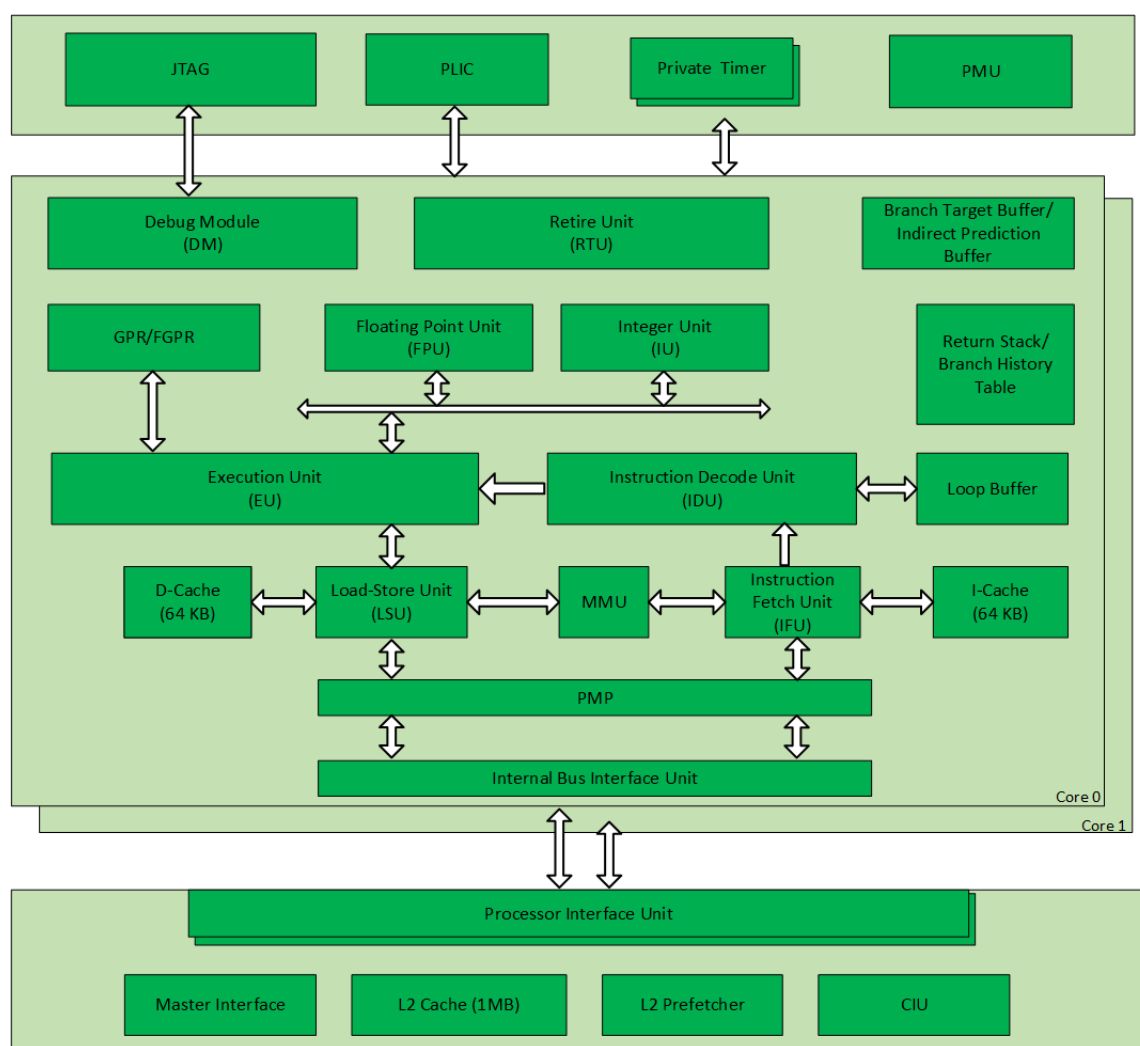


Fig. 2.1: C910/C920MP structure

instructions out of order to the next-level pipeline for execution. The IDU supports out-of-order scheduling and distribution of instructions. It speculatively issues instructions to mitigate performance loss caused by data correlation.

### 2.2.3 Execution units

Execution units include IUs and FPUs.

IUs include the arithmetic logic unit (ALU), multiplication (MULT) unit, division (DIV) unit, and branch/jump unit (BJU). The ALU performs 64-bit integer operations. The MULT unit supports 16\*16, 32\*32, and 64\*64 integer multiplication. The DIV unit is designed based on the radix-16 SRT algorithm. Its execution cycle varies with operands. The BJU can correct branch prediction errors within one cycle.

FPUs include the floating-point arithmetic logic unit (FALU), floating-point fused multiply-add unit (FMAU), and floating-point divide and square root unit (FDSU). FPUs support half-precision, single-precision operations. The FALU performs addition, subtraction, comparison, conversion, register data transmission, sign-injection, and classification operations. The FMAU performs common multiply and fused multiply-add operations. The FDSU performs floating-point divide and square root operations.

### 2.2.4 LSU

The LSU supports dual issue for scalar store/load instructions, single issue for vector store/load instructions, and full out-of-order execution for all store/load instructions. The LSU supports non-blocking access to caches. It supports byte, halfword, word, doubleword, and quadword store/load instructions, and supports sign/zero extension for byte and halfword load instructions. Store/load instructions can be executed in a pipeline so that only one data entry is accessed per cycle. The LSU supports 8-channel hardware prefetch. It can transfer data to the L1 D-Cache in advance. If the D-Cache is absent, the LSU supports parallel bus access.

### 2.2.5 RTU

The RTU consists of a re-order buffer and a physical register stack. The re-order buffer controls out-of-order recycling and in-order retirement of instructions. The physical register stack controls out-of-order recycling and transfer of results. The RTU improves the instruction retirement efficiency through parallel recycling and fast retirement of instructions. The RTU supports parallel retirement of up to three instructions per clock cycle and implements precise exceptions.

### 2.2.6 MMU

The MMU translates 39-bit virtual addresses to 40-bit physical addresses in compliance with the RISC-V SV39 standard. The MMU of C910/C920 provides extended software writeback methods and address attributes based on the hardware writeback criteria defined in SV39.

For more information, see [Memory Model](#) .

### 2.2.7 PMP

The PMP unit complies with the RISC-V standard and supports 8 or 16 entries, but does not support the NA4 mode. The minimum granularity supported by the PMP unit is 4 KB.

For more information, see [Memory Model](#) .

## 2.3 Multi-core subsystems

C910/C920 consists of the following multi-core subsystems: consistency interface unit (CIU), L2 cache, master device interface unit, platform-level interrupt controller (PLIC), timer, and custom multi-core single-port debug framework.

### 2.3.1 CIU

The CIU ensures data coherence between L1 D-Caches based on the MESI protocol. Two listening buffers are configured to parallel handle multiple listening requests, to fully utilize the listening bandwidth. The CIU adopts an efficient data bypassing mechanism. When a listening request hits an L1 D-Cache under listening, data is directly bypassed to the request initiation core. In addition, the CIU supports broadcasting of invalid TLB/I-Cache requests. This reduces the software costs of maintaining data coherence between TLB/I-Cache and D-Cache.

### 2.3.2 L2 cache

The L2 cache is tightly coupled to the CIU for synchronous access with L1 D-Caches. The L2 cache adopts a block-based pipelining architecture and can parallel handle two access requests within one cycle. It supports a maximum access bandwidth of 1024 bps. The operating frequency of the L2 cache is the same as that of C910/C920. The tag and data RAM access latency can be configured by using software.

### 2.3.3 Master device interface unit

The master device interface unit supports the AXI4 protocol and address access by keyword priority, and can work under different system clock to CPU clock ratios, for example, 1:1, 1:2, 1:3, 1:4, 1:5, 1:6, 1:7, and 1:8.

### 2.3.4 PLIC

The PLIC controls sampling and distribution of 1023 external interrupt sources. It supports level and pulse interrupts. You can set 32 interrupt priorities.

For more information, see *Interrupt Controllers*.

### 2.3.5 Timer

The multi-core system provides one shared 64-bit system timer. Each core has a private timer comparison value register. Values of the system timer are collected and compared with those in the private timer comparison value register to generate timer signals.

For more information, see *Interrupt Controllers*.

## 2.4 Interface overview

C910/C920 provides the following interfaces by feature: clock reset signal, bus system, interrupt system, debug system, low power system, DFT system, and CPU running monitoring signal. For more information, see Fig. 2.2.

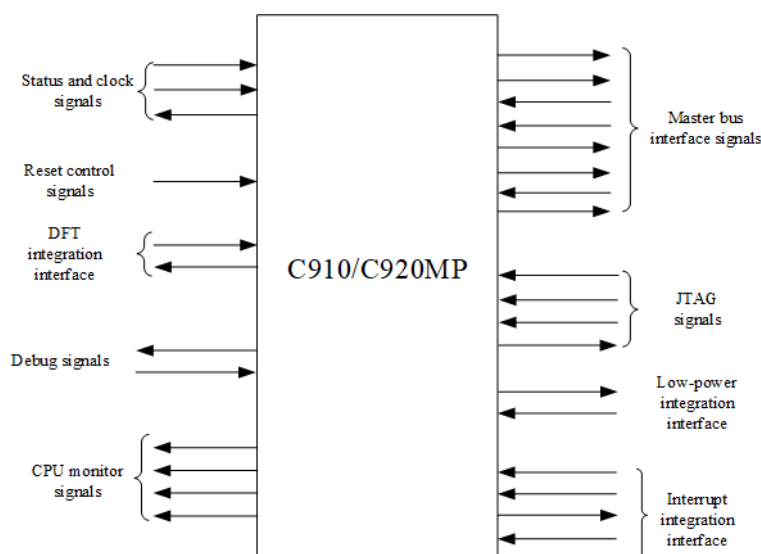


Fig. 2.2: C910/C920MP interfaces

This section describes the instruction sets implemented in C910/C920: RV base instruction sets and XuanTie extended instruction sets.

### 3.1 RV base instruction sets

#### 3.1.1 Integer instruction set (RV64I)

The integer instruction set includes instructions of the following types by feature:

- Add/Subtract instructions
- Logical operation instructions
- Shift instructions
- Compare instructions
- Data transmission instructions
- Branch and jump instructions
- Memory access instructions
- Control register operation instructions
- Low power instructions
- Exception-return instructions

- Special functional instructions

Table 3.1: RV64I instructions

Instruction	Description	Execution latency
<b>Add/Subtract instructions</b>		
ADD	A signed add instruction	1
ADDW	A signed add instruction that operates on the lower 32 bits	1
ADDI	A signed add immediate instruction	1
ADDIW	A signed add immediate instruction that operates on the lower 32 bits	1
SUB	A signed subtract instruction	1
SUBW	A signed subtract instruction that operates on the lower 32 bits	1
<b>Logic operation instructions</b>		
AND	A bitwise AND instruction.	1
ANDI	An immediate bitwise AND instruction	1
OR	A bitwise OR instruction	1
ORI	An immediate bitwise OR instruction	1
XOR	A bitwise XOR instruction.	1
XORI	An immediate bitwise XOR instruction	1
<b>Shift instructions</b>		
SLL	A logical left shift instruction	1
SLLW	A word logical left shift instruction that operates on the lower 32 bits	1
SLLI	An immediate logical left shift instruction	1
SLLIW	An immediate logical left shift instruction that operates on the lower 32 bits	1
SRL	A logical right shift instruction	1
SRLW	A logical right shift instruction that operates on the lower 32 bits	1
SRLI	An immediate logical right shift instruction	1
SRLIW	An immediate logical right shift instruction that operates on the lower 32 bits	1
SRA	An arithmetic right shift instruction	1
SRAW	An arithmetic right shift instruction that operates on the lower 32 bits	1
SRAI	An immediate arithmetic right shift instruction	1
SRAIW	An immediate arithmetic right shift instruction that operates on the lower 32 bits	1

Continued on next page



Table 3.1 – continued from previous page

Instruction	Description	Execution latency
<b>Compare instructions</b>		
SLT	A signed set-if-less-than instruction	1
SLTU	An unsigned set-if-less-than instruction	1
SLTI	A signed set-if-less-than-immediate instruction	1
SLTIU	An unsigned set-if-less-than-immediate instruction	1
<b>Data transmission instructions</b>		
LUI	A load upper immediate instruction	1
AUIPC	An add upper immediate to PC instruction	1
<b>Branch and jump instructions</b>		
BEQ	A branch-if-equal instruction	1
BNE	A branch-if-not-equal instruction	1
BLT	A signed branch-if-less-than instruction	1
BGE	A signed branch-if-greater-than-or-equal instruction	1
BLTU	An unsigned branch-if-less-than instruction	1
BGEU	An unsigned branch-if-greater-than-or-equal instruction	1
JAL	An instruction for directly jumping to a subroutine	1
JALR	An instruction for jumping to a subroutine by using an address in a register	1
<b>Memory access instructions</b>		
LB	A sign-extended byte load instruction	WEAK ORDER LOAD: $\geq 3$ STORE: 1 STRONG ORDER Aperiodic
LBU	An unsign-extended byte load instruction	Same as above
LH	A sign-extended halfword load instruction	Same as above
LHU	An unsign-extended halfword load instruction	Same as above
LW	A sign-extended word load instruction	Same as above
LWU	An unsign-extended word load instruction	Same as above
LD	A doubleword load instruction	Same as above
SB	A byte store instruction	Same as above
SH	A halfword store instruction	Same as above
SW	A word store instruction	Same as above
SD	A doubleword store instruction	Same as above
<b>Control register operation instructions</b>		
CSR <sub>RRW</sub>	A move instruction that reads/writes control registers	Blocked Aperiodic

Continued on next page

Table 3.1 – continued from previous page

Instruction	Description	Execution latency
CSRRS	A move instruction for setting control registers	Same as above
CSRRC	A move instruction that clears control register	Same as above
CSRRWI	A move instruction that reads/writes immediates in control registers	Same as above
CSRRSI	A move instruction for setting immediates in control registers	Same as above
CSRRCI	A move instruction that clears immediates in control registers	Same as above
<b>Low power instructions</b>		
WFI	An instruction for entering the low-power standby mode	Aperiodic
<b>Exception-return instructions</b>		
MRET	An instruction for returning from exceptions in machine mode (M-mode)	Block
SRET	An instruction for returning from exceptions in supervisor mode (S-mode)	Same as above
<b>Special functional instructions</b>		
FENCE	A memory synchronization instruction	Aperiodic
FENCE.I	An instruction stream synchronization instruction	Blocked
SFENCE.VMA	A virtual memory synchronization instruction	Same as above
EBREAK	A breakpoint instruction	1
ECALL	An environment call instruction	1

For more information, see [Appendix A-1 I instructions](#).

### 3.1.2 Multiply/Divide instruction set (RV64M)

Table 3.2: RV64M instructions

Name	Description	Execution latency
MUL	A signed multiply instruction	4
MULW	A signed multiply instruction that operates on the lower 32 bits	4
MULH	A signed multiply instruction that extracts upper bits	4
MULHS	A signed-unsigned multiply instruction that extracts upper bits	4
MULHU	An unsigned multiply instruction that extracts upper bits	4
DIV	A signed divide instruction.	3-20
DIVW	A signed divide instruction that operates on the lower 32 bits	3-12
DIVU	An unsigned divide instruction.	3-20
DIVUW	An unsigned divide instruction that operates on the lower 32 bits	3-12
REM	A signed remainder instruction	3-20
REMW	A signed remainder instruction that operates on the lower 32 bits	3-12

Continued on next page

Table 3.2 – continued from previous page

Name	Description	Execution latency
REMU	An unsigned remainder instruction.	3-20
REMUW	An unsigned remainder instruction that operates on the lower 32 bits	3-12

For more information, see [Appendix A-2 M instructions](#).

### 3.1.3 Atomic instruction set (RV64A)

Table 3.3: RV64A instructions

Instruction	Description	Execution latency
LR.W	A word load-reserved instruction.	This instruction is split into multiple atomic instructions for execution.
LR.D	A doubleword load-reserved instruction.	
SC.W	A word store-conditional instruction.	This instruction can be split into atomic instructions for blocked execution, but latency is not allowed.
SC.D	A doubleword store-conditional instruction.	
AMOSWAP.W	An atomic swap instruction that operates on the lower 32 bits.	
AMOSWAP.D	An atomic swap instruction.	
AMOADD.W	An atomic add instruction that operates on the lower 32 bits.	
AMOADD.D	An atomic add instruction.	
AMOXOR.W	An atomic bitwise XOR instruction that operates on the lower 32 bits.	
AMOXOR.D	An atomic bitwise XOR instruction.	
AMOAND.W	An atomic bitwise AND instruction that operates on the lower 32 bits.	
AMOAND.D	An atomic bitwise AND instruction.	
AMOOR.W	An atomic bitwise OR instruction that operates on the lower 32 bits.	
AMOOR.D	An atomic bitwise OR instruction	
AMOMIN.W	An atomic signed MIN instruction that operates on the lower 32 bits.	
AMOMIN.D	An atomic signed MIN instruction	
AMOMAX.W	An atomic signed MAX instruction that operates on the lower 32 bits.	
AMOMAX.D	An atomic signed MAX instruction.	
AMOMINU.W	An atomic unsigned MIN instruction that operates on the lower 32 bits.	
AMOMINU.D	An atomic unsigned MIN instruction.	

Continued on next page

Table 3.3 – continued from previous page

Instruction	Description	Execution latency
AMOMAXU.W	An atomic unsigned MAX instruction that operates on the lower 32 bits.	
AMOMAXU.D	An atomic unsigned MAX instruction.	

For more information, see [Appendix A-3 A instructions](#).

### 3.1.4 Single-precision floating-point instruction set (RV64F)

A single-precision floating-point instruction set includes instructions of the following types by feature:

- Operation instructions
- Sign injection instructions
- Data transmission instructions
- Compare instructions
- Data type conversion instructions
- Memory store instructions
- Floating-point classify instructions

Table 3.4: RV64F instructions

Instruction	Description	Latency
<b>Operation instructions</b>		
FADD.S	A single-precision floating-point add instruction.	3
FSUB.S	A single-precision floating-point subtract instruction.	3
FMUL.S	A single-precision floating-point multiply instruction	4
FMADD.S	A single-precision floating-point multiply-add instruction.	5
FMSUB.S	A single-precision floating-point multiply-subtract instruction.	5
FNMADD.S	A single-precision floating-point negate-(multiply-add) instruction.	5
FNMSUB.S	A single-precision floating-point negate -(multiply-subtract) instruction.	5
FDIV.S	A single-precision floating-point divide instruction.	4-10
FSQRT.S	A single-precision floating-point square-root instruction.	4-10
<b>Sign injection instructions</b>		
FSGNJ.S	A single-precision floating-point sign-injection instruction.	3
FSGNJN.S	A single-precision floating-point negate sign-injection instruction.	3

Continued on next page

Table 3.4 – continued from previous page

Instruction	Description	Latency
FSGNJX.S	A single-precision floating-point sign-injection XOR instruction.	3
<b>Data transmission instructions</b>		
FMV.X.W	A single-precision floating-point read move instruction	1+1 in split execution
FMV.W.X	A single-precision floating-point write move instruction.	1+1 in split execution
<b>Compare instructions</b>		
FMIN.S	A single-precision floating-point MIN instruction.	3
FMAX.S	A single-precision floating-point MAX instruction.	3
FEQ.S	A single-precision floating-point compare equal instruction.	3+1 in split execution
FLT.S	A single-precision floating-point compare less than instruction.	3+1 in split execution
FLE.S	A single-precision floating-point compare less than or equal to instruction.	3+1 in split execution
<b>Data type conversion instructions</b>		
FCVT.W.S	An instruction that converts a single-precision floating-point number into a signed integer.	3+1 in split execution
FCVT.W.U.S	An instruction that converts a single-precision floating-point number into an unsigned integer.	3+1 in split execution
FCVT.S.W	An instruction that converts a signed integer into a single-precision floating-point number.	3+1 in split execution
FCVT.S.WU	An instruction that converts an unsigned integer into a single-precision floating-point number.	3+1 in split execution
FCVT.L.S	An instruction that converts a single-precision floating-point number into a signed long integer.	3+1 in split execution
FCVT.L.U.S	An instruction that converts a single-precision floating-point number into an unsigned long integer.	3+1 in split execution
FCVT.S.L	An instruction that converts a signed long integer into a single-precision floating-point number.	1+3 in split execution
FCVT.S.LU	An instruction that converts an unsigned long integer into a single-precision floating-point number.	1+3 in split execution
<b>Memory store instructions</b>		

Continued on next page

Table 3.4 – continued from previous page

Instruction	Description	Latency
FLW	A single-precision floating-point load instruction.	WEAK ORDER LOAD: $\geq 3$ STORE: 1 STRONG ORDER Aperiodic
FSW	A single-precision floating-point store instruction.	Same as above
<b>Floating-point classify instructions</b>		
FCLASS.S	A single-precision floating-point classify instruction.	1+1

For more information, see [Appendix A-4 F instructions](#).

### 3.1.5 Compressed instruction set (RV64C)

The compressed instruction set includes instructions of the following types by feature:

- Add/Subtract instructions
- Logical operation instructions
- Shift instructions
- Data transmission instructions
- Branch and jump instructions
- Immediate offset access instructions

Table 3.5: RV64C instructions

Instruction	Description	Latency
<b>Add/Subtract instructions</b>		
C.ADD	A signed add instruction	1
C.ADDW	A signed add instruction that operates on the lower 32 bits.	1
C.ADDI	A signed add immediate instruction.	1
C.ADDIW	A signed add immediate instruction that operates on the lower 32 bits.	1
C.SUB	A compressed signed subtract instruction.	1
C.SUBW	A signed subtract instruction that operates on the lower 32 bits	1
C.ADDI16SP	An instruction that adds an immediate scaled by 16 to the stack pointer.	1

Continued on next page

Table 3.5 – continued from previous page

Instruction	Description	Latency
C.ADDI4SPN	An instruction that adds an immediate scaled by 4 to the stack pointer	1
<b>Logic operation instructions</b>		
C.AND	A bitwise AND instruction	1
C.ANDI	An immediate bitwise AND instruction	1
C.OR	A bitwise OR instruction	1
C.XOR	A bitwise XOR instruction	1
<b>Shift instructions</b>		
C.SLLI	An immediate logical left shift instruction.	1
C.SRLI	An immediate logical right shift instruction.	1
C.SRAI	An immediate arithmetic right shift instruction.	1
<b>Data transmission instructions</b>		
C.MV	A data move instruction	1
C.LI	An instruction for moving immediates in the lower bits	1
C.LUI	An instruction for moving immediates in the upper bits	1
<b>Branch and jump instructions</b>		
C.BEQZ	A branch-if-equal-to-zero instruction.	1
C.BNEZ	A branch- if-not-equal-to-zero instruction.	1
C.J	An unconditional jump instruction	1
C.JR	A register-based jump instruction	1
C.JALR	An instruction for jumping to a subroutine by using an address in a register	1
<b>Immediate offset access instructions</b>		
C.LW	A word load instruction	Weak order LOAD: $\geq 3$ STORE: 1 STRONG ORDER Aperiodic
C.SW	A word store instruction.	Same as above
C.LWSP	A word stack load instruction	Same as above
C.SWSP	A word stack store instruction	Same as above
C.LD	A doubleword load instruction.	Same as above
C.SD	A doubleword store instruction	Same as above
C.LDSP	A doubleword stack load instruction	Same as above
C.SDSP	A doubleword stack store instruction	Same as above
C.FLD	A double-precision load instruction.	Same as above
C.FSD	A double-precision store instruction.	Same as above

Continued on next page

Table 3.5 – continued from previous page

Instruction	Description	Latency
C.FLDSP	A double-precision stack store instruction.	Same as above
C.FSDSP	A double-precision stack load instruction.	Same as above
<b>Special instructions</b>		
C.NOP	A no-operation instruction	1
C.EBREAK	A breakpoint instruction	1

For more information, see [Appendix A-6 C Instructions](#).

## 3.2 XuanTie extended instruction sets

C910/C920 provides some extended custom instructions based on the RV64GC instruction sets. Extended half-precision floating-point instructions of C910/C920 can be directly used. All other extended instruction sets of C910/C920 must be enabled before they can be used; otherwise, illegal instruction errors will occur. To enable an extended instruction set, enable the THEADISAE bit in the MXSTATUS register.

### 3.2.1 Arithmetic operation instructions

Table 3.6: Arithmetic operation instructions

Instruction	Description	Execution Latency
<b>Add/Subtract instructions</b>		
ADDSL	An add register instruction that shifts registers	1
MULA	A multiply-add instruction	Additive numbers un-correlated: 4
MULS	A multiply-subtract instruction	Additive numbers un-correlated: 4
MULAW	A multiply-add instruction that operates on the lower 32 bits	Additive numbers correlated: 1
MULSW	A multiply-subtract instruction that operates on the lower 32 bits.	Additive numbers correlated: 1
MULAH	A multiply-add instruction that operates on the lower 16 bits	Additive numbers correlated: 1
MULSH	A multiply-subtract instruction that operates on the lower 16 bits.	Additive numbers correlated: 1
<b>Shift instructions</b>		
SRRI	A cyclic right shift instruction.	1

Continued on next page



Table 3.6 – continued from previous page

Instruction	Description	Execution Latency
SRRIW	A cyclic right shift instruction that operates on the lower 32 bits.	1
<b>Move instructions</b>		
MVEQZ	An instruction for moving values when the register value is 0	1
MVNEZ	An instruction for moving values when the register value is not 0	1

For more information, see [Appendix B-3 Arithmetic operation instructions](#).

### 3.2.2 Bit operation instructions

Table 3.7: Bit operation instructions

Instruction	Description	Execution latency
<b>Bit operation instructions</b>		
TST	An instruction for testing bits with the value of 0.	1
TSTNBZ	An instruction for testing bytes with the value of 0.	1
REV	An instruction for reversing the byte order.	1
REVV	An instruction for reversing the byte order in the lower 32 bits.	1
FF0	An instruction for fast finding the first bit with the value of 0 in a register.	1
FF1	An instruction for fast finding the first bit with the value of 1 in a register.	1
EXT	A signed extension instruction for extracting consecutive bits of a register.	1
EXTU	A zero extension instruction for extracting consecutive bits of a register.	1

For more information, see [Appendix B-4 Bitwise operation instructions](#).

### 3.2.3 Memory access instructions

Table 3.8: Memory access instructions

Instruction	Description	Execution latency
FLRD	A doubleword load instruction for shifting floating-point registers.	
FLRW	A word load instruction for shifting floating-point registers.	
FLURD	A doubleword load instruction for shifting the lower 32 bits in floating-point registers.	
FLURW	A word load instruction for shifting the lower 32 bits in floating-point registers.	
LRB	A byte load instruction for shifting registers and extending signed bits.	
LRH	A halfword load instruction for shifting registers and extending signed bits	
LRW	A halfword load instruction for shifting registers and extending signed bits	
LRD	A doubleword load instruction for shifting registers.	
LRBU	A byte load instruction for shifting registers and extending zero bits.	
LRHU	A halfword load instruction for shifting registers and extending zero bits.	
LRWU	A word load instruction for shifting registers and extending zero bits.	
LURB	A byte load instruction for shifting registers and extending signed bits.	
LURH	A halfword load instruction for shifting registers and extending signed bits.	
LURW	A word load instruction for shifting the lower 32 bits in registers and extending signed bits.	
LURD	A doubleword load instruction for shifting the lower 32 bits in floating-point registers.	
LURBU	A byte load instruction for shifting the lower 32 bits in registers and extending zero bits.	
LURHU	A halfword load instruction for shifting the lower 32 bits in registers and extending zero bits.	

Continued on next page

Table 3.8 – continued from previous page

Instruction	Description	Execution latency
LURWU	A word load instruction for shifting the lower 32 bits in registers and extending zero bits.	
LBIA	A base-address auto-increment instruction for loading bytes and extending signed bits.	This instruction is split into the load and ALU instructions for execution.
LBIB	A byte load instruction for auto-incrementing the base address and extending signed bits.	
LHIA	A base-address auto-increment instruction for loading halfwords and extending signed bits.	
LHIB	A halfword load instruction for auto-incrementing the base address and extending signed bits.	
LWIA	A base-address auto-increment instruction for loading words and extending signed bits.	
LWIB	The word load instruction for auto-incrementing the base address and extending signed bits.	
LDIA	A base-address auto-increment instruction for loading doublewords and extending signed bits.	
LDIB	A doubleword load instruction for auto-incrementing the base address and extending signed bits.	
LBUIA	A base-address auto-increment instruction for loading bytes and extending zero bits.	
LBUIB	A byte load instruction for auto-incrementing the base address and extending zero bits.	
LHUIA	An address auto-increment instruction for loading halfwords and extending zero bits.	
LHUIB	A halfword load instruction for auto-incrementing the base address and extending zero bits.	
LWUIA	An address auto-increment instruction for loading words and extending zero bits.	
LWUIB	A word load instruction for auto-incrementing the base address and extending zero bits.	
LDD	A double-register load instruction.	This instruction is split into two load instructions for execution.

Continued on next page

Table 3.8 – continued from previous page

Instruction	Description	Execution latency
LWD	A double-register word load instruction for extending signed bits.	
LWUD	A double-register word load instruction for extending zero bits.	
FSRD	A doubleword store instruction for shifting floating-point registers.	Weak order LOAD: $\geq 3$ STORE: 1 STRONG ORDER Aperiodic
FSRW	A word store instruction for shifting floating-point registers.	
FSURD	A doubleword store instruction for shifting the lower 32 bits in floating-point registers.	
FSURW	A word store instruction for shifting the lower 32 bits in floating-point registers.	
SRB	A byte store instruction for shifting registers.	
SRW	A word store instruction for shifting registers.	
SRD	A doubleword store instruction for shifting registers.	
SURB	A byte store instruction for shifting the lower 32 bits in registers.	
SURH	A halfword store instruction for shifting the lower 32 bits in registers.	
SURW	A word store instruction for shifting the lower 32 bits in registers.	
SURD	A doubleword store instruction for shifting the lower 32 bits in floating-point registers	
SBIA	A base-address auto-increment instruction for storing bytes	This instruction is split into the store and ALU instructions for execution.
SBIB	A byte store instruction for auto-incrementing the base address.	
SHIA	A base-address auto-increment instruction for storing halfwords.	
SHIB	A halfword store instruction for auto-incrementing the base address.	
SWIA	A base-address auto-increment instruction for storing words.	

Continued on next page

Table 3.8 – continued from previous page

Instruction	Description	Execution latency
SWIB	A word store instruction for auto-incrementing the base address.	
SDIA	A base-address auto-increment instruction for storing doublewords	
SDIB	A doubleword store instruction for auto-incrementing the base address.	
SDD	A double-register store instruction.	This instruction is split into two store instructions for execution.
SWD	An instruction for storing the lower 32 bits in double registers	

For more information, see [Appendix B-5 Storage instructions](#).

### 3.2.4 Cache instructions

Table 3.9: Cache instructions

Instruction	Description	**Execution latency(LMUL=1)**
DCACHE.CALL	An instruction that clears all dirty page table entries in the D-Cache.	Blocked Aperiodic
DCACHE.CIALL	An instruction that clears all dirty page table entries in the D-Cache and invalidates the entries.	
DCACHE.CIPA	An instruction that clears dirty page table entries that match the specified physical addresses in the D-Cache and invalidating the entries. (This instruction also acts on the L2 cache.)	
DCACHE.CISW	An instruction that clears dirty page table entries in the D-Cache based on the specified way and set and invalidates the entries.	
DCACHE.CIVA	An instruction that clears dirty page table entries that match the specified virtual addresses in the D-Cache and invalidates the entries. (This instruction also acts on the L2 cache.)	
DCACHE.CPA	An instruction that clears dirty page table entries that match the specified physical addresses in the D-Cache. (This instruction also acts on the L2 cache.)	

Continued on next page

Table 3.9 – continued from previous page

Instruction	Description	**Execution latency(LMUL=1)**
DCACHE.CPAL1	An instruction that clears dirty page table entries that match the specified physical addresses in the L1 D-Cache.	
DCACHE.CSW	An instruction that clears dirty page table entries in the D-Cache based on the specified way and set.	
DCACHE.CVA	An instruction that clears dirty page table entries that match the specified virtual addresses in the D-Cache. (This instruction also acts on the L2 cache.)	
DCACHE.CVAL1	An instruction that clears dirty page table entries that match the specified virtual addresses in the L1 D-Cache.	
DCACHE.IPA	An instruction that invalidates page table entries that match the specified physical addresses in the D-Cache. (This instruction also acts on the L2 cache.)	
DCACHE.ISW	An instruction that invalidates page table entries in the D-Cache based on the specified way and set.	
DCACHE.IVA	An instruction that invalidates page table entries that match the specified virtual addresses in the D-Cache. (This instruction also acts on the L2 cache.)	
DCACHE.IALL	An instruction that invalidates all page table entries in the D-Cache	Aperiodic
ICACHE.IALL	An instruction that invalidates all page table entries in the I-Cache	
ICACHE.IALLS	An instruction that invalidates all page table entries in the I-Cache through broadcasting	
ICACHE.IPA	An instruction that invalidates page table entries that match the specified physical addresses in the I-Cache.	
ICACHE.IVA	An instruction that invalidates page table entries that match the specified virtual addresses in the I-Cache.	

For more information, see [Appendix B-1 Cache instructions](#).

### 3.2.5 Multi-core synchronization instructions

Table 3.10: Multi-core synchronization instructions

Instruction	Description
SYNC	A synchronization instruction
SYNC.S	A synchronization broadcast instruction
SYNC.I	An instruction for synchronizing the clearing operation.
SYNC.IS	A broadcast instruction for synchronizing the clearing operation.

For more information, see [Appendix B-2 Multi-core synchronization instructions](#).

### 3.2.6 Half-precision floating-point instructions

Table 3.11: Half-precision floating-point instructions

Instruction	Description	Execution latency
<b>Operation instructions</b>		
FADD.H	A half-precision floating-point add instruction.	3
FSUB.H	A half-precision floating-point subtract instruction.	3
FMUL.H	A half-precision floating-point multiply instruction.	3
FMADD.H	A half-precision floating-point multiply-add instruction.	4
FMSUB.H	A half-precision floating-point multiply-subtract instruction.	4
FNMADD.H	A half-precision floating-point negate-(multiply-add) instruction.	4
FNMSUB.H	A half-precision floating-point negate -(multiply-subtract) instruction.	4
FDIV.H	A half-precision floating-point divide instruction.	4-7
FSQRT.H	A half-precision floating-point square-root instruction.	4-7
<b>Sign injection instructions</b>		
FSGNJ.H	A half-precision floating-point sign-injection instruction	3
FSGNJN.H	A half-precision floating-point negate sign-injection instruction	3
FSGNJX.H	A half-precision floating-point XOR sign-injection instruction	3
<b>Data transmission instructions</b>		
FMV.X.H	A half-precision floating-point read move instruction.	1+1

Continued on next page

Table 3.11 – continued from previous page

Instruction	Description	Execution latency
FMV.H.X	A half-precision floating-point write move instruction	1+1
<b>Compare instructions</b>		
FMIN.H	A half-precision floating-point MIN instruction	3
FMAX.H	A half-precision floating-point MAX instruction.	3
FEQ.H	A half-precision floating-point compare equal instruction.	3+1 in split execution
FLT.H	A half-precision floating-point compare less than instruction.	3+1 in split execution
FLE.H	A half-precision floating-point compare less than or equal to instruction.	3+1 in split execution
<b>Data type conversion instructions</b>		
FCVT.S.H	An instruction that converts a half-precision floating-point number into a single-precision floating-point number.	3
FCVT.H.S	An instruction that converts a single-precision floating-point number into a half-precision floating-point number.	3
FCVT.W.H	An instruction that converts a half-precision floating-point number into a signed integer.	3+1 in split execution
FCVT.WU.H	An instruction that converts a half-precision floating-point number into an unsigned integer.	3+1 in split execution
FCVT.H.W	An instruction that converts a signed integer into a half-precision floating-point number	3+1 in split execution
FCVT.H.WU	The instruction that converts an unsigned integer into a half-precision floating-point number.	3+1 in split execution
FCVT.L.H	An instruction that converts a half-precision floating-point number into a signed long integer.	3+1 in split execution
FCVT.LU.H	An instruction that converts a half-precision floating-point number into an unsigned long integer.	3+1 in split execution
FCVT.H.L	An instruction that converts a signed long integer into a half-precision floating-point number.	3+1 in split execution
FCVT.H.LU	An instruction that converts an unsigned long integer into a half-precision floating-point number.	3+1 in split execution
<b>Memory store instructions</b>		

Continued on next page



Table 3.11 – continued from previous page

Instruction	Description	Execution latency
FLH	A half-precision floating-point load instruction	Weak order LOAD: $\geq 3$ STORE: 1 STRONG ORDER
FSH	A half-precision floating-point store instruction.	Same as above
<b>Floating-point classify instructions</b>		
FCLASS.H	A single-precision floating-point classify instruction	1+1

For more information, see [Appendix B-6 Half-precision floating-point instructions](#) .

### 4.1 CPU modes

C910/C920 supports three RISC-V privilege modes : machine mode (M-mode), supervisor mode (S-mode), and user mode (U-mode). C910/C920 runs programs in M-mode after reset. The three modes correspond to different operation privileges and differ in the following aspects:

1. Register access
2. Use of privileged instructions
3. Memory access

**The U-mode provides the lowest privileges:**

User programs are allowed to access only the registers specific to the U-mode. This prevents user programs from accessing privileged information. The operating system manages and serves user programs by coordinating their behaviors.

**The S-mode provides higher privileges than the U-mode but lower privileges than the M-mode:**

Programs running in S-mode are not allowed to access control registers specific to the M-mode and are limited by physical memory protection (PMP). The page-based virtual memory acts as the core of the S-mode.

**The M-mode has the highest privileges:**

Programs running in M-mode have full access to memory, I/O resources, and underlying features required for starting and configuring the system. By default, the CPU switches to the M-mode to respond to exceptions and interrupts that occur in any mode unless the exceptions and interrupts are delegated.

Most instructions can run in all the three modes. However, some privileged instructions with major impact on the system are available only in S-mode or M-mode. For more information, see [Appendix A Standard Instructions](#) and [Appendix B T-Head Extended Instructions](#).

The privilege mode in which an exception occurs is different from that in which the CPU responds to the exception. The CPU switches to a higher privilege mode to respond to the exception, and switches back to the lower privilege mode after the exception is handled.

## 4.2 Register view

The register view of C910/C920 is shown in Fig. 4.1.

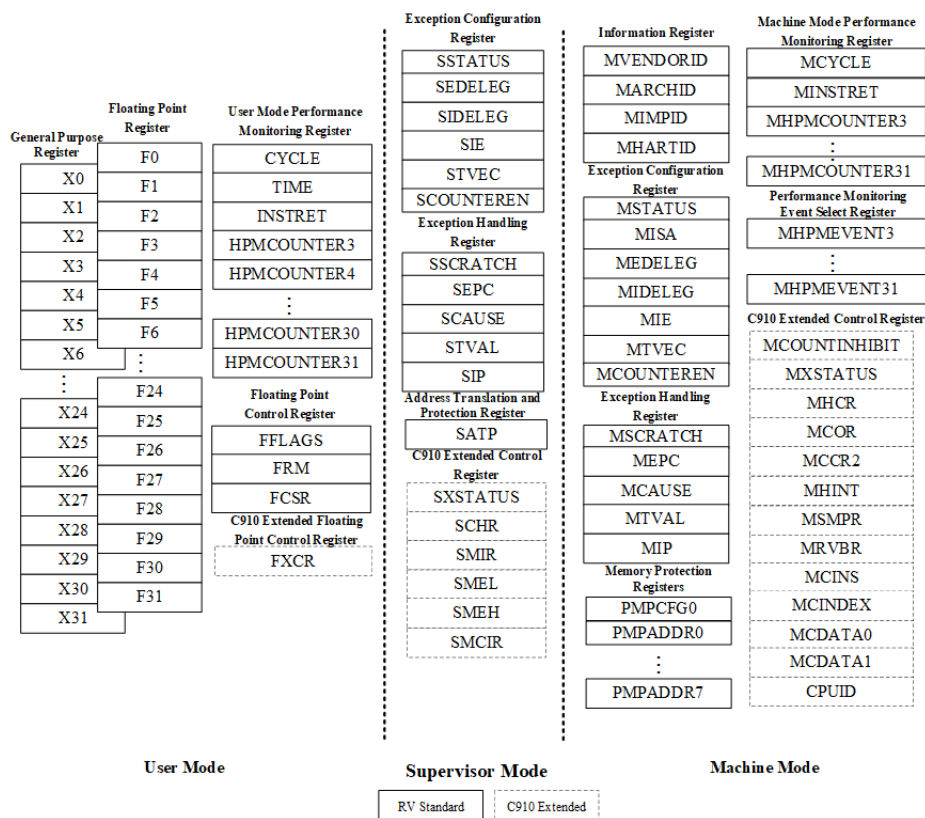


Fig. 4.1: Register view

## 4.3 General-purpose registers

C910/C920 provides thirty-two 64-bit general-purpose registers that have the same features as those defined in RISC-V. For more information, see [Table 4.1](#).

Table 4.1: General-purpose registers

Register	ABI name	Description
x0	zero	A hardwired zero register.
x1	ra	A return address register.
x2	sp	A stack pointer register.
x3	gp	A global pointer register.
x4	tp	A thread pointer register.
x5	t0	A temporary/standby link register.
x6-7	t1-2	Temporary registers.
x8	s0/fp	A reserved register/frame pointer register.
x9	s1	A reserved register.
x10-11	a0-1	Function argument/Return value registers.
x12-17	a2-7	Function argument registers.
x18-27	s2-11	Reserved registers.
x28-31	t3-6	Temporary registers.

The general-purpose registers are used to store instruction operands, instruction execution results, and address information.

## 4.4 Floating-point registers

In addition to standard RV64F instructions, C910/C920 also supports floating-point half-precision computing and provides 32 independent floating-point registers. These registers are accessible in U-mode, S-mode, and M-mode.

Table 4.2: Floating-point registers

Register	ABI name	Description
f0-7	ft0-7	Floating-point temporary registers.
f8-9	fs0-1	Floating-point reserved registers.
f10-11	fa0-1	Floating-point argument/return value registers.
f12-17	fa2-7	Floating-point argument registers.
f18-27	fs2-11	Floating-point reserved registers.
f28-31	ft8-11	Floating-point temporary registers.

Unlike x0, f0 is not hardwired to 0, but its bit values are variable like other floating-point registers. A single-precision floating-point number occupies only the lower 32 bits of a floating-point register, and the upper 32 bits must be set to 1; otherwise, the number will be considered nonnumeric. A half-precision floating-point number occupies only the lower 16 bits of a floating-point register, and the upper 48 bits must be set to 1; otherwise, the number will be considered nonnumeric.

The independent floating-point registers help increase the register capacity and bandwidth, improving performance of the CPU. Along with the floating-point registers, floating-point load and store instructions and instructions for transferring data between floating-point and general-purpose registers are added.

#### 4.4.1 Transmit data between floating-point and general-purpose registers

Data can be transmitted between floating-point and general-purpose registers through floating-point register move instructions. Floating-point register move instructions include:

- FMV.X.H/FMV.H.X: A half-precision data move instruction for floating-point registers.
- FMV.X.W/FMV.W.X: A single-precision data move instruction for floating-point registers.

When half-precision, single-precision data is transmitted from a general-purpose register to a floating-point register, the data format remains unchanged. Therefore, a program can directly use these registers without converting their types.

For more information, see [Appendix A-4 F instructions](#).

#### 4.4.2 Maintain consistency of register precision

Floating-point registers can store half-precision, single-precision and integer data. For example, the type of data stored in f1 depends on the last write operation, and may be any one of the four types.

Floating-point units (FPUs) do not detect data formats based on hardware. The hardware parses data formats in a floating-point register only based on the executed floating-point instruction, regardless of the data format in the last write operation in the register. In this case, the consistency of data precision in the register is ensured only by the compiler or program.

### 4.5 System control registers

#### 4.5.1 Standard control registers

This section describes RISC-V standard control registers implemented in C910/C920 by M-mode, S-mode, and U-mode.

The RISC-V standard M-mode control registers implemented in C910/C920 are described in [Table 4.3](#).

Table 4.3: RISC-V standard M-mode control registers

Register	Read/Write permission	ID	Description
<b>M-mode information registers</b>			
mvendorid	Read-only in M-mode	0xF11	A vendor ID register.

Continued on next page

Table 4.3 – continued from previous page

Register	Read/Write permission	ID	Description
marchid	Read-only in M-mode	0xF12	An architecture ID register.
mimpid	Read-only in M-mode	0xF13	An M-mode hardware implementation ID register.
mhartid	Read-only in M-mode	0xF14	An M-mode logical kernel ID register.
<b>M-mode exception configuration registers</b>			
mstatus	Read/Write in M-mode	0x300	An M-mode CPU status register.
misa	Read/Write in M-mode	0x301	An M-mode CPU instruction set attribute register.
medeleg	Read/Write in M-mode	0x302	An M-mode exception delegation control register.
mideleg	Read/Write in M-mode	0x303	An M-mode interrupt delegation control register.
mie	Read/Write in M-mode	0x304	An M-mode interrupt enable control register.
mtvec	Read/Write in M-mode	0x305	An M-mode vector base address register.
mcounteren	Read/Write in M-mode	0x306	An M-mode counter enable control register.
mcountinhibit	Read/Write in M-mode	0x320	An M-mode count inhibit register.
<b>M-mode exception handling registers</b>			
mscratch	Read/Write in M-mode	0x340	An M-mode temporary data backup register upon exceptions.
mepc	Read/Write in M-mode	0x341	An M-mode exception program counter.
mcause	Read/Write in M-mode	0x342	An M-mode exception event cause register.
mtval	Read/Write in M-mode	0x343	An M-mode exception event vector register.
mip	Read/Write in M-mode	0x344	An M-mode interrupt pending state register.
<b>M-mode memory protection registers</b>			
pmpcfg0	Read/Write in M-mode	0x3A0	Physical memory protection configuration register 0.
pmpcfg2	Read/Write in M-mode	0x3A2	Physical memory protection configuration register 2.
pmpaddr0	Read/Write in M-mode	0x3B0	Physical memory protection base address register 0.
pmpaddr15	Read/Write in M-mode	0x3BF	Physical memory protection base address register 7.
<b>M-mode counters/timers</b>			

Continued on next page

Table 4.3 – continued from previous page

Register	Read/Write permission	ID	Description
mcycle	Read/Write in M-mode	0xB00	An M-mode cycle counter.
minstret	Read/Write in M-mode	0xB02	An M-mode retired instruction counter.
mhpmcounter3	Read/Write in M-mode	0xB03	Machine-mode counter 3.
.....			
mhpmcounter31	Read/Write in M-mode	0xB1F	M-mode counter 31.
<b>M-mode counter configuration registers</b>			
mhpmevent3	Read/Write in M-mode	0x323	M-mode event select register 3.
.....			
mhpmevent31	Read/Write in M-mode	0x33F	M-mode event select register 31.

The RISC-V standard S-mode control registers implemented in C910/C920 are described in [Table 4.4](#).

Table 4.4: RISC-V standard S-mode control registers

Register	Read/Write permission	ID	Description
<b>S-mode exception configuration registers</b>			
sstatus	Read/Write in S-mode	0x100	An S-mode CPU status register.
sie	Read/Write in S-mode	0x104	An S-mode interrupt enable control register.
stvec	Read/Write in S-mode	0x105	An S-mode vector base address register.
scounteren	Read/Write in S-mode	0x106	An S-mode counter enable control register.
<b>S-mode exception handling registers</b>			
sscratch	Read/Write in S-mode	0x140	An S-mode temporary data backup register upon exceptions.
sepc	Read/Write in S-mode	0x141	An S-mode exception program counter.
scause	Read/Write in S-mode	0x142	An S-mode exception event cause register.
stval	Read/Write in S-mode	0x143	An S-mode exception event vector register.
sip	Read/Write in S-mode	0x144	An S-mode interrupt pending state register.
<b>S-mode address translation registers</b>			
satp	Read/Write in S-mode	0x180	An S-mode virtual address translation and protection register.

The RISC-V standard user-mode control registers implemented in C910/C920 are described in [Table 4.5](#).

Table 4.5: RISC-V standard U-mode control registers

Register	Read/Write permission	ID	Description
<b>U-mode floating-point control registers</b>			

Continued on next page

Table 4.5 – continued from previous page

Register	Read/Write permission	ID	Description
fflags	Read/Write in	0x001	A floating-point accrued exception status register.
frm	Read/Write in U-mode	0x002	A floating-point dynamic rounding mode control register.
fcsr	Read/Write in U-mode	0x003	A floating-point control and status register.
<b>U-mode counters/timers</b>			
cycle	Read/Write in U-mode	0xC00	A U-mode cycle counter.
time	Read/Write in U-mode	0xC01	A U-mode timer.
instret	Read/Write in U-mode	0xC02	A U-mode retired instruction counter.
hpmcounter3	Read/Write in U-mode	0xC03	A U-mode counter 3
... ..			
hpmcounter31	Read/Write in U-mode	0xC1F	U-mode counter 31.

## 4.5.2 Extended control registers

This section describes extended control registers implemented in C910/C920 by M-mode, S-mode, and U-mode.

The extended M-mode control registers of C910/C920 are described in [Table 4.6](#).

Table 4.6: Extended M-mode control registers of C910/C920

Register	Read/Write permission	ID	Description
<b>Extended M-mode CPU control and status registers</b>			
mxstatus	Read/Write in M-mode	0x7C0	An extended M-mode status register.
mhcr	Read/Write in M-mode	0x7C1	An M-mode hardware configuration register.
mcor	Read/Write in M-mode	0x7C2	An M-mode hardware operation register.
mccr2	Read/Write in M-mode	0x7C3	An M-mode L2 cache control register.
mhint	Read/Write in M-mode	0x7C5	An M-mode implicit operation register.

Continued on next page



Table 4.6 – continued from previous page

Register	Read/Write permission	ID	Description
mrvbr	Read-only in M-mode	0x7C7	An M-mode reset vector base address register.
mcounterwen	Read/Write in M-mode	0x7C9	An S-mode counter write enable register.
mcounterinten	Read/Write in M-mode	0x7CA	An M-mode event interrupt enable register.
mcounterof	Read/Write in M-mode	0x7CB	An M-mode overflow flag register.
<b>Extended M-mode cache access registers</b>			
mcins	Read/Write in M-mode	0x7D2	An M-mode cache instruction register.
mcindex	Read/Write in M-mode	0x7D3	An M-mode cache access index register.
mcdata0	Read/Write in M-mode	0x7D4	An M-mode cache data register 0.
mcdata1	Read/Write in M-mode	0x7D5	An M-mode cache data register 1.
<b>Extended M-mode CPU model registers</b>			
mcpuid	Read-only in M-mode	0xFC0	An M-mode CPU model register.
mapbaddr	Read-only in M-mode	0xFC1	An on-chip bus base address register.
<b>Extended multi-core registers</b>			
msmpr	Read/Write in M-mode	0x7F3	A snooping enable register.

For more information, see [Appendix C-1 M-mode control registers](#).

The extended S-mode control registers of C910/C920 are described in [Table 4.7](#).

Table 4.7: Extended S-mode control registers of C910/C920

Register	Read/Write permission	ID	Description
<b>Extended S-mode CPU control and status registers</b>			
sxtatus	Read/Write in S-mode	0x5C0	An extended S-mode status register.
shcr	Read/Write in S-mode	0x5C1	An S-mode hardware control register.
scounterinten	Read/Write in S-mode	0x5C4	An S-mode event interrupt enable register.
scounterof	Read/Write in S-mode	0x5C5	An S-mode event overflow flag register.
scycle	Read/Write in S-mode	0x5E0	An S-mode cycle counter.
.....			
shpmcounter31	Read/Write in S-mode	0x5FF	S-mode counter 31.
<b>Extended S-mode MMU registers</b>			
smir	Read/Write in S-mode	0x9C0	An S-mode MMU index register.
smel	Read/Write in S-mode	0x9C1	An S-mode MMU EntryLo register.
smeh	Read/Write in S-mode	0x9C2	An S-mode MMU EntryHi register.
smcir	Read/Write in S-mode	0x9C3	An S-mode MMU control register.

For more information, see [Appendix C-2 S-mode control registers](#).

The extended U-mode control registers of C910/C920 are described in [Table 4.8](#).

Table 4.8: Extended U-mode control registers of C910/C920

Register	Read/Write permission	ID	Description
<b>Extended U-mode floating-point control registers</b>			
fxcr	Read/Write in U-mode	0x800	An extended U-mode floating-point control register.

For more information, see [Appendix C-3 U-mode control registers](#).

## 4.6 Data formats

### 4.6.1 Integer data format

Values in a register are not distinguished by big-endian or little-endian type, but by signed or unsigned type. Values in a register are arranged from right to left with the least significant bit being the rightmost bit and the most significant bit being the leftmost bit, as shown in [Fig. 4.2](#).




Fig. 4.2: Integer data structure in registers

### 4.6.2 Floating-point data format

FPU of C910/C920 comply with the RISC-V standard and the ANSI/IEEE 754-2008 standard for floating-point arithmetic, and support half-precision, single-precision computation. The floating-point data format is shown in Fig. 4.3. Single-precision data occupies only the lower 32 bits of a floating-point register, and the upper 32 bits must be set to 1; otherwise, the data will be considered nonnumeric. Half-precision data occupies only the lower 16 bits of a floating-point register, and the upper 48 bits must be set to 1; otherwise, the data will be considered nonnumeric.

## 4.7 Big-endian and little-endian

The concepts of big-endian and little-endian are proposed with respect to the data storage formats of memories. In the big-endian scheme, the most significant byte of an address is stored to the lower bits in physical memory. In the little-endian scheme, the most significant byte of an address is stored to the upper bits in physical memory. The data formats are shown in Fig. 4.4.

C910/C920 supports only the little-endian scheme, and supports binary integers with standard complements. The length of each instruction operand can be explicitly encoded in programs (load/store instructions) or

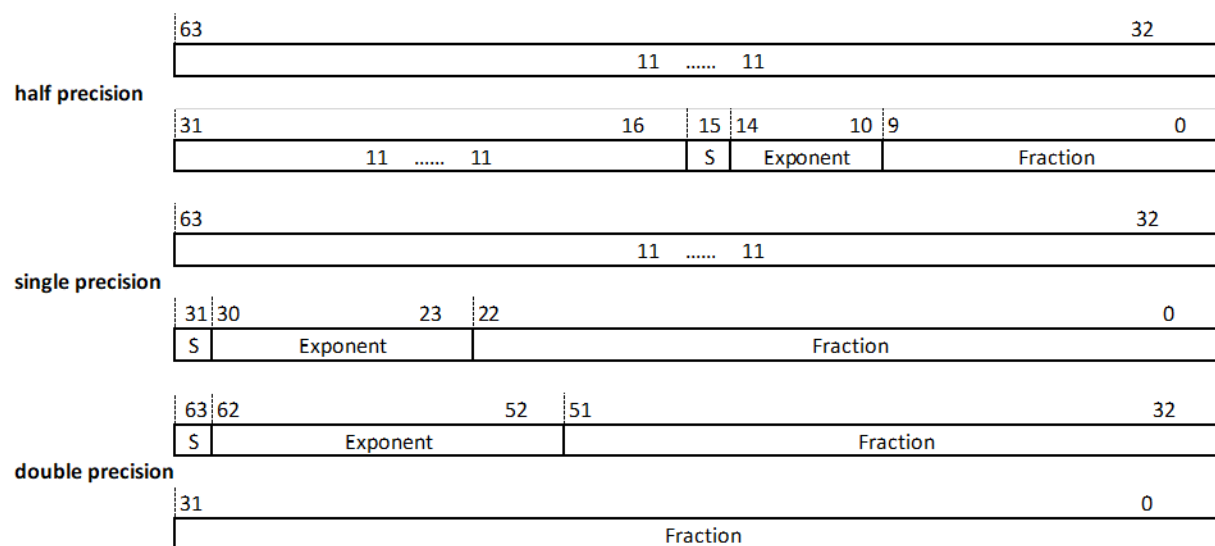


Fig. 4.3: Floating-point data structure in registers

A+7	A+6	A+5	A+4	A+3	A+2	A+1	A	
Byte7	Byte6	Byte5	Byte4	Byte3	Byte2	Byte1	Byte0	Double word at A
Byte7	Byte6	Byte5	Byte4	Byte3	Byte2	Byte1	Byte0	Word at A
Byte7	Byte6	Byte5	Byte4	Byte3	Byte2	Byte1	Byte0	Half word at A
Byte7	Byte6	Byte5	Byte4	Byte3	Byte2	Byte1	Byte0	Byte at A

little-endian scheme

A+4	A+5	A+6	A+7	A	A+1	A+2	A+3	
Byte3	Byte2	Byte1	Byte0	Byte7	Byte6	Byte5	Byte4	Double word at A
Byte4	Byte5	Byte6	Byte7	Byte3	Byte2	Byte1	Byte0	Word at A
Byte6	Byte7	Byte4	Byte5	Byte1	Byte0	Byte3	Byte2	Half word at A
Byte7	Byte6	Byte5	Byte4	Byte0	Byte1	Byte2	Byte3	Byte at A

big-endian V1 scheme

A+7	A+6	A+5	A+4	A+3	A+2	A+1	A	
Byte0	Byte1	Byte2	Byte3	Byte4	Byte5	Byte6	Byte7	Double word at A
Byte4	Byte5	Byte6	Byte7	Byte0	Byte1	Byte2	Byte3	Word at A
Byte6	Byte7	Byte4	Byte5	Byte2	Byte3	Byte0	Byte1	Half word at A
Byte7	Byte6	Byte5	Byte4	Byte3	Byte2	Byte1	Byte0	Byte at A

big-endian V2 scheme

Fig. 4.4: Data structure in memory

implicitly indicated in instruction operations (index operation and byte extraction) Usually, an instruction receives a 64-bit operand and generates a 64-bit result.

---

## Exceptions and Interrupts

---

### 5.1 Overview

Exception handling is a core feature of a CPU. Exceptions include instruction exceptions and external interrupts. When some exception events occur, the CPU is enabled to respond to these events. The events include hardware errors, instruction execution errors, and user program request services.

The key of exception handling is to save the operating status of the CPU when an exception occurs and resume the status when the CPU exits exception handling. Exceptions can be identified in all stages of the instruction pipeline. The CPU hardware ensures that subsequent instructions do not change the CPU status. Exceptions are handled at the boundary of an instruction. To be specific, the CPU responds to the exceptions when the instruction retires, and saves the address of the to-be-executed instruction when the CPU exits exception handling. Even if exceptions are identified before an instruction retires, the CPU does not handle the exceptions until the instruction retires. To ensure proper functioning of programs, the CPU does not repeatedly run the executed instructions after exception handling is completed.

In machine mode (M-mode), the CPU responds to an instruction exception or an external interrupt in the following procedure:

**Step 1:** Save the exception PC to the mepc register.

**Step 2:** Update the mcause and mtval registers based on the exception type.

**Step 3:** Save the machine interrupt-enable (MIE) bit in the mstatus register to the MPIE field, clear the MIE field, and prohibit responses to interrupts.

**Step 4:** Save the privilege mode applied before the exception occurs to the MPP field in the mstatus register, and switch to the M-mode.

**Step 5:** Obtain the entry address of exception program based on the base address and mode in the mtvec register, and run instructions of the exception program in sequence.

C910/C920 conforms to the exception vector table defined in RISC-V, as shown in Table 5.1.

Table 5.1: Exception and interrupt vector assignment

Interrupt flag	Exception vector ID	Description
1	0	Unavailable.
1	1	A software interrupt in supervisor mode (S-mode).
1	2	Reserved.
1	3	A software interrupt in M-mode.
1	4	Unavailable.
1	5	A timer interrupt in S-mode.
1	6	Reserved.
1	7	The timer interrupt in M-mode.
1	8	Unavailable.
1	9	An external interrupt in S-mode.
1	10	Reserved.
1	11	An external interrupt in M-mode.
1	17	A performance detection overflow interrupt.
1	Others	Reserved.
0	0	Unavailable.
0	1	A fetch instruction access error exception.
0	2	An illegal instruction exception.
0	3	A debug breakpoint exception.
0	4	A load instruction unaligned access exception.
0	5	A load instruction access error exception.
0	6	A store/atomic instruction unaligned access exception.
0	7	A store/atomic instruction access error exception.
0	8	A user-mode (U-mode) environment call exception.
0	9	An S-mode environment call exception.
0	10	Reserved.
0	11	An M-mode environment call exception.
0	12	An instruction fetch page error exception.
0	13	A load instruction page error exception.
0	14	Reserved.
0	15	A store/atomic instruction page error exception.
0	$\geq 16$	Reserved.

C910/C920 supports exception and interrupt delegation. When an exception or interrupt occurs in S-mode, the CPU switches to the M-mode for handling. This causes performance loss of the CPU. Delegation enables the CPU to respond to exceptions and interrupts in S-mode. Exceptions that occur in M-mode are not delegated, but still handled in M-mode. Interrupts that occur in M-mode can be delegated to the S-mode for handling, except the external interrupts, software interrupts, and timer interrupts that occur in M-mode. In M-mode, the CPU does not respond to delegated interrupts.

In S-mode and U-mode, the CPU can respond to all interrupts and exceptions that meet the specified criteria. The CPU responds to undelegated exceptions and interrupts in M-mode, and updates the machine-mode exception handling registers. The CPU responds to delegated exceptions and interrupts in S-mode, and updates the S-mode exception handling registers.

## 5.2 Exceptions

### 5.2.1 Exception handling

In M-mode, the CPU responds to illegal instruction or access error exceptions in the following procedure:

**Step 1:** Save the exception PC to the mepc register.

**Step 2:** Set the interrupt flag in the mcause register to 0, write the exception ID to the mcause register, and update the mtval register based on the rules defined in [Table 5.2](#).

**Step 3:** Save the machine interrupt-enable (MIE) bit in the mstatus register to the MPIE field, clear the MIE field, and prohibit responses to interrupts.

**Step 4:** Save the privilege mode applied before the exception occurs to the MPP field in the mstatus register, and switch to the M-mode.

**Step 5:** The PC fetches an instruction from the base address in the mtvec register and executes the instruction. The instruction is usually a jump instruction for jumping to the top-level handler. The handler analyzes the mcause register to obtain the exception ID and calls the handler corresponding to the exception ID.



Table 5.2: Updates to mtval when exceptions occur

Exception vector ID	Exception	mtval update
1	Fetch instruction access error exception	Virtual address accessed by the fetch instruction
2	Illegal instruction exception	Instruction code
3	Debug breakpoint exception	0
4	Load instruction unaligned access exception	Virtual address accessed by the load instruction
5	Load instruction access error exception	0
6	Store/Atomic instruction unaligned access exception	Virtual address accessed by the store/atomic instruction
7	Store/Atomic instruction access error exception	0
8	U-mode environment call exception	0
9	S-mode environment call exception	0
11	M-mode environment call exception	0
12	Fetch instruction page error exception	Virtual address accessed by the fetch instruction
13	Load instruction page access exception	Virtual address accessed by the load instruction
15	Store/Atomic instruction page error exception	Virtual address accessed by the store/atomic instruction

### 5.2.2 Return from exceptions

You can run the mret instruction to return from an exception. In this case, the CPU performs the following operations:

- Restore the mepc register to the PC. (The mepc register stores the PC applied when the exception occurs. You can adjust the mepc register to skip the exception instruction; otherwise, the exception instruction will be executed again.)
- Restore the value of the MPIE field in the mstatus register to the MIE field in the mstatus register.
- Restore the privilege mode applied before the exception occurs from the MPP field in the mstatus register.

### 5.2.3 Imprecise exceptions

In rare cases, the CPU may encounter imprecise exceptions. An imprecise exception means that the mepc register does not point to the instruction triggering the exception when the exception occurs. For example, the bus returns an error after the CPU executes a load instruction. An instruction can quickly retire in the pipeline, and the load instruction may have retired when the bus returns the error. Therefore, the mepc register points to a subsequent instruction instead of the load instruction.

However, imprecise exceptions rarely occur in practical systems. Once an imprecise exception occurs, the system may have encountered a fatal error.

## 5.3 Interrupts

### 5.3.1 Interrupt priorities

When receiving multiple interrupt requests, the CPU responds to them by their priorities (in descending order):

- L1 ECC interrupt
- M-mode external interrupt
- M-mode software interrupt
- M-mode timer interrupt
- S-mode external interrupt
- S-mode software interrupt
- S-mode timer interrupt
- PMU overflow interrupt
- L1 ECC interrupt (delegated)
- S-mode external interrupt (delegated)
- S-mode software interrupt (delegated)
- S-mode timer interrupt (delegated)
- PMU overflow interrupt (delegated)

### 5.3.2 Interrupt responses

In M-mode, the CPU responds to an interrupt in the following procedure:

**Step 1:** Execute the current instruction and save the PC of the next instruction to the mepc register.

**Step 2:** Set the interrupt flag in the mcause register to 1, write the interrupt ID to the mcause register, and update the mtval register to 0.

**Step 3:** Save the machine interrupt-enable (MIE) bit in the mstatus register to the MPIE field, clear the MIE field, and prohibit responses to interrupts.

**Step 4:** Save the privilege mode applied before the interrupt occurs to the MPP field in the mstatus register, and switch to the M-mode.

**Step 5 (The Mode field in the mtvec register is 0, indicating a direct interrupt):** The PC fetches an instruction from the base address in the mtvec register and executes the instruction. The instruction is usually a jump instruction for jumping to the top-level handler. The handler analyzes the mcause register to obtain the interrupt ID and calls the handler corresponding to the interrupt ID.

**Step 5 (The Mode field in the mtvec register is 1, indicating a vectored interrupt):** The PC fetches an instruction from the address calculated in (Base address in the mtvec register +  $4 \times$  Interrupt ID) and executes the instruction. The instruction is usually a jump instruction for jumping to the corresponding interrupt handler.

### 5.3.3 Return from interrupts

You can run the mret instruction to return from an interrupt. In this case, the CPU performs the following operations:

- Restore the mepc register to the PC. (The mepc register stores the PC of the next instruction and therefore does not need to be adjusted.)
- Restore the value of the MPIE field in the mstatus register to the MIE field in the mstatus register.
- Restore the privilege mode applied before the interrupt occurs from the MPP field in the mstatus register.

## 6.1 Overview

### 6.1.1 Memory attributes

C910/C920 supports two memory types: memory and device, which are distinguished by the SO bit. The memory supports speculative execution and out-of-order execution. It is further classified into cacheable memory and non-cacheable memory. The device supports only non-speculative in-order execution and therefore is non-cacheable. It is further classified into bufferable device and non-bufferable device. Bufferable indicates that a response to a write request can be quickly returned on an intermediate node. Non-bufferable indicates that a response to a write request is returned only after the end device completes writing.

To share data among multiple cores, C910/C920 allows you to set the shareable (SH) page attribute. A shareable page is shared among multiple cores, and the hardware maintains data coherence. A non-shareable page is exclusively occupied by a core, and the software, instead of hardware, maintains data coherence among multiple cores.

Table 6.1 describes the page attributes corresponding to each memory type.

Table 6.1: Memory types

Memory type	SO	C	B	SH	SEC
Cacheable memory	0	1	1	1	Reserved
Non-cacheable memory	0	0	1	1	Reserved
Bufferable device	1	0	1	1	Reserved
Non-bufferable device	1	0	0	1	Reserved

The CPU can obtain the page attribute of an address from the sysmap.h file or a page table entry (PTE). The two methods are described as follows:

1. Page attributes of addresses are determined by the sysmap.h file if virtual addresses are not translated into physical addresses, that is, the machine mode (M-mode) or MMU is disabled.
2. Page attributes of addresses depend on the MAEE field in the mxstatus register if virtual addresses are translated into physical addresses, that is, the CPU is not in M-mode and the MMU is enabled. If the MAEE field is enabled, page attributes of addresses are determined by page attributes extended in the corresponding PTEs. If the MAEE field is disabled, page attributes of addresses are determined by the sysmap.h file.

sysmap.h is an extended configuration file of C910/C920 that is open to users. You can define page attributes for different address ranges as required.

sysmap.h allows you to set page attributes for up to 8 address spaces. The largest address (non-inclusive) of address space  $i$  ( $i = 0$  to  $7$ ) is defined by the `SYSMAP_BASE_ADDRi` ( $i = 0$  to  $7$ ) macro. The smallest address (inclusive) is defined by the `SYSMAP_BASE_ADDR(i - 1)` macro. That is,

$$\text{SYSMAP\_BASE\_ADDR}(i - 1) \leq \text{Address of address space } i < \text{SYSMAP\_BASE\_ADDR}i.$$

The smallest address of address space 0 is 0x0. Page attributes of memory addresses beyond the eight address spaces defined in the sysmap.h file are cacheable/bufferable/shareable/security by default. The upper and lower boundaries of each address space is 4 KB aligned. Therefore, the `SYSMAP_BASE_ADDRi` macro defines the upper 28 bits of an address.

Page attributes of memory addresses within address space  $i$  ( $i = 0$  to  $7$ ) are defined by the `SYSMAP_FLAGi` ( $i = 0$  to  $7$ ) macro. The attribute layout is shown in Fig. 6.1.

4	3	2	1	0
Strong order	Cacheable	Bufferable	Shareable	Security

Fig. 6.1: Address attributes in the sysmap.h file

### 6.1.2 Memory ordering model

C910/C920MP adopts a weak memory ordering model, which is defined as follows:

- Ordering of access to the same address is maintained among multiple cores, including read after read (RAR), write after write (WAW), write after read (WAR), add read after write (RAW).
- Weak ordering of access to different addresses is allowed among multiple cores, including RAR, WAW, WAR, add RAW.
- Atomic other-multi-copy is ensured. When a core is able to obtain written data of another core, other cores must also be able to obtain the data. However, when a core is able to obtain its own written data, it is not required that other cores be able to obtain the data.

Weak memory ordering causes inconsistency between the actual read/write order among multiple cores and the access order defined by the program. Therefore, C910/C920 provides extended SYNC instructions to enforce memory access ordering in software.

SYNC instructions define the execution order of all instructions, ensuring that all instructions preceding a SYNC instruction are executed before the SYNC instruction. In addition, SYNC instructions can also be used to synchronize instruction memory. After instructions preceding a SYNC instruction are executed, the SYNC instruction clears the pipeline and re-fetches instructions. For more information, see [Table 6.2](#).

Table 6.2: SYNC instructions

Mnemonic	Description	Scope
SYNC.IS	Synchronize data and instruction memory	Shareable
SYNC.I	Synchronize data and instruction memory	Non-shareable
SYNC.S	Synchronize data memory	Shareable
SYNC	Synchronize data memory	Non-shareable

## 6.2 MMU

### 6.2.1 Overview

The memory management unit (MMU) of C910/C920 complies with the RISC-V SV39 standard. It provides the following features:

**Address translation:** Translates 39-bit virtual addresses to 40-bit physical addresses.

**Page protection:** Checks the read/write/execution permissions of page visitors.

**Page attribute management:** Extends address attribute bits and obtains page attributes based on access addresses for further processing by the system.

### 6.2.2 TLB

The MMU uses translation lookaside buffers (TLBs) to implement its features. A TLB stores virtual addresses used when the CPU accesses the memory. Before translating a virtual address, the MMU checks the

page attributes in the TLB and outputs a physical address corresponding to the virtual address.

The MMU of C910/C920 uses two levels of TLBs: the uTLB at level 1 and the jTLB at level 2. The uTLB includes the I-uTLB and the D-uTLB. After the CPU is reset, the hardware invalidates all entries in the uTLB and the jTLB, without the need of initializing software.

The I-uTLB provides 32 fully associative entries for storing pages in 4 KB, 2 MB, or 1 GB size. When an instruction fetch request hits the I-uTLB, the physical address and the corresponding permission attribute can be obtained in the current cycle.

The D-uTLB provides 17 fully associative entries for storing pages in 4 KB, 2 MB, or 1 GB size. When a load/store request hits the D-uTLB, the physical address and the corresponding permission attribute can be obtained in the current cycle.

The jTLB is a 4-way set-associative cache shared by instructions and data. It provides 1024 entries for storing pages in 4 KB, 2 MB, or 1 GB size. When a request misses the uTLB but hits the jTLB, the physical address and the corresponding permission attribute will be returned within at least three cycles.

### 6.2.3 Address translation process

The MMU is used to translate virtual addresses into physical addresses and check corresponding permissions. Specific address mappings and corresponding permissions are configured by the operating system and stored in page tables. C910/C920 implements address translation through indexing by at most three levels of page tables. The MMU accesses the L1 page table to obtain the base address of an L2 page table and the corresponding permission attributes, accesses the L2 page table to obtain the base address of an L3 page table and the corresponding permission attributes, and accesses the L3 page table to obtain the final physical address and the corresponding permission attributes. The MMU may obtain the final physical address, that is, a leaf table entry, at each level of access. The virtual page number (VPN) consists of 27 bits and is divided into three 9-bit VPN[i]. A part of the VPN is used for indexing in each access.

Content of leaf table entries is cached in the TLB to accelerate address translation. The content includes physical addresses translated from virtual addresses and corresponding permission attributes. If the uTLB is missed, the MMU accesses the jTLB. If the jTLB is missed, the MMU enables a hardware page table walk to access the memory to obtain the final address translation result.

A page table stores entry addresses of next-level page tables or physical information of the final page table. The page table structure is shown below:

#### Flags in bit [9:0]: page attributes

Features are described in *smel register*.

#### Flags in bit [63:59]: page attributes

The custom page attributes of C910/C920, which are available when the MAEE field in the mxstatus register is enabled. Features are described in *smel register*.

#### PPN: physical page number




Fig. 6.2: Page table structure

PPN[i] indicates the PPN corresponding to each level of page table.

The address translation process is described as follows:

If the TLB is hit when the CPU attempts to access a virtual address, the CPU directly obtains the physical address and the corresponding attributes from the TLB. If the TLB is missed, the MMU performs the following steps to translate the virtual address:

1. Obtain the access address {satp.PPN, VPN[2], 3' b0} of the L1 page table, and access the D-Cache/memory based on the address to obtain a 64-bit PTE of the L1 page table.
2. Check whether the PTE conforms to the physical memory protection (PMP) permission. If no, generate the corresponding access error exception. If yes, determine whether the X/W/R bit meets the condition of the leaf page table based on the rules shown in Table 6.4 . If yes, the final physical address has been found. Then go to step 3. If no, obtain the access address {PTE.PPN, next-level VPN, 3' b0} of the next-level page table, and access the D-Cache/memory again.
3. After the leaf page table is found, compare the X/W/R/L bit in the PMP register with the X/W/R bit in the PTE to obtain the minimum permissions, check the permissions, and write the content of the PTE back to jTLB.
4. If permission violation is found in any PMP check, generate the corresponding access error exception based on the access type.
5. Generate a page fault exception in the following three cases: the leaf page table is found but the access type does not conform to the setting of the A/D/X/W/R/U bit, no leaf page table is found after three accesses, or an access error is generated during access to the D-Cache/memory.
6. If the leaf page table is found in less than three accesses, a large page table has been obtained. In this case, check whether the PPN of the large page table is aligned based on the page size. If no, generate a page fault exception.

### 6.2.4 System control registers

In addition to the standard satp register, the MMU of C910/C920 provides the extended smir, smcir, supervisor-mode (S-mode) entry low (smel), and S-mode entry high (smeh) control registers. You can use the extended registers to directly read, write, probe, and invalidate the TLB.



#### 6.2.4.1 Supervisor address translation and protection (satp) register

The satp register is an MMU control register defined in the SV39 standard.




Fig. 6.3: Satp Register Descriptions

Mode: MMU address translation mode

Table 6.3: MMU address translation mode

RV64		
Value	Name	Description
0	Bare	No translation or protection
1-7	-	Reserved
8	Sv39	Page-based 39-bit virtual addressing
9	Sv48	Page-based 48-bit virtual addressing
10	Sv57	Reserved for page-based 57-bit virtual addressing
11	Sv64	Reserved for page-based 64-bit virtual addressing
12-15	-	Reserved

**ASID:** the current address space identifier (ASID)

Indicates the ASID of the current program.

**PPN: root PPN for hardware writeback**

Indicates the PPN used for L1 hardware writeback.

#### 6.2.4.2 smcir register

The smcir register enables you to probe, read, write, and invalidate the TLB.

**TLBP: TLB probe**

Indicates the operation of probing the TLB based on the smeh register.

When the TLB is hit, the value of the smir register is updated to the serial number of the TLB.

TLBR: TLB read

Indicates the operation of reading values of corresponding TLB entries based on indexes in the smir register, and updating the smeh and smel registers based on the values.




Fig. 6.4: Smcir Register Description

**TLBWI: TLB indexed write**

Indicates the operation of writing values of the smeh and smel registers to corresponding TLB entries based on indexes in the smir register.

**TLBWR: TLB random write**

Indicates the operation of writing values of the smeh and smel registers to corresponding TLB entries based on indexes in the random register.

**TLBIASID: TLB invalidation by ASID**

Indicates the operation of invalidating all TLB entries that match the specified ASID.

**TLBIALL: TLB initialization**

Indicates the operation of invalidating all TLB entries and initializing the TLB.

**TLBII: TLB invalidation by index**

Indicates the operation of invalidating all TLB entries that match the specified index in the smir register.

**TLBIAW: TLB invalidation by world**

Indicates the operation of invalidating all TLB entries corresponding to the trustable or non-trustable world.

This field is available only when trusted execution environment (TEE) extension is configured. It has not been implemented in C910/C920.

**ASID: the ASID used**

Indicates the ASID used for matching in the TLBIASID operation. The smcir register enables you to probe, read, write, and invalidate the TLB.

**6.2.4.3 smir register**

The smir register is used to index the TLB. In TLB probing, the index of a hit entry is updated. In TLB write indexing, the index field of the smir register is written to write the mapping to the corresponding index in the jTLB.

**P – Probe Failure**




Fig. 6.5: Smir Register Descriptions

0: indicates that TLB is hit when the TLBP instruction is executed.

1: indicates that TLB is missed when the TLBP instruction is executed.

#### Tfatal – Probe multiple

Specifies whether multiple matches occur when the TLBP instruction is executed.

0: indicates that no multiple matches occur.

1: indicates that multiple matches occur.

#### Index – TLB Index

1024-entry configuration: Index [9:8] is the way index, and index [7:0] is the set/entry index (4-way, 256 entries).

2048-entry configuration: Index [10:9] is the way index, and index [8:0] is the set/entry index (4-way, 512 entries).

#### 6.2.4.4 smeh register

The smeh register stores virtual addresses in TLB access and VPNs when TLB exceptions occur. The ASID indicates the process ID corresponding to the current page.




Fig. 6.6: Smeh Register Descriptions

**VPN: the virtual page number**

This field is updated by hardware when the TLB is read or a page error exception occurs. Software writes a value to this field before writing values to TLB entries.

#### Pagesize: the page size

The page size is indicated by using a one-hot, where 100 indicates a size of 4 KB, 010 indicates a size of 2 MB, and 001 indicates a size of 1 GB.

This field is updated by hardware when the TLB is read. Software writes a value to this field before writing values to TLB entries.

#### ASID: the ASID used

This field stores the ID of the current address space identified by the operating system. It is used to distinguish between processes.

This field is updated by hardware when the TLB is read. Software writes a value to this field before writing values to TLB entries.

#### 6.2.4.5 smel register

The smel register stores physical addresses in TLB access and page attributes.




Fig. 6.7: SMEL Register Descriptions

#### PPN: a 28-bit physical page number

##### SO – Strong Order

Indicates the access order required by memory.

1' b0: No strong order (Normal memory)

1' b1: Strong order(Device)

##### C – Cacheable

1' b0: Uncacheable

1' b1: Cacheable

##### B – Buffer

1' b0: Unbufferable

1' b1: Bufferable

### SH – Shareable

Indicates whether the page is shareable.

1' b0: Unshareable

1' b1: Shareable

### Sec (T – Trustable)

Indicates whether the page belongs to the trustable or non-trustable world. This bit is available only when TEE pro extension is configured.

1' b0: non-trustable

1' b1: trustable

### RSW – Reserved for Software

A bit reserved for software to implement custom page table features. The default value is 2' b0.

### D – Dirty

When the D bit is 1, it indicates whether data can be/has been written to the page.

1' b0: indicates that data has not been written/cannot be written to the page.

1' b1: indicates that data has been written/can be written to the page.

When the D bit is 0, a write operation to the page will trigger a page fault (store) exception. You can maintain the meanings of values of the D bit in the exception program through software.

### A – Accessed

When the A bit is 1, it indicates that the page is accessible. When the A bit is 0, it indicates that the page is inaccessible. Access to the page will trigger a page fault exception for the corresponding access type.

1' b0: indicates that the page is accessible.

1' b1: indicates that the page is accessible.

### G – Global

The global page ID, which indicates whether the page can be shared by multiple processes.

1' b0: indicates that the page is non-shareable and that the ASID is exclusive.

1' b1: indicates that the page is shareable.

### U – User

Indicates whether the page is accessible in user mode (U-mode).

1' b0: indicates that the page is inaccessible in U-mode. Access to the page in U-mode will trigger a page fault exception. **The default value is 1' b0.**

1' b1: indicates that the page is accessible in U-mode.

**X W R: executable, writable, readable**

Table 6.4: XWR permissions

X	W	R	Meaning
0	0	0	Pointer to next level of page table
0	0	1	Read-only page
0	1	0	Reserved for future use
0	1	1	Read-write page
1	0	0	Execute-only page
1	0	1	Read-execute page
1	1	0	Reserved for future page
1	1	1	Read-write-execute page

**V – Valid**

Indicates whether the physical page has been mapped to a virtual page. If the V bit of a page is 0, access to the page will cause a page fault exception.

1' b0: indicates that the physical page has not been mapped to a virtual page.

1' b1: indicates that the physical page has been mapped to a virtual page.

## 6.3 PMP

### 6.3.1 Overview

The PMP unit of C910/C920 complies with the RISC-V standard. The PMP unit checks the access permission on a physical address to determine whether the CPU has the read/write/execution permissions on the address in current mode.

The PMP unit of C910/C920 provides the following features:

- Supports 8/16 PMP entries, which are identified and indexed by 0 to 15.
- Supports the minimum address split granularity of 4 KB.
- Supports the OFF, top of range (TOR), and naturally aligned power-of-2 region (NAPOT) address matching modes, but not the naturally aligned four-byte region (NA4) mode.
- Supports three permissions: readable, writable, and executable.

- Supports software locks for PMP entries.

### 6.3.2 PMP control registers

A PMP entry consists of an 8-bit configuration register and a 64-bit address register. All PMP control registers are accessible in M-mode. Access to PMP control registers in other modes will trigger illegal instruction exceptions.

#### 6.3.2.1 Physical memory protection configuration (pmpcfg) register

The pmpcfg register supports permission configuration for 8 entries.

63	56	55	48	47	40	39	32	31	24	23	16	15	8	7	0	
entry7_cfg		entry6_cfg		entry5_cfg		entry4_cfg		entry3_cfg		entry2_cfg		entry1_cfg		entry0_cfg		pmpcfg0
8		8		8		8		8		8		8		8		

Fig. 6.8: Layout of the pmpcfg register

7	6	5	4	3	2	1	0
L (WARL)	O (WARL)		A (WARL)		X (WARL)	W (WARL)	R (WARL)
1	2		2		1	1	1

Fig. 6.9: pmpcfg register

For more information about the pmpcfg register, see [Table 6.5](#).

Table 6.5: Descriptions of the pmpcfg register

Bit	Name	Description
0	R	The readable attribute of the entry. <b>0</b> : indicates that the address matching the entry is non-readable. <b>1</b> : indicates that the address matching the entry is readable.
1	W	The writable attribute of the entry. <b>0</b> : indicates that the address matching the entry is non-writable. <b>1</b> : indicates that the address matching the entry is writable.
2	X	The executable attribute of the entry. <b>0</b> : indicates that the address matching the entry is non-executable. <b>1</b> : indicates that the address matching the entry is executable.
4:3	A	The address matching mode of the entry. <b>00</b> : indicates the OFF mode, in which the entry is invalid. <b>01</b> : indicates the TOR mode, in which the address of the adjacent entry is used as the matching range. <b>10</b> : indicates the NA4 mode, in which the matching range is 4 bytes. This mode is not supported. <b>11</b> : indicates the NAPOT mode, in which the matching range is a power of 2 and is at least 4 KB.
7	L	The lock enable bit of the entry. <b>0</b> : indicates that access in M-mode will succeed, and access results in S-mode/U-mode depend on the R/W/X settings. <b>1</b> : indicates that the entry is locked and cannot be modified. In TOR mode, the address register of the previous entry cannot be modified either. Access results in all modes depend on the R/W/X settings.

In TOR mode, assuming that the access address is A, the condition for hitting entry i is as follows:  $\text{pmpaddr}(i-1) \leq A < \text{pmpaddr}(i)$ . The lower boundary of entry 0 is 0.

Addresses and corresponding protection region sizes in NAPOT mode are shown in `Protection_region_code`.



Table 6.6: Protection region code :name: Protection\_region\_code

pmpaddr[37:9]	pmpcfg.A	Protection region size	Remarks
a_aaaa_aaaa_aaaa_aaaa_aaaa_aaaa_aaa0	NAPOT	4KB	Supported
a_aaaa_aaaa_aaaa_aaaa_aaaa_aaaa_aa01	NAPOT	8KB	Supported
a_aaaa_aaaa_aaaa_aaaa_aaaa_aaaa_a011	NAPOT	16KB	Supported
a_aaaa_aaaa_aaaa_aaaa_aaaa_aaaa_0111	NAPOT	32KB	Supported
a_aaaa_aaaa_aaaa_aaaa_aaaa_aaaa0_1111	NAPOT	64KB	Supported
a_aaaa_aaaa_aaaa_aaaa_aaaa_aa01_1111	NAPOT	128KB	Supported
a_aaaa_aaaa_aaaa_aaaa_aaaa_a011_1111	NAPOT	256KB	Supported
a_aaaa_aaaa_aaaa_aaaa_aaaa_0111_1111	NAPOT	512KB	Supported

The PMP unit of C910/C920 supports the minimum granularity of 4 KB in NAPOT mode, and does not support the NA4 mode.

### 6.3.2.2 Physical memory protection address (pmpaddr) register

The PMP unit provides pmpaddr 0 to pmpaddr 7/15 for storing physical addresses of entries.

As defined in the RISC-V standard, pmpaddr registers store bit [39:2] of physical addresses. The PMP unit of C910/C920 supports the minimum granularity of 4 KB. Therefore, bit[8:0] is not used for address authentication logic.

	63	38	37	9	8	0
	0		address [37:9] (WARL)		0 (WARL)	
Reset	0		0		0	

Fig. 6.10: pmpaddr registers

## 6.4 Memory access order

The following summarizes the processes of accessing an address space by C910/C920 in different scenarios.

Scenario 1: without VA-PA translation

To access a PA:

- Obtain the address attribute from the sysmap.h file.
- Perform PMP checks to determine whether the XWR permissions conform to the PMP settings.
- Access the address.

Scenario 2: with VA-PA translation

To access a VA:

- Translate the address by using the MMU to obtain the corresponding PTE.
- Obtain the following information from the PTE: the PA, address attribute (Note 1), and XWR permissions.
- Perform PMP checks to determine whether the XWR permissions conform to the PMP settings. (The minimum XWR permissions defined in the PMP register and PTE prevail.)
- Access the address.

(Note 1) When the MAEE field is 1, the address attribute comes from the PTE. When the MAEE field is 0, the address attribute comes from the sysmap.h file.

### 7.1 Memory Subsystem Overview

Each core of C910/C920 has its own I-Cache and D-Cache. Two cores share one L2 cache. Data coherence among multiple cores is maintained by hardware.

### 7.2 L1 I-Cache

#### 7.2.1 Overview

The L1 I-Cache provides the following features:

- 2-way set-associative, with a cache line size of 64 bytes;
- Virtually indexed, physically tagged (VIPT);
- Data width for access: 128 bits;
- First-in, first-out (FIFO);
- Invalidation to I-Cache or single cache line supported;
- Instruction prefetch supported;
- Way prediction supported;
- D-Cache snooping after a request misses the I-Cache (this feature can be enabled and disabled).

### 7.2.2 Way prediction

The C910/C920 I-Cache adopts the 2-way set-associative structure. To reduce power consumption in parallel access to two caches, C910/C920 implements I-Cache way prediction. When way prediction information is valid, access to invalid data ways is disabled, and the CPU accesses data only in the predicted way. You can configure the IWPE field in the mhint register to enable I-Cache way prediction.

Way prediction can be classified into the following two types by instruction fetch behavior:

**Sequential access:** When the CPU consecutively fetches instructions in a line, the CPU predicts way information of the current access based on the way hit information of the last access.

**Jump access:** A branch instruction obtains way prediction information of the target cache line along with the jump target address, and accesses one of the caches based on the information.

### 7.2.3 Loop acceleration buffer

C910/C920 provides a 32-byte loop acceleration buffer to cope with a large number of short loops in programs. When detecting a short-loop instruction sequence, the CPU loads it to the loop acceleration buffer. When a subsequent instruction fetch request hits the buffer, the CPU directly obtains the instruction and jump target address from the buffer and disables access to the I-Cache, branch history table, and branch and jump target predictor, reducing dynamic power consumption of instruction fetch. You can configure the LPE field in the mhint register to enable short-loop acceleration.

### 7.2.4 Branch history table

C910/C920 uses the branch history table to predict jump directions of conditional branch instructions. The branch history table is 64 KB in size. The bi-mode branch predictor predicts one branch result per cycle.

The branch history table consists of predictors and selectors. The predictors are classified into jump and non-jump predictors and are maintained in real time based on branch history information. The branch history table indexes ways based on branch history information and the address of the current branch instruction to predict the jump direction of the branch instruction.

The branch history table predicts jump directions of the following conditional branch instructions:

BEQ, BNE, BLT, BLTU, BGE, BGEU, C.BEQZ, and C.BNEZ

### 7.2.5 Branch and jump target predictor

The C910/C920 uses the branch and jump target predictor to predict jump target addresses of branch instructions. The branch and jump target predictor records the historical target addresses of branch instructions. If the current branch instruction hits the branch and jump target predictor, the recorded target address is used as the predicted target address of the current branch instruction.

The branch and jump target predictor provides the following features:

- Supports 1024 entries.
- Adopts the 2-way set-associative structure and supports selection and replacement based on the PC in the lower bits of a branch instruction.
- Maintains I-Cache way prediction information.
- Supports indexing by using a part of the PC of the current branch instruction.

The branch and jump target predictor predicts jump target addresses of the following branch instructions:

- BEQ, BNE, BLT, BLTU, BGE, BGEU, C.BEQZ, and C.BNEZ
- JAL and C.J

### 7.2.6 Indirect branch predictor

C910/C920 uses the indirect branch predictor to predict target addresses of indirect branch instructions. Indirect branch instructions obtain target addresses from registers. One indirect branch instruction can contain multiple branch target addresses, which cannot be predicted by using the conventional branch and jump target predictor. Therefore, C910/C920 uses the branch history-based indirect branch predictor to associate historical target addresses of an indirect branch instruction with its branch history information, and discretize different target addresses of one indirect branch instruction based on different branch history information. This makes it possible to predict multiple target addresses.

Indirect branch instructions include:

- JALR: except when the source register is x1 or x5
- C.JALR: except when the source register is x5
- C.JR: except when the source register is x1 or x5

### 7.2.7 Return address predictor

The return address predictor is used to quickly and accurately predict a return address when a function call ends. When the instruction fetch unit (IFU) obtains a valid function call instruction through decoding, it pushes a function return address to the return address predictor. When the IFU obtains a valid function return instruction through decoding, it pulls a function return address from the return address predictor. The return address predictor supports up to 12 nested function calls. If more than 12 function calls are nested, a target address prediction error will occur.

- Function call instructions include JAL, JALR, and C.JALR.
- Function return instructions include JALR, C.JR, and C.JALR.

For more information, see [Table 7.1](#).

Table 7.1: Instruction features

rd	rs1	rs1=rd	RAS action
!link	!link	-	none
!link	link	-	pop
link	!link	-	push
link	link	0	push and pop
link	link	1	push

### 7.2.8 Fast jump target predictor

To improve efficiency of the IFU in consecutive jumps, C910/C920 provides a fast jump target predictor at level 1 of the IFU. When the IFU jumps consecutively, the fast jump target predictor records the address of the second jump instruction and the jump target address. If an instruction fetch request hits the fast jump target predictor, the IFU starts to jump at level 1, reducing performance loss of at least one cycle.

The fast jump target predictor predicts jump target addresses of the following branch instructions:

- BEQ, BNE, BLT, BLTU, BGE, BGEU, C.BEQZ, and C.BNEZ
- JAL and C.J
- Function return instructions

## 7.3 L1 D-Cache

### 7.3.1 Overview

The L1 D-Cache provides the following features:

- 2-way set-associative, with a cache line size of 64 bytes;
- Physically indexed, physically tagged (PIPT);
- Maximum data width per read access: 128 bits, supporting byte, halfword, word, doubleword, and quadword access;
- Maximum data width per write access: 256 bits, supporting access to any combinations of bytes;
- Write policies: write-back with write-allocate, and write-back with write-no-allocate;
- First-in, first-out (FIFO);
- Invalidation and clearing by D-Cache or cache line supported;
- Multi-channel data prefetch for instructions.

### 7.3.2 Cache coherence

For requests with shareable and cacheable page attributes, data coherence between L1 D-Caches of different cores is maintained by hardware.

For requests with non-shareable and cacheable page attributes, the CPU does not maintain data coherence between L1 D-Caches. If non-shareable and cacheable pages need to be shared across cores, data coherence must be maintained by software.

C910/C920MP maintains data coherence between L1 D-Caches of different cores based on the MESI protocol. MESI indicates four states of each cache line in the D-Cache:

- M: indicates that the cache line is available only in this D-Cache and has been modified (UniqueDirty).
- E: indicates that the cache line is available only in this D-Cache and has not been modified (UniqueClean).
- S: indicates that the cache line may be available in multiple D-Caches and has not been modified (ShareClean).
- I: indicates that the cache line is not available in this D-Cache (Invalid).

### 7.3.3 Exclusive access

C910/C920 supports exclusive memory access instructions: LR and SC. You can use the two instructions to constitute a synchronization primitive such as an atomic lock to synchronize data between different processes of a core or between different cores. The LR instruction tags the address to be exclusively accessed. The SC instruction determines whether the tagged address is preempted by other processes. C910/C920 provides a local monitor in the L1 D-Cache and a global monitor in the L2 cache for each core. Each monitor consists of a state machine and an address buffer. The state machine has two states: IDLE and EXCLUSIVE.

Exclusive access to a cacheable page can be implemented with the local monitor. When the LR instruction is executed, it sets the state machine of the local monitor to the EXCLUSIVE state and stores the address to be accessed and the size to the buffer. When the SC instruction is executed, it reads the state of the local monitor, the address, and the size. If the state is EXCLUSIVE and the address exactly matches the size, the write operation is executed, a write success is returned, and the state machine is reset to the IDLE state. If the state or the address/size matching does not meet the requirement or the D-Cache is disabled, the write operation is not executed, a write failure is returned, and the state machine is reset to the IDLE state. When the write operation of another core performs matching against the local monitor at the same cache line address, the state machine is also reset to the IDLE state. The write operation in the current core or exclusive access to a different address does not affect the local monitor. In addition, the local monitor must be cleared when a process is switched.

Exclusive access to a non-cacheable page is implemented with both the local monitor and the global monitor. When the LR instruction is executed, it must set both the local monitor and the global monitor. After the local monitor passes the check, the SC instruction further checks the global monitor. If the global monitor

passes the check, the write operation is executed, a write success is returned, and the state of the state machine is cleared; otherwise, the write operation is not executed, a write failure is returned, and the state of the state machine is cleared. When the write operation of another core performs matching against a global monitor at an address, the state machine of the global monitor is reset to the IDLE state.

In C910/C920-based systems, we recommend that you use the LR and SC instructions to implement atomic locks. If the address attribute of an atomic lock is cacheable (either shareable or non-shareable), no special design is required for the SoC system. This is a typical case. If the address attribute of an atomic lock is non-cacheable, device, or strongly ordered, the system (for example, the slave client) must be integrated with an exclusive monitor. If an operation is performed in other ways, the response will be UNPREDICTABLE.

## 7.4 L2 Cache

### 7.4.1 Overview

The L2 cache provides the following features:

- 16-way set-associative, with a cache line size of 64 bytes;
- Strictly inclusive of the L1 D-Cache and L2 Cache. And non-strictly inclusive of the L1 I-Cache and L2 Cache;
- Physically indexed, physically tagged (PIPT);
- Maximum data width per access: 64 bytes;
- Write policies: write-back with write-allocate, and write-back with write-no-allocate;
- First-in, first-out (FIFO);
- Programmable RAM latency;
- Instruction prefetch and TLB prefetch supported;
- Block-based pipelining.

### 7.4.2 Cache coherence

The L2 cache of C910/C920MP maintains data coherence between D-Caches of different cores based on the MOESI protocol. MOESI indicates five states of each cache line in the D-Cache:

- M: indicates that the cache line is available only in this D-Cache and has been modified (UniqueDirty).
- O: indicates that the cache line may be available in multiple D-Caches and has been modified (ShareDirty).
- E: indicates that the cache line is available only in this D-Cache and has not been modified (UniqueClean).



- S: indicates that the cache line may be available in multiple D-Caches and has not been modified (ShareClean).
- I: indicates that the cache line is not available in this D-Cache (Invalid)

### 7.4.3 Structure

The L2 cache of C910/C920MP is built on a block-based pipelining architecture. Access addresses are discretized in two different blocks to allow parallel access and improve access efficiency.

The block mechanism is shown in Fig. 7.1 .

- The tag RAM is divided into two tag sub-blocks by PA[6]: tag bank 0 and tag bank 1, to handle two access requests in parallel within one clock cycle.
- Similarly, the data RAM is divided into two data sub-blocks by PA[6]: data bank 0 and data bank 1. Each data sub-block is further divided into four 128-bit micro blocks, to obtain one cache line in parallel.




Fig. 7.1: L2 Cache structure

### 7.4.4 RAM latency

The L2 cache has a long access latency because it is large in size. It usually takes multiple clock cycles to complete access to the L2 cache. C910/C920MP enables you to configure the access latency. You can set the setup time and latency of RAM in different processes. Detailed configurations are shown in Table 7.2 .

Table 7.2: RAM latency configurations

Item	Feature	Description
L2 TAG setup	L2 Cache Tag RAM setup: <b>1b0</b> : 0 cycles. Default value <b>1b1</b> : 1 cycle.	L2 Cache Tag The RAM setup affects only tags. The RAM access.
L2 TAG latency	L2 Cache Tag RAM latency: <b>3b000</b> : 1 cycle. Default value <b>3b001</b> : 2 cycles. <b>3b010</b> : 3 cycles. <b>3b011</b> : 4 cycles. <b>3b1xx</b> : 5 cycles.	
L2 DATA setup	L2 Cache Data RAM setup: <b>1b0</b> : 0 cycles. Default value <b>1b1</b> : 1 cycle.	L2 Cache Data The RAM setup affects only data. The RAM access.
L2 DATA latency	L2 Data RAM latency: <b>3b000</b> : 1 cycle. Default value <b>3b001</b> : 2 cycles. <b>3b010</b> : 3 cycles. <b>3b011</b> : 4 cycles. <b>3b100</b> : 5 cycles. <b>3b101</b> : 6 cycles. <b>3b110</b> : 7 cycles. <b>3b111</b> : 8 cycles.	

You can set the latency based on the time required for accessing the RAM. The default value of setup is 0. When the RAM setup time or winding length is long, you can modify setup to 1.

The number of access cycles with the preceding configurations is shown in [Table 7.3](#).

Table 7.3: Valid access latency of the tag RAM

Tag latency	Valid access latency of the tag RAM	
	TAG setup = 0	TAG setup = 1
/		
000	1	2
001	2	3
010	3	4
011	4	5
1xx	5	5

Table 7.4: Valid access latency of the data RAM

Tag latency	Valid access latency of the data RAM	
	TAG setup = 0	TAG setup = 1
/		
000	1	2
001	2	3
010	3	4
011	4	5
100	5	6
101	6	7
110	7	8
111	8	8

- The maximum valid L2 tag latency is 5 cycles.
- When tag setup is 1, one more cycle is required for access. Before the SRAM is accessed, the SRAM input signal will be flopped.
- The maximum valid L2 data latency is 8 cycles.
- When data setup is 1, one more cycle is required for access. Before the SRAM is accessed, the SRAM input signal will be flopped.

## 7.5 Accelerated memory access

This section describes the accelerated memory access features of C910/C920 L1 and L2 caches.

### 7.5.1 Instruction prefetch of the L1 I-Cache

The L1 I-Cache supports instruction prefetch. You can configure the IPLD field in the mhint register to enable this feature. When an instruction access request misses the current cache line, the next consecutive cache line is prefetched and stored to the prefetch buffer. When the instruction access request hits the prefetch buffer, the instruction is directly obtained from the prefetch buffer and written back to the I-Cache, reducing the instruction fetch latency.

This feature requires that the prefetched cache line and the current accessed cache line be on the same page, to ensure security of the instruction fetch address. In addition, you cannot allocate read-sensitive device address spaces to instruction spaces.

### 7.5.2 Multi-channel data prefetch of the L1 D-Cache

C910/C920 supports data prefetch to reduce the access latency of large-sized memory such as DDR SDRAMs. C910/C920 detects D-Cache misses to determine a fixed access mode through matching. Then the hardware

automatically prefetches cache lines and writes them back to the L1 D-Cache.

C910/C920 supports data prefetch through up to 8 channels and supports two prefetch methods: consecutive prefetch and strided prefetch (stride  $\leq 32$  cache lines).

C910/C920 also implements forward prefetch and backward prefetch (the stride is negative) to support various possible access modes.

Data prefetch is disabled when the CPU invalidates or clears the D-Cache.

You can configure the DPLD field in the mhint register to enable data prefetch and the DPLD\_DIS field to determine the number of cache lines to be prefetched at a time.

The following instructions support data prefetch:

- LB, LBU, LH, LHU, LW, LWU, and LD
- FLW and FLD
- LRB, LRH, LRW, LRD, LRBW, LRHW, LRWU, LURB, LURH, LURW, LURD, LURBU, LURHU, LURWU, LBI, LHI, LWI, LDI, LBUI, LHUI, LWUI, LDD, LWD, and LWUD

### 7.5.3 L1 adaptive write allocation mechanism

C910/C920 implements adaptive write allocation at L1. When the CPU detects consecutive memory write operations, the write allocation attribute of pages is automatically disabled.

You can configure the AMR field in the mhint register to enable L1 adaptive write allocation.

When the CPU invalidates or clears the D-Cache, adaptive write allocation is automatically disabled. After the invalidation or clearing is completed, the CPU detects consecutive memory write operations again.

The following instructions support adaptive write allocation:

- SB, SH, SW, and SD
- FSW and FSD
- SRB, SRH, SRW, SRD, SURB, SURH, SURW, SURD, SBI, SHI, SWI, SDI, SDD, and SWD

### 7.5.4 L2 prefetch mechanism

The L2 cache supports instruction prefetch and TLB prefetch. It supports the following prefetch features:

- The number of instructions prefetched at a time is software-configurable and can be 0, 1, 2, or 3. All prefetched instructions are written back to the L2 cache.
- Only one entry is prefetched from the TLB at a time.
- The prefetch range is a 4 KB page table, and addresses beyond the range will not be prefetched.

- You can use the machine-mode (M-mode) L2-cache control register (mccr2) to configure the prefetch mechanism.

## 7.6 L1/L2 cache operation instructions and registers

After the CPU is reset, the I-Cache and D-Cache are automatically invalidated and disabled by default.

Similarly, after the CPU is reset, the L2 cache is automatically invalidated. After the invalidation is completed, the L2 cache is automatically enabled and cannot be disabled. When the L1 cache is disabled, no data is written back to the L2 cache if the L2 cache is missed.

### 7.6.1 Extended registers of the L1 cache

Extended registers of the C910/C920 L1 cache are classified into the following types by feature:

- Cache enable and mode configuration: The M-mode hardware configuration register (mhcr) allows you to enable/disable the I-Cache/D-Cache and configure the write allocation and writeback modes. The supervisor-mode (S-mode) hardware configuration register (shcr) is a read-only register mapped to the mhcr register.
- Dirty page table entry clearing and invalidation: The M-mode cache operation register (mcor) allows you to clear and invalidate dirty page table entries in the I-Cache and the D-Cache.
- Cache read: The machine-mode cache access instruction register (mcins), M-mode cache access index register (mcindex), and M-mode cache access data register 0/1 (mcdata0/1) allow you to read data from the I-Cache and the D-Cache.

For more information, see *M-mode CPU control and status extension registers* and *M-mode cache access extension registers*.

### 7.6.2 Extended registers of the L2 cache

Extended registers of the C910/C920 L2 cache are classified into the following types by feature:

- L2 cache enable and latency configuration: The mccr2 register allows you to set the access latency of the L2 cache.
- L2 cache read: The mcins, mcindex, and mcdata0/1 registers allow you to read data from the L2 cache.

For more information, see *M-mode CPU control and status extension registers* and *M-mode cache access extension registers*.

### 7.6.3 L1/L2 cache operation instructions

C910/C920 provides extended L1/L2 cache operation instructions that invalidate page table entries by address, invalidate all page table entries, clear dirty page table entries by address, clear all dirty page table entries, clear and invalidate dirty page table entries by address, and clear and invalidate all dirty page table entries. For more information, see [Table 7.5](#).

Table 7.5: L1/L2 cache operation instructions

ICACHE.IALL	Invalidates all page table entries in the I-Cache.
ICACHE.IALLS	Invalidates all page table entries in the I-Cache through broadcasting.
ICACHE.IPA	Invalidates page table entries that match the specified physical addresses in the I-Cache.
ICACHE.IVA	Invalidates page table entries that match the specified virtual addresses in the I-Cache.
DCACHE.CALL	Clears all dirty page table entries in the D-Cache.
DCACHE.CIALL	Clears and invalidates all dirty page table entries in the D-Cache.
DCACHE.CIPA	Clears dirty page table entries that match the specified physical addresses in the D-Cache and invalidates the entries.
DCACHE.CISW	Clears dirty page table entries in the D-Cache based on the specified way and set and invalidates the entries.
DCACHE.CIVA	Clears dirty page table entries that match the specified virtual addresses in the D-Cache and invalidates the entries.
DCACHE.CPA	Clears dirty page table entries that match the specified physical addresses in the D-Cache.
DCACHE.CPAL1	Clears dirty page table entries that match the specified physical addresses in the L1 D-Cache.
DCACHE.CVA	Clears dirty page table entries that match the specified virtual addresses in the D-Cache.
DCACHE.CSW	Clears dirty page table entries in the D-Cache based on the specified way and set.
DCACHE.CVAL1	Clears dirty page table entries that match the specified virtual addresses in the L1 D-Cache.
DCACHE.IPA	Invalidates page table entries that match the specified physical addresses in the D-Cache.
DCACHE.ISW	Invalidates page table entries in the D-Cache based on the specified way and set.
DCACHE.IVA	Invalidates page table entries that match the specified virtual addresses in the D-Cache.
DCACHE.IALL	Invalidates all page table entries in the D-Cache.

For more information, see [Appendix B-1 Cache instructions](#).

---

## Vector Computations(not supportC910 )

---

C920 is compatible with RISC-V Vector Extension, Version 1.0-rc1-20210608 .

### 8.1 Vector programming model

C920 supports the following vector extension features:

- 32 independent vector registers from v0 to v31. Vector registers are 128 or 256 bits (VLEN=128/256), which depends on the vector computing capability option.
- Vector floating-point instructions support the FP16 and FP32 elements (SEW=16/32).
- Vector integer instructions support the INT8, INT16, INT32, and INT64 elements (SEW=8/16/32/64).
- Vector register groups are supported to improve the efficiency of vector computations. Four types of vector register groups are supported: 32, 16, 8, or 4 vector groups can be created, each of which contains 1, 2, 4, or 8 vector registers, respectively.

### 8.2 Vector control registers

Seven non-privileged control and status registers (CSRs) are added for C920:

- vstart

The vstart register specifies the position of the first element when a vector instruction is executed. After a vector instruction is executed, vstart is reset to zero. In most cases, software does not need to modify

vstart. In C920, only vector load/store instructions support non-zero vstart registers. All computational vector instructions support only vstart=0. Otherwise, instruction exceptions occur.

- vxsat

The vxsat register is valid only when the bit is set to 0. This register indicates whether the result of a fixed-point instruction is overflow.

- vxrm

The vxrm register provides four rounding modes: Round up, round to even, round towards zero, and round to odd.

- vcsr

Vector control core status register.

- vl

The vl register specifies the range of elements in the target register to be updated by a vector instruction. A vector instruction updates elements whose numbers are smaller than vl in the target register and resets elements whose numbers are greater than vl to zero. When vstart $\geq$ vl or vl=0, no element in the target register is updated.

- vtype

The vtype register defines basic data properties for vector computations, including: Invalid flag bits, element bits, and vector register groups. The vtype register also includes the EDIV bit. C920 does not support EDIV. Therefore, the EDIV bit is set to 0.

- vlenb (Vector Spec 0.8)

The vector bits of C920 are measured in bytes.

Therefore, C920 supports vector status maintenance (Vector Spec 0.8). The VS bit is defined in mstatus[10:9] to decide whether to save the vector-related register during context switching.

## 8.3 Vector exceptions

Vector instructions are classified into the following categories:

- Vector load
- Vector computation
- Vector store

Vector computation does not trigger exceptions. Vector store does not trigger exceptions because the bus ignores BRESP faults. Therefore, only vector load triggers exceptions. When an exception is triggered by vector load, the CPU discards the data that it reads and resets vstart to 0. The mepc points to the instruction. When an inexact exception occurs, mepc may point to subsequent instructions.



The CPU handles vector instruction interrupts in the same way as regular instructions. The CPU completes the current instruction and the mepc points to the next instruction. The remaining steps are the same as those in handling regular interrupts.

### 9.1 Security Requirements

This chapter describes software and hardware security design to meet trusted execution environment (TEE) requirements. System security requirements include:

- Support independent zones.
- Support zone isolation by code execution, memory access, external device, or I/O.
- Support isolation between applications and isolation between applications and the kernel within a zone.
- Support the multi-core SMP architecture.
- Support shared memory access among zones.
- Support the RISC-V 32-bit and 64-bit architectures.
- Supported trustworthy communication among zones.
- Support TEEs that comply with the GP specification.

### 9.2 Processor Security Model

The RISC-V ISA architecture supports the following privileged modes: M-mode, S-mode, and U-mode. These modes provide different execution and access permissions:

- In U-mode, only non-privileged instructions can be executed. In most cases, user applications are run in this mode.
- In S-mode, root user instructions can be executed and MMU management permissions are granted. In most cases, complex operating systems such as Linux are run in this mode.
- The M-mode provides the most execution and access privilege, including interrupt/exception handling and management, PMP management, and privileged access control management.

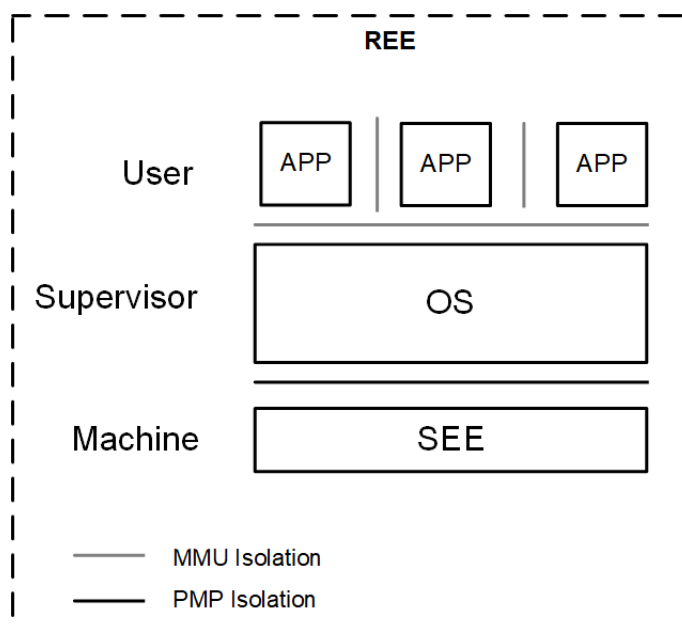


Fig. 9.1: RISC-V privileged modes

The S-mode and U-mode of RISC-V have no difference than other mainstream processor architectures, such as the supervisor and user modes of ARM. In U-mode, only non-privileged instructions can be executed. Applications running in U-mode must transition to S-mode through self-trap and access system resources under the management of the operating system. S-mode supports non-privileged instructions and privileged instructions, and provides permissions to access CSRs in S-mode. In addition, S-mode provides permissions to access MMUs. Memory protection and isolation in user mode and kernel mode are implemented through virtual memory management. The M-mode provides the most execution and access privilege. The RISC-V architecture adds privileged instructions that can be executed only in M-mode and system registers can be accessed only in M-mode, such as PMP. The most important feature of the M-mode is exception interception and handling. During exception handling, the processor transitions all exceptions to the M-mode through self-trap by default. The M-mode exception handler then forwards interrupts to the S-mode. The M-mode is usually used to run trusted firmware to adjust, allocate, and manage software and hardware resources.

To meet the isolation requirements for TEEs, security extensions are added to the Xuantie C series processors based on the RISC-V architecture. These processors can create multiple virtual zones based on software coordination. Fig. 9.2 shows the architecture. An operating system runs independently in each zone and applications run in the operating system. The operating system runs in S-mode and applications run in

U-mode. The processor can run in different zones. When the processor runs in a zone, it occupies the entire physical core in real time. The zone ID of the processor is changed to the ID of the current zone. Zone switching is completed by the trusted firmware in M-mode.

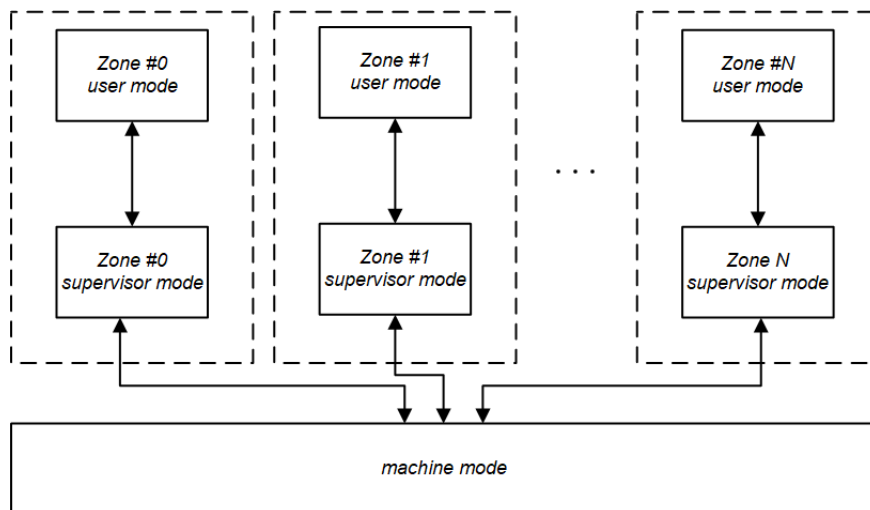


Fig. 9.2: Zones and privileged modes in Xuantie C series processors

## 9.3 System Security Architecture

### 9.3.1 Secure memory management

Each hardware thread can run in different zones through time-sharing. When a hardware thread runs in a zone, memory access is isolated in the zone. Other zones are not allowed to access the memory resources in the zone without authorization. In addition, the zone is not allowed to access memory resources in other zones without authorization. Different zones can exchange data through shared memory.

#### Physical Memory Protection (PMP)

The RISC-V architecture provides the PMP mechanism to isolate memory access in M-mode from memory access in S-mode and U-mode. PMP is configurable only in M-mode. PMP consists of multiple groups (8 to 16 groups in most cases) of address registers and configuration registers. The configuration registers can grant or revoke read, write, and execute permissions in S-mode and U-mode. PMP can also protect memory mapping I/O (MMIO). The M-mode trusted firmware can use PMP to limit the processor from accessing external device I/O.

When a hardware thread switches from one zone to another zone, the PMP configuration is also switched. The M-mode trusted firmware needs to save the PMP configuration in the current zone and loads the PMP configuration in the target zone to update the access permissions on memory and MMIO.

When multiple zones share memory, you can grant permissions to access the shared memory block to these zones by writing the access permissions to the PMP configuration table of each zone. The trusted firmware

will update the PMP table during zone switching. Fig. 9.3 shows the PMP configurations of multiple zones. The SHM area is a memory block shared by the zones.

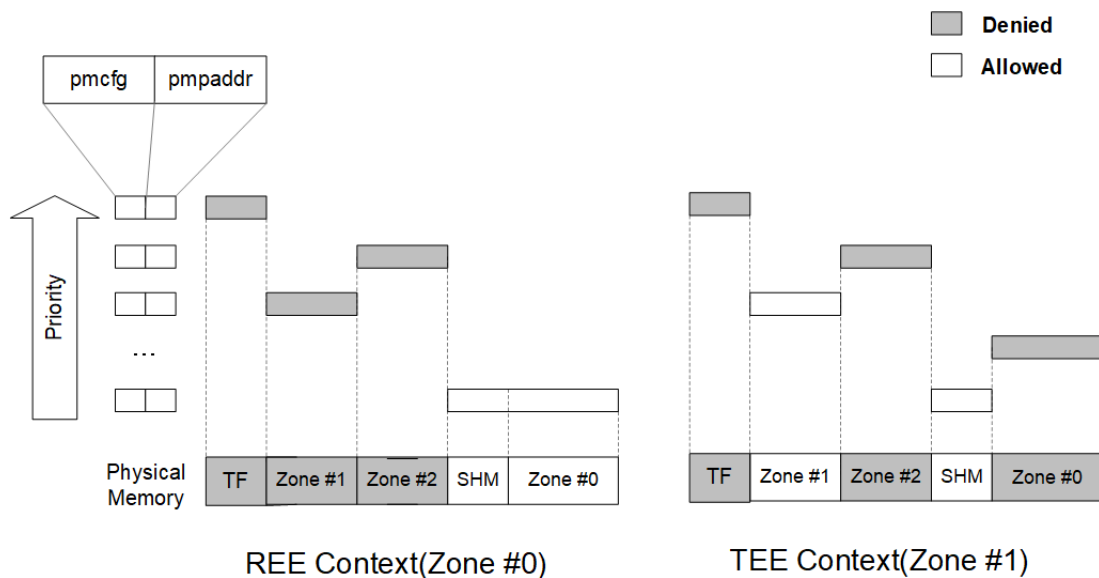


Fig. 9.3: PMP Configurations of Multiple Zones

### I/O Physical Memory Protection (IOPMP)

The RISC-V architecture provides a PMP mechanism to protect memory and MMIO access of RISC-V processors in different privileged modes.

Other master devices connected to the bus also require memory access protection: IOPMP. Same as PMP, IOPMP allows you to define access permissions. IOPMP checks whether the reads and writes sent through the bus meet the permission rules. Only legitimate reads and writes are transmitted to the target device. Typically, two methods are used to connect to an IOPMP:

1. Connect the requester to an IOPMP

Add an IOPMP between the bus and each master device, which is similar to PMP of RISC-V. Add IOPMPs for different master devices. These IOPMP are independent of each other. The design is simple but more flexible. However, the IOPMPs cannot be shared among master devices. As shown in Fig. 9.4 :




Fig. 9.4: Connect the requester to an IOPMP

2. Connect the destination device to an IOPMP

The IOPMP of the destination device needs to distinguish requests from different master devices. To do this, the requests sent by master devices must carry a master ID. As shown in Fig. 9.5:

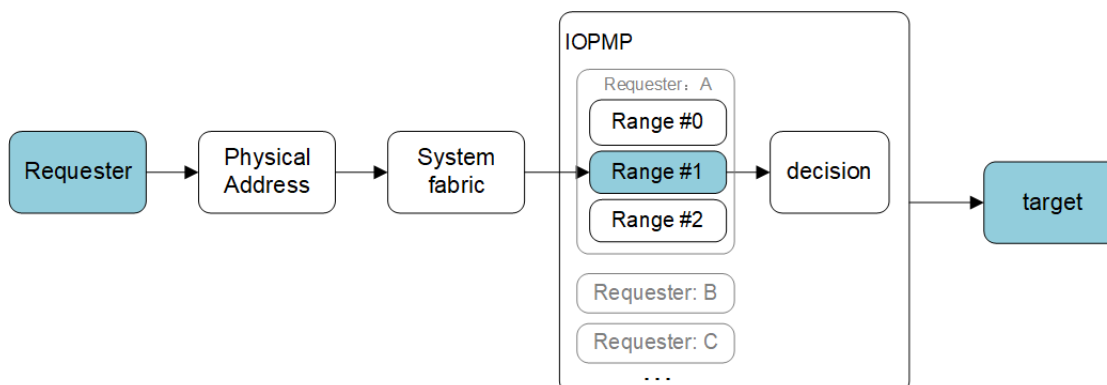


Fig. 9.5: Connect the destination device to an IOPMP

In Fig. 9.6 , the Xuantie processor mounts requesters to IOPMPs to build a secure SoC system framework.

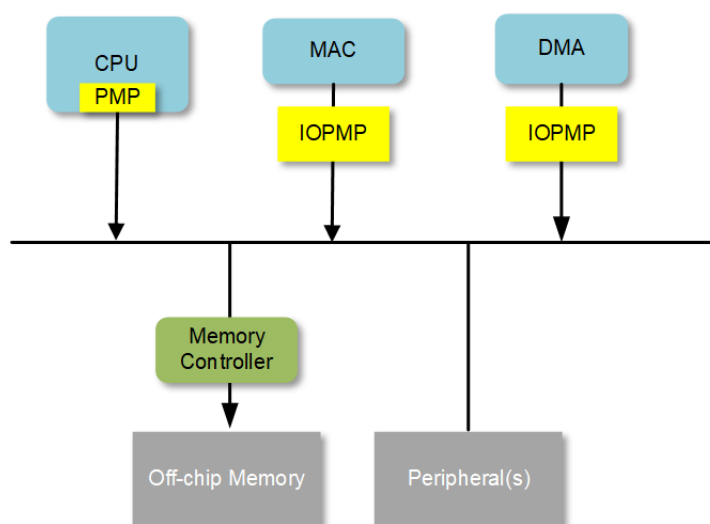


Fig. 9.6: SoC System Framework Based on PMP and IOPMP Isolation.

### Memory Management Units (MMUs)

MMUs are used to manage virtual memory in traditional operating systems. MMUs can be used to separate the user space and kernel space. The MMUs of a Xuantie processor integrates a configurable number of TLBs. Each TLB stores translations of virtual addresses to physical addresses and access permissions.

Different zones have separate TLBs to ensure that addresses in different zones are translated separately. Zone switching clears the corresponding TLB (sfence).

### Cache

When a processor runs in different zones, each zone has separate PMP configurations. PMP limits the zones from accessing physical memory MMIO and ensures that memory and I/O access among different zones are

not interfered.

In a Xuantie C series RISC-V processor, memory access that hits the cache is also protected by PMP. This means that all access to the cache must be verified by PMP. The access can reach the cache only after it passes the check. Multi-core cache coherence is also protected by PMP.

### Device Coherence Port (DCP)

Xuantie C910/C920 provides a DCP. DCP is an AXI slave interface of a processor. External master devices can access internal data with cache coherence through the DCP. This improves the efficiency of data exchange between the processor and external master device. Xuantie C910/C920 does not add protection to the DCP for access from external master devices. To protect external master devices, you need to mount the external master devices that are connected to the DCP to IOPMPs for protection.

## 9.3.2 Secure interrupts

In the PLIC specification of RISC-V, there are two modes of interrupt sources: M-mode interrupt sources and S-mode interrupt sources. M-mode interrupt sources are handled only by the M-mode. S-mode interrupt sources can be handled by the M-mode or S-mode. The M-mode has permissions to determine whether to send interrupts to the S-mode for handling. The M-mode of the RISC-V architecture provides interrupt interception to help isolate interrupts of different zones. [Table 9.1](#) describes how interrupts of different modes are handled.

Table 9.1: RISC-V interrupt handling model

Target mode of the interrupt source	Processor current mode	Delegation	Whether the interrupt is handled	Mode that handles the interrupt
M-mode	M-mode	Invalid	Yes	M-mode
	S-mode	Invalid	Yes	M-mode
	U-mode	Invalid	Yes	M-mode
S-mode S-mode	M-mode	0	Yes	M-mode
		1	No	
	M-mode	0	Yes	M-mode
		1	Yes	S-mode
	U-mode	0	Yes	M-mode
		1	Yes	S-mode

Interrupts are handled in the following ways based on the interrupt interception feature of the M-mode of RISC-V:

1. M-mode interrupt distribution
2. Interrupt groups

### M-mode interrupt distribution

The M-mode supports external interrupt interception. All external interrupts need to transition to the M-mode through self-trap. The M-mode trusted firmware will centrally manage all external interrupts, identify interrupt sources, and forward interrupts to different zones to handle these interrupts. This method meets the requirements for isolating interrupts among different zones. However, interrupts are forwarded by the trusted firmware. The trusted firmware needs to switching the context of the zone during interrupt forwarding, which delays interrupt handling.




Fig. 9.7: M-mode Interrupt Distribution in Xuantie C series processors

When this method is used, all external interrupts are sent to the M-mode trusted firmware. The trusted firmware first saves the fields of the current zone and reads the number of the external interrupt. Then, the trusted firmware selects a destination zone from the zone interrupt allocation table and obtains the entry of the interrupt handler of the destination zone. The trusted firmware can obtain the address of the interrupt entry by querying the stvec register. Before forwarding the interrupt to the entry, the trusted firmware needs to switch to the PMP configuration of the destination zone, checks the legitimacy of the address of the interrupt handler, and finally executes the mret instruction to switch to the destination zone. After the interrupt is handled, the interrupt handler needs to transition to the M-mode through an ecall. The M-mode trusted firmware will restore the fields of the interrupted zone and switch back to the zone.

### Interrupt groups

Interrupt forwarding through the M-mode severely delays interrupt handling. In addition, after an interrupt handler handles the interrupt, it needs to transition to the M-mode through an ecall. This results in incompatibility with the existing interrupt handlers (especially for Linux).

The PLIC provides separate control over each interrupt source and target. This means that the hardware thread to which the interrupts of an interrupt source are forwarded and the mode of the hardware thread can be separately configured. Currently, the execution environments of processors are classified into Rich



Execution Environments (REEs) and Trusted Execution Environments (TEEs). Secure interrupts are handled in TEEs. Most hardware interrupts are regular interrupts. Only a few number of hardware interrupts, such as secure timers, are secure interrupts. To reduce the interrupt handling delay posed by the M-mode, this solution creates interrupt groups. Interrupts from interrupt sources in the current zone are handled in the zone. Interrupts from interrupt sources that do not belong to the current zone are handled in M-mode. Interrupt context scenarios:

- The REE generates regular interrupts.
- The REE generates secure interrupts.
- The TEE generates regular interrupts.
- The TEE generates secure interrupts.

#### **The REE generates regular interrupts or the REE generates secure interrupts**

When the processor runs in the REE (Zone #0), the trusted firmware needs to perform the following operations:

1. Enables the S-mode for the interrupt source of regular interrupts.
2. Enables the M-mode for the interrupt source of secure interrupts.
3. Resets the first bit (SSIE\_DELEG), fifth bit (STIE\_DELEG), and ninth bit (SEIE\_DELEG) of the mideleg register. Assume that software interrupts and clock interrupts are configured as regular interrupts.
4. Enables mstatus.MIE and mstatus.SIE, and enables mie.MEIE, mie.MSIE, mie.MTIE, mie.SEIE, mie.SSIE, and mie.STIE.

#### **The TEE generates regular interrupts or the TEE generates secure interrupts**

When the processor runs in the TEE (Zone #1), the trusted firmware needs to perform the following operations:

1. Enables the S-mode for the interrupt source of regular interrupts.
2. Enables the M-mode for the interrupt source of secure interrupts.
3. Resets the first bit (SSIE\_DELEG), fifth bit (STIE\_DELEG), and ninth bit (SEIE\_DELEG) of the mideleg register. Assume that software interrupts and clock interrupts are configured as regular interrupts.
4. Enables mstatus.MIE and mstatus.SIE, and enables mie.MEIE, mie.MSIE, mie.MTIE, mie.SEIE, mie.SSIE, and mie.STIE.

### **9.3.3 Secure Access Control**

The M-mode is the most privileged mode that a hardware thread can run in RISC-V. A hardware thread running in M-mode has full access permissions on memory, I/O, and underlying features that are required for




Fig. 9.8: Interrupt Handling Rules When the Processor Is Running in Zone #0




Fig. 9.9: Interrupt Handling Rules When the Processor Is Running in Zone #1

booting and configuring the operating system. The M-mode is a privileged mode that must be implemented by all standard RISC-V processors. Simple RISC-V microcontrollers support only the M-mode.

The most important feature of the M-mode is exception interception and handling. By default, when an exception occurs (regardless of the privileged mode), the control permissions are transferred to the exception handler in M-mode. However, most exceptions in Linux are handled in S-mode. The exception handler in M-mode can redirect exceptions to the S-mode but these operations will severely delay exception handling. To address this issue, RISC-V provides the exception delegation mechanism. This mechanism can selectively transfer interrupts and exceptions to the S-mode for handling and bypass the M-mode. The Machine Interrupt Delegation (mideleg) CSR controls the interrupts or exceptions that are transferred to the S-mode.

Take note that control permissions are not transferred to a mode with less privilege when an interrupt or exception occurs, regardless of the delegation settings. Interrupts and exceptions in M-mode are handled only in M-mode. Interrupts and exceptions in S-mode are handled in M-mode or S-mode depending on the delegation settings. These interrupts and exceptions are never handled in U-mode.

The M-mode is sufficient for simple embedded systems. However, it is applicable only if the entire code repository is trusted because the M-mode provides full access to the hardware platform. In most cases, not all application code can be trusted because it is difficult to verify the security of every application. Therefore, RISC-V provides this mechanism to protect systems against untrusted code and isolate untrusted processes. The

untrusted code must be limited to access only the authorized memory block. Processors that support the M-mode and S-mode/U-mode support PMP, which allows the M-mode to specify the memory addresses that the S-mode/U-mode can access. PMP can also limit MMIO access. With the help of PMP, the M-mode can limit untrusted users or super users from accessing the memory and external devices.

### 9.3.4 Secure Debugging

Currently, C910/C920 does not support individual zone debugging. Only global zone debugging can be enabled or disabled.

## 10.1 Core local interrupt (CLINT) controller

C910/C920 implements the CLINT controller. It is a memory address mapping module that handles software and timer interrupts.

### 10.1.1 CLINT register address mapping

The CLINT controller occupies a 64 KB memory space. Addresses in the upper 13 bits depend on the SoC hardware integration. Address mapping in the lower 27 bits is shown in [Table 10.1](#). All registers support only access to word-aligned addresses.

Table 10.1: Memory-mapped addresses in CLINT registers

Address	Name	Type	Initial value	Description
0x4000000	MSIP0	Read/Write	0x00000000	The machine-mode (M-mode) software interrupt pending register for core 0. The upper bits are tied to 0, and bit [0] is valid.

Continued on next page

Table 10.1 – continued from previous page

Address	Name	Type	Initial value	Description
0x4000004	MSIP1	Read/Write	0x00000000	The M-mode software interrupt pending register for core 1. The upper bits are tied to 0, and bit [0] is valid.
0x4000008	MSIP2	Read/Write	0x00000000	The M-mode software interrupt pending register for core 2. The upper bits are tied to 0, and bit [0] is valid.
0x400000C	MSIP3	Read/Write	0x00000000	The M-mode software interrupt pending register for core 3. The upper bits are tied to 0, and bit [0] is valid.
Reserved	-	-	-	-
0x4004000	MTIMECMPL0	Read/Write	0xFFFFFFFF	The M-mode clock timer compare value register (the lower 32 bits) for core 0.
0x4004004	MTIMECMPH0	Read/Write	0xFFFFFFFF	The M-mode clock timer compare value register (the upper 32 bits) for core 0.
0x4004008	MTIMECMPL1	Read/Write	0xFFFFFFFF	The M-mode clock timer compare value register (the lower 32 bits) for core 1.
0x400400C	MTIMECMPH1	Read/Write	0xFFFFFFFF	The M-mode clock timer compare value register (the upper 32 bits) for core 1.
0x4004010	MTIMECMPL2	Read/Write	0xFFFFFFFF	The M-mode clock timer compare value register (the lower 32 bits) for core 2.
0x4004014	MTIMECMPH2	Read/Write	0xFFFFFFFF	The M-mode clock timer compare value register (the upper 32 bits) for core 2.
0x4004018	MTIMECMPL3	Read/Write	0xFFFFFFFF	The M-mode clock timer compare value register (the lower 32 bits) for core 3.
0x400401C	MTIMECMPH3	Read/Write	0xFFFFFFFF	The M-mode clock timer compare value register (the upper 32 bits) for core 3.
Reserved	-	-	-	-

Continued on next page

Table 10.1 – continued from previous page

Address	Name	Type	Initial value	Description
0x400C000	SSIP0	Read/Write	0x00000000	The supervisor-mode (S-mode) software interrupt pending register for core 0. The upper bits are tied to 0, and bit [0] is valid.
0x400C004	SSIP1	Read/Write	0x00000000	The S-mode software interrupt pending register for core 1. The upper bits are tied to 0, and bit [0] is valid.
0x400C008	SSIP2	Read/Write	0x00000000	The S-mode software interrupt pending register for core 2. The upper bits are tied to 0, and bit [0] is valid.
0x400C00C	SSIP3	Read/Write	0x00000000	The S-mode software interrupt pending register for core 3. The upper bits are tied to 0, and bit [0] is valid.
Reserved	-	-	-	-
0x400D000	STIMECMPL0	Read/Write	0xFFFFFFFF	The S-mode clock timer compare value register (the lower 32 bits) for core 0.
0x400D004	STIMECMPH0	Read/Write	0xFFFFFFFF	The S-mode clock timer compare value register (the upper 32 bits) for core 0.
0x400D008	STIMECMPL1	Read/Write	0xFFFFFFFF	The S-mode clock timer compare value register (the lower 32 bits) for core 1.
0x400D00C	STIMECMPH1	Read/Write	0xFFFFFFFF	The S-mode clock timer compare value register (the upper 32 bits) for core 1.
0x400D010	STIMECMPL2	Read/Write	0xFFFFFFFF	The S-mode clock timer compare value register (the lower 32 bits) for core 2.
0x400D014	STIMECMPH2	Read/Write	0xFFFFFFFF	The S-mode clock timer compare value register (the upper 32 bits) for core 2.

Continued on next page

Table 10.1 – continued from previous page

Address	Name	Type	Initial value	Description
0x400D018	STIMECMPL3	Read/Write	0xFFFFFFFF	The S-mode clock timer compare value register (the lower 32 bits) for core 3.
0x400D01C	STIMECMPH3	Read/Write	0xFFFFFFFF	The S-mode clock timer compare value register (the upper 32 bits) for core 3.
Reserved	-	-	-	-

### 10.1.2 Software interrupts

The CLINT controller can generate software interrupts.

Software interrupts are controlled by the software interrupt pending registers configured with address mappings. M-mode software interrupts are controlled by the machine software interrupt pending (MSIP) register. S-mode software interrupts are controlled by the supervisor software interrupt pending (SSIP) register.

You can set the xSIP bit to 1 to generate software interrupts or reset it to 0 clear software interrupts. CLINT S-mode software interrupt requests are valid only when the CLINTEE bit is enabled for the corresponding core.

In M-mode, the CPU is allowed to access and modify all software interrupt registers. In S-mode, the CPU is allowed to access and modify only the SSIP register. In user mode (U-mode), the CPU has no access to software interrupt registers.

The two groups of registers have the same structure. The bit layout and definition of the registers are shown in Fig. 10.1 .





Fig. 10.1: MSIP register

#### MSIP: the machine software interrupt pending bit

This bit indicates the status of M-mode software interrupts.

- When the MSIP bit is 1, valid M-mode software interrupt requests are available.
- When the MSIP bit is 0, no valid M-mode software interrupt requests are available.

SSIP register



### SSIP: the supervisor software interrupt pending bit

This bit indicates the status of S-mode software interrupts.

- When the SSIP bit is 1, valid S-mode software interrupt requests are available.
- When the SSIP bit is 0, no valid S-mode software interrupt requests are available.

### 10.1.3 Timer interrupts

The CLINT controller can generate timer interrupts.

A multi-core system has only one 64-bit system timer, *mtime*. *mtime* must run in the always-on voltage domain. *mtime* cannot be written but can be reset. The current value of *mtime* can be read from the time register of the performance monitoring unit (PMU). *mtime* is used to provide a unified time reference for multiple cores.

Each core has a group of 64-bit *M-mode clock timer compare value registers* (*mtimecmp* and *mtimecmph*) and a group of 64-bit *S-mode clock timer compare value registers* (*stimecmp* and *stimecmph*). You can modify the upper or lower 32 bits of these registers through word-aligned address access.

The CLINT controller compares the value of {CMPH[31:0], CMPL[31:0]} with the current value of *mtime* to determine whether to generate a timer interrupt. When the value of {CMPH[31:0], CMPL[31:0]} is greater than the current value of *mtime*, the CLINT controller does not generate an interrupt. When the value of {CMPH[31:0], CMPL[31:0]} is less than or equal to the current value of *mtime*, the CLINT controller generates a corresponding timer interrupt. You can rewrite the value of the *mtimecmp*/*stimecmp* register to clear the corresponding timer interrupt. S-mode timer interrupt requests are valid only when the CLINTEE bit is enabled for the corresponding core.

In M-mode, the CPU is allowed to access and modify all timer interrupt registers. In S-mode, the CPU is allowed to access and modify only the *stimecmp* and *stimecmph* registers. In U-mode, the CPU has no access to timer interrupt registers.

The two groups of registers have the same structure. The bit layout and definition of the registers are shown in Fig. 10.2 .

#### **mtimecmph/mtimecmpl: the M-mode clock timer compare value registers for the upper bits and the lower bits**

These registers store timer compare values.

- *mtimecmph*: stores the upper 32 bits of timer compare values.



	63	32	31	0
	MTIMECMPH		MTIMECMPL	
Reset	0xffffffff		0xffffffff	

Fig. 10.2: mtimecmph/mtimecmpl registers

- mtimecmpl: stores the lower 32 bits of timer compare values.

	63	32	31	0
	STIMECMPH		STIMECMPL	
Reset	0xffffffff		0xffffffff	

Fig. 10.3: stimecmph/stimecmpl registers

**stimecmph/stimecmpl: the S-mode clock timer compare value registers for the upper bits and the lower bits**

These registers store timer compare values.

- stimecmph: stores the upper 32 bits of timer compare values.
- stimecmpl: stores the lower 32 bits of timer compare values.

## 10.2 Platform-level interrupt controller (PLIC)

The PLIC controls sampling, priority arbitration, and distribution of external interrupt sources.

In the PLIC model, the M-mode and S-mode of each core can act as valid interrupt targets.

The PLIC of C910/C920 provides the following features:

- Interrupt distribution for 4 cores and 8 interrupt targets;
- Sampling of 1023 interrupt sources, supporting level and pulse interrupts;
- 32 interrupt priorities;
- Independent enable for each interrupt target;
- Independent interrupt threshold for each interrupt target;
- Configurable access permissions on PLIC registers.

### 10.2.1 Interrupt arbitration

In the PLIC, only interrupt sources that meet the specified conditions are involved in arbitration on an interrupt target. The conditions include:

- The interrupt source is in the pending state ( $IP = 1$ ).
- The interrupt priority is greater than 0.
- The enable bit for the interrupt target is enabled.

When multiple interrupts for an interrupt target are in the pending state, the PLIC selects the interrupt with the highest priority through arbitration. In the PLIC of C910/C920, M-mode interrupts have higher priorities than S-mode interrupts. In the same privilege mode, a larger value of the priority configuration register indicates a higher priority. Interrupts with a priority of 0 are invalid. If multiple interrupts have the same priority, they will be handled in ascending order of IDs.

The PLIC stores interrupt IDs that are determined based on arbitration results to the interrupt claim/complete register of the corresponding interrupt target.

### 10.2.2 Interrupt request and response

When the PLIC has a valid interrupt request for an interrupt target and the interrupt priority is higher than the interrupt threshold of the interrupt target, the PLIC sends the interrupt request to the interrupt target. When receiving the interrupt request, the interrupt target sends an interrupt response message to the PLIC if it is able to respond to the interrupt request.

The interrupt response mechanism functions as follows:

- The interrupt target initiates a read operation to the corresponding interrupt claim/complete register. The read operation returns the interrupt ID determined by the PLIC. The interrupt target proceeds to further processing based on the interrupt ID. If the interrupt ID is 0, no valid interrupt request is available, and the interrupt target ends the interrupt handling process.
- After receiving the read operation initiated by the interrupt target and returning the interrupt ID, the PLIC resets the IP bit of the interrupt source corresponding to the interrupt ID, and blocks subsequent sampling on the interrupt source before the current interrupt is completed.

### 10.2.3 Interrupt completion

After interrupt handling is completed, the interrupt target sends an interrupt completion message to the PLIC. The interrupt completion mechanism functions as follows:

- The interrupt target initiates a write operation to the corresponding interrupt claim/complete register, to write the ID of the completed interrupt to the register. If the interrupt is a level interrupt, the external interrupt source must be cleared before the write operation is initiated.
- After receiving the interrupt completion message, the PLIC does not update the interrupt claim/complete register, but unblocks sampling on the interrupt source corresponding to the interrupt ID to end the interrupt handling process.

### 10.2.4 PLIC register address mapping

The PLIC occupies a 64 MB memory space. Addresses in the upper 13 bits depend on the SoC hardware integration. Address mapping in the lower 27 bits is shown in Table 10.2. All registers support only access to word-aligned addresses. *PLIC registers are accessible through the load word instruction. Access results are placed in the lower 32 bits of 64-bit GPRs.*

Note: Registers not supported by C910/C920 are marked as reserved.

Table 10.2: PLIC register address mapping

Address	Name	Type	Initial value	Described
0x0000000	-	-	-	-
0x0000004	PLIC_PRI01	R/W	0x0	The priority configuration register for interrupts 1 to 1023.
0x0000008	PLIC_PRI02	R/W	0x0	
0x000000C	PLIC_PRI03	R/W	0x0	
...	...	...	...	
0x0000FFC	PLIC_PRI01023	R/W	0x0	
0x0001000	PLIC_IP0	R/W	0x0	The interrupt pending register for interrupts 1 to 31.
0x0001004	PLIC_IP1	R/W	0x0	The interrupt pending register for interrupts 32 to 63.
...	...			
0x000107C	PLIC_IP31	R/W	0x0	The interrupt pending register for interrupts 992 to 1023.
Reserved	-	-	-	-
0x0002000	PLIC_H0_MIE0	R/W	0x0	The M-mode interrupt enable register for interrupts 1 to 31 in core 0.
0x0002004	PLIC_H0_MIE1	R/W	0x0	The M-mode interrupt enable register for interrupts 32 to 63 in core 0.
...	...			
0x000207C	PLIC_H0_MIE31	R/W	0x0	The M-mode interrupt enable register for interrupts 992 to 1023 in core 0.
0x0002080	PLIC_H0_SIE0	R/W	0x0	The S-mode interrupt enable register for interrupts 1 to 31 in core 0.
0x0002084	PLIC_H0_SIE1	R/W	0x0	The S-mode interrupt enable register for interrupts 32 to 63 in core 0.
...	...			
0x00020FC	PLIC_H0_SIE31	R/W	0x0	The S-mode interrupt enable register for interrupts 992 to 1023 in core 0.

Continued on next page

Table 10.2 – continued from previous page

Address	Name	Type	Initial value	Described
0x0002100	PLIC_H1_MIE0	R/W	0x0	The M-mode interrupt enable register for interrupts 1 to 31 in core 1.
0x0002104	PLIC_H1_MIE1	R/W	0x0	The M-mode interrupt enable register for interrupts 32 to 63 in core 1.
.....				
0x000217C	PLIC_H1_MIE31	R/W	0x0	The M-mode interrupt enable register for interrupts 992 to 1023 in core 1.
0x0002180	PLIC_H1_SIE0	R/W	0x0	The S-mode interrupt enable register for interrupts 1 to 31 in core 1.
0x0002184	PLIC_H1_SIE1	R/W	0x0	The S-mode interrupt enable register for interrupts 1 to 31 in core 1.
.....				
0x00021FC	PLIC_H1_SIE31	R/W	0x0	The S-mode interrupt enable register for interrupts 992 to 1023 in core 1.
Reserved	-	-	-	-
0x01FFFFC	PLIC_CTRL	R/W	0x0	The PLIC permission control register.
0x0200000	PLIC_H0_MTH	R/W	0x0	The M-mode interrupt threshold register for core 0.
0x0200004	PLIC_H0_MCLAIM	R/W	0x0	The M-mode interrupt claim/complete register for core 0.
Reserved	-	-	-	-
0x0201000	PLIC_H0_STH	R/W	0x0	The S-mode interrupt threshold register for core 0.
0x0201004	PLIC_H0_SCLAIM	R/W	0x0	The S-mode interrupt claim/complete register for core 0.
Reserved	-	-	-	-
0x0202000	PLIC_H1_MTH	R/W	0x0	The M-mode interrupt threshold register for core 0.
0x0202004	PLIC_H1_MCLAIM	R/W	0x0	The M-mode interrupt claim/complete register for core 0.
Reserved	-	-	-	-
0x0203000	PLIC_H1_STH	R/W	0x0	The S-mode interrupt threshold register for core 0.
0x0203004	PLIC_H1_SCLAIM	R/W	0x0	The S-mode interrupt claim/complete register for core 0.
Reserved	-	-	-	-

### 10.2.5 PLIC\_PRIO register

This register is used to set the priorities of interrupt sources. Register read and write permissions For more information, see the descriptions of the PLIC\_PER register. The bit layout and definition of the register are shown in Fig. 10.4 .




Fig. 10.4: PLIC\_PRIO register

## PRIO: the interrupt priority

The lower 5 bits of the PLIC\_PRIO register are writable. The PLIC\_PRIO register supports 32 interrupt priorities. Interrupts with a priority of 0 are invalid.

M-mode interrupts have higher priorities than S-mode interrupts in any conditions. In the same privilege mode, the priority 1 is the lowest priority, and the priority 31 is the highest priority. When multiple interrupts have the same priority, interrupt IDs are further compared. An interrupt with a smaller ID has a higher priority.

### 10.2.6 PLIC\_IP register

The PLIC can read the PLIC\_IP register to obtain the pending state of each interrupt. If the ID of an interrupt is N, the interrupt information is stored in IP y ( $y = N \bmod 32$ ) in the PLIC\_IP x ( $x = N/32$ ) register. Bit 0 of the PLIC\_IP0 register is tied to 0. Register read and write permissions For more information, see the descriptions of the PLIC\_PER register. The bit layout and definition of the register are shown in [Fig. 10.5](#).




Fig. 10.5: PLIC\_IP x register

- IP: the interrupt pending state bit

This bit indicates the interrupt pending state of the corresponding interrupt source.

- When the IP bit is 1, the interrupt source has pending interrupts. You can run a memory store instruction to set this bit to 1. When the sampling logic of the interrupt source detects valid level or pulse interrupts, this bit is also set to 1.
- When the IP bit is 0, the interrupt source has no pending interrupt. You can run a memory store instruction to reset this bit. After an interrupt is handled, PLIC clears the corresponding IP bit.

### 10.2.7 PLIC\_IE register

Each interrupt target has an interrupt enable bit for each interrupt source, to enable the corresponding interrupts. The M-mode interrupt enable register is used to enable M-mode external interrupts. The S-mode interrupt enable register is used to enable S-mode external interrupts.

If the ID of an interrupt is  $N$ , the interrupt enable information is stored in IE  $y$  ( $y = N \bmod 32$ ) in the PLIC\_IE  $x$  ( $x = N/32$ ) register. The IE bit corresponding to ID 0 is set to 0. For more information about the read and write permissions on the register, see the descriptions of the PLIC\_PER register.

The bit layout and definition of the register are shown in Fig. 10.6.




Fig. 10.6: PLIC\_IE  $x$  register

- **IE: the interrupt enable state bit**

This bit indicates the interrupt enable state of the corresponding interrupt source.

- When the IE bit is 1, the interrupt source is enabled for the interrupt target.
- When the IE bit is 0, the interrupt source is disabled for the interrupt target.

### 10.2.8 PLIC\_CTRL register

The PLIC\_CTRL register is used to control access permissions on PLIC registers in S-mode.

- **S\_PER: the access permission control bit**

When the S\_PER bit is 0, the CPU has access to all PLIC registers only in M-mode. In S-mode, the CPU has access only to the S-mode PLIC\_TH register and S-mode PLIC\_CLAIM register, but not to the PLIC\_PER, PLIC\_PRIO, PLIC\_IP, or PLIC\_IE register. In U-mode, the CPU has no access to any PLIC registers.




Fig. 10.7: PLIC\_CTRL register

When the S\_PER bit is 1, the CPU has access to all PLIC registers in M-mode, and has access to all PLIC registers except PLIC\_PER in S-mode. In U-mode, the CPU has no access to any PLIC registers.

### 10.2.9 PLIC\_TH register

Each interrupt target has a PLIC\_TH register. The PLIC initiates an interrupt request to an interrupt target only when the interrupt request is valid and the interrupt priority is higher than the interrupt threshold of the interrupt target. For more information about the read and write permissions on the register, see the descriptions of the PLIC\_PER register.

The bit layout and definition of the register are shown in Fig. 10.8 .




Fig. 10.8: PLIC\_TH register

#### PRIOTHRESHOLD: the priority threshold

This bit indicates the interrupt threshold of the current interrupt target. When the interrupt threshold is 0, all interrupts are allowed.

### 10.2.10 PLIC\_CLAIM register

Each interrupt target has a PLIC\_CLAIM register. When the PLIC completes arbitration, this register is updated to the interrupt ID obtained in the current arbitration. For more information about the read and write permissions on the register, see the descriptions of the PLIC\_PER register.

The bit layout and definition of the register are shown in Fig. 10.9 .

#### CLAIM\_ID: the interrupt request ID

A read operation to the register returns the ID currently stored in the register. The read operation indicates that the interrupt corresponding to the ID is in the process of handling. The PLIC starts the interrupt claim process.

	31	10	9	0
	-			CLAIM ID
Reset	0			0

Fig. 10.9: PLIC\_CLAIM register

A write operation to the register indicates that the interrupt corresponding to the ID to be written has been handled. The write operation does not update the PLIC\_CLAIM register. The PLIC starts the interrupt complete process.

## 10.3 Multi-core interrupts

This section describes two common multi-core interrupt scenarios.

### 10.3.1 Multiple cores respond to external interrupts in parallel

In the PLIC model, one interrupt source can be mapped to multiple cores. When the interrupt source generates an interrupt request, the interrupt request is in the pending state with respect to multiple cores. Different cores run in different states, and they respond to the interrupt successively and read the PLIC\_CLAIM register to obtain the interrupt ID. Design of the PLIC ensures that only the first core accessing the PLIC\_CLAIM register obtains the valid ID, and other cores obtain an invalid ID ( $ID = 0$ ) and therefore do not handle the interrupt. In this case, the interrupt is handled only once.

Mapping an interrupt to multiple cores reduces the overall interrupt response time because any one of the cores has an opportunity to handle the interrupt. However, bandwidth of the cores that obtain the invalid ID is consumed, wasting additional CPU resources.

### 10.3.2 Send software interrupts across cores

In the programming model of the CLINT controller, software interrupts are stored in dedicated registers:

- M-mode software interrupts are stored in the MSIP0, MSIP1, MSIP2, MSIP3 registers.
- S-mode software interrupts are stored in the SSIP0, SSIP1, SSIP2, SSIP3 registers.

Addresses of the preceding registers are unified and known to all cores. Each core can initiate write operations to the registers to send software interrupts to other cores or itself.



### 11.1 AXI master device interface

The master device interface of C910/C920MP supports the AMBA 4.0 AXI protocol. For more information, see *AMBA Specifications —AMBA® AXI™ and ACE™ Protocol Specification*.

#### 11.1.1 Features of the AXI master device interface

The AXI master device interface controls address accesses and data transmission between C910/C920 and the AXI bus. It provides the following features:

- Complies with the AMBA 4.0 AXI protocol.
- Supports a bus width of 128 bits.
- Supports different frequency ratios between the system clock and the CPU master clock.
- Supports flop-out of all output signals and flop of all input signals to obtain better timing.

#### 11.1.2 Outstanding capability of the AXI master device interface

This section describes the outstanding capability of the AXI master device interface provided by C910/C920.

Table 11.1: Outstanding capability of the AXI master device interface

Parameter	Value	Description
Read Issuing Capability	$8n+28$ $n$ = Number of cores	Each core can issue up to 8 non-cacheable and device read requests. All cores can issue up to 28 cacheable read requests.
Write Issuing Capability	$8n+32$ $n$ = Number of cores	Each core can issue up to 8 non-cacheable and device write requests. All cores can issue up to 32 cacheable write requests.

Table 11.2: ARID encoding of the AXI master device interface

ARID[7:0]	Scenario	Outstanding requests of each ID
{2' b10, 6' b??????}	Cacheable read requests	Each ID has no outstanding requests. All cacheable read requests are outstanding. A total of 28 outstanding requests are supported.
{1' b0, 2' b(coreid), 5' h18}	Non-cacheable and weak-ordered read requests	Non-cacheable read requests are outstanding. A total of 31 outstanding requests are supported.
{1' b0, 2' b(coreid), 5' h1d}	Non-cacheable and strong-ordered read requests	

Table 11.3: AWID encoding of the AXI master device interface

AWID[7:0]	Scenario	Outstanding requests of each ID
{3' b111, 5' b?????}	Cacheable write requests	Each ID has no outstanding requests. All cacheable write requests are outstanding. A total of 32 outstanding requests are supported.
{4' b0000, 4' b????}	Non-cacheable and weak-ordered write requests	Each ID has no outstanding requests. All non-cacheable and weak-ordered write requests are outstanding. A total of 16 outstanding requests are supported.
{1' b0, 2' b(coreid), 5' h1d}	Non-cacheable and strong-ordered write requests	Non-cacheable and strong-ordered write requests are outstanding. A total of 31 outstanding requests are supported.

Note: The ARID and AWID encoding may vary with evolution of the CPU version. Therefore, SoC integration should not depend on specific IDs, but should conform to general-purpose rules of the AXI protocol.

### 11.1.3 Supported transmission types

The AXI master device interface supports the following transmission types:

- Burst types: INCR and WRAP;
- Transmission lengths: 1 and 4;
- Exclusive access;
- Transmission sizes: quadword, doubleword, word, halfword, and byte;
- Read and write.

Note: The AXI master device interface of C910/C920 implements only a subset of all AXI transmission types. Therefore, SoC integration should not depend on specific transmission types, but should conform to general-purpose rules of the AXI protocol.

### 11.1.4 Supported response types

The AXI master device interface supports the following types of responses from slave devices:

- OKAY
- EXOKAY
- SLVERR
- DECERR

### 11.1.5 CPU behavior in different bus responses

CPU behavior in different bus responses is shown in [Table 11.4](#).

Table 11.4: Bus exception handling

RRESP/BRESP	Result
OKAY	Indicates that common transfer access succeeds or exclusive transfer access fails. If exclusive read transfer access fails, it indicates that the bus does not support exclusive transfer, and an access error exception is generated. If exclusive write transfer access fails, it indicates that lock preemption fails, and no exception is generated.
EXOKAY	Indicates that exclusive access succeeds.
SLVERR/DECERR	Indicates that an access error occurs. If this error occurs in read transfer, an exception is generated. If this error occurs in write transfer, it is ignored.

### 11.1.6 Signals supported by the AXI master device interface

Signals supported by the AXI4 master device interface are described in Table 11.5.

Table 11.5: Signals supported by AXI channels

Signal name	I/O	Reset	Definition
<b>Interface related to the address read channel</b>			
biu_pad_arid[7:0]	O	0	The address ID of the read request.
biu_pad_araddr[39:0]	O	0	The address of the read request.
biu_pad_arlen[1:0]	O	0	The burst length of the read request. 00: 1 transfer 11: 4 transfers
biu_pad_arsize[2:0]	O	0	The data bit width per cycle of the read request. 000: 1 byte 001: 2 bytes 010: 4 bytes 011: 8 bytes 100: 16 bytes
biu_pad_arburst[1:0]	O	0	The transfer type corresponding to the read request: 01: INCR 10: WRAP
biu_pad_arlock	O	0	The access method corresponding to the read request: 0: normal access 1: exclusive access
biu_pad_arcache[3:0]	O	0	The memory access type corresponding to the read request: 0000: device no n-bufferable (strong order) 0001: device bufferable (strong order) 0011: normal non-cacheable bufferable (weak order) 1111: cacheable

Continued on next page

Table 11.5 – continued from previous page

Signal name	I/O	Reset	Definition
biu_pad_arprot[2:0]	O	0	The protection type of the read request: 0   1 [2]: data   instruction [1]: secure   non-secure. The value is always 1. [0]: user   privileged
biu_pad_aruser[2:0]	O	0	The user-defined signal of the read request: [2]: L2 Cache prefetch request [1]: machine-mode (M-mode) request [0]: MMU writeback request
biu_pad_arvalid	O	0	The valid signal of the address read channel.
pad_biu_arready	I	-	The slave ready signal of the address read channel.
<b>Interface related to the data read channel</b>			
pad_biu_rid[7:0]	I	-	The data ID of the read request.
pad_biu_rdata[127:0]	I	-	The data of the read request.
pad_biu_rresp[3:0]	I	-	The response information of the read request. [1:0]: 00: OKAY 01: EXOKAY 10: SLVERR 11: DECERR [2]: PASSDIRTY [3]: ISSHARED
pad_biu_rlast	I	-	Data of the last cycle of the read request.
pad_biu_rvalid	I	-	The valid signal of the data read request.
biu_pad_rready	O	1	The ready information of the data read channel.
<b>Interface related to the address write channel</b>			
biu_pad_awid[7:0]	O	0	The address ID of the write request.
biu_pad_awaddr[39:0]	O	0	The address of the write request.
biu_pad_awlen[1:0]	O	0	The burst length of the write request. 00: 1 cycle 11: 4 cycles

Continued on next page

Table 11.5 – continued from previous page

Signal name	I/O	Reset	Definition
biu_pad_awsiz[2:0]	O	0	The data bit width per cycle of the write request. 000: 1 byte 001: 2 bytes 010: 4 bytes 011: 8 bytes 100: 16 bytes
biu_pad_awburst[1:0]	O	0	The transfer type corresponding to the read request: 01: INCR 10: WRAP
biu_pad_awlock	O	0	The access method corresponding to the read request: 0: normal access 1: exclusive access
biu_pad_awcache[3:0]	O	0	The memory access type corresponding to the read request: 0000: device no n-bufferable (strong order) 0001: device bufferable (strong order) 0011: normal non-cacheable bufferable (weak order) 0111: write-back no-allocate 1111: write-back cacheable
biu_pad_awprot[2:0]	O	0	The protection type of the write request: 0   1 [2]: data   instruction [1]: secure   non-secure. The value is always 1. [0]: user   privileged
biu_pad_awvalid	O	0	The valid signal of the address write channel.
pad_biu_arready	I	-	The slave ready signal of the address read channel.
<b>Interface related to the data write channel</b>			
biu_pad_wdata[127:0]	O	0	The data of the write request.
biu_pad_wstrb[15:0]	O	0	The valid data fields of the data.
biu_pad_wlast	O	0	Data of the last cycle.

Continued on next page

Table 11.5 – continued from previous page

Signal name	I/O	Reset	Definition
biu_pad_wvalid	O	0	The valid signal of the data write channel.
biu_pad_wready	I	-	The slave ready signal of the address read channel.
<b>Interface related to the write response channel</b>			
pad_biu_bid[7:0]	I	-	The ID of the write response.
pad_biu_bresp[1:0]	I	-	The write response information, which is the response information of the write request. [1:0]: 00: OKAY 01: EXOKAY 10: SLVERR 11: DECERR
pad_biu_bvalid	I	-	The valid signal of the write response channel.
biu_pad_bready	O	1	The ready signal of the write response channel.

### 12.1 Features of the debug unit

The debug interface provides an interaction channel between software and the CPU. You can call the debug interface to obtain information stored in registers and memory of the CPU and information about other on-chip devices. You can also call the debug interface to download programs.

C910/C920MP supports the JTAG communication protocol (also known as JTAG5) compatible with the IEEE-1149.1 standard. It can be integrated with existing JTAG components or independent JTAG controllers.

The debug interface provides the following features:

- Supports debug by using standard JTAG interfaces.
- Supports synchronous and asynchronous debug, enabling the CPU to enter the debug mode in extreme conditions.
- Supports software breakpoints.
- Supports multiple memory breakpoints.
- Enables you to check and set the values of CPU registers.
- Enables you to check and modify memory values.
- Enables the CPU to run an instruction in a single step or multiple steps.
- Enables you to quickly download programs.



- Enables the CPU to enter the debug mode after it is reset.

Debug of C910/C920 is jointly completed by the debug software, debug proxy, debugger, and debug interface. The location of the debug interface in the CPU debug environment is shown in Fig. 12.1. The debug software is connected to the debug proxy over network. The debug proxy is connected to the debugger through a USB interface. The debugger communicates with the debug interface in JTAG mode.

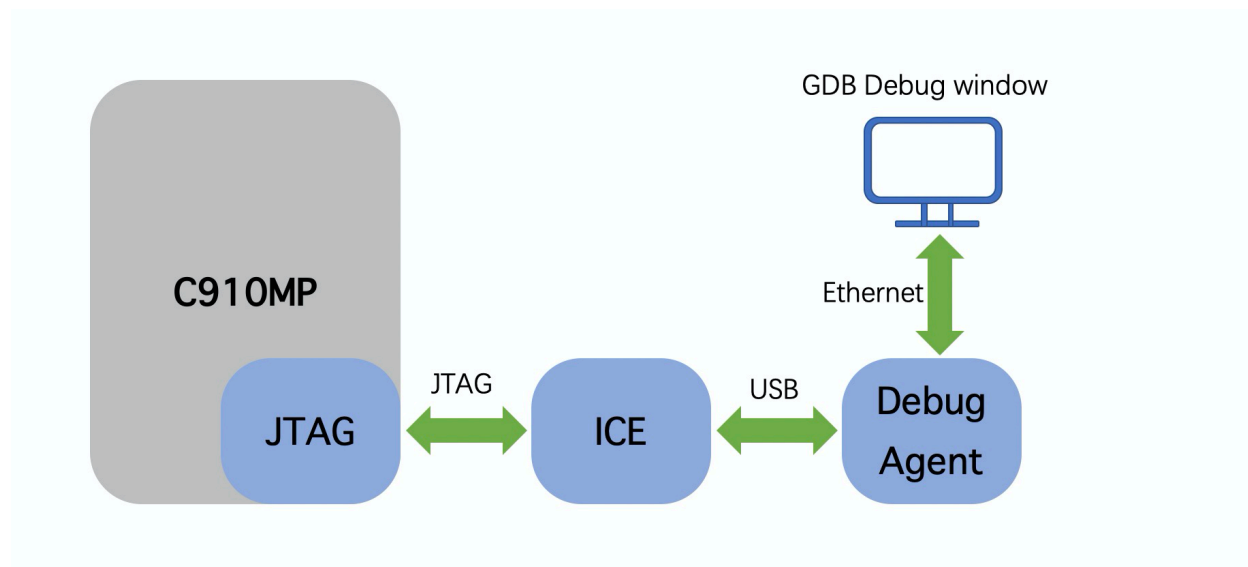


Fig. 12.1: Location of the debug interface in CPU debug environment

## 12.2 Connection between the debug unit and CPU cores

C910/C920MP adopts a multi-core single-port debug framework. It accesses the HAD unit of each core through a shared JTAG interface, and triggers the core to enter or exit the debug mode and access CPU resources. You can set the CORESEL field in the hacr register through the JTAG interface to specify a target core, and then configure HAD registers of the target core.

In the multi-core debug framework, when a core enters the debug mode, another one or more cores must also enter the debug mode; and when a core exits the debug mode, other one or more cores must also exit the debug mode. Therefore, C910/C920MP provides a unified event transmission module (ETM) to transmit debug events, including debug entry and exit events, between multiple cores. When a core receives a debug command from the ICE, it generates a debug entry or exit event and sends the event to the ETM. The ETM forwards the event to other cores to enable multiple cores to synchronously enter or exit the debug mode. The multi-core debug framework (with two cores) of C910/C920 is shown in Fig. 12.2.

- **Multiple cores enter the debug mode**

When a core enters the debug mode, it generates a DBG\_ENT signal. The EVENT\_OUTEN register of the core determines whether the DBG\_ENT signal can be sent to the ETM. When the EVENT\_OUTEN register is enabled, the DBG\_ENT signal is transmitted to other cores

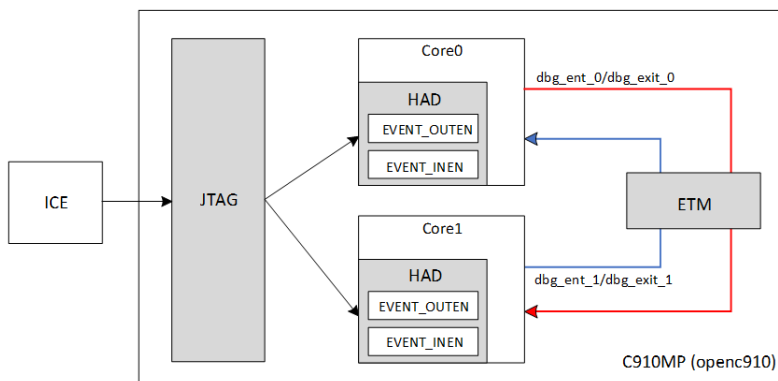


Fig. 12.2: Multi-core debug framework

through the ETM. Then the EVENT\_INEN registers of other cores determine whether the cores can enter the debug mode.

- **Multiple cores exit the debug mode**

When a core exits the debug mode, it generates a DBG\_EXIT signal. The EVENT\_OUTEN register of the core determines whether the DBG\_EXIT signal can be sent to the ETM. When the EVENT\_OUTEN register is enabled, the DBG\_EXIT signal is transmitted to other cores through the ETM. Then the EVENT\_INEN registers of the other cores determine whether the cores can exit the debug mode.

## 12.3 Debug interface signals

External interface signals of the debug unit include JTAG interface signals and debug interface signals. [Table 12.1](#) describes the debug interface signals.

Table 12.1: External interface signals of the debug unit

Signal name	Direction
corex_pad_halted	Output
pad_corex_dbgirq_b	Input
corex_pad_jdb_pm[1:0]	Output
pad_corex_dbg_mask	Input
pad_had_jtg_tclk	Input
pad_had_jtg_trst_b	Input
pad_had_jtg_tdi	Input
had_pad_jtg_tdo	Output
had_pad_jtg_tdo_en	Output
pad_had_jtg_tms	Input

**corex\_pad\_halted**

A high level indicates that the corresponding core is in debug mode.

**pad\_corex\_dbgrq\_b**

This signal is used to request multiple cores to synchronously enter the debug mode. It is valid at a low level. This signal is an external input signal of the core. It is synchronized by the system clock in the core and then transmitted to the HAD unit. The HAD unit uses this signal to request multiple cores to synchronously enter the debug mode. Setting the signal to 0 is equivalent to setting the DR bit in the hcr register of the HAD unit.

**corex\_pad\_jdb\_pm[1:0]**

This signal indicates the current operating mode of the corresponding core. You can determine whether the CPU has entered the debug mode based on this signal. For more information, see `Current_CPU_status_PM`.

Table 12.2: Current CPU status indicated by PM :name: `Current_CPU_status_PM`

<b>had_pad_jdb_pm[1:0]</b>	<b>Description</b>
00	Common mode
01	Low power mode
10	Debug mode
11	Reserved

**pad\_corex\_dbg\_mask**

A debug request masking signal driven by the SoC. It is used to mask debug requests for core x. This signal must be set to a high level in a power-off process.

**pad\_had\_jtg\_tclk**

A JTAG clock signal. This signal is an external input signal and is usually generated by the debugger. To ensure proper functioning between the debug unit and cores, the frequency of this signal must be lower than 1/2 that of the CPU clock signal.

**pad\_had\_jtg\_trst\_b**

A JTAG reset signal. It is used to reset the TAP state machine and other related control signals.

**JTAG5 signals**

- `pad_had_jtg_tdi`: A JTAG serial input signal of the HAD unit. The HAD unit performs sampling on this signal at the rising edge of JTAG TCLK, and the debugger sets this signal at the falling edge of JTAG TCLK.
- `pad_had_jtg_tms`: A JTAG mode selection signal. This signal is initiated by the debugger to control operation of the TAP state machine in the HAD unit.

- `had_pad_jtg_tdo`: A JTAG serial output signal of the HAD unit. The HAD unit sets this signal at the falling edge of JTAG TCLK, and the debugger performs sampling on this signal at the rising edge of JTAG TCLK.
- `had_pad_jtg_tdo_en`: A signal indicating that the `had_pad_jtg_tdo` signal is valid. The debugger usually monitors this signal to determine whether the `had_pad_jtg_tdo` signal is valid. Alternatively, the debugger can read the state of the TAP state machine in the debugger to determine whether to perform sampling on the `had_pad_jtg_tdo` signal. This signal is mainly used to determine the JTAG pin that generates output when multiple TAP state machines are implemented.

C910/C920 supports various power management features, including the support for multiple power domains, power-off of a single core, power-off of the cluster, and clearing of the L2 cache from external hardware interfaces. This chapter describes the power management features of C910/C920 in detail.

### 13.1 Power domain

C910/C920 can be divided into a maximum of five power domains:

- Each core is a power domain, including the computing unit, control logic, and cache RAM of the core.
- The L2 subsystem (also called top level) is a power domain, including the CIU, L2 Cache, Debug, PLIC, CLINT, and SYSIO submodules.

### 13.2 Overview of low-power modes

C910/C920 supports the following low-power modes:

- Normal mode: The cores and L2 are running properly.
- Core WFI mode: Some cores are in the wait for interrupt (WFI) mode.
- Individual-core power-off: Some cores are powered off.
- Cluster power-off: The entire cluster, including the cores and L2, is powered off.

## 13.3 Core WFI process

By executing the WFI low power instruction, a core enters WFI mode and outputs signal `core(x)_pad_lpm�_b[1:0]=2'b00`, which indicates that the core has entered WFI mode. The L2 subsystem will disable the global ICG of this core inside the cluster.

The core will be woken up and exit WFI mode upon the occurrence of the following events:

- Reset
- Interrupt request: external interrupt, software interrupt, or timer interrupt requests sent by the PLIC or CLINT submodules.
- Debug request

When one of the following events occurs, the core is temporarily woken up to process the event. It reenters low power mode after the event is processed. The core does not exit WFI mode during the entire process.

- Snoop request: Snoop requests sent by other cores.

## 13.4 Individual-core power-off process

The system can shut down the power of a core to completely terminate the static power of the core.

The process for powering off a core:

- Notifies SoC that the individual-core power-off process is to be executed. The implementation of this step is subject to the SoC design.
- Masks all interrupt requests, including external interrupts, software interrupts, and timer interrupts, and then disables the interrupt enable bit (MIE/SIE) of the mstatus/status register and the interrupt enable bit of the mie/sie register. If the power-off process is executed in M-mode, the interrupt enable bits of the mstatus and mie registers are disabled. If the power-off process is executed in S-mode, the interrupt bits of the status and sie registers are disabled.
- Disables data prefetch
- Executes D-Cache INV&CLR ALL to write dirty lines back to the L2 cache.
- Disables D-Cache (no store instruction allowed between the clear cache and disable cache operations).
- Disables the SMPEN bit to mask snoop requests.
- Executes the fence iorw, iorw instruction.
- Executes the WFI instruction to enter WFI mode.

The system performs the following operations:

- Detects a valid low-power output signal `core(x)_pad_lpm�_b` sent from the core.

- Sets `pad_core(x)_dbg_mask` to 1 to mask debug requests bound for the core to be powered off.
- Activates the output signal clamp bit of the core to be powered off.
- Sets the reset signal `pad_core(x)_rst_b` to 0 for the core to be powered off.
- Shuts down the power to the core.

A powered-off core can restart only by reset. The process of powering on a core again:

- The system detects a specific event and determines to wake up the core.
- The system sets the reset address of the core.
- The reset signal of the core is set to 0.
- The power is turned on and the reset signal remains unreleased.
- The output signal clamp bit of the core is released.
- The reset signal of the core is released.
- The core executes the initialization program, enables the SMPEN bit, or performs initialization operations, such as enabling the MMU or D-Cache.

## 13.5 Cluster power-off process (hardware clearing of the L2 cache)

Ensure that the power is shut down for all cores except the main core in the cluster. The main core is the last core to be powered off. It can be either of the two cores.

The main core performs the following operations:

- Notifies SoC that the cluster power-off process is to be executed. The implementation is subject to the SoC design.
- Masks all interrupt requests, including external interrupts, software interrupts, and timer interrupts, and then disables the interrupt enable bit (MIE/SIE) of the `mstatus/status` register and the interrupt enable bit of the `mie/sie` register.
- Disables data prefetch
- Executes the D-Cache INV&CLR ALL operation.
- Disables D-Cache (no store instruction allowed between the clear cache and disable cache operations).
- Disables the SMPEN bit.
- Executes the fence `iorw`, `iorw` instruction.
- Executes the WFI instruction to enter WFI mode.

The system performs the following operations:

- Detects a valid low-power output signal `core(x)_pad_lpm�_b` sent from the main core.

- Sets `pad_core(x)_dbg_mask` to 1 to mask debug requests bound for the main core.
- Activates the output signal clamp bit of the main core.
- Sets the reset signal `pad_core(x)_rst_b` to 0 for the main core.
- Shuts down the power of the main core.
- Sets `pad_cpu_l2cache_flush_req` to 1 to start clearing the L2 cache.
- Waits for C910/C920 to return `cpu_pad_l2cache_flush_done=1`.
- Sets `pad_cpu_l2cache_flush_req` to 0. (Then C910/C920 will set `cpu_pad_l2cache_flush_done` to 0.)
- Waits for C910/C920 to return `cpu_pad_no_op=1`.
- Activates the output signal clamp bit of the L2 subsystem.
- Sets the reset signal `pad_cpu_rst_b` of the L2 cache to 0.
- Shuts down the power of the L2 subsystem.

The cluster is powered on again by reset. The process of powering on the cluster again:

- The reset signal is set to 0 for all cores and the L2 subsystem in the cluster.
- The power is turned on, the reset signal remains unreleased, and the PLL is stable.
- The output signal clamp bits of the cores and the L2 subsystem are released.
- The reset signals of the cores and the L2 subsystem are released.
- The reset exception service program is executed to recover the CPU.

## 13.6 Cluster power-off process (software clearing of the L2 cache)

Ensure that the power is shut down for all cores except the main core in the cluster. In this scenario, it is recommended that SoC distinguishes the main core and secondary core. Core 0 can function as the main core.

The main core performs the following operations:

- Notifies SoC that the cluster power-off process is to be executed. The implementation is subject to the SoC design.
- Masks all interrupt requests, including external interrupts, software interrupts, and timer interrupts, and then disables the interrupt enable bit (MIE/SIE) of the `mstatus/status` register and the interrupt enable bit of the `mie/sie` register.
- Disables data prefetch
- Executes the D-Cache INV&CLR ALL operation.



- Disables D-Cache (no store instruction allowed between the clear cache and disable cache operations).
- Disables the SMPEN bit.
- Executes the fence iorw, iorw instruction.
- Executes the CLR & INV L2 Cache operation.
- Executes the fence iorw, iorw instruction.
- Executes the WFI instruction to enter WFI mode.

The system performs the following operations:

- Detects a valid low-power output signal `core(x)_pad_lpm�_b` sent from the main core.
- Sets `pad_core(x)_dbg_mask` to 1 to mask debug requests bound for the main core.
- Activates the output signal clamp bit of the main core.
- Sets the reset signal `pad_core(x)_rst_b` to 0 for the main core.
- Shuts down the power of the main core.
- Waits for C910/C920 to return `cpu_pad_no_op=1`.
- Activates the output signal clamp bit of the L2 subsystem.
- Sets the reset signal `pad_cpu_rst_b` of the L2 cache to 0.
- Shuts down the power of the L2 subsystem.

## 13.7 Simplified scenario: overall cluster power-off process (hardware clearing of the L2 cache)

In some systems, SoC designers may take a simple way to divide power domains. That is, take the entire C910/C920 cluster (two cores and one L2 subsystem) as a power domain and power off the cluster as a whole, instead of powering off each core separately. The cluster can be powered off (hardware clearing of the L2 cache) through the following steps:

The system performs the following operations:

- Notifies SoC that the overall cluster power-off process is to be executed. The implementation is subject to the SoC design.

The core (no need to distinguish the main core and secondary core, as the process is the same for them) performs the following operations:

- Masks all interrupt requests including external interrupts, software interrupts, and timer interrupts, and disables the interrupt enable bit (MIE/SIE) of the `mstatus/sstatus` register, as well as the interrupt enable bit of the `mie/sie` register.
- Disables data prefetch

- Executes INV&CLR D-Cache ALL to write dirty lines back to the L2 cache.
- Disables D-Cache (no store instruction allowed between the clear cache and disable cache operations).
- Disables the SMPEN bit to mask snoop requests.
- Executes the fence iorw, iorw instruction.
- Executes the WFI instruction.

The system performs the following operations:

- Waits for `core(x)_pad_lpmdb[1:0]==2'b00`, which means all CPU cores enter the low power state.
- Sets `pad_core(x)_dbg_mask` to 1 for all cores to mask debug requests.
- Sets `pad_cpu_l2cache_flush_req` to 1 to start hardware clearing for the L2 cache.
- Waits for C910/C920 to return `cpu_pad_l2cache_flush_done=1`, which means the L2 cache is cleared.
- Sets `pad_cpu_l2cache_flush_req` to 0. (Then, C910/C920 will set `cpu_pad_l2cache_flush_done` to 0.)
- Waits for `cpu_pad_no_op==1'b1`, which means the L2 cache enters the idle state. (All CPU cores are still in the low power state.)
- Activates the output signal clamp of the cluster.
- Asserts all reset signals.
- Powers off the entire cluster.

## 13.8 Simplified scenario: overall cluster power-off process (software clearing of the L2 cache)

The simple power domain division also applies. That is, the entire C910/C920 cluster (two cores and the L2 subsystem) is taken as a power domain and is powered off as a whole. The cluster can be powered off (software clearing of the L2 cache) through the following steps:

The system performs the following operations:

- Notifies SoC that the overall cluster power-off process is to be executed. The implementation is subject to the SoC design.

The secondary core (for example, CPU1) performs the following operations:

- Masks all interrupt requests including external interrupts, software interrupts, and timer interrupts, and disables the interrupt enable bit (MIE/SIE) of the `mstatus/sstatus` register, as well as the interrupt enable bit of the `mie/sie` register.
- Disables data prefetch
- Executes INV&CLR D-Cache ALL to write dirty lines back to the L2 cache.

- Disables D-Cache (no store instruction allowed between the clear cache and disable cache operations).
- Disables the SMPEN bit to mask snoop requests.
- Executes the fence iorw, iorw instruction.
- Executes the WFI instruction.

The main core (for example, CPU0) performs the following operations:

- Masks all interrupt requests including external interrupts, software interrupts, and timer interrupts, and disables the interrupt enable bit (MIE/SIE) of the mstatus/sstatus register, as well as the interrupt enable bit of the mie/sie register.
- Disables data prefetch
- Executes INV&CLR D-Cache ALL to write dirty lines back to the L2 cache.
- Disables D-Cache (no store instruction allowed between the clear cache and disable cache operations).
- Disables the SMPEN bit to mask snoop requests.
- Executes the fence iorw, iorw instruction.
- Waits for all secondary cores to enter WFI mode. (The implementation is subject to the SoC design.)
- Executes INV&CLR L2 Cache ALL to clear the L2 cache.
- Executes the fence iorw, iorw instruction.
- Executes the WFI instruction.

The system performs the following operations:

- Waits for `core(x)_pad_lpmdb[1:0]==2'b00` and `cpu_pad_no_op==1'b1`, which means all CPU cores enter the low power state and the L2 cache enters the idle state.
- Sets `pad_core(x)_dbg_mask` to 1 for all cores to mask debug requests.
- Activates the output signal clamp of the cluster.
- Asserts all reset signals.
- Powers off the entire cluster.

## 13.9 Low power consumption related programming models and interface signals

### 13.9.1 Programming models

#### M-mode snoop enable register (MSMPR)

This register is 64 bits wide. Only bit [0] has a definition (=SMPEN) and its default value is 0. This register controls whether cores can accept snoop requests.

- When MSMPR.SMPEN is 0, the cores cannot process snoop requests, and the L2 subsystem masks snoop requests bound for the cores.
- When MSMPR.SMPEN is 1, the cores can process snoop requests, and the L2 subsystem sends snoop requests to the cores.

The corresponding SMPEN bit must be set to 0 before a core is powered off. After the core is powered on, SMPEN must be set to 1 before the software enables the D-Cache and MMU. When a core is in normal running mode, its SMPEN bit must be set to 1.

### M-mode reset vector base address register (mrivr)

Each core has an mrivr register for determining the restart address of the core. The access permission for mrivr registers is MRO. The initial value of the mrivr register of a core is determined by the hardware signal pad\_core(x)\_rvba[39:1]. Note: pad\_core(x)\_rvba[0] is always 0.

## 13.9.2 Interface signals

### Reset signal control

C910/C920 top level has three reset signals: pad\_core0\_rst\_b, pad\_core1\_rst\_b, and pad\_cpu\_rst\_b. SoC can use the preceding signals to control the reset of Core 0, Core 1, and L2.

C910/C920 communicates with the power management unit of SoC by using the following signals:

- core(x)\_pad\_lpmd\_b: Indicates whether a core is in WFI mode. 2' b11 indicates normal mode, and 2' b00 indicates WFI mode.
- cpu\_pad\_no\_op: Indicates whether the L2 cache is idle. This signal is valid (a high level) when all cores enter low power mode and the L2 cache finishes all transmissions.
- pad\_cpu\_l2cache\_flush\_req and cpu\_pad\_l2cache\_flush\_done: Clear the L2 cache under the control of SoC. These signals are used in the cluster power-off process. The req signal is driven by SoC, and the done signal is driven by C910/C920. First, SoC sets the req signal to 1 to start the L2 cache clearing process > C910/C920 finishes clearing the L2 cache and returns done=1 > SoC sets the req signal to 0 > C910/C920 sets the done signal to 0.

### 14.1 PMU overview

The performance monitoring unit (PMU) of C910/C920 complies with the RISC-V standard and collects software and hardware information during a program operation for software developers to optimize their programs.

The software and hardware information collected by the PMU includes the following:

- Number of running clocks and the time
- Instruction statistics
- Statistics of key components of the CPU

### 14.2 PMU programming model

#### 14.2.1 PMU functions

Basic functions of the PMU are:

- Prohibits the counting of all events by using the `mcountinhibit` register.
- Resets the PMU counters, including `mcycle`, `minstret`, and `mhpmmcounter3` to `mhpmmcounter31`.
- Configures the corresponding events for each PMU counter. In C910/C920, the mappings between events and counters are fixed. Therefore, events must be configured for the PMU counters based on

a fixed pattern. For example, 0x1 must be written to mhpmevent3, which means that mhpcounter3 counts the number of 0x1 events (L1 ICache access count), and 0x2 must be written to mhpmevent4, which means that mhpmevent4 counts the number of 0x2 events (L1 ICache miss count), and so forth.

- Grants access permissions. The mcounteren register determines whether PMU counters can be accessed in S-mode, and scounteren determines whether PMU counters can be accessed in U-mode.
- Discharges the prohibition by using the mcountinhibit register and starts counting.

For more information, see *PMU setting example*.

14.2.2 PMU event overflow interrupt


C910/C920 implements the M-mode event overflow mark register (mcounteren) and M-mode event interrupt enable register (mcounteren). For more information about register functions and read/write permissions, see appendix C-1 M-mode control register. In the mcounteren register, the bits and event counters are in one-to-one correspondence, indicating whether the event counters overflow. In the mcounteren register, the bits and event counters are in one-to-one correspondence, indicating whether to initiate an interrupt request when an event counter overflows.

The unified interrupt vector number of overflow interrupts initiated by the PMU is 17. The interrupt enabling and processing process is the same as that of common interrupts. For more information, see *Exceptions and Interrupts*.

14.3 PMU related control registers

14.3.1 M-mode counter access enable register (mcounteren)

The mcounteren register determines whether U-mode counters can be accessed in S-mode.




M-mode counter access enable register (mcounteren)

Table 14.1: Description of the M-mode counter access enable register

Bit	Read/Write	Name	Description
31:3	Read/Write	HPM $n$	The access bit of the hpmcountern register in S-mode. 0: An illegal instruction exception will occur for accesses to the hpmcountern register in S-mode. 1: The hpmcountern register can be normally accessed in S-mode.
2	Read/Write	IR	The access bit of the minstret register in S-mode. 0: An instruction exception will occur for accesses to the minstret register in S-mode. 1: The minstret register can be normally accessed in S-mode.
1	Read/Write	TM	The access bit of the time register in S-mode. 0: An illegal instruction exception will occur for accesses to the time register in S-mode. 1: When the corresponding bit of the mcounteren register is 1, the time register can be normally accessed in S-mode. Otherwise, an illegal instruction exception will occur.
0	Read/Write	CY	The ccess bit of the mcycle register in S-mode. 0: An illegal instruction exception will occur for accesses to the cycle register in S-mode. 1: The cycle register can be normally accessed in S-mode.

### 14.3.2 S-mode counter access enable register (scounteren)

The scounteren register determines whether U-mode counters can be accessed in U-mode.




S-mode counter access enable register (scounteren)

Table 14.2: Description of the scounteren register

Bit	Read/Write	Name	Description
31:3	Read/Write	HPM $n$	The access bit of the hpmcounter $n$ register in U-mode. 0: An illegal instruction exception will occur for accesses to the hpmcounter $n$ register in U-mode. 1: When the corresponding bit of the mcounteren register is 1, the hpmcounter register can be normally accessed in U-mode. Otherwise, an illegal instruction exception will occur.
2	Read/Write	IR	The access bit of the instret register in U-mode. 0: An illegal instruction exception will occur for accesses to the instret register in U-mode. 1: When the corresponding bit of the mcounteren register is 1, the instret register can be normally accessed in U-mode. Otherwise, an illegal instruction exception will occur.
1	Read/Write	TM	The access bit of the time register in U-mode. 0: An illegal instruction exception will occur for accesses to the time register in U-mode. 1: When the corresponding bit of the mcounteren register is 1, the time register can be normally accessed in U-mode. Otherwise, an illegal instruction exception will occur.
0	Read/Write	CY	The access bit of the cycle register in U-mode. 0: An illegal instruction exception will occur for accesses to the cycle register in U-mode. 1: When the corresponding bit of the mcounteren register is 1, the cycle register can be normally accessed in U-mode. Otherwise, an illegal instruction exception will occur.

### 14.3.3 M-mode count inhibit register (mcountinhibit)

The mcountinhibit register inhibits counting of M-mode counters. When performance analysis is not required, counters can be disabled to reduce the power consumption of the CPU.



M-mode count inhibit register (mcountinhibit)



Table 14.3: Description of the M-mode count inhibit register


Bit	Read/Write	Name	Description
31:3	Read/Write	MHPM $n$	$n$ Count inhibit bit of the mhpcounter register 0: normal counting 1: counting inhibited
2	Read/Write	MIR	Count inhibit bit of the minstret register 0: normal counting 1: counting inhibited
1	-	-	-
0	Read/Write	MCY	Count inhibit bit of the mcycle register 0: normal counting 1: counting inhibited

#### 14.3.4 S-mode write enable register (mcounterwen)

The mcounterwen register determines whether S-mode event counters can be written in S-mode. This register is an M-mode extension register. For the register description, see [Appendix C-1 M-mode control registers](#).

#### 14.3.5 Performance monitoring event select register (mhpmevent3-31)

The mhpmevent3-31 register selects the counting event corresponding to a counter. In C910/C920, a counter corresponds to an event, which cannot be modified. Therefore, only the corresponding event ID can be written to each event selector. An event counter performs counting normally only after the index value of the corresponding event is written to the event selector, and the event counter is initialized by using the csrw instruction.



M-mode performance monitoring event select register (mhpmevent)

Table 14.4 describes the M-mode performance monitoring event select register.

Table 14.4: Description of the M-mode performance monitoring event select register

Bit	Read/Write	Name	Description
63:0	Read/Write	Event index	Performance monitoring event index 0: no event 0x1 to 0x1A: performance monitoring events implemented by hardware. For more information, see: numref:Counter_event_list. >0x1A: performance monitoring events that are not defined by hardware. These events are customized for use by the software.

Counter\_event\_list describes the correspondence between event selectors, events, and counters.

Table 14.5: List of correspondence between counters and events

:name: Counter\_event\_list :widths: 25 18 40 25

Index	Event	Counter
0x1	L1 ICache Access Counter	mhpmcounter3
0x2	L1 ICache Miss Counter	mhpmcounter4
0x3	I-UTLB Miss Counter	mhpmcounter5
0x4	D-UTLB Miss Counter	mhpmcounter6
0x5	JTLB Miss Counter	mhpmcounter7
0x6	Conditional Branch Mispredict Counter	mhpmcounter8
0x7	Reserved	mhpmcounter9
0x8	Indirect Branch Mispredict Counter	mhpmcounter10
0x9	Indirect Branch Instruction Counter	mhpmcounter11
0xA	LSU Spec Fail Counter	mhpmcounter12
0xB	Store Instruction Counter	mhpmcounter13
0xC	L1 DCache Read Access Counter	mhpmcounter14
0xD	L1 DCache Read Miss Counter	mhpmcounter15
0xE	L1 DCache Write Access Counter	mhpmcounter16
0xF	L1 DCache Write Miss Counter	mhpmcounter17
0x10	L2 Cache Read Access Counter	mhpmcounter18
0x11	L2 Cache Read Miss Counter	mhpmcounter19
0x12	L2 Cache Write Access Counter	mhpmcounter20
0x13	L2 Cache Write Miss Counter	mhpmcounter21
0x14	RF Launch Fail Counter	mhpmcounter22
0x15	RF Reg Launch Fail Counter	mhpmcounter23
0x16	RF Instruction Counter	mhpmcounter24
0x17	LSU Cross 4K Stall Counter	mhpmcounter25
0x18	LSU Other Stall Counter	mhpmcounter26
0x19	LSU SQ Discard Counter	mhpmcounter27
0x1A	LSU SQ Data Discard Counter	mhpmcounter28
>= 0x1B	Not Defined	mhpmcounter29~31

### 14.3.6 Event counters

Event counters are divided into three groups: M-mode event counters, U-mode event counters, and S-mode event counters (extended in C910/C920). For more information, see: [numref:mac\\_counter\\_list](#).

Table 14.6: M-mode event counter list

Name	Index	Read/Write	Initial value	Description
MCYCLE	0xB00	MRW	0x0	The cycle counter.
MINSTRET	0xB02	MRW	0x0	The instructions-retired counter.
MHPMCOUNTER3	0xB03	MRW	0x0	A performance-monitoring counter.
MHPMCOUNTER4	0xB04	MRW	0x0	A performance-monitoring counter.
...	...	...	...	...
MHPMCOUNTER31	0xB1F	MRW	0x0	A performance-monitoring counter.

Table 14.7 lists the U-mode event counters.

Table 14.7: U-mode event counter list

Name	Index	Read/Write	Initial value	Description
CYCLE	0xC00	URO	0x0	The cycle counter.
TIME	0xC01	URO	0x0	The timer.
INSTRET	0xC02	URO	0x0	The instructions-retired counter.
HPMCOUNTER3	0xC03	URO	0x0	A performance-monitoring counter.
HPMCOUNTER4	0xC04	URO	0x0	A performance-monitoring counter.
...	...	...	...	...
HPMCOUNTER31	0xC1F	URO	0x0	A performance-monitoring counter.

Table 14.8: S-mode event counter list

Name	Index	Read/Write	Initial value	Description
SCYCLE	0x5E0	SRO	0x0	The cycle counter.
SINSTRET	0x5E2	SRO	0x0	The instructions-retired counter.
SHPMCOUNTER3	0x5E3	SRO	0x0	A performance-monitoring counter.
SHPMCOUNTER4	0x5E4	SRO	0x0	A performance-monitoring counter.
...	...	...	...	...
SHPMCOUNTER31	0x5FF	SRO	0x0	A performance-monitoring counter.

The U-mode CYCLE, INSTRET, and HPMCOUNTER<sub>n</sub> counters are read-only mappings of the corresponding M-mode event counters. The timer is the read-only mapping of the MTIME register.

The S-mode SCYCLE, SINSTRET, and SHPMCOUNTER<sub>n</sub> counters are mappings of corresponding M-mode event counters.

This chapter describes various program examples, including the MMU setting example, PMP setting example, cache setting example, multi-core startup example, synchronization primitive example, PLIC setting example, and PMU setting example.

### 15.1 Optimal CPU performance configuration

The optimal performance of C910/C920 can be achieved by using the following configurations:

- MHCR = 0x11ff
- MHINT = 0x6e30c
- MCCR2 = 0xe0000009 (Note: MCCR2 contains RAM delay settings. In this example, all delays are 0. Customers need to set a proper RAM delay based on the actual situation.)
- MXSTATUS = 0x638000
- MSMPR = 0x1

```
# mhcr
li x3, 0x11ff
csrs mhcr,x3

#mhint
li x3, 0x6e30c
```

(continues on next page)

(continued from previous page)

```

csrs mhint,x3

# mxstatus
li x3, 0x638000
csrs mxstatus,x3

# msmpr
csrsi msmpr,0x1

# mCCR2
li x3, 0xe0000009
csrs mCCR2,x3

```

## 15.2 MMU setting example

```

/*****

* Function: An example of setting C910/C920MP MMU.
* Memory space: Virtual address <-> physical address.
*
* Pagesize 4K: vpn: {vpn2,vpn1,vpn0} <-> ppn: {ppn2,ppn1,ppn0}
* Pagesize 2M: vpn: {vpn2,vpn1} <-> ppn:{ppn2,ppn1}
* Pagesize 1G: vpn: {vnp2} <-> ppn: {ppn2}
*
*****/

/*C910/C920 will invalidate all MMU TLB entries automatically when reset*/
/*You can use sfence.vma to invalid all MMU TLB entries if necessary*/
sfence.vma x0, x0

/* Pagesize 4K: vpn: {vpn2, vpn1, vpn0} <-> ppn: {ppn2, ppn1, ppn0}*/
/* First-level page addr base: PPN (defined in satp)*/
/* Second-level page addr base: BASE2 (self define)*/
/* Third-level page addr base: BASE3 (self define)*/
/* 1. Get first-level page addr base: PPN and vpn*/
/* Get PPN*/
csrr x3, satp
li x4, 0xffffffff

```

(continues on next page)

(continued from previous page)

```

and x3, x3, x4

/*2. Config first-level page*/
/*First-level page addr: {PPN, vpn2, 3' b0}, first-level page pte:{ 44' b BASE2, 10' b1}
→ */
/*Get first-level page addr*/
slli x3, x3, 12
/*Get vpn2*/
li x4, VPN
li x5, 0x7fc0000
and x4, x4, x5
srli x4, x4, 15
and x5, x3, x4
/*Store pte at first-level page addr*/
li x6, {44' b BASE2, 10' b1}
sd x6, 0(x5)

/*3. Config second-level page*/
/*Second-level page addr: {BASE2, vpn1, 3' b0}, second-level page pte:{ 44' b BASE3, 10'
b1} */
/*Get second-level page addr*/
/* VPN1*/
li x4, VPN
li x5, 0x3fe00
and x4, x4, x5
srli x4, x4, 9
/*BASE2*/
li x5, BASE2
srli x5, x5, 12
and x5, x5, x4
/*Store pte at second-level page addr*/
li x6, {44' b BASE3, 10' b1}
sd x6, 0(x5)

/*4. Config third-level page*/
/*Third-level page addr: {BASE3, vpn0, 3' b0}, third-level page pte:{
theadflag, ppn2, ppn1, ppn0, 9' b flags,1' b1} */
/*Get second-level page addr*/
/* VPN0*/
li x4, VPN
li x5, 0x1ff

```

(continues on next page)

(continued from previous page)

```

and x4, x4, x5
srli x4, x4, 3
/*BASE3*/
li x5, BASE3
srli x5, x5, 12
and x5, x5, x4
/*Store pte at second-level page addr*/
li x6, { theadflag, ppn2, ppn1, ppn0, 9' b flags, 1' b1}
sd x6, 0(x5)

/* Pagesize 2M: vpn: {vpn2, vpn1} <-> ppn: {ppn2, ppn1}*/
/*First-level page addr base: PPN (defined in satp)*/
/*Second-level page addr base: BASE2 (self define)*/

/*1. Get first-level page addr base: PPN and vpn*/
/* Get PPN*/
csrr x3, satp
li x4, 0xffffffff
and x3, x3, x4

/*2. Config first-level page*/
/*First-level page addr: {PPN, vpn2, 3' b0}, first-level page pte:{ 44' b
BASE2, 10' b1}*/
/*Get first-level page addr*/
slli x3, x3, 12
/*Get vpn2*/
li x4, VPN
li x5, 0x7fc0000
and x4, x4, x5
srli x4, x4, 15
and x5, x3, x4
/*Store pte at first-level page addr*/
li x6, {44' b BASE2, 10' b1}
sd x6, 0(x5)

/*3. Config second-level page*/
/*Second-level page addr: {BASE2, vpn1, 3' b0}, second-level page pte:{
theadflag, ppn2, ppn1, 9' b0, 9' b flags, 1' b1} */
/*Get second-level page addr*/

```

(continues on next page)



(continued from previous page)

```

/*VPN1*/
li x4, VPN
li x5, 0x3fe00
and x4, x4, x5
srli x4, x4, 9
/*BASE2*/
li x5, BASE2
srli x5, x5, 12
and x5, x5, x4
/*Store pte at second-level page addr*/
li x6, { theadflag, ppn2, ppn1, 9' b0, 9' b flags, 1' b1}
sd x6, 0(x5)

/* Pagesize 1G: vpn: {vpn2} <-> ppn: {ppn2}*/
/*First-level page addr base: PPN (defined in satp)*/
/*1. Get first-level page addr base: PPN and vpn*/
/* Get PPN*/
csrr x3, satp
li x4, 0xfffffffffff
and x3, x3, x4

/*2. Config first-level page*/
/*First-level page addr: {PPN, vpn2, 3' b0}, first-level page pte:{
theadflag, ppn2, 9' b0, 9' b0, 9' b flags, 1' b1}*/
/*Get first-level page addr*/
slli x3, x3, 12
/*Get vpn2*/
li x4, VPN
li x5, 0x7fc0000
and x4, x4, x5
srli x4, x4, 15
and x5, x3, x4
/*Store pte at first-level page addr*/
li x6, { theadflag, ppn2, 9' b0, 9' b0, 9' b flags, 1' b1}
sd x6, 0(x5)

```

## 15.3 PMP setting example

```

/*****
* Function: An example of setting C910/C920MP PMP.
* 0x0 ~ 0xf0000000, TOR mode, RWX
* 0xf0000000 ~ 0xf8000000, NAPOT mode, RW
* 0xffff73000 ~ 0xffff74000, NAPOT mode, RW
* 0xfffc0000 ~ 0xfffc2000, NAPOT mode, RW
* Different execution permissions are configured for the preceding four regions. PMP must
  ↳ be configured to prevent the CPU from executing instructions to unsupported address
  ↳ regions in different modes, especially in M-mode where the CPU has full execution
  ↳ permissions by default. Specifically, after you configure address regions that require
  ↳ execution permissions, no permission should be configured for the rest address regions.
  ↳ For more information, see the following example.
*****/

# pmpaddr0, 0x0 0xf0000000, TOR mode, read and write permissions
li x3, (0xf0000000 >> 2)
csrw pmpaddr0, x3
# pmpaddr1, 0xf0000000 0xf8000000, NAPOT mode, read and write permissions
li x3, ( 0xf0000000 >> 2 | (0x8000000-1) >> 3))
csrw pmpaddr1, x3
# pmpaddr2, 0xffff73000 0xffff74000, NAPOT mode, read and write permissions
li x3, ( 0xffff73000 >> 2 | (0x1000-1) >> 3))
csrw pmpaddr2, x3
# pmpaddr3, 0xfffc0000 0xfffc2000, NAPOT mode, read and write permissions
li x3, ( 0xfffc0000 >> 2 | (0x2000-1) >> 3))
csrw pmpaddr3, x3
# pmpaddr4, 0xf0000000 0x100000000, NAPOT mode, no permissions
li x3, ( 0xf0000000 >> 2 | (0x100000000-1) >> 3))
csrw pmpaddr4, x3
# pmpaddr5, 0x100000000 0xffffffff, TOR mode, no permissions
li x3, (0xffffffff >> 2)
csrw pmpaddr5, x3
# PMPCFG0 configures the execution permission, mode, and lock bit of entries.
When lock is 1, it is valid only in M-mode.
li x3, 0x88989b9b9b8f
csrw pmpcfg0, x3
# pmpaddr5, 0x100000000 0xffffffff: In TOR mode, when 0x100000000 <= addr <
0xffffffff, pmpaddr5 will be hit. However, pmpaddr5 cannot be hit in the address
↳ range 0xffffffff000 ~

```

(continues on next page)

(continued from previous page)

0xffffffff (the minimum PMP granularity is 4 KB in C910/C920). An NAPOT entry must be configured to mask the last 4 KB space of a 1 TB space.

## 15.4 Cache examples

### 15.4.1 Cache enabling example

```
/*C910/C920 will invalidate all I-cache automatically when reset*/
/*You can invalidate I-cache by yourself if necessary*/
/*Invalidate I-cache*/
li x3, 0x33
csrc mcor, x3
li x3, 0x11
csrs mcor, x3
// You can also use icache instructions to replace the invalidate sequence
// if theadisaee is enabled.
//icache.iall
//sync.is

/*Enable I-cache*/
li x3, 0x1
csrs mhcr, x3

/*C910/C920 will invalidate all D-cache automatically when reset*/
/*You can invalidate D-cache by yourself if necessary*/
/*Invalidate D-cache*/
li x3, 0x33
csrc mcor, x3
li x3, 0x12
csrs mcor, x3

// You can also use dcache instructions to replace the invalidate sequence
// if theadisaee is enabled.
// dcache.iall
// sync.is

/*Enable D-cache*/
li x3, 0x2
```

(continues on next page)

(continued from previous page)

```

csrs mhcr, x3

/*C910/C920 will invalidate all L2 cache automatically when reset*/
/*You can invalidate L2 by yourself if necessary*/
/*Invalidate L2-cache if theadisaee is enabled*/
l2cache.iall
sync.is

/*Enable L2-cache*/
li x3, 8
csrs mCCR2, x3

```

### 15.4.2 Example of synchronization between the instruction and data caches

#### CPU0

```

sd x3,0(x4) // a new instruction defined in x3
              // is stored to program memory address defined in x4.
dcache.cval1 r0 // clean the new instruction to the shared L2 cache.
sync.s          // ensure completion of clean operation.
              // the dcache clean is not necessarily if INSDC is not enabled.
icache.iva r0 // invalid icache according to shareable configuration.
sync.s/fence.i // ensure completion in all CPUs.
sd x5,0(x6)    // set flag to signal operation completion.
sync.is
jr x4 // jmp to new code

```

#### CPU1 CPU3

```

WAIT_FINISH:
ld x7,0(x6)
bne x7,x5, WAIT_FINISH // wait CPU0 modification finish.
sync.is
jr x4                  // jmp to new code

```

### 15.4.3 Example of synchronization between the TLB and the data cache

#### CPU0

```
sd x4,0(x3) // update a new translation table entry
sync.is/fence.i // ensure completion of update operation.
sfence.vma x5,x0 // invalid the TLB by va
sync.is/fence.i // ensure completion of TLB invalidation and
                // synchronises context
```

## 15.5 Synchronization primitive examples

### CPU0

```
li x1, 0x1
li x6, 0x0

ACQUIRE_LOCK:           // (x3) is the lock address. 0: Free; 1: Busy.
lr x4, 0(x3)             // Read lock
bnez x4, ACQUIRE_LOCK   // Try again if the lock is in use
sc x5, x1, 0(x3)         // Attempt to store new value
bne x6, x5, ACQUIRE_LOCK // Try again if fail
sync.s

...                      // Critical section code
```

### CPU1

```
sync.s/fence.i // Ensure all operations are observed before clearing the lock.
sd x0, 0(x3)    // Clear the lock.
```

## 15.6 PLIC setting example

```
//Init id 1 machine mode int for hart 0
/*1. Set hart threshold if needed*/
li x3, (plic_base_addr + 0x200000) // h0 mthreshold addr
li x4, 0xa //threshold value
sw x4,0x0(x3) // set hart0 threshold as 0xa

/*2. Set priority for int id 1*/
li x3, (plic_base_addr + 0x0) // int id 1 prio addr
li x4, 0x1f // prio value
```

(continues on next page)

(continued from previous page)

```

sw x4,0x4(x3) // init id1 priority as 0x1f

/*3. Enable m-mode int id1 to hart*/
li x3, (plic_base_addr + 0x2000) // h0 mie0 addr
li x4, 0x2
sw x4,0x0(x3) // enable int id1 to hart0

/*4. Set ip or wait external int*/
/*following code set ip*/
li x3, (plic_base_addr + 0x1000) // h0 mthreshold addr
li x4, 0x2 // id 1 pending
sw x4, 0x0(x3) // set int id1 pending

/*5. Core enters interrupt handler, read PLIC_CLAIM and get ID*/

/*6. Core takes interrupt*/

/*7. Core needs to clear external interrupt source if LEVEL(not PULSE)
configured, then core writes ID to PLIC_CLAIM and exits interrupt*/

```

## 15.7 PMU setting example

```

/*1. Inhibit counters counting*/
li x3, 0xffffffff
csrw mcountinhibit, x3

/*2. C910/C920 will initial all pmu counters when reset*/
/*you can initial pmu counters manually if necessarily*/
csrw mcycle, x0
csrw minstret, x0
csrw mhpmpcounter3, x0
.....
csrw mhpmpcounter31, x0

/*3. Configure mhpmevent*/
li x3, 0x1
csrw mhpmevent3, x3 // mhpmpcounter3 count event: L1 ICache Access Counter
li x3, 0x2

```

(continues on next page)

(continued from previous page)

```
csrw mhpmevent4, x3 // mhpcounter4 count event: L1 ICache Miss Counter
.....
li x3, 0x13
csrw mhpmevent21, x3 // mhpcounter21 count event: L2 Cache write miss Counter

/*4. Configure mcounteren and scounteren*/
li x3, 0xffffffff
csrw mcounteren, x3 // enable super mode to read hpmcounter
li x3, 0xffffffff
csrw scounteren, x3 // enable user mode to read hpmcounter

/*5. Enable counters to count when you want*/
csrw mcountinhibit, x0
```

---

## Appendix A Standard Instructions

---

C910MP implements the RV64IMAFDC instruction set architecture. The instructions are described in the following sections by instruction set.

### 16.1 Appendix A-1 I instructions

The following describes the RISC-V I instructions implemented by C910/C920. The instructions are sorted in alphabetic order.

The instructions are 32 bits wide by default. However, in specific cases, the system assembles some instructions into 16-bit compressed instructions. For more information about compressed instructions, see [Appendix A-6 C Instructions](#).

#### 16.1.1 ADD: a signed add instruction

**Syntax:**

add rd, rs1, rs2

**Operation:**

$rd \leftarrow rs1 + rs2$

**Permission:**

Machine mode (M-mode)/Supervisor mode (S-mode)/User mode (U-mode)



**Exception:**

None

**Instruction format:**

31	25	24	20	19	15	14	12	11	7	6	0
0000000				rs2		rs1		000	rd		0110011

**16.1.2 ADDI: a signed add immediate instruction****Syntax:**

addi rd, rs1, imm12

**Operation:** $rd \leftarrow rs1 + \text{sign\_extend}(imm12)$ **Permission:**

M mode/S mode/U mode

**Exception:**

None

**Instruction format:**

31	20	19	15	14	12	11	7	6	0
imm12[11:0]				rs1		000	rd		0010011

**16.1.3 ADDIW: a signed add immediate instruction that operates on the lower 32 bits****Syntax:**

addiw rd, rs1, imm12

**Operation:** $\text{tmp}[31:0] \leftarrow rs1[31:0] + \text{sign\_extend}(imm12)[31:0]$  $rd \leftarrow \text{sign\_extend}(\text{tmp}[31:0])$ **Permission:**

M mode/S mode/U mode

**Exception:**

None

**Instruction format:**

31	20	19	15	14	12	11	7	6	0
imm12[11:0]				rs1		000	rd		0011011

### 16.1.4 ADDW: a signed add instruction that operates on the lower 32 bits

**Syntax:**

addw rd, rs1, rs2

**Operation:**

$\text{tmp}[31:0] \leftarrow \text{rs1}[31:0] + \text{rs2}[31:0]$

$\text{rd} \leftarrow \text{sign\_extend}(\text{tmp}[31:0])$

**Permission:**

M mode/S mode/U mode

**Exception:**

None

**Instruction format:**

31	25	24	20	19	15	14	12	11	7	6	0
0000000				rs2		rs1		000	rd		0111011

### 16.1.5 AND: a bitwise AND instruction

**Syntax:**

and rd, rs1, rs2

**Operation:**

$\text{rd} \leftarrow \text{rs1} \& \text{rs2}$

**Permission:**

M mode/S mode/U mode

**Exception:**

None

**Instruction format:**

31	25	24	20	19	15	14	12	11	7	6	0
0000000				rs2		rs1		111	rd		0110011

### 16.1.6 ANDI: an immediate bitwise AND instruction

**Syntax:**

```
andi rd, rs1, imm12
```

**Operation:**

$$rd \leftarrow rs1 \& \text{sign\_extend}(imm12)$$
**Permission:**

M mode/S mode/U mode

**Exception:**

None

**Instruction format:**

31	20	19	15	14	12	11	7	6	0
<div>imm12[11:0]</div> <div>rs1</div> <div>111</div> <div>rd</div> <div>0010011</div>									

### 16.1.7 AUIPC: an instruction that adds the immediate in the upper bits to the PC

**Syntax:**

```
auipc rd, imm20
```

**Operation:**

$$rd \leftarrow \text{current pc} + \text{sign\_extend}(imm20 \ll 12)$$
**Permission:**

M mode/S mode/U mode

**Exception:**

None

**Instruction format:**

31	12	11	7	6	0
<div>imm20[19:0]</div> <div>rd</div> <div>0010111</div>					

### 16.1.8 BEQ: a branch-if-equal instruction

**Syntax:**

```
beq rs1, rs2, label
```

**Operation:**

```

if (rs1 == rs2)
    next pc = current pc + sign_extend(imm12 << 1)
else
    next pc = current pc + 4

```

**Permission:**

M mode/S mode/U mode


**Exception:**

None

**Notes:**

The compiler calculates immediate 12 based on the label.

The jump range of the instruction is  $\pm 4$  KB address space.

**Instruction format:**

### 16.1.9 BGE: a signed branch-if-greater-than-or-equal instruction

**Syntax:**

```
bge rs1, rs2, label
```

**Operation:**

```

if (rs1 >= rs2)
    next pc = current pc + sign_extend(imm12 << 1)
else
    next pc = current pc + 4

```

**Permission:**

M mode/S mode/U mode


**Exception:**

None

**Notes:**

The compiler calculates immediate 12 based on the label.

The jump range of the instruction is  $\pm 4$  KB address space.

**Instruction format:****16.1.10 BGEU: an unsigned branch-if-greater-than-or-equal instruction****Syntax:**

```
bgeu rs1, rs2, label
```

**Operation:**

```

if (rs1 >= rs2)
    next pc = current pc + sign_extend(imm12<<1)
else
    next pc = current pc + 4
  
```

**Permission:**

M mode/S mode/U mode


**Exception:**

None

**Notes:**

The compiler calculates immediate 12 based on the label.

The jump range of the instruction is  $\pm 4$  KB address space.

**Instruction format:**

### 16.1.11 BLT: a signed branch-if-less-than instruction

**Syntax:**

```
blt rs1, rs2, label
```

**Operation:**

```
if (rs1 < rs2)
```

```
    next pc = current pc + sign_extend(imm12<<1)
```

```
else
```

```
    next pc = current pc + 4
```

**Permission:**

M mode/S mode/U mode


**Exception:**

None

**Notes:**

The compiler calculates immediate 12 based on the label.

The jump range of the instruction is  $\pm 4$  KB address space.

**Instruction format:**


### 16.1.12 BLTU: an unsigned branch-if-less-than instruction

**Syntax:**

```
bltu rs1, rs2, label
```

**Operation:**

```
if (rs1 < rs2)
```

```
    next pc = current pc + sign_extend(imm12<<1)
```

```
else
```

```
    next pc = current pc + 4
```

**Permission:**

M mode/S mode/U mode


**Exception:**

None

**Notes:**

The compiler calculates immediate 12 based on the label.

The jump range of the instruction is  $\pm 4$  KB address space.

**Instruction format:****16.1.13 BNE: a branch-if-not-equal instruction****Syntax:**

```
bne rs1, rs2, label
```

**Operation:**

```
if (rs1 != rs2)
```

```
    next pc = current pc + sign_extend(imm12<<1)
```

```
else
```

```
    next pc = current pc + 4
```

**Permission:**

M mode/S mode/U mode


**Exception:**

None

**Notes:**

The compiler calculates immediate 12 based on the label.

The jump range of the instruction is  $\pm 4$  KB address space.

**Instruction format:**

### 16.1.14 CSRRC: a move instruction that clears control registers

**Syntax:**

```
csrrc rd, csr, rs1
```

**Operation:**

$$rd \leftarrow csr$$

$$csr \leftarrow csr \& (\sim rs1)$$
**Permission:**

M mode/S mode/U mode

**Exception:**

Illegal instruction.

**Notes:**

Accessible control registers vary under different privileges. For more information, see the descriptions of control registers.

When  $rs1 = x0$ , this instruction does not initiate write operations and therefore does not cause write-related exceptions.

**Instruction format:**

31	20	19	15	14	12	11	7	6	0
csr				rs1		011	rd		1110011

### 16.1.15 CSRRCI: a move instruction that clears immediates in control registers

**Syntax:**

```
csrrci rd, csr, imm5
```

**Operation:**

$$rd \leftarrow csr$$

$$csr \leftarrow csr \& \sim \text{zero\_extend}(\text{imm5})$$
**Permission:**

M mode/S mode/U mode

**Exception:**

Illegal instruction.

**Notes:**



Accessible control registers vary under different privileges. For more information, see the descriptions of control registers.

When  $rs1 = x0$ , this instruction does not initiate write operations and therefore does not cause write-related exceptions.

**Instruction format:**

31	20	19	15	14	12	11	7	6	0	
csr			imm5		111		rd		1110011	

### 16.1.16 CSRRS: a move instruction for setting control registers

**Syntax:**

`csrrs rd, csr, rs1`

**Operation:**

$rd \leftarrow csr$

$csr \leftarrow csr \mid rs1$

**Permission:**

M mode/S mode/U mode

**Exception:**

Illegal instruction.

**Notes:**

Accessible control registers vary under different privileges. For more information, see the descriptions of control registers.

When  $rs1 = x0$ , this instruction does not initiate write operations and therefore does not cause write-related exceptions.

**Instruction format:**

31	20	19	15	14	12	11	7	6	0	
csr			rs1		010		rd		1110011	

### 16.1.17 CSRRSI: a move instruction for setting immediates in control registers

**Syntax:**

`csrrsi rd, csr, imm5`

**Operation:**

$$rd \leftarrow csr$$

$$csr \leftarrow csr \mid \text{zero\_extend}(\text{imm5})$$
**Permission:**

M mode/S mode/U mode

**Exception:**

Illegal instruction.

**Notes:**

Accessible control registers vary under different privileges. For more information, see the descriptions of control registers.

When  $rs1 = x0$ , this instruction does not initiate write operations and therefore does not cause write-related exceptions.

**Instruction format:**

31	20	19	15	14	12	11	7	6	0
csr				imm5		110	rd		1110011

**16.1.18 CSRRW: a move instruction that reads/writes control registers****Syntax:**

$$\text{csrrw } rd, csr, rs1$$
**Operation:**

$$rd \leftarrow csr$$

$$csr \leftarrow rs1$$
**Permission:**

M mode/S mode/U mode

**Exception:**

Illegal instruction.

**Notes:**

Accessible control registers vary under different privileges. For more information, see the descriptions of control registers.

When  $rs1 = x0$ , this instruction does not initiate write operations and therefore does not cause write-related exceptions.

**Instruction format:**

31	20	19	15	14	12	11	7	6	0	
csr			rs1		001		rd		1110011	

### 16.1.19 CSRRWI: a move instruction that reads/writes immediates in control registers

**Syntax:**

csrrwi rd, csr, imm5

**Operation:**

$rd \leftarrow csr$

$csr[4:0] \leftarrow imm5$

$csr[63:5] \leftarrow csr[63:5]$

**Permission:**

M mode/S mode/U mode

**Exception:**

Illegal instruction.

**Notes:**

Accessible control registers vary under different privileges. For more information, see the descriptions of control registers.

When  $rs1 = x0$ , this instruction does not initiate write operations and therefore does not cause write-related exceptions.

**Instruction format:**

31	20	19	15	14	12	11	7	6	0
csr			imm5		101	rd		1110011	

### 16.1.20 EBREAK: a breakpoint instruction

**Syntax:**

ebreak

**Operation:**

Generates breakpoint exceptions or enables the core to enter the debug mode.

**Permission:**

M mode/S mode/U mode

**Exception:**

Breakpoint exceptions

**Instruction format:**

31	20	19	15	14	12	11	7	6	0			
000000000001					00000		000		00000		1110011	

**16.1.21 ECALL: an environment call instruction****Syntax:**

ecall

**Operation:**

Generates environment call exceptions.

**Permission:**

M mode/S mode/U mode

**Exception:**

U-mode, S-mode, and M-mode environment call exceptions

**Instruction format:**

31	20	19	15	14	12	11	7	6	0	
000000000000			00000		000		00000		1110011	

**16.1.22 FENCE: a memory synchronization instruction****Syntax:**

fence iorw, iorw

**Operation:**

Ensures that all memory or device read/write instructions before this instruction are observed earlier than those after this instruction.

**Permission:**

M mode/S mode/U mode

**Exception:**

None

**Notes:**

When the PI and SO bits are both 1, the instruction syntax is fence i,o, and so on.

Instruction format:

31	28	27	26	25	24	23	22	21	20	19	15	14	12	11	7	6	0
0000	pi	po	pr	pw	si	so	sr	sw	00000		000		00000				0001111

### 16.1.23 FENCE.I: an instruction stream synchronization instruction

Syntax:

fence.i

Operation:

Clears the I-Cache to ensure that the data access results before this instruction can be accessed by fetch operations after the instruction.

Permission:

M mode/S mode/U mode

Exception:

None

Instruction format:

31	28	27	24	23	20	19	15	14	12	11	7	6	0
0000		0000		0000		00000		001		00000			0001111

### 16.1.24 JAL: an instruction for directly jumping to a subroutine

Syntax:

jal rd, label

Operation:

$\text{next pc} \leftarrow \text{current pc} + \text{sign\_extend}(\text{imm20} \ll 1)$

$\text{rd} \leftarrow \text{current pc} + 4$

Permission:

M mode/S mode/U mode

Exception:


None

Notes:

The compiler calculates immediate 20 based on the label.

The jump range of the instruction is  $\pm 1$  MB address space.

**Instruction format:**



### 16.1.25 JALR: an instruction for jumping to a subroutine by using an address in a register

**Syntax:**

jalr rd, rs1, imm12

**Operation:**

$\text{next pc} \leftarrow (\text{rs1} + \text{sign\_extend}(\text{imm12}) \& 64' \text{ hfffffffffffffe})$

$\text{rd} \leftarrow \text{current pc} + 4$

**Permission:**

M mode/S mode/U mode

**Exception:**


None

**Notes:**

When the CPU runs in M-mode or the MMU is disabled, the jump range of the instruction is the entire 1 TB address space.

When the CPU does not run in M-mode and the MMU is enabled, the jump range of the instruction is the entire 512 GB address space.

**Instruction format:**



### 16.1.26 LB: a sign-extended byte load instruction

**Syntax:**

lb rd, imm12(rs1)

**Operation:**

$\text{address} \leftarrow \text{rs1} + \text{sign\_extend}(\text{imm12})$

$\text{rd} \leftarrow \text{sign\_extend}(\text{mem}[\text{address}])$

**Permission:**

M mode/S mode/U mode

**Exception:**

Unaligned access exceptions, access error exceptions, and page error exceptions on load instructions

**Instruction format:**

31	20	19	15	14	12	11	7	6	0
imm12[11:0]				rs1		000	rd		0000011

**16.1.27 LBU: an unsign-extended byte load instruction****Syntax:**

lbu rd, imm12(rs1)

**Operation:** $\text{address} \leftarrow \text{rs1} + \text{sign\_extend}(\text{imm12})$  $\text{rd} \leftarrow \text{zero\_extend}(\text{mem}[\text{address}])$ **Permission:**

M mode/S mode/U mode

**Exception:**

Unaligned access exceptions, access error exceptions, and page error exceptions on load instructions

**Instruction format:**

31	20	19	15	14	12	11	7	6	0
imm12[11:0]				rs1		100	rd		0000011

**16.1.28 LD: a doubleword load instruction****Syntax:**

ld rd, imm12(rs1)

**Operation:** $\text{address} \leftarrow \text{rs1} + \text{sign\_extend}(\text{imm12})$  $\text{rd} \leftarrow \text{mem}[(\text{address}+7):\text{address}]$ **Permission:**

M mode/S mode/U mode

**Exception:**

Unaligned access exceptions, access error exceptions, and page error exceptions on load instructions

**Instruction format:**

31	20	19	15	14	12	11	7	6	0
imm12[11:0]				rs1	011	rd		0000011	

**16.1.29 LH: a sign-extended halfword load instruction****Syntax:**

lh rd, imm12(rs1)

**Operation:**

address  $\leftarrow$  rs1 + sign\_extend(imm12)

rd  $\leftarrow$  sign\_extend(mem[(address+1):address])

**Permission:**

M mode/S mode/U mode

**Exception:**

Unaligned access exceptions, access error exceptions, and page error exceptions on load instructions

**Instruction format:**

31	20	19	15	14	12	11	7	6	0
imm12[11:0]				rs1	001	rd		0000011	

**16.1.30 LHU: an unsign-extended halfword load instruction****Syntax:**

lhu rd, imm12(rs1)

**Operation:**

address  $\leftarrow$  rs1 + sign\_extend(imm12)

rd  $\leftarrow$  zero\_extend(mem[(address+1):address])

**Permission:**

M mode/S mode/U mode

**Exception:**

Unaligned access exceptions, access error exceptions, and page error exceptions on load instructions

**Instruction format:**



31	20	19	15	14	12	11	7	6	0
imm12[11:0]				rs1	101	rd		0000011	

### 16.1.31 LUI: an instruction for loading the immediate in the upper bits

**Syntax:**

lui rd, imm20

**Operation:**

$rd \leftarrow \text{sign\_extend}(\text{imm20} \ll 12)$

**Permission:**

M mode/S mode/U mode

**Exception:**

None

**Instruction format:**

31	12	11	7	6	0
imm20[19:0]				rd	0110111

### 16.1.32 LW: a sign-extended word load instruction

**Syntax:**

lw rd, imm12(rs1)

**Operation:**

$\text{address} \leftarrow \text{rs1} + \text{sign\_extend}(\text{imm12})$

$rd \leftarrow \text{sign\_extend}(\text{mem}[(\text{address}+3):\text{address}])$

**Permission:**

M mode/S mode/U mode

**Exception:**

Unaligned access exceptions, access error exceptions, and page error exceptions on load instructions

**Instruction format:**

31	20	19	15	14	12	11	7	6	0
imm12[11:0]				rs1	010	rd		0000011	

### 16.1.33 LWU: an unsign-extended word load instruction

**Syntax:**

```
lwu rd, imm12(rs1)
```

**Operation:**

$$\text{address} \leftarrow \text{rs1} + \text{sign\_extend}(\text{imm12})$$

$$\text{rd} \leftarrow \text{zero\_extend}(\text{mem}[(\text{address}+3):\text{address}])$$
**Permission:**

M mode/S mode/U mode

**Exception:**

Unaligned access exceptions, access error exceptions, and page error exceptions on load instructions

**Instruction format:**

31	20	19	15	14	12	11	7	6	0
imm12[11:0]									
rs1									
110									
rd									
0000011									

### 16.1.34 MRET: an instruction for returning from exceptions in M-mode

**Syntax:**

```
mret
```

**Operation:**

$$\text{next pc} \leftarrow \text{mepc}$$

$$\text{mstatus.mie} \leftarrow \text{mstatus.mpie}$$

$$\text{mstatus.mpie} \leftarrow 1$$
**Permission:**

M mode

**Exception:**

Illegal instruction.

**Instruction format:**

31	25	24	20	19	15	14	12	11	7	6	0
0011000										1110011	
00010										00000	
00000										000	
00000										00000	

### 16.1.35 OR: a bitwise OR instruction

**Syntax:**

$$\text{or rd, rs1, rs2}$$
**Operation:**

$$\text{rd} \leftarrow \text{rs1} \mid \text{rs2}$$
**Permission:**

M mode/S mode/U mode

**Exception:**

None

**Instruction format:**

31	25	24	20	19	15	14	12	11	7	6	0
0000000				rs2		rs1		110	rd		0110011

### 16.1.36 ORI: an immediate bitwise OR instruction

**Syntax:**

$$\text{ori rd, rs1, imm12}$$
**Operation:**

$$\text{rd} \leftarrow \text{rs1} \mid \text{sign\_extend}(\text{imm12})$$
**Permission:**

M mode/S mode/U mode

**Exception:**

None

**Instruction format:**

31	20	19	15	14	12	11	7	6	0
imm12[11:0]				rs1		110	rd		0010011

### 16.1.37 SB: a byte store instruction

**Syntax:**

$$\text{sb rs2, imm12(rs1)}$$
**Operation:**

$$\text{address} \leftarrow \text{rs1} + \text{sign\_extend}(\text{imm12})$$

$$\text{mem}[:, \text{address}] \leftarrow \text{rs2}[7:0]$$
**Permission:**

M mode/S mode/U mode

**Exception:**

Unaligned access exceptions, access error exceptions, and page error exceptions on store instructions

**Instruction format:**

31	25	24	20	19	15	14	12	11	7	6	0
imm12[11:5]				rs2		rs1		000		imm12[4:0]	0100011

**16.1.38 SD: a doubleword store instruction****Syntax:**

sd rs2, imm12(rs1)

**Operation:**

$$\text{address} \leftarrow \text{rs1} + \text{sign\_extend}(\text{imm12})$$

$$\text{mem}[(\text{address}+7):\text{address}] \leftarrow \text{rs2}$$
**Permission:**

M mode/S mode/U mode

**Exception:**

Unaligned access exceptions, access error exceptions, and page error exceptions on store instructions

**Instruction format:**

31	25	24	20	19	15	14	12	11	7	6	0
imm12[11:5]				rs2		rs1		011		imm12[4:0]	0100011

**16.1.39 SFENCE.VMA: a virtual memory synchronization instruction****Syntax:**

sfence.vma rs1,rs2

**Operation:**

Invalidates and synchronizes virtual memory.

**Permission:**

M mode/S mode

**Exception:**

Illegal instruction.

**Notes:**

When the TVM bit in the mstatus is 1, running this instruction in S-mode will trigger an illegal instruction exception.

rs1 is the virtual address, and rs2 is the address space identifier (ASID).

- When rs1 and rs2 are both x0, all TLB entries are invalidated.
- When rs1! and rs2 are both x0, all TLB entries that hit the virtual address specified by rs1 are invalidated.
- When rs1 and rs2! are both x0, all TLB entries that hit the process ID specified by rs2 are invalidated.
- When rs1! and rs2! are both x0, all TLB entries that hit the virtual address specified by rs1 and the process ID specified by rs2 are invalidated.

**Instruction format:**

31	25	24	20	19	15	14	12	11	7	6	0
0001001	rs2			rs1		000	00000			1110011	

**16.1.40 SH: a halfword store instruction****Syntax:**

sh rs2, imm12(rs1)

**Operation:**

address ← rs1 + sign\_extend(imm12)

mem[(address+1):address] ← rs2[15:0]

**Permission:**

M mode/S mode/U mode

**Exception:**

Unaligned access exceptions, access error exceptions, and page error exceptions on store instructions

**Instruction format:**

31	25	24	20	19	15	14	12	11	7	6	0
imm12[11:5]				rs2		rs1		001	imm12[4:0]		0100011

#### 16.1.41 SLL: a logical left shift instruction

##### Syntax:

sll rd, rs1, rs2

##### Operation:

$rd \leftarrow rs1 \ll rs2[5:0]$

##### Permission:

M mode/S mode/U mode

##### Exception:

None

##### Instruction format:

31	25	24	20	19	15	14	12	11	7	6	0
0000000				rs2		rs1		001	rd		0110011

#### 16.1.42 SLLI: an immediate logical left shift instruction

##### Syntax:

slli rd, rs1, shamt6

##### Operation:

$rd \leftarrow rs1 \ll shamt6$

##### Permission:

M mode/S mode/U mode

##### Exception:

None

**Instruction format:**

31	26	25	20	19	15	14	12	11	7	6	0
000000		shamt6		rs1		001		rd		0010011	

### 16.1.43 SLLIW: an immediate logical left shift instruction that operates on the lower 32 bits

**Syntax:**

sllw rd, rs1, shamt5

**Operation:**

$$\text{tmp}[31:0] \leftarrow (\text{rs1}[31:0] \ll \text{shamt5})[31:0]$$

$$\text{rd} \leftarrow \text{sign\_extend}(\text{tmp}[31:0])$$
**Permission:**

M mode/S mode/U mode

**Exception:**

None

**Instruction format:**

31	25	24	20	19	15	14	12	11	7	6	0
0000000		shamt5		rs1		001		rd		0011011	

### 16.1.44 SLLW: a logical left shift instruction that operates on the lower 32 bits

**Syntax:**

sllw rd, rs1, rs2

**Operation:**

$$\text{tmp}[31:0] \leftarrow (\text{rs1}[31:0] \ll \text{rs2}[4:0])[31:0]$$

$$\text{rd} \leftarrow \text{sign\_extend}(\text{tmp}[31:0])$$
**Permission:**

M mode/S mode/U mode

**Exception:**

None

**Instruction format:**

31	25	24	20	19	15	14	12	11	7	6	0
0000000	rs2	rs1	001	rd	0111011						

### 16.1.45 SLT: a signed set-if-less-than instruction

#### Syntax:

slt rd, rs1, rs2

#### Operation:

if (rs1 < rs2)

rd ← 1

else

rd ← 0

#### Permission:

M mode/S mode/U mode

#### Exception:

None

#### Instruction format:

31	25	24	20	19	15	14	12	11	7	6	0
0000000	rs2	rs1	010	rd	0110011						

### 16.1.46 SLTI: a signed set-if-less-than-immediate instruction

#### Syntax:

slti rd, rs1, imm12

#### Operation:

if (rs1 < sign\_extend(imm12))

rd ← 1

else

rd ← 0

#### Permission:

M mode/S mode/U mode



**Exception:**

None

**Instruction format:**

31	20	19	15	14	12	11	7	6	0	
imm12[11:0]			rs1		010		rd		0010011	

**16.1.47 SLTIU: an unsigned set-if-less-than-immediate instruction****Syntax:**

sltiu rd, rs1, imm12

**Operation:**

if (rs1 &lt; zero\_extend(imm12))

rd ← 1

else

rd ← 0

**Permission:**

M mode/S mode/U mode

**Exception:**

None

**Instruction format:**

31	20	19	15	14	12	11	7	6	0	
imm12[11:0]			rs1		011		rd		0010011	

**16.1.48 SLTU: an unsigned set-if-less-than instruction****Syntax:**

sltu rd, rs1, rs2

**Operation:**

if (rs1 &lt; rs2)

rd ← 1

else

rd ← 0

**Permission:**

M mode/S mode/U mode

**Exception:**

None

**Instruction format:**

31	25	24	20	19	15	14	12	11	7	6	0
0000000		rs2		rs1		011		rd		0110011	

**16.1.49 SRA: an arithmetic right shift instruction****Syntax:**

sra rd, rs1, rs2

**Operation:** $rd \leftarrow rs1 \ggg rs2[5:0]$ **Permission:**

M mode/S mode/U mode

**Exception:**

None

**Instruction format:**

31	25	24	20	19	15	14	12	11	7	6	0
0100000		rs2		rs1		101		rd		0110011	

**16.1.50 SRAI: an immediate arithmetic right shift instruction****Syntax:**

srai rd, rs1, shamt6

**Operation:** $rd \leftarrow rs1 \ggg shamt6$ **Permission:**

M mode/S mode/U mode

**Exception:**

None

**Instruction format:**

31	26	25	20	19	15	14	12	11	7	6	0
010000	shamt6	rs1	101	rd	0010011						

### 16.1.51 SLLIW: an immediate arithmetic right shift instruction that operates on the lower 32 bits

#### Syntax:

sraiw rd, rs1, shamt5

#### Operation:

$\text{tmp}[31:0] \leftarrow (\text{rs1}[31:0] \gg \text{shamt5})[31:0]$

$\text{rd} \leftarrow \text{sign\_extend}(\text{tmp}[31:0])$

#### Permission:

M mode/S mode/U mode

#### Exception:

None

#### Instruction format:

31	25	24	20	19	15	14	12	11	7	6	0
0100000	shamt5	rs1	101	rd	0011011						

### 16.1.52 SRAW: an arithmetic right shift instruction that operates on the lower 32 bits

#### Syntax:

sraw rd, rs1, rs2

#### Operation:

$\text{tmp} \leftarrow (\text{rs1}[31:0] \gg \text{rs2}[4:0])[31:0]$

$\text{rd} \leftarrow \text{sign\_extend}(\text{tmp})$

#### Permission:

M mode/S mode/U mode

#### Exception:

None

#### Instruction format:

31	25	24	20	19	15	14	12	11	7	6	0
0100000	rs2	rs1	101	rd	0111011						

### 16.1.53 SRET: an instruction for returning from exceptions in S-mode

**Syntax:**

sret

**Operation:**

next pc  $\leftarrow$  sepc

sstatus.sie  $\leftarrow$  sstatus.spie

sstatus.spie  $\leftarrow$  1

**Permission:**

S mode

**Exception:**

Illegal instruction.

**Instruction format:**

31	25	24	20	19	15	14	12	11	7	6	0
0001000	00010	00000	000	00000	1110011						

### 16.1.54 SRL: a logical right shift instruction

**Syntax:**

srl rd, rs1, rs2

**Operation:**

rd  $\leftarrow$  rs1  $\gg$  rs2[5:0]

**Permission:**

M mode/S mode/U mode

**Exception:**

None

**Instruction format:**

31	25	24	20	19	15	14	12	11	7	6	0
0000000	rs2	rs1	101	rd	0110011						

### 16.1.55 SRLI: an immediate logical right shift instruction

**Syntax:**

srli rd, rs1, shamt6

**Operation:**

$$rd \leftarrow rs1 \gg \text{shamt6}$$
**Permission:**

M mode/S mode/U mode

**Exception:**

None

**Instruction format:**

31	26	25	20	19	15	14	12	11	7	6	0
000000			shamt6		rs1		101		rd		0010011

### 16.1.56 SRLIW: an immediate logical right shift instruction that operates on the lower 32 bits

**Syntax:**

$$\text{srlw } rd, rs1, \text{shamt5}$$
**Operation:**

$$\text{tmp}[31:0] \leftarrow (rs1[31:0] \gg \text{shamt5})[31:0]$$

$$rd \leftarrow \text{sign\_extend}(\text{tmp}[31:0])$$
**Permission:**

M mode/S mode/U mode

**Exception:**

None

**Instruction format:**

31	25	24	20	19	15	14	12	11	7	6	0
0000000			shamt5		rs1		101		rd		0011011

### 16.1.57 SRLW: a logical right shift instruction that operates on the lower 32 bits

**Syntax:**

$$\text{srlw } rd, rs1, rs2$$
**Operation:**

$$\text{tmp} \leftarrow (rs1[31:0] \gg rs2[4:0])[31:0]$$

$$rd \leftarrow \text{sign\_extend}(\text{tmp})$$

**Permission:**

M mode/S mode/U mode

**Exception:**

None

**Instruction format:**

31	25	24	20	19	15	14	12	11	7	6	0
0000000		rs2		rs1		101		rd		0111011	

**16.1.58 SUB: a signed subtract instruction****Syntax:**

sub rd, rs1, rs2

**Operation:** $rd \leftarrow rs1 - rs2$ **Permission:**

M mode/S mode/U mode

**Exception:**

None

**Instruction format:**

31	25	24	20	19	15	14	12	11	7	6	0
0100000		rs2		rs1		000		rd		0110011	

**16.1.59 SUBW: a signed subtract instruction that operates on the lower 32 bits****Syntax:**

subw rd, rs1, rs2

**Operation:** $tmp[31:0] \leftarrow rs1[31:0] - rs2[31:0]$  $rd \leftarrow \text{sign\_extend}(tmp[31:0])$ **Permission:**

M mode/S mode/U mode

**Exception:**

None

Instruction format:

31	25	24	20	19	15	14	12	11	7	6	0
0100000		rs2		rs1		000		rd		0111011	

### 16.1.60 SW: a word store instruction

Syntax:

sw rs2, imm12(rs1)

Operation:

address  $\leftarrow$  rs1 + sign\_extend(imm12)

mem[(address+3):address]  $\leftarrow$  rs2[31:0]

Permission:

M mode/S mode/U mode

Exception:

Unaligned access exceptions, access error exceptions, and page error exceptions on store instructions

Instruction format:

31	25	24	20	19	15	14	12	11	7	6	0
imm12[11:5]		rs2		rs1		010		imm12[4:0]		0100011	

### 16.1.61 WFI: an instruction for entering the low power mode

Syntax:

wfi

Operation:

Triggers the CPU to enter the low power mode. In this mode, the CPU clock and most device clocks are disabled.

Permission:

M mode/S mode/U mode

Exception:

None

Instruction format:

31	25	24	20	19	15	14	12	11	7	6	0
0001000		00101		00000		000		00000		1110011	

### 16.1.62 XOR: a bitwise XOR instruction

**Syntax:**

```
xor rd, rs1, rs2
```

**Operation:**

$$rd \leftarrow rs1 \oplus rs2$$
**Permission:**

M mode/S mode/U mode

**Exception:**

None

**Instruction format:**

31	25	24	20	19	15	14	12	11	7	6	0
0000000				rs2		rs1		100		rd	
										0110011	



### 16.2.1 DIV: a signed divide instruction

**Syntax:**

div rd, rs1, rs2

**Operation:**

$rd \leftarrow rs1 / rs2$

**Permission:**

Machine mode (M-mode)/Supervisor mode (S-mode)/User mode (U-mode)

**Exception:**

None

**Notes:**

When the divisor is 0, the division result is 0xffffffffffff.

When overflow occurs, the division result is 0x8000000000000000.

**Instruction format:**

31	25	24	20	19	15	14	12	11	7	6	0
0000001		rs2		rs1		100		rd		0110011	

### 16.2.2 DIVU: an unsigned divide instruction

**Syntax:**

divu rd, rs1, rs2

**Operation:**

$rd \leftarrow rs1 / rs2$

**Permission:**

M mode/S mode/U mode

**Exception:**

None

**Notes:**

When the divisor is 0, the division result is 0xffffffffffff.

**Instruction format:**

31	25	24	20	19	15	14	12	11	7	6	0
0000001		rs2		rs1		101		rd		0110011	

### 16.2.3 DIVUW: an unsigned divide instruction that operates on the lower 32 bits

**Syntax:**

```
divuw rd, rs1, rs2
```

**Operation:**

$$\text{tmp}[31:0] \leftarrow (\text{rs1}[31:0] / \text{rs2}[31:0])[31:0]$$

$$\text{rd} \leftarrow \text{sign\_extend}(\text{tmp}[31:0])$$
**Permission:**

M mode/S mode/U mode

**Exception:**

None

**Notes:**

When the divisor is 0, the division result is 0xffffffffffff.

**Instruction format:**

31	25	24	20	19	15	14	12	11	7	6	0
0000001		rs2		rs1		101		rd		0111011	

### 16.2.4 DIVW: a signed divide instruction that operates on the lower 32 bits

**Syntax:**

```
divw rd, rs1, rs2
```

**Operation:**

$$\text{tmp}[31:0] \leftarrow (\text{rs1}[31:0] / \text{rs2}[31:0])[31:0]$$

$$\text{rd} \leftarrow \text{sign\_extend}(\text{tmp}[31:0])$$
**Permission:**

M mode/S mode/U mode

**Exception:**

None

**Notes:**

When the divisor is 0, the division result is 0xffffffffffff.

When overflow occurs, the division result is 0xffffffff80000000.

**Instruction format:**

31	25	24	20	19	15	14	12	11	7	6	0
0000001		rs2		rs1		100		rd		0111011	

### 16.2.5 MUL: a signed multiply instruction

**Syntax:**

mul rd, rs1, rs2

**Operation:**

$rd \leftarrow (rs1 * rs2)[63:0]$

**Permission:**

M mode/S mode/U mode

**Exception:**

None

**Instruction format:**

31	25	24	20	19	15	14	12	11	7	6	0
0000001		rs2		rs1		000		rd		0110011	

### 16.2.6 MULH: a signed multiply instruction that extracts the upper bits

**Syntax:**

mulh rd, rs1, rs2

**Operation:**

$rd \leftarrow (rs1 * rs2)[127:64]$

**Permission:**

M mode/S mode/U mode

**Exception:**

None

**Instruction format:**

31	25	24	20	19	15	14	12	11	7	6	0
0000001		rs2		rs1		001		rd		0110011	

### 16.2.7 MULHSU: a signed-unsigned multiply instruction that extracts the upper bits

**Syntax:**

mulusu rd, rs1, rs2

**Operation:**

$rd \leftarrow (rs1 * rs2)[127:64]$

**Permission:**

M mode/S mode/U mode

**Exception:**

None

**Notes:**

rs1 indicates a signed number, and rs2 indicates an unsigned number.

**Instruction format:**

31	25	24	20	19	15	14	12	11	7	6	0
0000001		rs2		rs1		010		rd		0110011	

### 16.2.8 MULHU: an unsigned multiply instruction that extracts the upper bits

**Syntax:**

mulhu rd, rs1, rs2

**Operation:**

$rd \leftarrow (rs1 * rs2)[127:64]$

**Permission:**

M mode/S mode/U mode

**Exception:**

None

**Instruction format:**

31	25	24	20	19	15	14	12	11	7	6	0
0000001		rs2		rs1		011		rd		0110011	

### 16.2.9 MULW: a signed multiply instruction that operates on the lower 32 bits

**Syntax:**

mulw rd, rs1, rs2

**Operation:**

$$\text{tmp} \leftarrow (\text{rs1}[31:0] * \text{rs2}[31:0])[31:0]$$

$$\text{rd} \leftarrow \text{sign\_extend}(\text{tmp}[31:0])$$
**Permission:**

M mode/S mode/U mode

**Exception:**

None

**Instruction format:**

31	25	24	20	19	15	14	12	11	7	6	0
0000001		rs2		rs1		000		rd		0111011	

**16.2.10 REM: a signed remainder instruction****Syntax:**

rem rd, rs1, rs2

**Operation:**
$$\text{rd} \leftarrow \text{rs1} \% \text{rs2}$$
**Permission:**

M mode/S mode/U mode

**Exception:**

None

**Notes:**

When the divisor is 0, the remainder operation result is the dividend.

When overflow occurs, the remainder operation result is 0x0.

**Instruction format:**

31	25	24	20	19	15	14	12	11	7	6	0
0000001		rs2		rs1		110		rd		0110011	

**16.2.11 REMU: an unsigned remainder instruction****Syntax:**

remu rd, rs1, rs2

**Operation:**
$$\text{rd} \leftarrow \text{rs1} \% \text{rs2}$$

**Permission:**

M mode/S mode/U mode

**Exception:**

None

**Notes:**

When the divisor is 0, the remainder operation result is the dividend.

**Instruction format:**

31	25	24	20	19	15	14	12	11	7	6	0
0000001		rs2		rs1		111		rd		0110011	

### 16.2.12 REMUW: an unsigned remainder instruction that operates on the lower 32 bits

**Syntax:**

remw rd, rs1, rs2

**Operation:**

$$\text{tmp} \leftarrow (\text{rs1}[31:0] \% \text{rs2}[31:0])[31:0]$$

$$\text{rd} \leftarrow \text{sign\_extend}(\text{tmp})$$
**Permission:**

M mode/S mode/U mode

**Exception:**

None

**Notes:**

When the divisor is 0, the remainder operation result is obtained by extending the signed bit [31] of the dividend.

**Instruction format:**

31	25	24	20	19	15	14	12	11	7	6	0
0000001		rs2		rs1		111		rd		0111011	

### 16.2.13 REMW: a signed remainder instruction that operates on the lower 32 bits

**Syntax:**

remw rd, rs1, rs2

**Operation:**

$$\text{tmp}[31:0] \leftarrow (\text{rs1}[31:0] \% \text{rs2}[31:0])[31:0]$$

$$\text{rd} \leftarrow \text{sign\_extend}(\text{tmp}[31:0])$$
**Permission:**

M mode/S mode/U mode

**Exception:**

None

**Notes:**

When the divisor is 0, the remainder operation result is obtained by extending the signed bit [31] of the dividend.

When overflow occurs, the remainder operation result is 0x0.

**Instruction format:**

31	25	24	20	19	15	14	12	11	7	6	0
0000001		rs2		rs1		110		rd		0111011	

## 16.3 Appendix A-3 A instructions

The following describes the RISC-V A instructions implemented by C910/C920. The instructions are 32 bits wide and sorted in alphabetic order.

### 16.3.1 AMOADD.D: an atomic add instruction

**Syntax:**

$$\text{amoadd.d.aqrl rd, rs2, (rs1)}$$
**Operation:**

$$\text{rd} \leftarrow \text{mem}[\text{rs1}+7:\text{rs1}]$$

$$\text{mem}[\text{rs1}+7:\text{rs1}] \leftarrow \text{mem}[\text{rs1}+7:\text{rs1}] + \text{rs2}$$
**Permission:**

M mode/S mode/U mode

**Exception:**

Unaligned access exceptions, access error exceptions, and page error exceptions on atomic instructions

**Affected flag bits:**

None

**Notes:**

The aq and rl bits determine the sequences of executing the memory access instructions before and after this instruction.

- When aq and rl are both 0, the corresponding assembler instruction is `amoadd.d rd, rs2, (rs1)`.
- When aq is 0 and rl is 1, the corresponding assembler instruction is `amoadd.d.rl rd, rs2, (rs1)`. Results of all memory access instructions before this instruction must be observed before this instruction is executed.
- When aq is 1 and rl is 0, the corresponding assembler instruction is `amoadd.d.aq rd, rs2, (rs1)`. All memory access instructions after this instruction can be executed only after execution of this instruction is completed.
- When aq and rl are both 1, the corresponding assembler instruction is `amoadd.d.aqrl rd, rs2, (rs1)`. Results of all memory access instructions before this instruction must be observed before this instruction is executed, and all memory access instructions after this instruction can be executed only after execution of this instruction is completed.

**Instruction format:**

31	27	26	25	24	20	19	15	14	12	11	7	6	0
00000	aq	rl		rs2		rs1		011		rd		0101111	

**16.3.2 AMOADD.W: an atomic add instruction that operates on the lower 32 bits****Syntax:**

`amoadd.w.aqrl rd, rs2, (rs1)`

**Operation:**

$rd \leftarrow \text{sign\_extend}(\text{mem}[\text{rs1}+3:\text{rs1}])$

$\text{mem}[\text{rs1}+3:\text{rs1}] \leftarrow \text{mem}[\text{rs1}+3:\text{rs1}] + \text{rs2}[31:0]$

**Permission:**

M mode/S mode/U mode

**Exception:**

Unaligned access exceptions, access error exceptions, and page error exceptions on atomic instructions

**Affected flag bits:**

None

**Notes:**

The aq and rl bits determine the sequences of executing the memory access instructions before and after this instruction.



- When `aq` and `rl` are both 0, the corresponding assembler instruction is `amoadd.w rd, rs2, (rs1)`.
- When `aq` is 0 and `rl` is 1, the corresponding assembler instruction is `amoadd.w.rl rd, rs2, (rs1)`. Results of all memory access instructions before this instruction must be observed before this instruction is executed.
- When `aq` is 1 and `rl` is 0, the corresponding assembler instruction is `amoadd.w.aq rd, rs2, (rs1)`. All memory access instructions after this instruction can be executed only after execution of this instruction is completed.
- When `aq` and `rl` are both 1, the corresponding assembler instruction is `amoadd.w.aqrl rd, rs2, (rs1)`. Results of all memory access instructions before this instruction must be observed before this instruction is executed, and all memory access instructions after this instruction can be executed only after execution of this instruction is completed.

**Instruction format:**

31	27	26	25	24	20	19	15	14	12	11	7	6	0
00000	aq	rl		rs2		rs1		010		rd		0101111	

**16.3.3 AMOAND.D: an atomic bitwise AND instruction****Syntax:**

```
amoand.d.aqrl rd, rs2, (rs1)
```

**Operation:**

$$rd \leftarrow \text{mem}[rs1+7:rs1]$$

$$\text{mem}[rs1+7:rs1] \leftarrow \text{mem}[rs1+7:rs1] \& rs2$$
**Permission:** M mode/S mode/U mode**Exception:**

Unaligned access exceptions, access error exceptions, and page error exceptions on atomic instructions

**Affected flag bits:**

None

**Notes:**

The `aq` and `rl` bits determine the sequences of executing the memory access instructions before and after this instruction.

- When `aq` and `rl` are both 0, the corresponding assembler instruction is `amoand.d rd, rs2, (rs1)`.
- When `aq` is 0 and `rl` is 1, the corresponding assembler instruction is `amoand.d.rl rd, rs2, (rs1)`. Results of all memory access instructions before this instruction must be observed before this instruction is executed.

- When `aq` is 1 and `rl` is 0, the corresponding assembler instruction is `amoand.d.aq rd, rs2, (rs1)`. All memory access instructions after this instruction can be executed only after execution of this instruction is completed.
- When `aq` and `rl` are both 1, the corresponding assembler instruction is `amoand.d.aqrl rd, rs2, (rs1)`. Results of all memory access instructions before this instruction must be observed before this instruction is executed, and all memory access instructions after this instruction can be executed only after execution of this instruction is completed.

**Instruction format:**

31	27	26	25	24	20	19	15	14	12	11	7	6	0
01100	aq	rl		rs2		rs1		011		rd		0101111	

### 16.3.4 AMOAND.W: an atomic bitwise AND instruction that operates on the lower 32 bits

**Syntax:**

```
amoand.w.aqrl rd, rs2, (rs1)
```

**Operation:**

$$rd \leftarrow \text{sign\_extend}(\text{mem}[\text{rs1}+3:\text{rs1}])$$

$$\text{mem}[\text{rs1}+3:\text{rs1}] \leftarrow \text{mem}[\text{rs1}+3:\text{rs1}] \& \text{rs2}[31:0]$$
**Permission:**

M mode/S mode/U mode

**Exception:**

Unaligned access exceptions, access error exceptions, and page error exceptions on atomic instructions

**Affected flag bits:**

None

**Notes:**

The `aq` and `rl` bits determine the sequences of executing the memory access instructions before and after this instruction.

- When `aq` and `rl` are both 0, the corresponding assembler instruction is `amoand.w rd, rs2, (rs1)`.
- When `aq` is 0 and `rl` is 1, the corresponding assembler instruction is `amoand.w.rl rd, rs2, (rs1)`. Results of all memory access instructions before this instruction must be observed before this instruction is executed.

- When aq is 1 and rl is 0, the corresponding assembler instruction is amoand.w.aq rd, rs2, (rs1). All memory access instructions after this instruction can be executed only after execution of this instruction is completed.
- When aq and rl are both 1, the corresponding assembler instruction is amoand.w.aqrl rd, rs2, (rs1). Results of all memory access instructions before this instruction must be observed before this instruction is executed, and all memory access instructions after this instruction can be executed only after execution of this instruction is completed.

**Instruction format:**

31	27	26	25	24	20	19	15	14	12	11	7	6	0
01100	aq	rl		rs2		rs1		010		rd		0101111	

**16.3.5 AMOMAX.D: an atomic signed MAX instruction****Syntax:**

amomax.d.aqrl rd, rs2, (rs1)

**Operation:**

$rd \leftarrow \text{mem}[rs1+7:rs1]$

$\text{mem}[rs1+7:rs1] \leftarrow \max(\text{mem}[rs1+7:rs1], rs2)$

**Permission:**

M mode/S mode/U mode

**Exception:**

Unaligned access exceptions, access error exceptions, and page error exceptions on atomic instructions

**Affected flag bits:**

None

**Notes:**

The aq and rl bits determine the sequences of executing the memory access instructions before and after this instruction.

- When aq and rl are both 0, the corresponding assembler instruction is amomax.d rd, rs2, (rs1).
- When aq is 0 and rl is 1, the corresponding assembler instruction is amomax.d.rl rd, rs2, (rs1). Results of all memory access instructions before this instruction must be observed before this instruction is executed.
- When aq is 1 and rl is 0, the corresponding assembler instruction is amomax.d.aq rd, rs2, (rs1). All memory access instructions after this instruction can be executed only after execution of this instruction is completed.

- When aq and rl are both 1, the corresponding assembler instruction is amomax.d.aqrl rd, rs2, (rs1). Results of all memory access instructions before this instruction must be observed before this instruction is executed, and all memory access instructions after this instruction can be executed only after execution of this instruction is completed.

**Instruction format:**

31	27	26	25	24	20	19	15	14	12	11	7	6	0
10100	aq	rl		rs2		rs1		011		rd		0101111	

### 16.3.6 AMOMAX.W: an atomic signed MAX instruction that operates on the lower 32 bits

**Syntax:**

amomax.w.aqrl rd, rs2, (rs1)

**Operation:**

$rd \leftarrow \text{sign\_extend}(\text{mem}[\text{rs1}+3:\text{rs1}])$

$\text{mem}[\text{rs1}+3:\text{rs1}] \leftarrow \max(\text{mem}[\text{rs1}+3:\text{rs1}], \text{rs2}[31:0])$

**Permission:**

M mode/S mode/U mode

**Exception:**

Unaligned access exceptions, access error exceptions, and page error exceptions on atomic instructions

**Affected flag bits:**

None

**Notes:**

The aq and rl bits determine the sequences of executing the memory access instructions before and after this instruction.

- When aq and rl are both 0, the corresponding assembler instruction is amomax.w rd, rs2, (rs1).
- When aq is 0 and rl is 1, the corresponding assembler instruction is amomax.w.rl rd, rs2, (rs1). Results of all memory access instructions before this instruction must be observed before this instruction is executed.
- When aq is 1 and rl is 0, the corresponding assembler instruction is amomax.w.aq rd, rs2, (rs1). All memory access instructions after this instruction can be executed only after execution of this instruction is completed.
- When aq and rl are both 1, the corresponding assembler instruction is amomax.w.aqrl rd, rs2, (rs1). Results of all memory access instructions before this instruction must be observed before this instruction is executed.

tion is executed, and all memory access instructions after this instruction can be executed only after execution of this instruction is completed.

#### Instruction format:

31	27	26	25	24	20	19	15	14	12	11	7	6	0
10100	aq	rl		rs2		rs1		010		rd		0101111	

### 16.3.7 MOMAXU.DA: an atomic unsigned MAX instruction

#### Syntax:

amomaxu.d.aqrl rd, rs2, (rs1)

#### Operation:

$rd \leftarrow \text{mem}[rs1+7:rs1]$

$\text{mem}[rs1+7:rs1] \leftarrow \max(\text{mem}[rs1+7:rs1], rs2)$

#### Permission:

M mode/S mode/U mode

#### Exception:

Unaligned access exceptions, access error exceptions, and page error exceptions on atomic instructions

#### Affected flag bits:

None

#### Notes:

The aq and rl bits determine the sequences of executing the memory access instructions before and after this instruction.

- When aq and rl are both 0, the corresponding assembler instruction is amomaxu.d rd, rs2, (rs1).
- When aq is 0 and rl is 1, the corresponding assembler instruction is amomaxu.d.rl rd, rs2, (rs1). Results of all memory access instructions before this instruction must be observed before this instruction is executed.
- When aq is 1 and rl is 0, the corresponding assembler instruction is amomaxu.d.aq rd, rs2, (rs1). All memory access instructions after this instruction can be executed only after execution of this instruction is completed.
- When aq and rl are both 1, the corresponding assembler instruction is amomaxu.d.aqrl rd, rs2, (rs1). Results of all memory access instructions before this instruction must be observed before this instruction is executed, and all memory access instructions after this instruction can be executed only after execution of this instruction is completed.

#### Instruction format:

31	27	26	25	24	20	19	15	14	12	11	7	6	0
11100	aq	rl	rs2	rs1	011	rd	0101111						

### 16.3.8 AMOMAXU.W: an atomic unsigned MAX instruction that operates on the lower 32 bits.

#### Syntax:

amomaxu.w.aqrl rd, rs2, (rs1)

#### Operation:

$rd \leftarrow \text{zero\_extend}(\text{mem}[\text{rs1}+3:\text{rs1}])$

$\text{mem}[\text{rs1}+3:\text{rs1}] \leftarrow \text{max}(\text{mem}[\text{rs1}+3:\text{rs1}], \text{rs2}[31:0])$

#### Permission:

M mode/S mode/U mode

#### Exception:

Unaligned access exceptions, access error exceptions, and page error exceptions on atomic instructions

#### Affected flag bits:

None

#### Notes:

The aq and rl bits determine the sequences of executing the memory access instructions before and after this instruction.

- When aq and rl are both 0, the corresponding assembler instruction is amomaxu.w rd, rs2, (rs1).
- When aq is 0 and rl is 1, the corresponding assembler instruction is amomaxu.w.rl rd, rs2, (rs1). Results of all memory access instructions before this instruction must be observed before this instruction is executed.
- When aq is 1 and rl is 0, the corresponding assembler instruction is amomaxu.w.aq rd, rs2, (rs1). All memory access instructions after this instruction can be executed only after execution of this instruction is completed.
- When aq and rl are both 1, the corresponding assembler instruction is amomaxu.w.aqrl rd, rs2, (rs1). Results of all memory access instructions before this instruction must be observed before this instruction is executed, and all memory access instructions after this instruction can be executed only after execution of this instruction is completed.

#### Instruction format:

31	27	26	25	24	20	19	15	14	12	11	7	6	0
11100	aq	rl	rs2	rs1	010	rd	0101111						

### 16.3.9 AMOMIN.D: an atomic signed MIN instruction

#### Syntax:

```
amomin.d.aqrl rd, rs2, (rs1)
```

#### Operation:

$$rd \leftarrow \text{mem}[rs1+7:rs1]$$

$$\text{mem}[rs1+7:rs1] \leftarrow \min(\text{mem}[rs1+7:rs1], rs2)$$

#### Permission:

M mode/S mode/U mode

#### Exception:

Unaligned access exceptions, access error exceptions, and page error exceptions on atomic instructions

#### Affected flag bits:

None

#### Notes:

The aq and rl bits determine the sequences of executing the memory access instructions before and after this instruction.

- When aq and rl are both 0, the corresponding assembler instruction is amomin.d rd, rs2, (rs1).
- When aq is 0 and rl is 1, the corresponding assembler instruction is amomin.d.rl rd, rs2, (rs1). Results of all memory access instructions before this instruction must be observed before this instruction is executed.
- When aq is 1 and rl is 0, the corresponding assembler instruction is amomin.d.aq rd, rs2, (rs1). All memory access instructions after this instruction can be executed only after execution of this instruction is completed.
- When aq and rl are both 1, the corresponding assembler instruction is amomin.d.aqrl rd, rs2, (rs1). Results of all memory access instructions before this instruction must be observed before this instruction is executed, and all memory access instructions after this instruction can be executed only after execution of this instruction is completed.

#### Instruction format:

31	27	26	25	24	20	19	15	14	12	11	7	6	0
10000	aq	rl		rs2		rs1		011		rd		0101111	

### 16.3.10 AMOMIN.W: an atomic signed MIN instruction that operates on the lower 32 bits

#### Syntax:

```
amomin.w.aqrl rd, rs2, (rs1)
```

#### Operation:

```
rd ← sign_extend(mem[rs1+3: rs1])
mem[rs1+3:rs1] ← min(mem[rs1+3:rs1], rs2[31:0])
```

#### Permission:

M mode/S mode/U mode

#### Exception:

Unaligned access exceptions, access error exceptions, and page error exceptions on atomic instructions

#### Affected flag bits:

None

#### Notes:

The aq and rl bits determine the sequences of executing the memory access instructions before and after this instruction.

- When aq and rl are both 0, the corresponding assembler instruction is amomin.w rd, rs2, (rs1).
- When aq is 0 and rl is 1, the corresponding assembler instruction is amomin.w.rl rd, rs2, (rs1). Results of all memory access instructions before this instruction must be observed before this instruction is executed.
- When aq is 1 and rl is 0, the corresponding assembler instruction is amomin.w.aq rd, rs2, (rs1). All memory access instructions after this instruction can be executed only after execution of this instruction is completed.
- When aq and rl are both 1, the corresponding assembler instruction is amomin.w.aqrl rd, rs2, (rs1). Results of all memory access instructions before this instruction must be observed before this instruction is executed, and all memory access instructions after this instruction can be executed only after execution of this instruction is completed.

#### Instruction format:

31	27	26	25	24	20	19	15	14	12	11	7	6	0
10000	aq	rl		rs2		rs1		010		rd		0101111	



### 16.3.11 AMOMINU.D: an atomic unsigned MIN instruction

**Syntax:**

$$\text{amominu.d.aqrl rd, rs2, (rs1)}$$
**Operation:**

$$\text{rd} \leftarrow \text{mem}[\text{rs1}+7:\text{rs1}]$$

$$\text{mem}[\text{rs1}+7:\text{rs1}] \leftarrow \min(\text{mem}[\text{rs1}+7:\text{rs1}], \text{rs2})$$
**Permission:**

$$\text{M mode/S mode/U mode}$$
**Exception:**

Unaligned access exceptions, access error exceptions, and page error exceptions on atomic instructions

**Affected flag bits:**

None

**Notes:**

The aq and rl bits determine the sequences of executing the memory access instructions before and after this instruction.

- When aq and rl are both 0, the corresponding assembler instruction is `amominu.d rd, rs2, (rs1)`.
- When aq is 0 and rl is 1, the corresponding assembler instruction is `amominu.d.rl rd, rs2, (rs1)`. Results of all memory access instructions before this instruction must be observed before this instruction is executed.
- When aq is 1 and rl is 0, the corresponding assembler instruction is `amominu.d.aq rd, rs2, (rs1)`. All memory access instructions after this instruction can be executed only after execution of this instruction is completed.
- When aq and rl are both 1, the corresponding assembler instruction is `amominu.d.aqrl rd, rs2, (rs1)`. Results of all memory access instructions before this instruction must be observed before this instruction is executed, and all memory access instructions after this instruction can be executed only after execution of this instruction is completed.

**Instruction format:**

31	27	26	25	24	20	19	15	14	12	11	7	6	0
11000	aq	rl		rs2		rs1		011		rd		0101111	

### 16.3.12 AMOMINU.W: an atomic unsigned MIN instruction that operates on the lower 32 bits

#### Syntax:

```
amominu.w.aqrl rd, rs2, (rs1)
```

#### Operation:

```
rd ← sign_extend(mem[rs1+3: rs1])
mem[rs1+3:rs1] ← min(mem[rs1+3:rs1], rs2[31:0])
```

#### Permission:

M mode/S mode/U mode

#### Exception:

Unaligned access exceptions, access error exceptions, and page error exceptions on atomic instructions

#### Affected flag bits:

None

#### Notes:

The aq and rl bits determine the sequences of executing the memory access instructions before and after this instruction.

- When aq and rl are both 0, the corresponding assembler instruction is amominu.w rd, rs2, (rs1).
- When aq is 0 and rl is 1, the corresponding assembler instruction is amominu.w.rl rd, rs2, (rs1). Results of all memory access instructions before this instruction must be observed before this instruction is executed.
- When aq is 1 and rl is 0, the corresponding assembler instruction is amominu.w.aq rd, rs2, (rs1). All memory access instructions after this instruction can be executed only after execution of this instruction is completed.
- When aq and rl are both 1, the corresponding assembler instruction is amominu.w.aqrl rd, rs2, (rs1). Results of all memory access instructions before this instruction must be observed before this instruction is executed, and all memory access instructions after this instruction can be executed only after execution of this instruction is completed.

#### Instruction format:

31	27	26	25	24	20	19	15	14	12	11	7	6	0
11000	aq	rl		rs2		rs1		010		rd		0101111	

### 16.3.13 AMOOR.D: an atomic bitwise OR instruction.

#### Syntax:

```
amoor.d.aqrl rd, rs2, (rs1)
```

#### Operation:

$$rd \leftarrow \text{mem}[rs1+7:rs1]$$

$$\text{mem}[rs1+7:rs1] \leftarrow \text{mem}[rs1+7:rs1] \mid rs2$$

#### Permission:

M mode/S mode/U mode

#### Exception:

Unaligned access exceptions, access error exceptions, and page error exceptions on atomic instructions

#### Affected flag bits:

None

#### Notes:

The aq and rl bits determine the sequences of executing the memory access instructions before and after this instruction.

- When aq and rl are both 0, the corresponding assembler instruction is amoor.d rd, rs2, (rs1).
- When aq is 0 and rl is 1, the corresponding assembler instruction is amoor.d.rl rd, rs2, (rs1). Results of all memory access instructions before this instruction must be observed before this instruction is executed.
- When aq is 1 and rl is 0, the corresponding assembler instruction is amoor.d.aq rd, rs2, (rs1). All memory access instructions after this instruction can be executed only after execution of this instruction is completed.
- When aq and rl are both 1, the corresponding assembler instruction is amoor.d.aqrl rd, rs2, (rs1). Results of all memory access instructions before this instruction must be observed before this instruction is executed, and all memory access instructions after this instruction can be executed only after execution of this instruction is completed.

#### Instruction format:

31	27	26	25	24	20	19	15	14	12	11	7	6	0
01000	aq	rl		rs2		rs1		011		rd		0101111	

### 16.3.14 AMOOR.W: an atomic bitwise OR instruction that operates on the lower 32 bits

#### Syntax:

```
amoor.w.aqrl rd, rs2, (rs1)
```

**Operation:**

$$rd \leftarrow \text{sign\_extend}(\text{mem}[\text{rs1}+3:\text{rs1}])$$

$$\text{mem}[\text{rs1}+3:\text{rs1}] \leftarrow \text{mem}[\text{rs1}+3:\text{rs1}] \mid \text{rs2}[31:0]$$
**Permission:**

M mode/S mode/U mode

**Exception:**

Unaligned access exceptions, access error exceptions, and page error exceptions on atomic instructions

**Affected flag bits:**

None

**Notes:**

The aq and rl bits determine the sequences of executing the memory access instructions before and after this instruction.

- When aq and rl are both 0, the corresponding assembler instruction is amoor.w rd, rs2, (rs1).
- When aq is 0 and rl is 1, the corresponding assembler instruction is amoor.w.rl rd, rs2, (rs1). Results of all memory access instructions before this instruction must be observed before this instruction is executed.
- When aq is 1 and rl is 0, the corresponding assembler instruction is amoor.w.aq rd, rs2, (rs1). All memory access instructions after this instruction can be executed only after execution of this instruction is completed.
- When aq and rl are both 1, the corresponding assembler instruction is amoor.w.aqrl rd, rs2, (rs1). Results of all memory access instructions before this instruction must be observed before this instruction is executed, and all memory access instructions after this instruction can be executed only after execution of this instruction is completed.

**Instruction format:**

31	27	26	25	24	20	19	15	14	12	11	7	6	0
01000	aq	rl		rs2		rs1		010		rd		0101111	

**16.3.15 AMOSWAP.D: an atomic swap instruction****Syntax:**

```
amoswap.d.aqrl rd, rs2, (rs1)
```

**Operation:**

$$rd \leftarrow \text{mem}[rs1+7: rs1]$$

$$\text{mem}[rs1+7:rs1] \leftarrow rs2$$
**Permission:**

M mode/S mode/U mode

**Exception:**

Unaligned access exceptions, access error exceptions, and page error exceptions on atomic instructions

**Affected flag bits:** None**Notes:**

The aq and rl bits determine the sequences of executing the memory access instructions before and after this instruction.

- When aq and rl are both 0, the corresponding assembler instruction is amoswap.d rd, rs2, (rs1).
- When aq is 0 and rl is 1, the corresponding assembler instruction is amoswap.d.rl rd, rs2, (rs1). Results of all memory access instructions before this instruction must be observed before this instruction is executed.
- When aq is 1 and rl is 0, the corresponding assembler instruction is amoswap.d.aq rd, rs2, (rs1). All memory access instructions after this instruction can be executed only after execution of this instruction is completed.
- When aq and rl are both 1, the corresponding assembler instruction is amoswap.d.aqrl rd, rs2, (rs1). Results of all memory access instructions before this instruction must be observed before this instruction is executed, and all memory access instructions after this instruction can be executed only after execution of this instruction is completed.

**Instruction format:**

31	27	26	25	24	20	19	15	14	12	11	7	6	0
00001	aq	rl		rs2		rs1		011		rd		0101111	

**16.3.16 AMOSWAP.W: an atomic swap instruction that operates on the lower 32 bits****Syntax:**

amoswap.w.aqrl rd, rs2, (rs1)

**Operation:**

$$rd \leftarrow \text{sign\_extend}(\text{mem}[rs1+3: rs1])$$

$$\text{mem}[rs1+3:rs1] \leftarrow rs2[31:0]$$
**Permission:**

M mode/S mode/U mode

**Exception:**

Unaligned access exceptions, access error exceptions, and page error exceptions on atomic instructions

**Affected flag bits:** None

**Notes:**

The aq and rl bits determine the sequences of executing the memory access instructions before and after this instruction.

- When aq and rl are both 0, the corresponding assembler instruction is amoswap.w rd, rs2, (rs1).
- When aq is 0 and rl is 1, the corresponding assembler instruction is amoswap.w.rl rd, rs2, (rs1). Results of all memory access instructions before this instruction must be observed before this instruction is executed.
- When aq is 1 and rl is 0, the corresponding assembler instruction is amoswap.w.aq rd, rs2, (rs1). All memory access instructions after this instruction can be executed only after execution of this instruction is completed.
- When aq and rl are both 1, the corresponding assembler instruction is amoswap.w.aqrl rd, rs2, (rs1). Results of all memory access instructions before this instruction must be observed before this instruction is executed, and all memory access instructions after this instruction can be executed only after execution of this instruction is completed.

**Instruction format:**

31	27	26	25	24	20	19	15	14	12	11	7	6	0
00001	aq	rl		rs2		rs1		010		rd		0101111	

**16.3.17 AMOXOR.D: an atomic bitwise XOR instruction****Syntax:**

amoxor.d.aqrl rd, rs2, (rs1)

**Operation:**

$rd \leftarrow \text{mem}[rs1+7:rs1]$

$\text{mem}[rs1+7:rs1] \leftarrow \text{mem}[rs1+7:rs1] \wedge rs2$

**Permission:**

M mode/S mode/U mode

**Exception:**

Unaligned access exceptions, access error exceptions, and page error exceptions on atomic instructions

**Affected flag bits:**

None

**Notes:**

The aq and rl bits determine the sequences of executing the memory access instructions before and after this instruction.

- When aq and rl are both 0, the corresponding assembler instruction is amoxor.d rd, rs2, (rs1).
- When aq is 0 and rl is 1, the corresponding assembler instruction is amoxor.d.rl rd, rs2, (rs1). Results of all memory access instructions before this instruction must be observed before this instruction is executed.
- When aq is 1 and rl is 0, the corresponding assembler instruction is amoxor.d.aq rd, rs2, (rs1). All memory access instructions after this instruction can be executed only after execution of this instruction is completed.
- When aq and rl are both 1, the corresponding assembler instruction is amoxor.d.aqrl rd, rs2, (rs1). Results of all memory access instructions before this instruction must be observed before this instruction is executed, and all memory access instructions after this instruction can be executed only after execution of this instruction is completed.

**Instruction format:**

31	27	26	25	24	20	19	15	14	12	11	7	6	0
00100	aq	rl		rs2		rs1		011		rd		0101111	

### 16.3.18 AMOXOR.W: an atomic bitwise XOR instruction that operates on the lower 32 bits

**Syntax:**

```
amoxor.w.aqrl rd, rs2, (rs1)
```

**Operation:**

```
rd ← sign_extend(mem[rs1+3: rs1])
mem[rs1+3:rs1] ← mem[rs1+3:rs1] ^ rs2[31:0]
```

**Permission:**

M mode/S mode/U mode

**Exception:**

Unaligned access exceptions, access error exceptions, and page error exceptions on atomic instructions

**Affected flag bits:**

None

**Notes:**

The aq and rl bits determine the sequences of executing the memory access instructions before and after this instruction.

- When aq and rl are both 0, the corresponding assembler instruction is amoxor.w rd, rs2, (rs1).
- When aq is 0 and rl is 1, the corresponding assembler instruction is amoxor.w.rl rd, rs2, (rs1). Results of all memory access instructions before this instruction must be observed before this instruction is executed.
- When aq is 1 and rl is 0, the corresponding assembler instruction is amoxor.w.aq rd, rs2, (rs1). All memory access instructions after this instruction can be executed only after execution of this instruction is completed.
- When aq and rl are both 1, the corresponding assembler instruction is amoxor.w.aqrl rd, rs2, (rs1). Results of all memory access instructions before this instruction must be observed before this instruction is executed, and all memory access instructions after this instruction can be executed only after execution of this instruction is completed.

**Instruction format:**

31	27	26	25	24	20	19	15	14	12	11	7	6	0
00100	aq	rl		rs2		rs1		010		rd		0101111	

### 16.3.19 LR.D: a doubleword load-reserved instruction

**Syntax:**

lr.d.aqrl rd, (rs1)

**Operation:**

$rd \leftarrow \text{mem}[rs1+7: rs1]$

$\text{mem}[rs1+7:rs1]$  is reserved

**Permission:**

M mode/S mode/U mode

**Exception:**

Unaligned access exceptions, access error exceptions, and page error exceptions on atomic instructions

**Affected flag bits:**

None

**Notes:**

The aq and rl bits determine the sequences of executing the memory access instructions before and after this instruction.

- When aq and rl are both 0, the corresponding assembler instruction is lr.d rd, (rs1).



- When aq is 0 and rl is 1, the corresponding assembler instruction is `lr.d.rl rd, (rs1)`. Results of all memory access instructions before this instruction must be observed before this instruction is executed.
- When aq is 1 and rl is 0, the corresponding assembler instruction is `lr.d.aq rd, (rs1)`. All memory access instructions after this instruction can be executed only after execution of this instruction is completed.
- When aq and rl are both 1, the corresponding assembler instruction is `lr.d.aqrl rd, (rs1)`. Results of all memory access instructions before this instruction must be observed before this instruction is executed, and all memory access instructions after this instruction can be executed only after execution of this instruction is completed.

**Instruction format:**

31	27	26	25	24	20	19	15	14	12	11	7	6	0
00010	aq	rl	00000	rs1	011	rd	0101111						

**16.3.20 LR.W: a word load-reserved instruction****Syntax:**

`lr.w.aqrl rd, (rs1)`

**Operation:**

$rd \leftarrow \text{sign\_extend}(\text{mem}[\text{rs1}+3: \text{rs1}])$

`mem[rs1+3:rs1]` is reserved

**Permission:**

M mode/S mode/U mode

**Exception:** Unaligned access exceptions, access error exceptions, and page error exceptions on atomic instructions

**Affected flag bits:** None

**Notes:** The aq and rl bits determine the sequences of executing the memory access instructions before and after this instruction.

- When aq and rl are both 0, the corresponding assembler instruction is `lr.w rd, (rs1)`.
- When aq is 0 and rl is 1, the corresponding assembler instruction is `lr.w.rl rd, (rs1)`. Results of all memory access instructions before this instruction must be observed before this instruction is executed.
- When aq is 1 and rl is 0, the corresponding assembler instruction is `lr.w.aq rd, (rs1)`. All memory access instructions after this instruction can be executed only after execution of this instruction is completed.
- When aq and rl are both 1, the corresponding assembler instruction is `lr.w.aqrl rd, (rs1)`. Results of all memory access instructions before this instruction must be observed before this instruction is executed, and all memory access instructions after this instruction can be executed only after execution of this instruction is completed.

Instruction format:

31	27	26	25	24	20	19	15	14	12	11	7	6	0
00010	aq	rl	00000	rs1	010	rd	0101111						

### 16.3.21 SC.D: a doubleword store-conditional instruction

Syntax:

sc.d.aqrl rd, rs2, (rs1)

Operation:

If(mem[rs1+7:rs1] is reserved)

mem[rs1+7: rs1]  $\leftarrow$  rs2

rd  $\leftarrow$  0

else

rd  $\leftarrow$  1

Permission:

M mode/S mode/U mode

Exception:

Unaligned access exceptions, access error exceptions, and page error exceptions on atomic instructions

Affected flag bits:

None

Notes:

The aq and rl bits determine the sequences of executing the memory access instructions before and after this instruction.

- When aq and rl are both 0, the corresponding assembler instruction is sc.d rd, rs2, (rs1).
- When aq is 0 and rl is 1, the corresponding assembler instruction is sc.d.rl rd, rs2, (rs1). Results of all memory access instructions before this instruction must be observed before this instruction is executed.
- When aq is 1 and rl is 0, the corresponding assembler instruction is sc.d.aq rd, rs2, (rs1). All memory access instructions after this instruction can be executed only after execution of this instruction is completed.
- When aq and rl are both 1, the corresponding assembler instruction is sc.d.aqrl rd, rs2, (rs1). Results of all memory access instructions before this instruction must be observed before this instruction is executed, and all memory access instructions after this instruction can be executed only after execution of this instruction is completed.

Instruction format:

31	27	26	25	24	20	19	15	14	12	11	7	6	0
00011	aq	rl			rs2		rs1		011		rd		0101111

### 16.3.22 SC.W: a word store-conditional instruction

Syntax:

sc.w.aqrl rd, rs2, (rs1)

Operation:

if(mem[rs1+3:rs1] is reserved)

mem[rs1+3:rs1]  $\leftarrow$  rs2[31:0]

rd  $\leftarrow$  0

else

rd  $\leftarrow$  1

Permission:

M mode/S mode/U mode

Exception:

Unaligned access exceptions, access error exceptions, and page error exceptions on atomic instructions

Affected flag bits:

None

Notes:

The aq and rl bits determine the sequences of executing the memory access instructions before and after this instruction.

- When aq and rl are both 0, the corresponding assembler instruction is sc.w rd, rs2, (rs1).
- When aq is 0 and rl is 1, the corresponding assembler instruction is sc.w.rl rd, rs2, (rs1). Results of all memory access instructions before this instruction must be observed before this instruction is executed.
- When aq is 1 and rl is 0, the corresponding assembler instruction is sc.w.aq rd, rs2, (rs1). All memory access instructions after this instruction can be executed only after execution of this instruction is completed.
- When aq and rl are both 1, the corresponding assembler instruction is sc.w.aqrl rd, rs2, (rs1). Results of all memory access instructions before this instruction must be observed before this instruction is executed, and all memory access instructions after this instruction can be executed only after execution of this instruction is completed.

Instruction format:

31	27	26	25	24	20	19	15	14	12	11	7	6	0
00011	aq	rl		rs2		rs1		010		rd		0101111	

## 16.4 Appendix A-4 F instructions

The following describes the RISC-V F instructions implemented by C910/C920. The instructions are 32 bits wide and sorted in alphabetic order.

For single-precision floating-point instructions, if the upper 32 bits in the source register are not all 1, the single-precision data is treated as qNaN.

When the fs bit in the mstatus register is 2' b00, running any instruction listed in this appendix will trigger an illegal instruction exception. When the fs bit in the mstatus register is not 2' b00, it is set to 2' b11 after any instruction listed in this appendix is executed.

### 16.4.1 FADD.S: a single-precision floating-point add instruction

**Syntax:**

fadd.s fd, fs1, fs2, rm

**Operation:**

$frd \leftarrow fs1 + fs2$

**Permission:**

M mode/S mode/U mode

**Exception:**

Illegal instruction.

**Affected flag bits:**

Floating-point status bits NV, OF, and NX

**Notes:**

RM determines the round-off mode:

- 3' b000: Rounds off to the nearest even number. The corresponding assembler instruction is fadd.s fd, fs1, fs2, rne.
- 3' b001: Rounds off to zero. The corresponding assembler instruction is fadd.s fd, fs1, fs2, rtz.
- 3' b010: Rounds off to negative infinity. The corresponding assembler instruction is fadd.s fd, fs1, fs2, rdn.

- 3' b011: Rounds off to positive infinity. The corresponding assembler instruction is `fadd.s fd, fs1, fs2, rup`.
- 3' b100: Rounds off to the nearest large value. The corresponding assembler instruction is `fadd.s fd, fs1, fs2, rmm`.
- 3' b101: This code is reserved and not used.
- 3' b110: This code is reserved and not used.
- 3' b111: Dynamically rounds off based on the `rm` bit of the floating-point control and status register (FCSR), `fcsr`. The corresponding assembler instruction is `fadd.s fd, fs1, fs2`.

**Instruction format:**

31	25	24	20	19	15	14	12	11	7	6	0
0000000				fs2		fs1		rm	fd		1010011

**16.4.2 FCLASS.S: a single-precision floating-point classify instruction****Syntax:**

```
fclass.s rd, fs1
```

**Operation:**

```
if ( fs1 = -inf)
```

```
    rd ← 64' h1
```

```
if ( fs1 = -norm)
```

```
    rd ← 64' h2
```

```
if ( fs1 = -subnorm)
```

```
    rd ← 64' h4
```

```
if ( fs1 = -zero)
```

```
    rd ← 64' h8
```

```
if ( fs1 = +zero)
```

```
    rd ← 64' h10
```

```
if ( fs1 = +subnorm)
```

```
    rd ← 64' h20
```

```
if ( fs1 = +norm)
```

```
    rd ← 64' h40
```

```
if ( fs1 = +Inf)
```

```

rd ← 64' h80
if ( fs1 = sNaN)
    rd ← 64' h100
if ( fs1 = qNaN)
    rd ← 64' h200

```

**Permission:**

M mode/S mode/U mode

**Exception:**

Illegal instruction.

**Affected flag bits:**

None

**Instruction format:**

31	25	24	20	19	15	14	12	11	7	6	0
1110000		00000		fs1		001		rd		1010011	

### 16.4.3 FCVT.L.S: an instruction that converts a single-precision floating-point number into a signed long integer

**Syntax:**

```
fcvt.l.s rd, fs1, rm
```

**Operation:**

```
rd ← single_convert_to_signed_long(fs1)
```

**Permission:**

M mode/S mode/U mode

**Exception:**

Illegal instruction.

**Affected flag bits:**

Floating-point status bits NV and NX

**Notes:**

RM determines the round-off mode:

- 3' b000: Rounds off to the nearest even number. The corresponding assembler instruction is fcvt.l.s rd, fs1, rne.

- 3' b001: Rounds off to zero. The corresponding assembler instruction is `fcvt.l.s rd, fs1, rtz`.
- 3' b010: Rounds off to negative infinity. The corresponding assembler instruction is `fcvt.l.s rd, fs1, rdn`.
- 3' b011: Rounds off to positive infinity. The corresponding assembler instruction is `fcvt.l.s rd, fs1, rup`.
- 3' b100: Rounds off to the nearest large value. The corresponding assembler instruction is `fcvt.l.s rd, fs1, rmm`.
- 3' b101: This code is reserved and not used.
- 3' b110: This code is reserved and not used.
- 3' b111: Dynamically rounds off based on the `rm` bit of the `fcsr` register. The corresponding assembler instruction is `fcvt.l.s rd, fs1`.

**Instruction format:**

31	25	24	20	19	15	14	12	11	7	6	0
1100000				00010		fs1		rm	rd		1010011

### 16.4.4 FCVT.LU.S: an instruction that converts a single-precision floating-point number into an unsigned long integer

**Syntax:**

```
fcvt.lu.s rd, fs1, rm
```

**Operation:**

```
rd ← single_convert_to_unsigned_long(fs1)
```

**Permission:**

M mode/S mode/U mode

**Exception:**

Illegal instruction.

**Affected flag bits:**

Floating-point status bits NV and NX

**Notes:**

RM determines the round-off mode:

- 3' b000: Rounds off to the nearest even number. The corresponding assembler instruction is `fcvt.lu.s rd, fs1, rne`.
- 3' b001: Rounds off to zero. The corresponding assembler instruction is `fcvt.lu.s rd, fs1, rtz`.

- 3' b010: Rounds off to negative infinity. The corresponding assembler instruction is `fcvt.lu.s rd, fs1, rdn`.
- 3' b011: Rounds off to positive infinity. The corresponding assembler instruction is `fcvt.lu.s rd, fs1, rup`.
- 3' b100: Rounds off to the nearest large value. The corresponding assembler instruction is `fcvt.lu.s rd, fs1, rmm`.
- 3' b101: This code is reserved and not used.
- 3' b110: This code is reserved and not used.
- 3' b111: Dynamically rounds off based on the `rm` bit of the `fcsr` register. The corresponding assembler instruction is `fcvt.lu.s rd, fs1`.

**Instruction format:**

31	25	24	20	19	15	14	12	11	7	6	0		
1100000				00011		fs1		rm		rd		1010011	

### 16.4.5 FCVT.S.L: an instruction that converts a signed long integer into a single-precision floating-point number

**Syntax:**

```
fcvt.s.l fd, rs1, rm
```

**Operation:**

```
fd ← signed_long_convert_to_single(fs1)
```

**Permission:**

M mode/S mode/U mode

**Exception:**

Illegal instruction.

**Affected flag bits:**

Floating-point status bit NX

**Notes:**

RM determines the round-off mode:

- 3' b000: Rounds off to the nearest even number. The corresponding assembler instruction is `fcvt.s.l fd, rs1, rne`.
- 3' b001: Rounds off to zero. The corresponding assembler instruction is `fcvt.s.l fd, rs1, rtz`.



- 3' b010: Rounds off to negative infinity. The corresponding assembler instruction is `fcvt.s.l fd, fs1, rdn`.
- 3' b011: Rounds off to positive infinity. The corresponding assembler instruction is `fcvt.s.l fd, fs1, rup`.
- 3' b100: Rounds off to the nearest large value. The corresponding assembler instruction is `fcvt.s.l fd, fs1, rmm`.
- 3' b101: This code is reserved and not used.
- 3' b110: This code is reserved and not used.
- 3' b111: Dynamically rounds off based on the `rm` bit of the `fcsr` register. The corresponding assembler instruction is `fcvt.s.l fd, fs1`.

**Instruction format:**

31	25	24	20	19	15	14	12	11	7	6	0	
1101000				00010		rs1		rm		fd		1010011

### 16.4.6 FCVT.S.LU: an instruction that converts an unsigned long integer into a single-precision floating-point number

**Syntax:**

`fcvt.s.l fd, fs1, rm`

**Operation:**

`fd` ← `unsigned_long_convert_to_single_fp(fs1)`

**Permission:**

M mode/S mode/U mode

**Exception:**

Illegal instruction.

**Affected flag bits:**

Floating-point status bit NX

**Notes:**

RM determines the round-off mode:

- 3' b000: Rounds off to the nearest even number. The corresponding assembler instruction is `fcvt.s.lu fd, fs1, rne`.
- 3' b001: Rounds off to zero. The corresponding assembler instruction is `fcvt.s.lu fd, fs1, rtz`.
- 3' b010: Rounds off to negative infinity. The corresponding assembler instruction is `fcvt.s.lu fd, fs1, rdn`.

- 3' b011: Rounds off to positive infinity. The corresponding assembler instruction is `fcvt.s.lu fd, fs1, rup`.
- 3' b100: Rounds off to the nearest large value. The corresponding assembler instruction is `fcvt.s.lu fd, fs1, rmm`.
- 3' b101: This code is reserved and not used.
- 3' b110: This code is reserved and not used.
- 3' b111: Dynamically rounds off based on the `rm` bit of the `fcsr` register. The corresponding assembler instruction is `fcvt.s.lu fd, fs1`.

**Instruction format:**

31	25	24	20	19	15	14	12	11	7	6	0
1101000		00011		rs1		rm		fd		1010011	

### 16.4.7 FCVT.S.W: an instruction that converts a signed integer into a single-precision floating-point number

**Syntax:**

```
fcvt.s.w fd, rs1, rm
```

**Operation:**

$$fd \leftarrow \text{signed\_int\_convert\_to\_single}(fs1)$$
**Permission:**

M mode/S mode/U mode

**Exception:**

Illegal instruction.

**Affected flag bits:**

Floating-point status bit `NX`

**Notes:**

RM determines the round-off mode:

- 3' b000: Rounds off to the nearest even number. The corresponding assembler instruction is `fcvt.s.w fd, rs1, rne`.
- 3' b001: Rounds off to zero. The corresponding assembler instruction is `fcvt.s.w fd, rs1, rtz`.
- 3' b010: Rounds off to negative infinity. The corresponding assembler instruction is `fcvt.s.w fd, rs1, rdn`.

- 3' b011: Rounds off to positive infinity. The corresponding assembler instruction is `fcvt.s.w fd, rs1, rup`.
- 3' b100: Rounds off to the nearest large value. The corresponding assembler instruction is `fcvt.s.w fd, rs1, rmm`.
- 3' b101: This code is reserved and not used.
- 3' b110: This code is reserved and not used.
- 3' b111: Dynamically rounds off based on the `rm` bit of the `fcsr` register. The corresponding assembler instruction is `fcvt.s.w fd, rs1`.

**Instruction format:**

31	25	24	20	19	15	14	12	11	7	6	0
1101000				00000				rs1		rm	
								fd		1010011	

### 16.4.8 FCVT.S.WU: an instruction that converts an unsigned integer into a single-precision floating-point number

**Syntax:**

`fcvt.s.wu fd, rs1, rm`

**Operation:**

$fd \leftarrow \text{unsigned\_int\_convert\_to\_single\_fp}(fs1)$

**Permission:**

M mode/S mode/U mode

**Exception:**

Illegal instruction.

**Affected flag bits:**

Floating-point status bit NX

**Notes:**

RM determines the round-off mode:

- 3' b000: Rounds off to the nearest even number. The corresponding assembler instruction is `fcvt.s.wu fd, rs1, rne`.
- 3' b001: Rounds off to zero. The corresponding assembler instruction is `fcvt.s.wu fd, rs1, rtz`.
- 3' b010: Rounds off to negative infinity. The corresponding assembler instruction is `fcvt.s.wu fd, rs1, rdn`.

- 3' b011: Rounds off to positive infinity. The corresponding assembler instruction is `fcvt.s.wu fd, rs1, rup`.
- 3' b100: Rounds off to the nearest large value. The corresponding assembler instruction is `fcvt.s.wu fd, rs1, rmm`.
- 3' b101: This code is reserved and not used.
- 3' b110: This code is reserved and not used.
- 3' b111: Dynamically rounds off based on the `rm` bit of the `fcsr` register. The corresponding assembler instruction is `fcvt.s.wu fd, rs1`.

**Instruction format:**

31	25	24	20	19	15	14	12	11	7	6	0
1101000	00001	rs1	rm	fd	1010011						

### 16.4.9 FCVT.W.S: an instruction that converts a single-precision floating-point number into a signed integer

**Syntax:**

`fcvt.w.s rd, fs1, rm`

**Operation:**

```
tmp ← single_convert_to_signed_int(fs1)
rd ← sign_extend(tmp)
```

**Permission:**

M mode/S mode/U mode

**Exception:**

Illegal instruction.

**Affected flag bits:**

Floating-point status bits NV and NX

**Notes:**

RM determines the round-off mode:

- 3' b000: Rounds off to the nearest even number. The corresponding assembler instruction is `fcvt.w.s rd, fs1, rne`.
- 3' b001: Rounds off to zero. The corresponding assembler instruction is `fcvt.w.s rd, fs1, rtz`.
- 3' b010: Rounds off to negative infinity. The corresponding assembler instruction is `fcvt.w.s rd, fs1, rdn`.

- 3' b011: Rounds off to positive infinity. The corresponding assembler instruction is `fcvt.w.s rd, fs1, rup`.
- 3' b100: Rounds off to the nearest large value. The corresponding assembler instruction is `fcvt.w.s rd, fs1, rmm`.
- 3' b101: This code is reserved and not used.
- 3' b110: This code is reserved and not used.
- 3' b111: Dynamically rounds off based on the `rm` bit of the `fcsr` register. The corresponding assembler instruction is `fcvt.w.s rd, fs1`.

**Instruction format:**

31	25	24	20	19	15	14	12	11	7	6	0
1100000				00000		fs1		rm	rd	1010011	

### 16.4.10 FCVT.WU.S: an instruction that converts a single-precision floating-point number into an unsigned integer

**Syntax:**

```
fcvt.wu.s rd, fs1, rm
```

**Operation:**

```
tmp ← single_convert_to_unsigned_int(fs1)
rd ← sign_extend(tmp)
```

**Permission:**

M mode/S mode/U mode

**Exception:**

Illegal instruction.

**Affected flag bits:**

Floating-point status bits NV and NX

**Notes:**

RM determines the round-off mode:

- 3' b000: Rounds off to the nearest even number. The corresponding assembler instruction is `fcvt.wu.s rd, fs1, rne`.
- 3' b001: Rounds off to zero. The corresponding assembler instruction is `fcvt.wu.s rd, fs1, rtz`.
- 3' b010: Rounds off to negative infinity. The corresponding assembler instruction is `fcvt.wu.s rd, fs1, rdn`.

- 3' b011: Rounds off to positive infinity. The corresponding assembler instruction is `fcvt.wu.s rd, fs1, rup`.
- 3' b100: Rounds off to the nearest large value. The corresponding assembler instruction is `fcvt.wu.s rd, fs1, rmm`.
- 3' b101: This code is reserved and not used.
- 3' b110: This code is reserved and not used.
- 3' b111: Dynamically rounds off based on the `rm` bit of the `fcsr` register. The corresponding assembler instruction is `fcvt.wu.s rd, fs1`.

**Instruction format:**

31	25	24	20	19	15	14	12	11	7	6	0
1100000		00001		fs1		rm		rd		1010011	

**16.4.11 FDIV.S: a single-precision floating-point divide instruction****Syntax:**

`fdiv.s fd, fs1, fs2, rm`

**Operation:**

$fd \leftarrow fs1 / fs2$

**Permission:**

M mode/S mode/U mode

**Exception:**

Illegal instruction.

**Affected flag bits:**

Floating-point status bits NV, DZ, OF, UF, and NX

**Notes:**

RM determines the round-off mode:

- 3' b000: Rounds off to the nearest even number. The corresponding assembler instruction is `fdiv.s fd, fs1, fs2, rne`.
- 3' b001: Rounds off to zero. The corresponding assembler instruction is `fdiv.s fd, fs1, fs2, rtz`.
- 3' b010: Rounds off to negative infinity. The corresponding assembler instruction is `fdiv.s fd, fs1, fs2, rdn`.
- 3' b011: Rounds off to positive infinity. The corresponding assembler instruction is `fdiv.s fd, fs1, fs2, rup`.

- 3' b100: Rounds off to the nearest large value. The corresponding assembler instruction is `fdiv.s fd, fs1, fs2, rmm`.
- 3' b101: This code is reserved and not used.
- 3' b110: This code is reserved and not used.
- 3' b111: Dynamically rounds off based on the `rm` bit of the `fcsr` register. The corresponding assembler instruction is `fdiv.s fd, fs1, fs2`.

**Instruction format:**

31	25	24	20	19	15	14	12	11	7	6	0
0001100		fs1		fs2		rm		fd		1010011	

### 16.4.12 FEQ.S: a single-precision floating-point compare equal instruction

**Syntax:**

`feq.s rd, fs1, fs2`

**Operation:**

if(`fs1 == fs2`)

`rd`  $\leftarrow$  1

else

`rd`  $\leftarrow$  0

**Permission:**

M mode/S mode/U mode

**Exception:**

Illegal instruction.

**Affected flag bits:**

Floating-point status bit NV

**Instruction format:**

31	25	24	20	19	15	14	12	11	7	6	0
1010000		fs2		fs1		010		rd		1010011	

### 16.4.13 FLE.S: a single-precision floating-point compare less than or equal to instruction

**Syntax:**

`fle.s rd, fs1, fs2`

**Operation:**

```

if(fs1 <= fs2)
    rd ← 1
else
    rd ← 0

```

**Permission:**

M mode/S mode/U mode

**Exception:**

Illegal instruction.

**Affected flag bits:**

Floating-point status bit NV

**Instruction format:**

31	25	24	20	19	15	14	12	11	7	6	0
1010000				fs2	fs1		000	rd		1010011	

### 16.4.14 FLT.S: a single-precision floating-point compare less than instruction

**Syntax:**

```
flt.s rd, fs1, fs2
```

**Operation:**

```

if(fs1 < fs2)
    rd ← 1
else
    rd ← 0

```

**Permission:**

M mode/S mode/U mode

**Exception:**

Illegal instruction.

**Affected flag bits:**

Floating-point status bit NV

**Instruction format:**



31	25	24	20	19	15	14	12	11	7	6	0
1010000		fs2		fs1		001		rd		1010011	

### 16.4.15 FLW: a single-precision floating-point load instruction

**Syntax:**

flw fd, imm12(rs1)

**Operation:**

address  $\leftarrow$  rs1 + sign\_extend(imm12)

fd[31:0]  $\leftarrow$  mem[(address+3):address]

fd[63:32]  $\leftarrow$  32' hfffffff

**Permission:**

M mode/S mode/U mode

**Exception:**

Unaligned access, access error, page error, or Illegal instruction.

**Affected flag bits:**

None

**Instruction format:**

31	20	19	15	14	12	11	7	6	0
imm12[11:0]		rs1		010		fd		0000111	

### 16.4.16 FMADD.S: a single-precision floating-point multiply-add instruction

**Syntax:**

fmadd.s fd, fs1, fs2, fs3, rm

**Operation:**

rd  $\leftarrow$  fs1 \* fs2 + fs3

**Permission:**

M mode/S mode/U mode

**Exception:**

Illegal instruction.

**Affected flag bits:**

Floating-point status bits NV, OF, UF, and IX

**Notes:**

RM determines the round-off mode:

- 3' b000: Rounds off to the nearest even number. The corresponding assembler instruction is `fmadd.s fd, fs1, fs2, fs3, rne`.
- 3' b001: Rounds off to zero. The corresponding assembler instruction is `fmadd.s fd, fs1, fs2, fs3, rtz`.
- 3' b010: Rounds off to negative infinity. The corresponding assembler instruction is `fmadd.s fd, fs1, fs2, fs3, rdn`.
- 3' b011: Rounds off to positive infinity. The corresponding assembler instruction is `fmadd.s fd, fs1, fs2, fs3, rup`.
- 3' b100: Rounds off to the nearest large value. The corresponding assembler instruction is `fmadd.s fd, fs1, fs2, fs3, rmm`.
- 3' b101: This code is reserved and not used.
- 3' b110: This code is reserved and not used.
- 3' b111: Dynamically rounds off based on the `rm` bit of the `fcsr` register. The corresponding assembler instruction is `fmadd.s fd, fs1, fs2, fs3`.

**Instruction format:**

31	27	26	25	24	20	19	15	14	12	11	7	6	0
fs3				00	fs2				fs1		rm	fd	1000011

**16.4.17 FMAX.S: a single-precision floating-point MAX instruction****Syntax:**

```
fmmax.s fd, fs1, fs2
```

**Operation:**

```
if(fs1 >= fs2)
```

```
    fd ← fs1
```

```
else
```

```
    fd ← fs2
```

**Permission:**

M mode/S mode/U mode

**Exception:**

Illegal instruction.

**Affected flag bits:**

Floating-point status bit NV

**Instruction format:**

31	25	24	20	19	15	14	12	11	7	6	0
0010100		fs2		fs1		001		fd		1010011	

### 16.4.18 FMIN.S: a single-precision floating-point MIN instruction

**Syntax:**

fmin.s fd, fs1, fs2

**Operation:**

if(fs1 >= fs2)

fd ← fs2

else

fd ← fs1

**Permission:**

M mode/S mode/U mode

**Exception:**

Illegal instruction.

**Affected flag bits:**

Floating-point status bit NV

**Instruction format:**

31	25	24	20	19	15	14	12	11	7	6	0
0010100		fs2		fs1		000		fd		1010011	

### 16.4.19 FMSUB.S: a single-precision floating-point multiply-subtract instruction

**Syntax:**

fmsub.s fd, fs1, fs2, fs3, rm

**Operation:**

fd ← fs1\*fs2 - fs3

**Permission:**

M mode/S mode/U mode

**Exception:**

Illegal instruction.

**Affected flag bits:**

Floating-point status bits NV, OF, UF, and IX

**Notes:**

RM determines the round-off mode:

- 3' b000: Rounds off to the nearest even number. The corresponding assembler instruction is `fmsub.s fd, fs1, fs2, fs3, rne`.
- 3' b001: Rounds off to zero. The corresponding assembler instruction is `fmsub.s fd, fs1, fs2, fs3, rtz`.
- 3' b010: Rounds off to negative infinity. The corresponding assembler instruction is `fmsub.s fd, fs1, fs2, fs3, rdn`.
- 3' b011: Rounds off to positive infinity. The corresponding assembler instruction is `fmsub.s fd, fs1, fs2, fs3, rup`.
- 3' b100: Rounds off to the nearest large value. The corresponding assembler instruction is `fmsub.s fd, fs1, fs2, fs3, rmm`.
- 3' b101: This code is reserved and not used.
- 3' b110: This code is reserved and not used.
- 3' b111: Dynamically rounds off based on the `rm` bit of the `fcsr` register. The corresponding assembler instruction is `fmsub.s fd, fs1, fs2, fs3`.

**Instruction format:**

31	27	26	25	24	20	19	15	14	12	11	7	6	0
fs3				00	fs2				fs1		rm	fd	1000111

**16.4.20 FMUL.S: a single-precision floating-point multiply instruction****Syntax:**

`fmul.s fd, fs1, fs2, rm`

**Operation:**

$fd \leftarrow fs1 * fs2$

**Permission:**

M mode/S mode/U mode

**Exception:**

Illegal instruction.

**Affected flag bits:**

Floating-point status bits NV, OF, UF, and NX

**Notes:**

RM determines the round-off mode:

- 3' b000: Rounds off to the nearest even number. The corresponding assembler instruction is `fmul.s fd, fs1, fs2, rne`.
- 3' b001: Rounds off to zero. The corresponding assembler instruction is `fmul.s fd, fs1, fs2, rtz`.
- 3' b010: Rounds off to negative infinity. The corresponding assembler instruction is `fmul.s fd, fs1, fs2, rdn`.
- 3' b011: Rounds off to positive infinity. The corresponding assembler instruction is `fmul.s fd, fs1, fs2, rup`.
- 3' b100: Rounds off to the nearest large value. The corresponding assembler instruction is `fmul.s fd, fs1, fs2, rmm`.
- 3' b101: This code is reserved and not used.
- 3' b110: This code is reserved and not used.
- 3' b111: Dynamically rounds off based on the `rm` bit of the `fcsr` register. The corresponding assembler instruction is `fmul.s fs1, fs2`.

**Instruction format:**

31	25	24	20	19	15	14	12	11	7	6	0		
0001000				fs2		fs1		rm		fd		1010011	

**16.4.21 FMV.W.X: a single-precision floating-point write move instruction****Syntax:**

`fmv.w.x fd, rs1`

**Operation:**

$fd[31:0] \leftarrow rs[31:0]$

$fd[63:32] \leftarrow 32' \text{ hfffffff}$

**Permission:**

M mode/S mode/U mode

**Exception:**

Illegal instruction.

**Affected flag bits:**

None

**Instruction format:**

31	25	24	20	19	15	14	12	11	7	6	0
1111000		00000		rs1		000		fd		1010011	

### 16.4.22 FMV.X.H: a single-precision floating-point read move instruction

**Syntax:**

fmv.x.w rd, fs1

**Operation:**

$\text{tmp}[31:0] \leftarrow \text{fs1}[31:0]$

$\text{rd} \leftarrow \text{sign\_extend}(\text{tmp}[31:0])$

**Permission:**

M mode/S mode/U mode

**Exception:**

Illegal instruction.

**Affected flag bits:**

None

**Instruction format:**

31	25	24	20	19	15	14	12	11	7	6	0
1110000		00000		fs1		000		rd		1010011	

### 16.4.23 FNMADD.S: a single-precision floating-point negate-(multiply-add) instruction

**Syntax:**

fnmadd.s fd, fs1, fs2, fs3, rm

**Operation:**

$\text{fd} \leftarrow -(\text{fs1} * \text{fs2} + \text{fs3})$

**Permission:**

M mode/S mode/U mode

**Exception:**

Illegal instruction.

**Affected flag bits:**

Floating-point status bits NV, OF, UF, and IX

**Notes:**

RM determines the round-off mode:

- 3' b000: Rounds off to the nearest even number. The corresponding assembler instruction is `fnmadd.s fd, fs1, fs2, fs3, rne`.
- 3' b001: Rounds off to zero. The corresponding assembler instruction is `fnmadd.s fd, fs1, fs2, fs3, rtz`.
- 3' b010: Rounds off to negative infinity. The corresponding assembler instruction is `fnmadd.s fd, fs1, fs2, fs3, rdn`.
- 3' b011: Rounds off to positive infinity. The corresponding assembler instruction is `fnmadd.s fd, fs1, fs2, fs3, rup`.
- 3' b100: Rounds off to the nearest large value. The corresponding assembler instruction is `fnmadd.s fd, fs1, fs2, fs3, rmm`.
- 3' b101: This code is reserved and not used.
- 3' b110: This code is reserved and not used.
- 3' b111: Dynamically rounds off based on the `rm` bit of the `fcsr` register. The corresponding assembler instruction is `fnmadd.s fd, fs1, fs2, fs3`.

**Instruction format:**

31	27	26	25	24	20	19	15	14	12	11	7	6	0
fs3				00	fs2				fs1		rm	fd	1001111

### 16.4.24 FNMSUB.S: a single-precision floating-point negate-(multiply-subtract) instruction

**Syntax:**

```
fnmsub.s fd, fs1, fs2, fs3, rm
```

**Operation:**

$$fd \leftarrow -(fs1 * fs2 - fs3)$$
**Permission:**

M mode/S mode/U mode

**Exception:**

Illegal instruction.

**Affected flag bits:**

Floating-point status bits NV, OF, UF, and IX

#### Notes:

RM determines the round-off mode:

- 3' b000: Rounds off to the nearest even number. The corresponding assembler instruction is `fnmsub.s fd, fs1, fs2, fs3, rne`.
- 3' b001: Rounds off to zero. The corresponding assembler instruction is `fnmsub.s fd, fs1, fs2, fs3, rtz`.
- 3' b010: Rounds off to negative infinity. The corresponding assembler instruction is `fnmsub.s fd, fs1, fs2, fs3, rdn`.
- 3' b011: Rounds off to positive infinity. The corresponding assembler instruction is `fnmsub.s fd, fs1, fs2, fs3, rup`.
- 3' b100: Rounds off to the nearest large value. The corresponding assembler instruction is `fnmsub.s fd, fs1, fs2, fs3, rmm`.
- 3' b101: This code is reserved and not used.
- 3' b110: This code is reserved and not used.
- 3' b111: Dynamically rounds off based on the `rm` bit of the `fcsr` register. The corresponding assembler instruction is `fnmsub.s fd, fs1, fs2, fs3`.

#### Instruction format:

31	27	26	25	24	20	19	15	14	12	11	7	6	0	
fs3				00	fs2				fs1		rm		fd	1001011

### 16.4.25 FSGNJ.S: a single-precision floating-point sign-injection instruction

#### Syntax:

`fsgnj.s fd, fs1, fs2`

#### Operation:

$fd[30:0] \leftarrow fs1[30:0]$

$fd[31] \leftarrow fs2[31]$

$fd[63:32] \leftarrow 32' \text{ hfffffff}$

#### Permission:

M mode/S mode/U mode

#### Exception:

Illegal instruction.

#### Affected flag bits:



None

**Instruction format:**

31	25	24	20	19	15	14	12	11	7	6	0
0010000		fs2		fs1		000		fd		1010011	

### 16.4.26 FSGNJN.S: a single-precision floating-point negate sign-injection instruction

**Syntax:**

fsgnfn.s fd, fs1, fs2

**Operation:**

$fd[30:0] \leftarrow fs1[30:0]$

$fd[31] \leftarrow ! fs2[31]$

$fd[63:32] \leftarrow 32' \text{ hfffffff}$

**Permission:**

M mode/S mode/U mode

**Exception:**

Illegal instruction.

**Affected flag bits:**

None

**Instruction format:**

31	25	24	20	19	15	14	12	11	7	6	0
0010000		fs2		fs1		001		fd		1010011	

### 16.4.27 FSGNJX.S: a single-precision floating-point XOR sign-injection instruction

**Syntax:**

fsgnjx.s fd, fs1, fs2

**Operation:**

$fd[30:0] \leftarrow fs1[30:0]$

$fd[31] \leftarrow fs1[31] \wedge fs2[31]$

$fd[63:32] \leftarrow 32' \text{ hfffffff}$

**Permission:**

M mode/S mode/U mode

**Exception:**

Illegal instruction.

**Affected flag bits:**

None

**Instruction format:**

31	25	24	20	19	15	14	12	11	7	6	0
0010000		fs2		fs1		010		fd		1010011	

### 16.4.28 FSQRT.S: a single-precision floating-point square-root instruction

**Syntax:**

fsqrt.s fd, fs1, rm

**Operation:**

$fd \leftarrow \text{sqrt}(fs1)$

**Permission:**

M mode/S mode/U mode

**Exception:**

Illegal instruction.

**Affected flag bits:**

Floating-point status bits NV and NX

**Notes:**

RM determines the round-off mode:

- 3' b000: Rounds off to the nearest even number. The corresponding assembler instruction is fsqrt.s fd, fs1, rne.
- 3' b001: Rounds off to zero. The corresponding assembler instruction is fsqrt.s fd, fs1, rtz.
- 3' b010: Rounds off to negative infinity. The corresponding assembler instruction is fsqrt.s fd, fs1, rdn.
- 3' b011: Rounds off to positive infinity. The corresponding assembler instruction is fsqrt.s fd, fs1, rup.
- 3' b100: Rounds off to the nearest large value. The corresponding assembler instruction is fsqrt.s fd, fs1, rmm.
- 3' b101: This code is reserved and not used.
- 3' b110: This code is reserved and not used.

- 3' b111: Dynamically rounds off based on the rm bit of the fcsr register. The corresponding assembler instruction is fsqrt.s fd, fs1.

**Instruction format:**

31	25	24	20	19	15	14	12	11	7	6	0
0101100	00000	fs1	rm	fd	1010011						

### 16.4.29 FSUB.S: a single-precision floating-point subtract instruction

**Syntax:**

fsub.s fd, fs1, fs2, rm

**Operation:**

$fd \leftarrow fs1 - fs2$

**Permission:**

M mode/S mode/U mode

**Exception:**

Illegal instruction.

**Affected flag bits:**

Floating-point status bits NV, OF, and NX

**Notes:**

RM determines the round-off mode:

- 3' b000: Rounds off to the nearest even number. The corresponding assembler instruction is fsub.f, fs1, fs2, rne.
- 3' b001: Rounds off to zero. The corresponding assembler instruction is fsub.s fd, fs1, fs2, rtz.
- 3' b010: Rounds off to negative infinity. The corresponding assembler instruction is fsub.s fd, fs1, fs2, rdn.
- 3' b011: Rounds off to positive infinity. The corresponding assembler instruction is fsub.s fd, fs1, fs2, rup.
- 3' b100: Rounds off to the nearest large value. The corresponding assembler instruction is fsub.s fd, fs1, fs2, rmm.
- 3' b101: This code is reserved and not used.
- 3' b110: This code is reserved and not used.
- 3' b111: Dynamically rounds off based on the rm bit of the fcsr register. The corresponding assembler instruction is fsub.s fd, fs1, fs2.

Instruction format:

31	25	24	20	19	15	14	12	11	7	6	0
0000100				fs2		fs1		rm		fd	
										1010011	

### 16.4.30 FSW: a single-precision floating-point store instruction

Syntax:

fs2 fs2, imm12(rs1)

Operation:

address  $\leftarrow$  rs1 + sign\_extend(imm12)

mem[(address+31):address]  $\leftarrow$  fs2[31:0]

Permission:

M mode/S mode/U mode

Exception:

Unaligned access, access error, page error, or illegal instruction.

Instruction format:

31	25	24	20	19	15	14	12	11	7	6	0
imm12[11:5]				fs2		rs1		010		imm12[4:0]	
										0100111	

## 16.5 Appendix A-5 D instructions

The following describes the RISC-V D instructions implemented by C910/C920. The instructions are 32 bits wide and sorted in alphabetic order.

When the fs bit in the mstatus register is 2' b00, running any instruction listed in this appendix will trigger an illegal instruction exception. When the fs bit in the mstatus register is not 2' b00, it is set to 2' b11 after any instruction listed in this appendix is executed.

### 16.5.1 FADD.D: a double-precision floating-point add instruction

Syntax:

fadd.d fd, fs1, fs2, rm

Operation:

fd  $\leftarrow$  fs1 + fs2

Permission:

M mode/S mode/U mode

**Exception:**

Illegal instruction.

**Affected flag bits:**

Floating-point status bits NV, OF, and NX

**Notes:**

RM determines the round-off mode:

- 3' b000: Rounds off to the nearest even number. The corresponding assembler instruction is `fadd.d fd, fs1, fs2, rne`.
- 3' b001: Rounds off to zero. The corresponding assembler instruction is `fadd.d fd, fs1, fs2, rtz`.
- 3' b010: Rounds off to negative infinity. The corresponding assembler instruction is `fadd.d fd, fs1, fs2, rdn`.
- 3' b011: Rounds off to positive infinity. The corresponding assembler instruction is `fadd.d fd, fs1, fs2, rup`.
- 3' b100: Rounds off to the nearest large value. The corresponding assembler instruction is `fadd.d fd, fs1, fs2, rmm`.
- 3' b101: This code is reserved and not used.
- 3' b110: This code is reserved and not used.
- 3' b111: Dynamically rounds off based on the `rm` bit of the `fcsr` register. The corresponding assembler instruction is `fadd.d fd, fs1, fs2`.

**Instruction format:**

31	25	24	20	19	15	14	12	11	7	6	0
0000001				fs2		fs1		rm	fd		1010011

### 16.5.2 FCLASS.D: a double-precision floating-point classify instruction

**Syntax:**

`fclass.d rd, fs1`

**Operation:** if ( `fs1` = -Inf)

$rd \leftarrow 64' \text{ h1}$

if ( `fs1` = -norm)

$rd \leftarrow 64' \text{ h2}$

if ( `fs1` = -subnorm)

```

rd ← 64' h4
if ( fs1 = -zero)
    fd ← 64' h8
if ( fs1 = +Zero)
    rd ← 64' h10
if ( fs1 = +subnorm)
    rd ← 64' h20
if ( fs1 = +norm)
    rd ← 64' h40
if ( fs1 = +Inf)
    rd ← 64' h80
if ( fs1 = sNaN)
    rd ← 64' h100
if ( fs1 = qNaN)
    rd ← 64' h200

```

**Permission:**

M mode/S mode/U mode

**Exception:**

Illegal instruction.

**Affected flag bits:**

None

**Instruction format:**

31	25	24	20	19	15	14	12	11	7	6	0
1110001	00000	fs1	001	rd	1010011						

### 16.5.3 FCVT.D.L: an instruction that converts a signed long integer into a double-precision floating-point number

**Syntax:**

```
fcvt.d.l fd, rs1, rm
```

**Operation:**

```
fd ← signed_long_convert_to_double(fs1)
```

**Permission:**

M mode/S mode/U mode

**Exception:**

Illegal instruction.

**Affected flag bits:**

Floating-point status bit NX

**Notes:**

RM determines the round-off mode:

- 3' b000: Rounds off to the nearest even number. The corresponding assembler instruction is `fcvt.d.l fd, rs1, rne`.
- 3' b001: Rounds off to zero. The corresponding assembler instruction is `fcvt.d.l fd, rs1, rtz`.
- 3' b010: Rounds off to negative infinity. The corresponding assembler instruction is `fcvt.d.l fd, rs1, rdn`.
- 3' b011: Rounds off to positive infinity. The corresponding assembler instruction is `fcvt.d.l fd, rs1, rup`.
- 3' b100: Rounds off to the nearest large value. The corresponding assembler instruction is `fcvt.d.l fd, rs1, rmm`.
- 3' b101: This code is reserved and not used.
- 3' b110: This code is reserved and not used.
- 3' b111: Dynamically rounds off based on the `rm` bit of the `fcsr` register. The corresponding assembler instruction is `fcvt.d.l fd, rs1`.

**Instruction format:**

31	25	24	20	19	15	14	12	11	7	6	0	
1101001				00010		rs1		rm		fd		1010011

### 16.5.4 FCVT.D.LU: an instruction that converts an unsigned long integer into a double-precision floating-point number

**Syntax:**`fcvt.d.lu fd, rs1, rm`**Operation:** $fd \leftarrow \text{unsigned\_long\_convert\_to\_double}(fs1)$ **Permission:**

M mode/S mode/U mode

**Exception:**

Illegal instruction.

**Affected flag bits:**

Floating-point status bit NX

**Notes:**

RM determines the round-off mode:

- 3' b000: Rounds off to the nearest even number. The corresponding assembler instruction is `fcvt.d.lu fd, rs1, rne`.
- 3' b001: Rounds off to zero. The corresponding assembler instruction is `fcvt.d.lu fd, rs1, rtz`.
- 3' b010: Rounds off to negative infinity. The corresponding assembler instruction is `fcvt.d.lu fd, rs1, rdn`.
- 3' b011: Rounds off to positive infinity. The corresponding assembler instruction is `fcvt.d.lu fd, rs1, rup`.
- 3' b100: Rounds off to the nearest large value. The corresponding assembler instruction is `fcvt.d.lu fd, rs1, rmm`.
- 3' b101: This code is reserved and not used.
- 3' b110: This code is reserved and not used.
- 3' b111: Dynamically rounds off based on the `rm` bit of the `fcsr` register. The corresponding assembler instruction is `fcvt.d.lu fd, rs1`.

**Instruction format:**

31	25	24	20	19	15	14	12	11	7	6	0
1101001	00011	rs1	rm	fd	1010011						

### 16.5.5 FCVT.D.S: an instruction that converts a single-precision floating-point number into a double-precision floating-point number

**Syntax:**

`fcvt.d.s fd, fs1`

**Operation:**

$fd \leftarrow \text{single\_convert\_to\_double}(fs1)$

**Permission:**

M mode/S mode/U mode



**Exception:**

Illegal instruction.

**Affected flag bits:**

None

**Instruction format:**

31	25	24	20	19	15	14	12	11	7	6	0
0100001	00000	fs1	000	fd	1010011						

### 16.5.6 FCVT.D.W: an instruction that converts a signed integer into a double-precision floating-point number

**Syntax:**

fcvt.d.w fd, rs1

**Operation:** $fd \leftarrow \text{signed\_int\_convert\_to\_double}(fs1)$ **Permission:**

M mode/S mode/U mode

**Exception:**

Illegal instruction.

**Affected flag bits:**

None

**Instruction format:**

31	25	24	20	19	15	14	12	11	7	6	0
1101001	00000	rs1	000	fd	1010011						

### 16.5.7 FCVT.D.WU: an instruction that converts an unsigned integer into a double-precision floating-point number

**Syntax:**

fcvt.d.wu fd, rs1

**Operation:** $fd \leftarrow \text{unsigned\_int\_convert\_to\_double}(fs1)$ **Permission:**

M mode/S mode/U mode

**Exception:**

Illegal instruction.

**Affected flag bits:**

None

**Instruction format:**

31	25	24	20	19	15	14	12	11	7	6	0
1101001		00001		rs1		000		fd		1010011	

### 16.5.8 FCVT.L.D: an instruction that converts a double-precision floating-point number into a signed long integer

**Syntax:**

`fcvt.l.d rd, fs1, rm`

**Operation:**

$rd \leftarrow \text{double\_convert\_to\_signed\_long}(fs1)$

**Permission:**

M mode/S mode/U mode

**Exception:**

Illegal instruction.

**Affected flag bits:**

Floating-point status bits NV and NX

**Notes:**

RM determines the round-off mode:

- 3' b000: Rounds off to the nearest even number. The corresponding assembler instruction is `fcvt.l.d rd, fs1, rne`.
- 3' b001: Rounds off to zero. The corresponding assembler instruction is `fcvt.l.d rd, fs1, rtz`.
- 3' b010: Rounds off to negative infinity. The corresponding assembler instruction is `fcvt.l.d rd, fs1, rdn`.
- 3' b011: Rounds off to positive infinity. The corresponding assembler instruction is `fcvt.l.d rd, fs1, rup`.
- 3' b100: Rounds off to the nearest large value. The corresponding assembler instruction is `fcvt.l.d rd, fs1, rmm`.

- 3' b101: This code is reserved and not used.
- 3' b110: This code is reserved and not used.
- 3' b111: Dynamically rounds off based on the rm bit of the fcsr register. The corresponding assembler instruction is `fcvt.l.d rd, fs1`.

**Instruction format:**

31	25	24	20	19	15	14	12	11	7	6	0
1100001		00010		fs1		rm		rd		1010011	

### 16.5.9 FCVT.LU.D: an instruction that converts a double-precision floating-point number into an unsigned long integer

**Syntax:**

`fcvt.lu.d rd, fs1, rm`

**Operation:**

`rd ← double_convert_to_unsigned_long(fs1)`

**Permission:**

M mode/S mode/U mode

**Exception:**

Illegal instruction.

**Affected flag bits:**

Floating-point status bits NV and NX

**Notes:**

RM determines the round-off mode:

- 3' b000: Rounds off to the nearest even number. The corresponding assembler instruction is `fcvt.lu.d rd, fs1, rne`.
- 3' b001: Rounds off to zero. The corresponding assembler instruction is `fcvt.lu.d rd, fs1, rtz`.
- 3' b010: Rounds off to negative infinity. The corresponding assembler instruction is `fcvt.lu.d rd, fs1, rdn`.
- 3' b011: Rounds off to positive infinity. The corresponding assembler instruction is `fcvt.lu.d rd, fs1, rup`.
- 3' b100: Rounds off to the nearest large value. The corresponding assembler instruction is `fcvt.lu.d rd, fs1, rmm`.
- 3' b101: This code is reserved and not used.

- 3' b110: This code is reserved and not used.
- 3' b111: Dynamically rounds off based on the rm bit of the fcsr register. The corresponding assembler instruction is `fcvt.lu.d rd, fs1`.

**Instruction format:**

31	25	24	20	19	15	14	12	11	7	6	0		
1100001				00011		fs1		rm		rd		1010011	

### 16.5.10 FCVT.S.D: an instruction that converts a double-precision floating-point number into a single-precision floating-point number

**Syntax:**

```
fcvt.s.d fd, fs1, rm
```

**Operation:**

```
fd ← double_convert_to_single(fs1)
```

**Permission:**

M mode/S mode/U mode

**Exception:**

Illegal instruction.

**Affected flag bits:**

Floating-point status bits NV, OF, UF, and NX

**Notes:**

RM determines the round-off mode:

- 3' b000: Rounds off to the nearest even number. The corresponding assembler instruction is `fcvt.s.d fd, fs1, rne`.
- 3' b001: Rounds off to zero. The corresponding assembler instruction is `fcvt.s.d fd, fs1, rtz`.
- 3' b010: Rounds off to negative infinity. The corresponding assembler instruction is `fcvt.s.d fd, fs1, rdn`.
- 3' b011: Rounds off to positive infinity. The corresponding assembler instruction is `fcvt.s.d fd, fs1, rup`.
- 3' b100: Rounds off to the nearest large value. The corresponding assembler instruction is `fcvt.s.d fd, fs1, rmm`.
- 3' b101: This code is reserved and not used.
- 3' b110: This code is reserved and not used.

- 3' b111: Dynamically rounds off based on the rm bit of the fcsr register. The corresponding assembler instruction is `fcvt.s.d fd, fs1`.

**Instruction format:**

31	25	24	20	19	15	14	12	11	7	6	0
0100000	00001	fs1	rm	fd	1010011						

### 16.5.11 FCVT.W.D: an instruction that converts a double-precision floating-point number into a signed integer

**Syntax:**

`fcvt.w.d rd, fs1, rm`

**Operation:**

`tmp ← double_convert_to_signed_int(fs1)`

`rd ← sign_extend(tmp)`

**Permission:**

M mode/S mode/U mode

**Exception:**

Illegal instruction.

**Affected flag bits:**

Floating-point status bits NV and NX

**Notes:**

RM determines the round-off mode:

- 3' b000: Rounds off to the nearest even number. The corresponding assembler instruction is `fcvt.w.d rd, fs1, rne`.
- 3' b001: Rounds off to zero. The corresponding assembler instruction is `fcvt.w.d rd, fs1, rtz`.
- 3' b010: Rounds off to negative infinity. The corresponding assembler instruction is `fcvt.w.d rd, fs1, rdn`.
- 3' b011: Rounds off to positive infinity. The corresponding assembler instruction is `fcvt.w.d rd, fs1, rup`.
- 3' b100: Rounds off to the nearest large value. The corresponding assembler instruction is `fcvt.w.d rd, fs1, rmm`.
- 3' b101: This code is reserved and not used.
- 3' b110: This code is reserved and not used.

- 3' b111: Dynamically rounds off based on the rm bit of the fcsr register. The corresponding assembler instruction is `fcvt.w.d rd, fs1`.

**Instruction format:**

31	25	24	20	19	15	14	12	11	7	6	0
1100001	00000	fs1	rm	rd	1010011						

### 16.5.12 FCVT.WU.D: an instruction that converts a double-precision floating-point number into an unsigned integer

**Syntax:**

`fcvt.wu.d rd, fs1, rm`

**Operation:**

`tmp ← double_convert_to_unsigned_int(fs1)`

`rd ← sign_extend(tmp)`

**Permission:**

M mode/S mode/U mode

**Exception:**

Illegal instruction.

**Affected flag bits:**

Floating-point status bits NV and NX

**Notes:**

RM determines the round-off mode:

- 3' b000: Rounds off to the nearest even number. The corresponding assembler instruction is `fcvt.wu.d rd, fs1, rne`.
- 3' b001: Rounds off to zero. The corresponding assembler instruction is `fcvt.wu.d rd, fs1, rtz`.
- 3' b010: Rounds off to negative infinity. The corresponding assembler instruction is `fcvt.wu.d rd, fs1, rdn`.
- 3' b011: Rounds off to positive infinity. The corresponding assembler instruction is `fcvt.wu.d rd, fs1, rup`.
- 3' b100: Rounds off to the nearest large value. The corresponding assembler instruction is `fcvt.wu.d rd, fs1, rmm`.
- 3' b101: This code is reserved and not used.
- 3' b110: This code is reserved and not used.

- 3' b111: Dynamically rounds off based on the rm bit of the fcsr register. The corresponding assembler instruction is `fcvt.wu.d rd, fs1`.

**Instruction format:**

31	25	24	20	19	15	14	12	11	7	6	0
1100001	00001	fs1	rm	rd	1010011						

### 16.5.13 FDIV.D: a double-precision floating-point divide instruction

**Syntax:**

`fdiv.d fd, fs1, fs2, rm`

**Operation:**

$fd \leftarrow fs1 / fs2$

**Permission:**

M mode/S mode/U mode

**Exception:**

Illegal instruction.

**Affected flag bits:**

Floating-point status bits NV, DZ, OF, UF, and NX

**Notes:**

RM determines the round-off mode:

- 3' b000: Rounds off to the nearest even number. The corresponding assembler instruction is `fdiv.d fd, fs1, fs2, rne`.
- 3' b001: Rounds off to zero. The corresponding assembler instruction is `fdiv.d fd fs1, fs2, rtz`.
- 3' b010: Rounds off to negative infinity. The corresponding assembler instruction is `fdiv.d fd, fs1, fs2, rdn`.
- 3' b011: Rounds off to positive infinity. The corresponding assembler instruction is `fdiv.d fd, fs1, fs2, rup`.
- 3' b100: Rounds off to the nearest large value. The corresponding assembler instruction is `fdiv.d fd, fs1, fs2, rmm`.
- 3' b101: This code is reserved and not used.
- 3' b110: This code is reserved and not used.
- 3' b111: Dynamically rounds off based on the rm bit of the fcsr register. The corresponding assembler instruction is `fdiv.d fd, fs1, fs2`.

Instruction format:

31	25	24	20	19	15	14	12	11	7	6	0
0001101			fs2		fs1		rm		fd		1010011

### 16.5.14 FEQ.D: a double-precision floating-point compare equal instruction

Syntax:

feq.d rd, fs1, fs2

Operation:

if(fs1 == fs2)

rd  $\leftarrow$  1

else

rd  $\leftarrow$  0

Permission:

M mode/S mode/U mode

Exception:

Illegal instruction.

Affected flag bits:

Floating-point status bit NV

Instruction format:

31	25	24	20	19	15	14	12	11	7	6	0
1010001			fs2		fs1		010		rd		1010011

### 16.5.15 FLD: a double-precision floating-point load instruction

Syntax:

fld fd, imm12(rs1)

Operation:

address  $\leftarrow$  rs1 + sign\_extend(imm12)

fd[63:0]  $\leftarrow$  mem[(address+7):address]

Permission:

M mode/S mode/U mode



**Exception:**

Unaligned access, access error, page error, or illegal instruction.

**Affected flag bits:**

None

**Instruction format:**

31	20	19	15	14	12	11	7	6	0		
imm12[11:0]				rs1		011		fd		0000111	

### 16.5.16 FLE.D: a double-precision floating-point compare less than or equal to instruction

**Syntax:**

fle.d rd, fs1, fs2

**Operation:**

if(fs1 <= fs2)

rd ← 1

else

rd ← 0

**Permission:**

M mode/S mode/U mode

**Exception:**

Illegal instruction.

**Affected flag bits:**

Floating-point status bit NV

**Instruction format:**

31	25	24	20	19	15	14	12	11	7	6	0
1010001		fs2		fs1		000		rd		1010011	

### 16.5.17 FLT.D: a double-precision floating-point compare less than instruction

**Syntax:**

flt.d rd, fs1, fs2

**Operation:**

```

if(fs1 < fs2)
    rd ← 1
else
    rd ← 0

```

**Permission:**

M mode/S mode/U mode

**Exception:**

Illegal instruction.

**Affected flag bits:**

Floating-point status bit NV

**Instruction format:**

31	25	24	20	19	15	14	12	11	7	6	0
1010001				fs2	fs1		001	rd		1010011	

### 16.5.18 FMADD.D: a double-precision floating-point multiply-add instruction

**Syntax:**

```
fmadd.d fd, fs1, fs2, fs3, rm
```

**Operation:**

$$fd \leftarrow fs1 * fs2 + fs3$$
**Permission:**

M mode/S mode/U mode

**Exception:**

Illegal instruction.

**Affected flag bits:**

Floating-point status bits NV, OF, UF, and IX

**Notes:**

RM determines the round-off mode:

- 3' b000: Rounds off to the nearest even number. The corresponding assembler instruction is `fmadd.d fd, fs1, fs2, fs3, rne`.
- 3' b001: Rounds off to zero. The corresponding assembler instruction is `fmadd.d fd, fs1, fs2, fs3, rtz`.

- 3' b010: Rounds off to negative infinity. The corresponding assembler instruction is `fmadd.d fd, fs1, fs2, fs3, rdn`.
- 3' b011: Rounds off to positive infinity. The corresponding assembler instruction is `fmadd.d fd, fs1, fs2, fs3, rdn`.
- 3' b100: Rounds off to the nearest large value. The corresponding assembler instruction is `fmadd.d fd, fs1, fs2, fs3, rmm`.
- 3' b101: This code is reserved and not used.
- 3' b110: This code is reserved and not used.
- 3' b111: Dynamically rounds off based on the `rm` bit of the `fcsr` register. The corresponding assembler instruction is `fmadd.d fd, fs1, fs2, fs3`.

**Instruction format:**

31	27	26	25	24	20	19	15	14	12	11	7	6	0			
fs3				01	fs2				fs1				rm	fd	1000011	

**16.5.19 FMAX.D: a double-precision floating-point MAX instruction****Syntax:**

`fmax.d fd, fs1, fs2`

**Operation:**

`if(fs1 >= fs2)`

`fd ← fs1`

`else`

`fd ← fs2`

**Permission:**

M mode/S mode/U mode

**Exception:**

Illegal instruction.

**Affected flag bits:**

Floating-point status bit NV

**Instruction format:**

31	25	24	20	19	15	14	12	11	7	6	0
0010101				fs2		fs1		001	fd		1010011

### 16.5.20 FMIN.D: a double-precision floating-point MIN instruction

**Syntax:**

```
fmin.d fd, fs1, fs2
```

**Operation:**

```
if(fs1 >= fs2)
```

```
    fd ← fs2
```

```
else
```

```
    fd ← fs1
```

**Permission:**

M mode/S mode/U mode

**Exception:**

Illegal instruction.

**Affected flag bits:**

Floating-point status bit NV

**Instruction format:**

31	25	24	20	19	15	14	12	11	7	6	0
0010101		fs2		fs1		000		fd		1010011	

### 16.5.21 FMSUB.D: a double-precision floating-point multiply-subtract instruction

**Syntax:**

```
fmsub.d fd, fs1, fs2, fs3, rm
```

**Operation:**

```
fd ← fs1*fs2 - fs3
```

**Permission:**

M mode/S mode/U mode

**Exception:**

Illegal instruction.

**Affected flag bits:**

Floating-point status bits NV, OF, UF, and IX

**Notes:**

RM determines the round-off mode:

- 3' b000: Rounds off to the nearest even number. The corresponding assembler instruction is `fmsub.d fd, fs1, fs2, fs3, rne`.
- 3' b001: Rounds off to zero. The corresponding assembler instruction is `fmsub.d fd, fs1, fs2, fs3, rtz`.
- 3' b010: Rounds off to negative infinity. The corresponding assembler instruction is `fmsub.d fd, fs1, fs2, fs3, rdn`.
- 3' b011: Rounds off to positive infinity. The corresponding assembler instruction is `fmsub.d fd, fs1, fs2, fs3, rup`.
- 3' b100: Rounds off to the nearest large value. The corresponding assembler instruction is `fmsub.d fd, fs1, fs2, fs3, rmm`.
- 3' b101: This code is reserved and not used.
- 3' b110: This code is reserved and not used.
- 3' b111: Dynamically rounds off based on the `rm` bit of the `fcsr` register. The corresponding assembler instruction is `fmsub.d fd, fs1, fs2, fs3`.

**Instruction format:**

31	27	26	25	24	20	19	15	14	12	11	7	6	0
fs3				01	fs2			fs1		rm	fd		1000111

### 16.5.22 FMUL.D: a double-precision floating-point multiply instruction

**Syntax:**

`fmul.d fd, fs1, fs2, rm`

**Operation:**

$fd \leftarrow fs1 * fs2$

**Permission:**

M mode/S mode/U mode

**Exception:**

Illegal instruction.

**Affected flag bits:**

Floating-point status bits NV, OF, UF, and NX

**Notes:**

RM determines the round-off mode:

- 3' b000: Rounds off to the nearest even number. The corresponding assembler instruction is `fmul.d fd, fs1, fs2, rne`.
- 3' b001: Rounds off to zero. The corresponding assembler instruction is `fmul.d fd, fs1, fs2, rtz`.
- 3' b010: Rounds off to negative infinity. The corresponding assembler instruction is `fmul.d fd, fs1, fs2, rdn`.
- 3' b011: Rounds off to positive infinity. The corresponding assembler instruction is `fmul.d fd, fs1, fs2, rup`.
- 3' b100: Rounds off to the nearest large value. The corresponding assembler instruction is `fmul.d fd, fs1, fs2, rmm`.
- 3' b101: This code is reserved and not used.
- 3' b110: This code is reserved and not used.
- 3' b111: Dynamically rounds off based on the `rm` bit of the `fcsr` register. The corresponding assembler instruction is `fmul. fd, fs1, fs2`.

**Instruction format:**

31	25	24	20	19	15	14	12	11	7	6	0
0001001				fs2		fs1		rm	fd		1010011

**16.5.23 FMV.D.X: a double-precision floating-point write move instruction****Syntax:**

```
fmv.d.x fd, rs1
```

**Operation:**

$$fd \leftarrow rs1$$
**Permission:**

M mode/S mode/U mode

**Exception:**

Illegal instruction.

**Affected flag bits:**

None

**Notes:**

This instruction moves data from an integer register to a floating-point register.

**Instruction format:**

31	25	24	20	19	15	14	12	11	7	6	0
1111001	00000	rs1	000	fd	1010011						

### 16.5.24 FMV.X.D: a double-precision floating-point read move instruction

**Syntax:**

fmv.x.d rd, fs1

**Operation:**

$rd \leftarrow fs1$

**Permission:**

M mode/S mode/U mode

**Exception:**

Illegal instruction.

**Affected flag bits:**

None

**Notes:**

This instruction moves data from a floating-point register to an integer register.

**Instruction format:**

31	25	24	20	19	15	14	12	11	7	6	0
1110001	00000	fs1	000	rd	1010011						

### 16.5.25 FNMADD.D: a double-precision floating-point negate-(multiply-add) instruction

**Syntax:**

fnmadd.d fd, fs1, fs2, fs3, rm

**Operation:**

$fd \leftarrow -(fs1 * fs2 + fs3)$

**Permission:**

M mode/S mode/U mode

**Exception:**

Illegal instruction.

**Affected flag bits:**

Floating-point status bits NV, OF, UF, and IX

#### Notes:

RM determines the round-off mode:

- 3' b000: Rounds off to the nearest even number. The corresponding assembler instruction is `fnmadd.d fd, fs1, fs2, fs3, rne`.
- 3' b001: Rounds off to zero. The corresponding assembler instruction is `fnmadd.d fd, fs1, fs2, fs3, rtz`.
- 3' b010: Rounds off to negative infinity. The corresponding assembler instruction is `fnmadd.d fd, fs1, fs2, fs3, rdn`.
- 3' b011: Rounds off to positive infinity. The corresponding assembler instruction is `fnmadd.d fd, fs1, fs2, fs3, rup`.
- 3' b100: Rounds off to the nearest large value. The corresponding assembler instruction is `fnmadd.d fd, fs1, fs2, fs3, rmm`.
- 3' b101: This code is reserved and not used.
- 3' b110: This code is reserved and not used.
- 3' b111: Dynamically rounds off based on the `rm` bit of the `fcsr` register. The corresponding assembler instruction is `fnmadd.d fd, fs1, fs2, fs3`.

#### Instruction format:

31	27	26	25	24	20	19	15	14	12	11	7	6	0
fs3				01	fs2				fs1	rm	fd	1001111	

### 16.5.26 FNMSUB.D: a double-precision floating-point negate-(multiply-subtract) instruction

#### Syntax:

`fnmsub.d fd, fs1, fs2, fs3, rm`

#### Operation:

$fd \leftarrow -(fs1 * fs2 - fs3)$

#### Permission:

M mode/S mode/U mode

#### Exception:

Illegal instruction.

#### Affected flag bits:

Floating-point status bits NV, OF, UF, and IX



**Notes:**

RM determines the round-off mode:

- 3' b000: Rounds off to the nearest even number. The corresponding assembler instruction is `fnmsub.d fd, fs1, fs2, fs3, rne`.
- 3' b001: Rounds off to zero. The corresponding assembler instruction is `fnmsub.d fd, fs1, fs2, fs3, rtz`.
- 3' b010: Rounds off to negative infinity. The corresponding assembler instruction is `fnmsub.d fd, fs1, fs2, fs3, rdn`.
- 3' b011: Rounds off to positive infinity. The corresponding assembler instruction is `fnmsub.d fd, fs1, fs2, fs3, rup`.
- 3' b100: Rounds off to the nearest large value. The corresponding assembler instruction is `fnmsub.d fd, fs1, fs2, fs3, rmm`.
- 3' b101: This code is reserved and not used.
- 3' b110: This code is reserved and not used.
- 3' b111: Dynamically rounds off based on the `rm` bit of the `fcsr` register. The corresponding assembler instruction is `fnmsub.d fd, fs1, fs2, fs3`.

**Instruction format:**

31	27	26	25	24	20	19	15	14	12	11	7	6	0	
fs3				01	fs2				fs1		rm		fd	1001011

**16.5.27 FSD: a double-precision floating-point store instruction****Syntax:**

`fsd fs2, imm12(rs1)`

**Operation:**

$\text{address} \leftarrow \text{rs1} + \text{sign\_extend}(\text{imm12})$

$\text{mem}[(\text{address}+63):\text{address}] \leftarrow \text{fs2}[63:0]$

**Permission:**

M mode/S mode/U mode

**Exception:**

Unaligned access, access error, page error, or illegal instruction.

**Instruction format:**

31	25	24	20	19	15	14	12	11	7	6	0
imm12[11:5]				fs2		rs1		011	imm12[4:0]		0100111

### 16.5.28 FSGNJ.D: a double-precision floating-point sign-injection instruction

**Syntax:**

```
fsgnj.d fd, fs1, fs2
```

**Operation:**

$$fd[62:0] \leftarrow fs1[62:0]$$

$$fd[63] \leftarrow fs2[63]$$
**Permission:**

M mode/S mode/U mode

**Exception:**

Illegal instruction.

**Affected flag bits:**

None

**Instruction format:**

31	25	24	20	19	15	14	12	11	7	6	0
0010001	fs2			fs1			000	fd		1010011	

### 16.5.29 FSGNJN.D: a double-precision floating-point negate sign-injection instruction

**Syntax:**

```
fsgnjn.d fd, fs1, fs2
```

**Operation:**

$$fd[62:0] \leftarrow fs1[62:0]$$

$$fd[63] \leftarrow !fs2[63]$$
**Permission:**

M mode/S mode/U mode

**Exception:**

Illegal instruction.

**Affected flag bits:**

None

**Instruction format:**

31	25	24	20	19	15	14	12	11	7	6	0
0010001	fs2			fs1			001	fd		1010011	

### 16.5.30 FSGNJX.D: a double-precision floating-point XOR sign-injection instruction

**Syntax:**

$$\text{fsgnjx.d fd, fs1, fs2}$$
**Operation:**

$$\text{fd}[62:0] \leftarrow \text{fs1}[62:0]$$

$$\text{fd}[63] \leftarrow \text{fs1}[63] \wedge \text{fs2}[63]$$
**Permission:**

M mode/S mode/U mode

**Exception:**

Illegal instruction.

**Affected flag bits:**

None

**Instruction format:**

31	25	24	20	19	15	14	12	11	7	6	0
0010001	fs2			fs1			010	fd		1010011	

### 16.5.31 FSQRT.D: a double-precision floating-point square-root instruction

**Syntax:**

$$\text{fsqrt.d fd, fs1, rm}$$
**Operation:**

$$\text{fd} \leftarrow \text{sqrt}(\text{fs1})$$
**Permission:**

M mode/S mode/U mode

**Exception:**

Illegal instruction.

**Affected flag bits:**

Floating-point status bits NV and NX

**Notes:**

RM determines the round-off mode:

- 3' b000: Rounds off to the nearest even number. The corresponding assembler instruction is `fsqrt.d fd, fs1, rne`.
- 3' b001: Rounds off to zero. The corresponding assembler instruction is `fsqrt.d fd, fs1, rtz`.
- 3' b010: Rounds off to negative infinity. The corresponding assembler instruction is `fsqrt.d fd, fs1, rdn`.
- 3' b011: Rounds off to positive infinity. The corresponding assembler instruction is `fsqrt.d fd, fs1, rup`.
- 3' b100: Rounds off to the nearest large value. The corresponding assembler instruction is `fsqrt.d fd, fs1, rmm`.
- 3' b101: This code is reserved and not used.
- 3' b110: This code is reserved and not used.
- 3' b111: Dynamically rounds off based on the `rm` bit of the `fcsr` register. The corresponding assembler instruction is `fsqrt.d fd, fs1`.

**Instruction format:**

31	25	24	20	19	15	14	12	11	7	6	0
0101101	00000	fs1	rm	fd	1010011						

### 16.5.32 FSUB.D: a double-precision floating-point subtract instruction

**Syntax:**

`fsub.d fd, fs1, fs2, rm`

**Operation:**

$fd \leftarrow fs1 - fs2$

**Permission:**

M mode/S mode/U mode

**Exception:**

Illegal instruction.

**Affected flag bits:**

Floating-point status bits NV, OF, and NX

**Notes:**

RM determines the round-off mode:

- 3' b000: Rounds off to the nearest even number. The corresponding assembler instruction is `fsub.d fd, fs1, fs2, rne`.
- 3' b001: Rounds off to zero. The corresponding assembler instruction is `fsub.d fd, fs1, fs2, rtz`.

- 3' b010: Rounds off to negative infinity. The corresponding assembler instruction is `fsub.d fd, fs1, fs2, rdn`.
- 3' b011: Rounds off to positive infinity. The corresponding assembler instruction is `fsub.d fd, fs1, fs2, rup`.
- 3' b100: Rounds off to the nearest large value. The corresponding assembler instruction is `fsub.d fd, fs1, fs2, rmm`.
- 3' b101: This code is reserved and not used.
- 3' b110: This code is reserved and not used.
- 3' b111: Dynamically rounds off based on the `rm` bit of the `fcsr` register. The corresponding assembler instruction is `fsub.dfd, fs1, fs2`.

**Instruction format:**

31	25	24	20	19	15	14	12	11	7	6	0
0000101		fs2		fs1		rm		fd		1010011	

## 16.6 Appendix A-6 C Instructions

This section describes RISC-V C instructions implemented by C910/C920. The instructions are 16 bits wide and sorted in alphabetic order.

### 16.6.1 C.ADD: a signed add instruction

**Syntax:**

`c.add rd, rs2`

**Operation:**

$rd \leftarrow rs1 + rs2$

**Permission:**

M mode/S mode/U mode

**Exception:**

None

**Notes:**

$rs1 = rd \neq 0$

$rs2 \neq 0$

**Instruction format:**

15	13	12	11	7	6	2	1	0
100	1	rs1/rd			rs2		10	

### 16.6.2 C.ADDI: a signed add immediate instruction

**Syntax:**

c.addi rd, nzimm6

**Operation:**

$rd \leftarrow rs1 + \text{sign\_extend}(nzimm6)$

**Permission:**

M mode/S mode/U mode

**Exception:**

None

**Notes:**

$rs1 = rd \neq 0$

$nzimm6 \neq 0$

**Instruction format:**

15	13	12	11	7	6	2	1	0
000		rs1/rd			nzimm6[4:0]		01	

└─── nzimm6[5]

### 16.6.3 C.ADDIW: an add immediate instruction that operates on the lower 32 bits

**Syntax:**

c.addiw rd, imm6

**Operation:**

$\text{tmp}[31:0] \leftarrow rs1[31:0] + \text{sign\_extend}(\text{imm6})$

$rd \leftarrow \text{sign\_extend}(\text{tmp}[31:0])$


**Permission:**

M mode/S mode/U mode

**Exception:**

None

**Notes:**

$$rs1 = rd \neq 0$$
**Instruction format:**

### 16.6.4 C.ADDI4SPN: an instruction that adds an immediate scaled by 4 to the stack pointer

**Syntax:**

$$c.addi4spn\ rd,\ sp,\ nzuimm8 \ll 2$$
**Operation:**

$$rd \leftarrow sp + \text{zero\_extend}(nzuimm8 \ll 2)$$
**Permission:**

M mode/S mode/U mode

**Exception:**


None

**Notes:**

$$nzuimm8 \neq 0$$

Typical rd code registers are:

- 000 x8
- 001 x9
- 010 x10
- 011 x11
- 100 x12
- 101 x13
- 110 x14
- 111 x15

**Instruction format:**

### 16.6.5 C.ADDI16SP: an instruction that adds an immediate scaled by 16 to the stack pointer

**Syntax:**

```
c.addi16sp sp, nzuimm6<<4
```


**Operation:**

$$sp \leftarrow sp + \text{sign\_extend}(nzuimm6 \ll 4)$$
**Permission:**

M mode/S mode/U mode

**Exception:**

None

**Instruction format:**


### 16.6.6 C.ADDW: a signed add instruction that operates on the lower 32 bits

**Syntax:**

```
c.addw rd, rs2
```

**Operation:**

$$\text{tmp}[31:0] \leftarrow \text{rs1}[31:0] + \text{rs2}[31:0]$$

$$\text{rd} \leftarrow \text{sign\_extend}(\text{tmp}[31:0])$$
**Permission:**

M mode/S mode/U mode

**Exception:**

None

**Notes:**

rs1 = rd

Typical rd/rs1 and rs2 code registers are:

- 000: x8
- 001: x9



- 010: x10
- 011: x11
- 100: x12
- 101: x13
- 110: x14
- 111: x15

**Instruction format:**

15	13	12	11	10	9	7	6	5	4	2	1	0
100	1	11	rs1/rd	01	rs2	01						

### 16.6.7 C.AND: a bitwise AND instruction

**Syntax:**

c.and rd, rs2

**Operation:**

$rd \leftarrow rs1 \& rs2$

**Permission:**

M mode/S mode/U mode

**Exception:**

None

**Notes:**

$rs1 = rd$

Typical rd/rs1 and rs2 code registers are:

- 000: x8
- 001: x9
- 010: x10
- 011: x11
- 100: x12
- 101: x13
- 110: x14
- 111: x15

Instruction format:

15	13	12	11	10	9	7	6	5	4	2	1	0
100	0	11	rs1/rd	11	rs2	01						

### 16.6.8 C.ANDI: an immediate bitwise AND instruction

Syntax:

c.andi rd, imm6

Operation:

$rd \leftarrow rs1 \& \text{sign\_extend}(\text{imm6})$

Permission:

M mode/S mode/U mode

Exception:

None

Notes:

rs1 = rd

Typical rd/rs1 code registers are:

- 000: x8
- 001: x9
- 010: x10
- 011: x11
- 100: x12
- 101: x13
- 110: x14
- 111: x15

Instruction format:

15	13	12	11	10	9	7	6	5	4	3	2	1	0
100			10	rs1/rd									01

└── imm6[5]

### 16.6.9 C.BEQZ: a branch-if-equal-to-zero instruction

#### Syntax:

c.beqz rs1, label

#### Operation:

if (rs1 == 0)

next pc = current pc + imm8<<1;

else

next pc = current pc + 2;

#### Permission:

M mode/S mode/U mode

#### Exception:

None

#### Notes:


Typical rs1 code registers are:

- 000: x8
- 001: x9
- 010: x10
- 011: x11
- 100: x12
- 101: x13
- 110: x14
- 111: x15

The compiler calculates immediate 8 based on the label.

The jump range of the instruction is  $\pm 256$  B address space.

#### Instruction format:



### 16.6.10 C.BNEZ: a branch-if-not-equal-to-zero instruction

#### Syntax:

c.bnez rs1, label

#### Operation:

if (rs1 != 0)

next pc = current pc + imm8<<1;

else

next pc = current pc + 2;

#### Permission:

M mode/S mode/U mode

#### Exception:

None

#### Notes:


Typical rs1 code registers are:

- 000: x8
- 001: x9
- 010: x10
- 011: x11
- 100: x12
- 101: x13
- 110: x14
- 111: x15

The compiler calculates immediate 12 based on the label.

The jump range of the instruction is  $\pm 256$  B address space.

#### Instruction format:



### 16.6.11 C.EBREAK: a break instruction

**Syntax:**

c.ebreak

**Operation:**

Generates breakpoint exceptions or enables the core to enter the debug mode.

**Permission:**

M mode/S mode/U mode

**Exception:**

Breakpoint exceptions

**Instruction format:**

15	13	12	11	7	6	2	1	0
100	1	00000	00000	10				

### 16.6.12 C.FLD: a floating-point load doubleword instruction

**Syntax:**

c.fld fd, uimm5<<3(rs1)

**Operation:**

address  $\leftarrow$  rs1 + zero\_extend(uimm5<<3)

fd  $\leftarrow$  mem[address+7:address]

**Permission:**

M mode/S mode/U mode

**Exception:**

Unaligned access exceptions, access error exceptions, and page error exceptions on load instructions

**Notes:**

Typical rs1 code registers are:


- 000: x8
- 001: x9
- 010: x10
- 011: x11
- 100: x12

- 101: x13
- 110: x14
- 111: x15

Typical fd code registers are:

- 000: f8
- 001: f9
- 010: f10
- 011: f11
- 100: f12
- 101: f13
- 110: f14
- 111: f15

**Instruction format:**



### 16.6.13 C.FLDSP: a floating-point doubleword load stack instruction

**Syntax:**

`c.fldsp fd, uimm6<<3(sp)`

**Operation:**

$\text{address} \leftarrow \text{sp} + \text{zero\_extend}(\text{uimm6} \ll 3)$

$\text{fd} \leftarrow \text{mem}[\text{address}+7:\text{address}]$


**Permission:**

M mode/S mode/U mode

**Exception:**

Unaligned access exceptions, access error exceptions, and page error exceptions on load instructions

**Instruction format:**



### 16.6.14 C.FSD: a floating-point store doubleword instruction

#### Syntax:

c.fsd fs2, uimm5<<3(rs1)

#### Operation:

address  $\leftarrow$  rs1 + zero\_extend(uimm5<<3)

mem[address+7:address]  $\leftarrow$  fs2

#### Permission:

M mode/S mode/U mode

#### Exception:

Unaligned access exceptions, access error exceptions, and page error exceptions on store instructions

#### Notes:

Typical fs1 code registers are:


- 000: x8
- 001: x9
- 010: x10
- 011: x11
- 100: x12
- 101: x13
- 110: x14
- 111: x15

Typical rs2 code registers are:

- 000: f8
- 001: f9
- 010: f10
- 011: f11
- 100: f12

- 101: f13
- 110: f14
- 111: f15

**Instruction format:**



### 16.6.15 C.FSDSP: a floating-point store doubleword stack pointer instruction

**Syntax:**

`c.fsdsp fs2, uimm6<<3(sp)`

**Operation:**

$\text{address} \leftarrow \text{sp} + \text{zero\_extend}(\text{uimm6} \ll 3)$

$\text{mem}[\text{address}+7:\text{address}] \leftarrow \text{fs2}$


**Permission:**

M mode/S mode/U mode

**Exception:**

Unaligned access exceptions, access error exceptions, and page error exceptions on store instructions

**Instruction format:**



### 16.6.16 C.J: a unconditional jump instruction

**Syntax:**

`c.j label`

**Operation:**

$\text{next pc} \leftarrow \text{current pc} + \text{sign\_extend}(\text{imm} \ll 1);$

**Permission:**

M mode/S mode/U mode



**Exception:**

None

**Notes:**

The compiler calculates immediate 11 based on the label.

The jump range of the instruction is  $\pm 2$  KB address space.

**Instruction format:**

15	13	12								2	1	0
101	imm11[10 3 8:7 9 5 6 2:0 4]										01	

**16.6.17 C.JALR: a jump and link register instruction****Syntax:**

```
c.jalr rs1
```

**Operation:**

```
next pc  $\leftarrow$  rs1;
```

```
x1  $\leftarrow$  current pc + 2;
```

**Permission:**

```
M mode/S mode/U mode
```

**Exception:**

None

**Notes:**

```
rs1  $\neq$  0.
```

When MMU is enabled, the jump range is the entire 512 GB address space.

When MMU is disabled, the jump range is the entire 1 TB address space.

**Instruction format:**

15	13	12	11			7	6				2	1	0
100	1	rs1					00000					10	

**16.6.18 C.JR: a jump register instruction****Syntax:**

```
c.jr rs1
```

**Operation:**

$$\text{next pc} = \text{rs1};$$
**Permission:**

M mode/S mode/U mode

**Exception:**

None

**Notes:**

$$\text{rs1} \neq 0.$$

When MMU is enabled, the jump range is the entire 512 GB address space.

When MMU is disabled, the jump range is the entire 1 TB address space.

**Instruction format:**

15	13	12	11	7	6	2	1	0
100	0		rs1		00000		10	

**16.6.19 C.LD: a load doubleword instruction****Syntax:**

$$\text{c.ld rd, uimm5} < < 3(\text{rs1})$$
**Operation:**

$$\text{address} \leftarrow \text{rs1} + \text{zero\_extend}(\text{uimm5} < < 3)$$

$$\text{rd} \leftarrow \text{mem}[\text{address} + 7 : \text{address}]$$
**Permission:**

M mode/S mode/U mode

**Exception:**

Unaligned access exceptions, access error exceptions, and page error exceptions on load instructions


**Notes:**

Typical rs1/rd code registers are:

- 000: x8
- 001: x9
- 010: x10
- 011: x11

- 100: x12
- 101: x13
- 110: x14
- 111: x15

**Instruction format:**



### 16.6.20 C.LDSP: a load doubleword instruction

**Syntax:**

`c.ldsp rd, uimm6<<3(sp)`

**Operation:**

$\text{address} \leftarrow \text{sp} + \text{zero\_extend}(\text{uimm6} \ll 3)$

$\text{rd} \leftarrow \text{mem}[\text{address}+7:\text{address}]$

**Permission:**

M mode/S mode/U mode


**Exception:**

Unaligned access exceptions, access error exceptions, and page error exceptions on load instructions

**Notes:**

$\text{rd} \neq 0$

**Instruction format:**



### 16.6.21 C.LI: a load immediate instruction

**Syntax:**

`c.li rd, imm6`

**Operation:**

$$rd \leftarrow \text{sign\_extend}(\text{imm6})$$
**Permission:**


M mode/S mode/U mode

**Exception:**

None

**Notes:**

rd != 0

**Instruction format:****16.6.22 C.LUI: a load upper immediate instruction****Syntax:**

c.lui rd, nzimm6

**Operation:**

$$rd \leftarrow \text{sign\_extend}(\text{nzimm6} \ll 12)$$
**Permission:**

M mode/S mode/U mode


**Exception:**

None

**Notes:**

rd != 0。

Nzimm6 != 0。

**Instruction format:**

### 16.6.23 C.LW: a load word instruction

**Syntax:**

```
c.lw rd, uimm5<<2(rs1)
```

**Operation:**

$$\text{address} \leftarrow \text{rs1} + \text{zero\_extend}(\text{uimm5} \ll 2)$$

$$\text{tmp}[31:0] \leftarrow \text{mem}[\text{address}+3:\text{address}]$$

$$\text{rd} \leftarrow \text{sign\_extend}(\text{tmp}[31:0])$$
**Permission:**

M mode/S mode/U mode


**Exception:**

Unaligned access exceptions, access error exceptions, and page error exceptions on load instructions

**Notes:**

Typical rs1/rd code registers are:

- 000: x8
- 001: x9
- 010: x10
- 011: x11
- 100: x12
- 101: x13
- 110: x14
- 111: x15

**Instruction format:**


### 16.6.24 C.LWSP: a load word stack pointer instruction

**Syntax:**

```
c.lwsp rd, uimm6<<2(sp)
```

**Operation:**

$$\text{address} \leftarrow \text{sp} + \text{zero\_extend}(\text{uimm6} \ll 2)$$

$$\text{tmp}[31:0] \leftarrow \text{mem}[\text{address}+3:\text{address}]$$

$$\text{rd} \leftarrow \text{sign\_extend}(\text{tmp}[31:0])$$
**Permission:**


M mode/S mode/U mode

**Exception:**

Unaligned access exceptions, access error exceptions, and page error exceptions on load instructions

**Notes:**

$\text{rd} \neq 0$

**Instruction format:****16.6.25 C.MV: an instruction that copies the value in rs to rd****Syntax:**

c.mv rd, rs2

**Operation:**

$\text{rd} \leftarrow \text{rs2};$

**Permission:**


M mode/S mode/U mode

**Exception:**

None

**Notes:**

$\text{rs2} \neq 0, \text{rd} \neq 0。$

**Instruction format:**

### 16.6.26 C.NOP: a no-operation instruction

**Syntax:**

c.nop

**Operation:**

No operations

**Permission:**

M mode/S mode/U mode

**Exception:**

None

**Instruction format:**

15	13	12	11	7	6	2	1	0
000	0	00000	00000	01				

### 16.6.27 C.OR: a bitwise OR instruction

**Syntax:**

c.or rd, rs2

**Operation:**

$rd \leftarrow rs1 \mid rs2$

**Permission:**

M mode/S mode/U mode

**Exception:**

None

**Notes:**

$rs1 = rd$

Typical rd/rs1 code registers are:

- 000: x8
- 001: x9
- 010: x10
- 011: x11
- 100: x12

- 101: x13
- 110: x14
- 111: x15

**Instruction format:**

15	13	12	11	10	9	7	6	5	4	2	1	0
100	0	11	rs1/rd			10	rs2			01		

### 16.6.28 C.SD: a store doubleword instruction

**Syntax:**

`c.sd rs2, uimm5<<3(rs1)`

**Operation:**

$\text{address} \leftarrow \text{rs1} + \text{zero\_extend}(\text{uimm5} \ll 3)$

$\text{mem}[\text{address}+7:\text{address}] \leftarrow \text{rs2}$

**Permission:**

M mode/S mode/U mode

**Exception:**

Unaligned access exceptions, access error exceptions, and page error exceptions on store instructions


**Notes:**

Typical rs1/rd code registers are:

- 000: x8
- 001: x9
- 010: x10
- 011: x11
- 100: x12
- 101: x13
- 110: x14
- 111: x15

**Instruction format:**





### 16.6.29 C.SDSP: a store doubleword stack pointer instruction

**Syntax:**

`c.fsdsp rs2, uimm6<<3(sp)`

**Operation:**

$\text{address} \leftarrow \text{sp} + \text{zero\_extend}(\text{uimm6} \ll 3)$

$\text{mem}[\text{address}+7:\text{address}] \leftarrow \text{rs2}$


**Permission:**

M mode/S mode/U mode

**Exception:**

Unaligned access exceptions, access error exceptions, and page error exceptions on store instructions

**Instruction format:**



### 16.6.30 C.SLLI: an immediate logical left shift instruction

**Syntax:**

`c.slli rd, nzuimm6`

**Operation:**

$\text{rd} \leftarrow \text{rs1} \ll \text{nzuimm6}$

**Permission:**

M mode/S mode/U mode

**Exception:**


None

**Notes:**

$\text{rs1} == \text{rd}$

$\text{rd/rs1} \neq 0, \text{nzuimm6} \neq 0$

Instruction format:



### 16.6.31 C.SRAI: a right shift arithmetic immediate instruction

Syntax:

c.srli rd, nzuimm6

Operation:

$rd \leftarrow rs1 \ggg nzuimm6$

Permission:

M mode/S mode/U mode

Exception:

None

Notes:


nzuimm6 != 0

rs1 == rd

Typical rs1/rd code registers are:

- 000: x8
- 001: x9
- 010: x10
- 011: x11
- 100: x12
- 101: x13
- 110: x14
- 111: x15

Instruction format:



### 16.6.32 C.SRLI: an immediate right shift instruction

**Syntax:**

c.srli rd, nzuimm6

**Operation:**

$rd \leftarrow rs1 \gg nzuimm6$

**Permission:**

M mode/S mode/U mode

**Exception:**

None

**Notes:**


$nzuimm6 \neq 0$

$rs1 == rd$

Typical rs1/rd code registers are:

- 000: x8
- 001: x9
- 010: x10
- 011: x11
- 100: x12
- 101: x13
- 110: x14
- 111: x15

**Instruction format:**



### 16.6.33 C.SW: a store word instruction

**Syntax:**

c.sw rs2, uimm5<<2(rs1)

**Operation:**

$$\text{address} \leftarrow \text{rs1} + \text{zero\_extend}(\text{uimm5} \ll 2)$$

$$\text{mem}[\text{address}+3:\text{address}] \leftarrow \text{rs2}$$
**Permission:**

M mode/S mode/U mode


**Exception:**

Unaligned access exceptions, access error exceptions, and page error exceptions on store instructions

**Notes:**

Typical rs1/rs2 code registers are:

- 000: x8
- 001: x9
- 010: x10
- 011: x11
- 100: x12
- 101: x13
- 110: x14
- 111: x15

**Instruction format:****16.6.34 C.SWSP: a store word stack pointer instruction****Syntax:**

$$\text{c.swsp rs2, uimm6} \ll 2(\text{sp})$$
**Operation:**

$$\text{address} \leftarrow \text{sp} + \text{zero\_extend}(\text{uimm6} \ll 2)$$

$$\text{mem}[\text{address}+3:\text{address}] \leftarrow \text{rs2}$$
**Permission:**

M mode/S mode/U mode

**Exception:**

Unaligned access exceptions, access error exceptions, and page error exceptions on store instructions

**Instruction format:**

15	13	12	7	6	2	1	0
110	uimm6[3:0 5:4]			rs2	10		

**16.6.35 C.SUB: a signed subtract instruction****Syntax:**

c.sub rd, rs2

**Operation:**

$rd \leftarrow rs1 - rs2$

**Permission:**

M mode/S mode/U mode

**Exception:**

None

**Notes:**

$rs1 == rd$

Typical rs1/rd code registers are:

- 000: x8
- 001: x9
- 010: x10
- 011: x11
- 100: x12
- 101: x13
- 110: x14
- 111: x15

**Instruction format:**

15	13	12	11	10	9	7	6	5	4	2	1	0
100	0	11	rs1/rd		00	rs2		01				

**16.6.36 C.SUBW: a signed subtract instruction that operates on the lower 32 bits****Syntax:**

```
c.subw rd, rs2
```

**Operation:**

$$\text{tmp}[31:0] \leftarrow \text{rs1}[31:0] - \text{rs2}[31:0]$$

$$\text{rd} \leftarrow \text{sign\_extend}(\text{tmp})$$
**Permission:**

M mode/S mode/U mode

**Exception:**

None

**Notes:**

$\text{rs1} == \text{rd}$

Typical rs1/rd code registers are:

- 000: x8
- 001: x9
- 010: x10
- 011: x11
- 100: x12
- 101: x13
- 110: x14
- 111: x15

**Instruction format:**

15	13	12	11	10	9	7	6	5	4	2	1	0
100	1	11	rs1/rd			00	rs2			01		

**16.6.37 C.XOR: a bitwise XOR instruction****Syntax:**

```
c.xor rd, rs2
```

**Operation:**

$$\text{rd} \leftarrow \text{rs1} \wedge \text{rs2}$$

**Permission:**

M mode/S mode/U mode

**Exception:**

None

**Notes:**

rs1 == rd

Typical rs1/rd code registers are:

- 000: x8
- 001: x9
- 010: x10
- 011: x11
- 100: x12
- 101: x13
- 110: x14
- 111: x15

**Instruction format:**

15	13	12	11	10	9	7	6	5	4	2	1	0
100	0	11	rs1/rd	01	rs2	01						

## 16.7 Appendix A-8 Pseudo instructions

RISC-V implements a series of pseudo instructions. The instructions listed in this section are for reference only and are sorted in alphabetic order.

Pseudo instruction	Base instruction	Meaning
beqz rs, offset	beq rs, x0, offset	Takes the branch if rs is zero.
bnez rs, offset	bne rs, x0, offset	Takes the branch if rs is not zero.
blez rs, offset	bge x0,rs, offset	Takes the branch if rs is less than or equal to zero.
bgez rs, offset	bge rs, x0, offset	Takes the branch if rs is greater than or equal to zero.
bltz rs, offset	blt rs, x0, offset	Takes the branch if rs is less than zero.

Continued on next page

Table 16.1 – continued from previous page

Pseudo instruction	Base instruction	Meaning
bgtz rs, offset	blt x0, xs, offset	Takes the branch if rs is greater than zero.
bgt rs, rt, offset	blt rt, rs, offset	Takes the branch if rs is greater than rt.
ble rs, rt, offset	bge rt, rs, offset	Takes the branch if rs is less than or equal to rt.
bgtu rs, rt, offset	bltu rt, rs, offset	Takes the branch if rs is greater than rt, using unsigned comparison.
bleu rs, rt, offset	bgeu rt, rs, offset	Takes the branch if rs is less than or equal to rt, using unsigned comparison.
call offset	auipc x6, offset[31:12] jalr x1, x6, offset[11:0]	Calls far-away subroutine.
csrrc csr, rs	csrrc x0, csr, rs	Clears bits in the control/status register (CSR).
csrci csr, imm	csrrci x0, csr, imm	Clears bits in the CSR, immediate.
csrs csr, rs	csrrs x0, csr, rs	Sets bits in the CSR.
csrsi csr, imm	csrrsi x0, csr, imm	Sets bits in the CSR, immediate
csrw csr, rs	csrrw x0, csr, rs	Writes the CSR.
csrwi csr, imm	csrrwi x0, csr, imm	Writes the CSR, immediate.
fabs.d rd, rs	fsgnjx.d rd, rs, rs	Calculates the double-precision floating point (FP) absolute value.
fabs.s rd, rs	fsgnjx.s rd, rs, rs	Calculates the single-precision FP absolute value.
fence	fence iorw, iorw	Fences on all memory and I/O.
fl{w d} rd, symbol, rt	auipc rt, symbol[31:12] fl{w d} rd, symbol[11:0](rt)	An FP load global instruction.
fmv.d rd, rs	fsgnj.d rd, rs, rs	A double-precision FP copy instruction.
fmv.s rd, rs	fsgnj.s rd, rs, rs	A single-precision FP copy instruction.
fneg.d rd, rs	fsgnjd rd, rs, rs	A double-precision FP negate instruction.
fneg.s rd, rs	fsgnjd.s rd, rs, rs	A single-precision FP negate instruction.

Continued on next page



Table 16.1 – continued from previous page

Pseudo instruction	Base instruction	Meaning
fcsr rd	csrrs x0, fcsr, x0	Reads FP CSR.
frflags rd	csrrs rd, fflags, x0	Reads FP exception flags.
frfm rd	csrrs rd, frm, x0	Reads FP rounding mode.
fscsr rs	csrrw x0, fcsr, rs	Writes FP CSR.
fscsr rd, rs	csrrs rd, fcsr, rs	Swaps FP CSR.
fsflags rs	csrrw x0, fflags, rs	Writes FP exception flags.
fsflags rd, rs	csrrs rd, fflags, rs	Swaps FP exception flags.
fsflagsi imm	csrrwi x0, fflags, imm	Writes FP exception flags, immediate.
fsflagsi rd, imm	csrrwi rd, fflags, imm	Swaps FP exception flags, immediate.
fsrm rs	csrrw x0, frm, rs	Writes FP rounding mode.
fsrm rd, rs	csrrs rd, frm, rs	Swaps FP rounding mode.
fsrmi imm	csrrwi x0, frm, imm	Writes FP rounding mode, immediate.
fsrmi rd, imm	csrrwi rd, frm, imm	Swaps FP rounding mode, immediate.
fs{w d} rd, symbol,rt	auipc rt,symbol[31:12] fs{w d} rd, symbol[11:0](rt)	An FP store global instruction.
j offset	jal x0, offset	A jump instruction.
jal offset	jal x1, offset	Jumps to subroutine and link.
jalr rs	jalr x1, rs, 0	Jumps to subroutine and links register.
jr rs	jalr x0, rs, 0	A jump register instruction.
la rd, symbol	auipc rd, symbol[31:12] addi rd, rd, symbol[11:0]	A load address instruction.
li rd, immediate	Split into multiple instructions based on the size of the immediate	A load immediate instruction
l{b h w d} rd,symbol, rt	auipc rt,  symbol[31:12] l{b h w d} rd,symbol[11:0](rt)	A load global instruction.
mv rd, rs	addi rd, rs, 0	A instruction that copies the value in rs to rd.
neg rd, rs	sub rd, x0, rs	A register negate instruction.
negw rd, rs	subw rd, x0, rs	Negates the lower 32 bits of registers.
nop	addi x0,x0,0	A no operation instruction.
not rd, rs	xori rd, rs, -1	A register NOT instruction.

Continued on next page

Table 16.1 – continued from previous page

Pseudo instruction	Base instruction	Meaning
rdcycle[h] rd	csrrs rd, cycle[h], x0	A read cycle counter instruction.
rdinstret[h] rd	csrrs rd, instret[h], x0	Reads instructions-retired counter.
rdtime[h] rd	csrrs rd, time[h], x0	Reads real-time clock.
ret	jalr x0, x1, 0	Returns from subroutine.
s{b h w d} rd, symbol, rt	auipc rt, symbol[31:12] s{b h w d} rd, symbol[11:0](rt)	A store global instruction.
seqz rd, rs	sltiu rd, rs, 1	Sets 0 in registers to 1.
sextw rd, rs	addiw rd, rs, 0	A sign extend word instruction.
sgtz rd, rs	slt rd, rs, x0, rs	Sets rd to 1 if rs is greater than zero.
sltz rd, rs	slt rd, rs, rs, x0	Sets rd to 1 if rs is less than zero.
snez rd, rs	sltu rd, rs, x0, rs	Sets rd to 1 if rs is not equal to zero.
tail offset	auipc x6, offset[31:12] jalr x0, x6, offset[11:0]	Tail call far-away subroutine.

---

## Appendix B T-Head Extended Instructions

---

Apart from the GC instruction sets defined in the standard, C910 provides custom instruction sets, including the cache instruction set, synchronization instruction set, arithmetic operation instruction set, bitwise operation instruction set, storage instruction set, and half-precision floating-point instruction set.

Among these instruction sets, the cache instructions, synchronization instructions, arithmetic operation instructions, bitwise operation instructions, and storage instructions can be executed only when the value of `mxstatus.theadisaee` is 1. Otherwise, an instruction exception will occur. Half-precision floating-point instructions can be executed only when the value of `mstatus.fs` is 2'b00. Otherwise, an illegal instruction exception will occur. The following describes each instruction in these instruction sets.

### 17.1 Appendix B-1 Cache instructions

You can use the cache instruction set to manage caches. Each instruction has 32 bits.

Arithmetic operation instructions in this instruction set are described in alphabetical order.

#### 17.1.1 DCACHE.CALL: an instruction that clears all dirty page table entries in the D-Cache

**Syntax:**

`dcache.call`

**Operation:**

Clears all page table entries in the L1 D-Cache and writes all dirty page table entries back into the next-level storage. You can perform this operation only on the current core.

**Permission:**

M mode/S mode

**Exception:**

Illegal instruction.

**Notes:**

If the value of `mxstatus.theadisaee` is 0, executing this instruction causes an exception of illegal instruction.

If the value of `mxstatus.theadisaee` is 1, executing this instruction in U mode causes an exception of illegal instruction.

**Instruction format:**

31	25	24	20	19	15	14	12	11	7	6	0
0000000			00001		00000		000		00000		0001011

### 17.1.2 DCACHE.CIALL: an instruction that clears all dirty page table entries in the D-Cache and invalidates the D-Cache

**Syntax:**

`dcache.ciall`

**Operation:**

Writes all dirty page table entries in the L1 D-Cache back into the next-level storage and invalidates all these page table entries.

**Permission:**

M mode/S mode

**Exception:**

Illegal instruction.

**Notes:**

If the value of `mxstatus.theadisaee` is 0, executing this instruction causes an exception of illegal instruction.

If the value of `mxstatus.theadisaee` is 1, executing this instruction in U mode causes an exception of illegal instruction.

**Instruction format:**

31	25	24	20	19	15	14	12	11	7	6	0
0000000			00011		00000		000		00000		0001011

### 17.1.3 DCACHE.CIPA: clears dirty page table entries that match the specified physical addresses from the D-Cache and invalidates the the D-Cache

#### Syntax:

```
dcache.cipa rs1
```

#### Operation:

Writes page table entries that match the specified physical addresses of the D-Cache or L2 Cache of rs1 back into the next-level storage and invalidates these page table entries. You can perform this operation on all cores and the L2 Cache.

#### Permission:

M mode/S mode

#### Exception:

Illegal instruction.

#### Notes:

If the value of mxstatus.theadisaee is 0, executing this instruction causes an exception of illegal instruction.

If the value of mxstatus.theadisaee is 1, executing this instruction in U mode causes an exception of illegal instruction.

#### Instruction format:

31	25	24	20	19	15	14	12	11	7	6	0
0000001	01011	rs1	000	00000	0001011						

### 17.1.4 DCACHE.CISW: an instruction that clears dirty page table entries in the D-Cache based on the specified way and set and invalidates the D-Cache

#### Syntax:

```
dcache.cisw rs1
```

#### Operation:

Writes the dirty page table entry that matches the specified way and set from the L1 Cache of rs1 back into the next-level storage and invalidates this page table entry. You can perform this operation only on the current core.

#### Permission:

M mode/S mode

#### Exception:

Illegal instruction.

**Notes:**

C910/C920 D-Cache is a 2-way set-associative cache. `rs1[31]` specifies the way and `rs1[s:6]` specifies the set. When the size of the D-Cache is 32 KB, `w` denotes 13. When the size of the D-Cache is 64 KB, `w` denotes 14, and so forth.

- If the value of `mxstatus.theadisaee` is 0, executing this instruction causes an exception of illegal instruction.
- If the value of `mxstatus.theadisaee` is 1, executing this instruction in U mode causes an exception of illegal instruction.

**Instruction format:**

31	25	24	20	19	15	14	12	11	7	6	0
0000001	00011		rs1		000		00000		0001011		

### 17.1.5 DCACHE.CIVA: an instruction that clears dirty page table entries that match the specified virtual addresses in the D-Cache and invalidates the D-Cache

**Syntax:**

```
dcache.civa rs1
```

**Operation:**

Writes the page table entry that matches the specified virtual address from the D-Cache or L2 Cache of `rs1` back into the next-level storage and invalidates this page table entry. You can perform this operation on the current core and the L2 Cache. The sharing attribute of the virtual address determines whether you can perform this operation on other cores.

**Permission:**

M mode/S mode/U mode

**Exception:**

Illegal instruction or error page during instruction loading.

**Notes:**

- If the value of `mxstatus.theadisaee` is 0, executing this instruction causes an exception of illegal instruction.
- If the value of `mxstatus.theadisaee` is 1 and the value of `mxstatus.ucme` is 1, this instruction can be executed in U mode.
- If the value of `mxstatus.theadisaee` is 1 and the value of `mxstatus.ucme` is 0, executing this instruction in U mode causes an exception of illegal instruction.

**Instruction format:**

31	25	24	20	19	15	14	12	11	7	6	0
0000001	00111	rs1	000	00000	0001011						

### 17.1.6 DCACHE.CPA: an instruction that clears dirty page table entries that match the specified physical addresses from the D-Cache

#### Syntax:

dcache.cpa rs1

#### Operation:

Writes the page table entry that matches the specified physical address from the D-Cache or L2 Cache of rs1 back into the next-level storage. You can perform this operation on all cores and the L2 Cache.

#### Permission:

M mode/S mode

#### Exception:

Illegal instruction.

#### Notes:

If the value of mxstatus.theadisae is 0, executing this instruction causes an exception of illegal instruction.

If the value of mxstatus.theadisae is 1, executing this instruction in U mode causes an exception of illegal instruction.

#### Instruction format:

31	25	24	20	19	15	14	12	11	7	6	0
0000001	01001	rs1	000	00000	0001011						

### 17.1.7 DCACHE.CPAL1: an instruction that clears dirty page table entries that match the specified physical addresses from the L1 D-Cache

#### Syntax:

dcache.cpal1 rs1

**Operation:** Writes the page table entry that matches the specified physical address from the D-Cache of rs1 back into the next-level storage. You can perform this operation on all cores and the L1 Cache.

#### Permission:

M mode/S mode

**Exception:**

Illegal instruction.

**Notes:**

If the value of `mxstatus.theadisaee` is 0, executing this instruction causes an exception of illegal instruction.

If the value of `mxstatus.theadisaee` is 1, executing this instruction in U mode causes an exception of illegal instruction.

**Instruction format:**

31	25	24	20	19	15	14	12	11	7	6	0
0000001	01000	rs1	000	00000	0001011						

### 17.1.8 DCACHE.CVA: an instruction that clears dirty page table entries that match the specified virtual addresses in the D-Cache

**Syntax:**

`dcache.cva rs1`

**Operation:**

Writes the page table entry that matches the specified virtual address from the D-Cache or L2 Cache of `rs1` back into the next-level storage. You can perform this operation on the current core and the L2 Cache. The sharing attribute of the virtual address determines whether you can perform this operation on other cores.

**Permission:**

M mode/S mode

**Exception:**

Illegal instruction or error page during instruction loading.

**Notes:**

If the value of `mxstatus.theadisaee` is 0, executing this instruction causes an exception of illegal instruction.

If the value of `mxstatus.theadisaee` is 1, executing this instruction in U mode causes an exception of illegal instruction.

**Instruction format:**

31	25	24	20	19	15	14	12	11	7	6	0
0000001	00101	rs1	000	00000	0001011						



### 17.1.9 DCACHE.CVAL1: an instruction that clears dirty page table entries that match the specified virtual addresses in the L1 D-Cache

#### Syntax:

```
dcache.cval1 rs1
```

#### Operation:

Writes the page table entry that matches the specified virtual address from the D-Cache of s1 back into the next-level storage. You can perform this operation on all cores and the L1 Cache.

#### Permission:

M mode/S mode/U mode

#### Exception:

Illegal instruction or error page during instruction loading.

#### Notes:

If the value of mxstatus.theadisaee is 0, executing this instruction causes an exception of illegal instruction.

If the value of mxstatus.theadisaee is 1 and the value of mxstatus.ucme is 0, executing this instruction in U mode causes an exception of illegal instruction.

#### Instruction format:

31	25	24	20	19	15	14	12	11	7	6	0
0000001	00100	rs1	000	00000	0001011						

### 17.1.10 DCACHE.IPA: an instruction that invalidates page table entries that match the specified physical addresses in the D-Cache

#### Syntax:

```
dcache.ipa rs1
```

#### Operation:

Invalidates the page table entry that matches the specified physical address in the D-Cache or L2 Cache of rs1. You can perform this operation on all cores and the L2 Cache.

#### Permission:

M mode/S mode

#### Exception:

Illegal instruction.

#### Notes:

- If the value of `mxstatus.theadisaee` is 0, executing this instruction causes an exception of illegal instruction.
- If the value of `mxstatus.theadisaee` is 1, executing this instruction in U mode causes an exception of illegal instruction.

**Instruction format:**

31	25	24	20	19	15	14	12	11	7	6	0
0000001		01010		rs1		000		00000		0001011	

### 17.1.11 DCACHE.ISW: an instruction that invalidates page table entries in the D-Cache based on the specified way and set and invalidates the D-Cache

**Syntax:**

```
dcache.isw rs1
```

**Operation:**

Invalidates the page table entry in the D-Cache based on the specified set and way. You can perform this operation only on the current core.

**Permission:**

M mode/S mode

**Exception:**

Illegal instruction.

**Notes:**

C910/C920 D-Cache is a 2-way set-associative cache. `rs1[31]` specifies the way and `rs1[s:6]` specifies the set. When the size of the D-Cache is 32 KB, `w` denotes 13. When the size of the D-Cache is 64 KB, `w` denotes 14, and so forth.

- If the value of `mxstatus.theadisaee` is 0, executing this instruction causes an exception of illegal instruction.
- If the value of `mxstatus.theadisaee` is 1, executing this instruction in U mode causes an exception of illegal instruction.

**Instruction format:**

31	25	24	20	19	15	14	12	11	7	6	0
0000001		00010		rs1		000		00000		0001011	

### 17.1.12 DCACHE.IVA: an instruction that invalidates the D-Cache based on the specified virtual address

#### Syntax:

```
dcache.iva rs1
```

#### Operation:

Invalidates the page table entry that matches the specified virtual address from the D-Cache or L2 Cache of rs1. You can perform this operation on the current core and the L2 Cache. The sharing attribute of the virtual address determines whether you can perform this operation on other cores.

#### Permission:

M mode/S mode

#### Exception:

Illegal instruction or error page during instruction loading.

#### Notes:

- If the value of mxstatus.theadisaee is 0, executing this instruction causes an exception of illegal instruction.
- If the value of mxstatus.theadisaee is 1, executing this instruction in U mode causes an exception of illegal instruction.

#### Instruction format:

31	25	24	20	19	15	14	12	11	7	6	0
0000001	00110	rs1	000	00000	0001011						

### 17.1.13 DCACHE.IALL: an instruction that invalidates all page table entries in the D-Cache.

#### Syntax:

```
dcache.iall
```

#### Operation:

Invalidates all page table entries in the L1 Cache. You can perform this operation only on the current core.

#### Permission:

M mode/S mode

#### Exception:

Illegal instruction.

#### Notes:

- If the value of `mxstatus.theadisaee` is 0, executing this instruction causes an exception of illegal instruction.
- If the value of `mxstatus.theadisaee` is 1, executing this instruction in U mode causes an exception of illegal instruction.

**Instruction format:**

31	25	24	20	19	15	14	12	11	7	6	0
0000000		00010		00000		000		00000		0001011	

### 17.1.14 ICACHE.IALL: an instruction that invalidates all page table entries in the I-Cache

**Syntax:**

```
icache.iall
```

**Operation:**

Invalidates all page table entries in the I-Cache. You can perform this operation only on the current core.

**Permission:**

M mode/S mode

**Exception:**

Illegal instruction.

**Notes:**

If the value of `mxstatus.theadisaee` is 0, executing this instruction causes an exception of illegal instruction.

If the value of `mxstatus.theadisaee` is 1, executing this instruction in U mode causes an exception of illegal instruction.

**Instruction format:**

31	25	24	20	19	15	14	12	11	7	6	0
0000000		10000		00000		000		00000		0001011	

### 17.1.15 ICACHE.IALLS: an instruction that invalidates all page table entries in the I-Cache through broadcasting

**Syntax:**

```
icache.ialls
```

**Operation:**

Invalidates all page table entries in the I-Cache and invalidates all page table entries in the I-Cache of other cores through broadcasting. You can perform this operation on all cores.

**Permission:**

M mode/S mode

**Exception:**

Illegal instruction.

**Notes:**

If the value of `mxstatus.theadisaee` is 0, executing this instruction causes an exception of illegal instruction.

If the value of `mxstatus.theadisaee` is 1, executing this instruction in U mode causes an exception of illegal instruction.

**Instruction format:**

31	25	24	20	19	15	14	12	11	7	6	0
0000000	10001	00000	000	00000	0001011						

### 17.1.16 ICACHE.IPA: an instruction that invalidates page table entries that match the specified physical addresses in the I-Cache

**Syntax:**

`icache.ipa rs1`

**Operation:**

Invalidates the page table entry that matches the specified physical address in the I-Cache of `rs1`. You can perform this operation on all cores.

**Permission:**

M mode/S mode

**Exception:**

Illegal instruction.

**Notes:**

If the value of `mxstatus.theadisaee` is 0, executing this instruction causes an exception of illegal instruction.

If the value of `mxstatus.theadisaee` is 1, executing this instruction in U mode causes an exception of illegal instruction.

**Instruction format:**

31	25	24	20	19	15	14	12	11	7	6	0
0000001	11000	rs1	000	00000	0001011						

### 17.1.17 ICACHE.IVA: an instruction that invalidates page table entries that match the specified virtual addresses in the I-Cache

**Syntax:**

```
icache.iva rs1
```

**Operation:**

Invalidates the page table entry that matches the specified virtual address in the I-Cache of rs1. You can perform this operation only on the current core. The sharing attribute of the virtual address determines whether you can perform this operation on other cores.

**Permission:**

M mode/S mode/U mode

**Exception:**

Illegal instruction or error page during instruction loading.

**Notes:**

If the value of mxstatus.theadisaee is 0, executing this instruction causes an exception of illegal instruction.

If the value of mxstatus.theadisaee is 1 and the value of mxstatus.ucme is 1, this instruction can be executed in U mode.

If the value of mxstatus.theadisaee is 1 and the value of mxstatus.ucme is 0, executing this instruction in U mode causes an exception of illegal instruction.

**Instruction format:**

31	25	24	20	19	15	14	12	11	7	6	0
0000001			10000		rs1		000		00000		0001011

### 17.1.18 L2CACHE.CALL: an instruction that clears all dirty page table entries in the L2 Cache

**Syntax:**

```
l2cache.call
```

**Operation:**

Writes all dirty page table entries from the L2 Cache back into the next-level storage.

**Permission:**

M mode/S mode

**Exception:**

Illegal instruction.

**Notes:**

- If the value of `mxstatus.theadisaee` is 0, executing this instruction causes an exception of illegal instruction.
- If the value of `mxstatus.theadisaee` is 1, executing this instruction in U mode causes an exception of illegal instruction.

**Instruction format:**

31	25	24	20	19	15	14	12	11	7	6	0
0000000	10101	00000	000	00000	0001011						

### 17.1.19 L2CACHE.CIALL: an instruction that clears all dirty page table entries in the L2 Cache and invalidates the L2 Cache

**Syntax:**

```
l2cache.ciall
```

**Operation:**

Writes all dirty page table entries from the L2 Cache back into the next-level storage and invalidates all page table entries in the L2 Cache.

**Permission:**

M mode/S mode

**Exception:**

Illegal instruction.

**Notes:**

- If the value of `mxstatus.theadisaee` is 0, executing this instruction causes an exception of illegal instruction.
- If the value of `mxstatus.theadisaee` is 1, executing this instruction in U mode causes an exception of illegal instruction.

**Instruction format:**

31	25	24	20	19	15	14	12	11	7	6	0
0000000	10111	00000	000	00000	0001011						

### 17.1.20 L2CACHE.IALL: an instruction that invalidates the L2 Cache

**Syntax:**

```
l2cache.iall
```

**Operation:**

Invalidates all page table entries in the L2 Cache.

**Permission:**

M mode/S mode

**Exception:**

Illegal instruction.

**Notes:**

- If the value of `mxstatus.cskisayee` is 0, executing this instruction causes an exception of illegal instruction.
- If the value of `mxstatus.cskisayee` is 1, executing this instruction in U mode causes an exception of illegal instruction.

**Instruction format:**

31	25	24	20	19	15	14	12	11	7	6	0
0000000		10110		00000		000		00000		0001011	

### 17.1.21 DCACHE.CSW: an instruction that clears dirty page table entries in the D-Cache based on the specified set and way

**Syntax:**

`dcache.csw rs1`

**Operation:**

Writes the dirty page table entry from the D-Cache back into the next-level storage device based on the specified set and way.

**Permission:**

M mode/S mode

**Exception:**

Illegal instruction.

**Notes:**

C910/C920 D-Cache is a 2-way set-associative cache. `rs1[31]` specifies the way and `rs1[s:6]` specifies the set. When the size of the D-Cache is 32 KB, `w` denotes 13. When the size of the D-Cache is 64 KB, `w` denotes 14, and so forth.

If the value of `mxstatus.theadisaee` is 0, executing this instruction causes an exception of illegal instruction.



If the value of `mxstatus.theadisaee` is 1, executing this instruction in U mode causes an exception of illegal instruction.

**Instruction format:**




Fig. 17.1: DCACHE.CSW

## 17.2 Appendix B-2 Multi-core synchronization instructions

This synchronization instruction set extends multi-core synchronization instructions. Each instruction has 32 bits. Synchronization instructions in this instruction set are described in alphabetical order.

### 17.2.1 SFENCE.VMAS: a broadcast instruction that synchronizes the virtual memory address

**Syntax:**

`sfence.vmas rs1,rs2`

**Operation:**

Invalidates and synchronizes page table entries in the virtual memory and broadcasts them to other cores in the cluster.

**Permission:**

M mode/S mode

**Exception:**

Illegal instruction.

**Notes:**

rs1 is the virtual address, and rs2 is the address space identifier (ASID).

- If the value of rs1 is x0 and the value of rs2 is x0, invalidate all page table entries in the TLB and broadcast them to other cores in the cluster.
- When rs1! and rs2 are both x0, all TLB entries that hit the virtual address specified by rs1 are invalidated and broadcast to other cores in the cluster.

- When rs1 and rs2! are both x0, all TLB entries that hit the process ID specified by rs2 are invalidated and broadcast to other cores in the cluster.
- When rs1! and rs2! are both x0, all TLB entries that hit the virtual address specified by rs1 and the process ID specified by rs2 are invalidated and broadcast to other cores in the cluster.

If the value of mxstatus.theadisaee is 0, executing this instruction causes an exception of illegal instruction.

If the value of mxstatus.theadisaee is 1, executing this instruction in U mode causes an exception of illegal instruction.

#### Instruction format:

31	25	24	20	19	15	14	12	11	7	6	0
0000010				rs2		rs1		000	00000		0001011

### 17.2.2 SYNC: an instruction that performs the synchronization operation

#### Syntax:

sync

#### Operation:

Ensures that all preceding instructions retire earlier than this instruction and all subsequent instructions retire later than this instruction.

#### Permission:

M mode/S mode/U mode

#### Exception:

Illegal instruction.

#### Instruction format:

31	25	24	20	19	15	14	12	11	7	6	0
0000000				11000		00000		000	00000		0001011

### 17.2.3 SYNC.I: an instruction that synchronizes the clearing operation.

#### Syntax:

sync.i

#### Operation:

Ensures that all preceding instructions retire earlier than this instruction and all subsequent instructions retire later than this instruction, and clears the pipeline when this instruction retires.

#### Permission:

M mode/S mode/U mode

**Exception:**

Illegal instruction.

**Instruction format:**

31	25	24	20	19	15	14	12	11	7	6	0
0000000	11010	00000	000	00000	0001011						

### 17.2.4 SYNC.IS: a broadcast instruction that synchronizes the clearing operation

**Syntax:**

sync.is

**Operation:**

Ensures that all preceding instructions retire earlier than this instruction and all subsequent instructions retire later than this instruction. Clears the pipeline when this instruction retires and broadcasts the request to other cores.

**Permission:**

M mode/S mode/U mode

**Exception:**

Illegal instruction.

**Instruction format:**

31	25	24	20	19	15	14	12	11	7	6	0
0000000	11011	00000	000	00000	0001011						

### 17.2.5 SYNC.S: a broadcast instruction that performs a synchronization operation

**Syntax:**

sync.s

**Operation:**

Ensures that all preceding instructions retire earlier than this instruction and all subsequent instructions retire later than this instruction, and broadcasts the request to other cores.

**Permission:**

M mode/S mode/U mode

**Exception:**

Illegal instruction.

Instruction format:

31	25	24	20	19	15	14	12	11	7	6	0
0000000	11001	00000	000	00000	0001011						

## 17.3 Appendix B-3 Arithmetic operation instructions

The arithmetic operation instruction set extends arithmetic operation instructions. Each instruction has 32 bits.

Arithmetic operation instructions in this instruction set are described in alphabetical order.

### 17.3.1 ADDSL: an add register instruction that shifts registers

**Syntax:**

addsl rd rs1, rs2, imm2

**Operation:**

$rd \leftarrow rs1 + rs2 \ll imm2$

**Permission:**

M mode/S mode/U mode

**Exception:**

Illegal instruction.

Instruction format:

31	27	26	25	24	20	19	15	14	12	11	7	6	0
00000	imm2	rs2	rs1	001	rd	0001011							

### 17.3.2 MULA: a multiply-add instruction

**Syntax:**

mula rd, rs1, rs2

**Operation:**

$rd \leftarrow rd + (rs1 * rs2)[63:0]$

**Permission:**

M mode/S mode/U mode

**Exception:**

Illegal instruction.

**Instruction format:**

31	27	26	25	24	20	19	15	14	12	11	7	6	0
00100	00	rs2			rs1			001	rd			0001011	

**17.3.3 MULA: a multiply-add instruction that operates on the lower 16 bits****Syntax:**

mulah rd, rs1, rs2

**Operation:**

$\text{tmp}[31:0] \leftarrow \text{rd}[31:0] + (\text{rs1}[15:0] * \text{rs2}[15:0])$

$\text{rd} \leftarrow \text{sign\_extend}(\text{tmp}[31:0])$

**Permission:**

M mode/S mode/U mode

**Exception:**

Illegal instruction.

**Instruction format:**

31	27	26	25	24	20	19	15	14	12	11	7	6	0
00101	00	rs2			rs1			001	rd			0001011	

**17.3.4 MULAW: a multiply-add instruction that operates on the lower 32 bits****Syntax:**

mulaw rd, rs1, rs2

**Operation:**

$\text{tmp}[31:0] \leftarrow \text{rd}[31:0] + (\text{rs1}[31:0] * \text{rs2}[31:0])[31:0]$

$\text{rd} \leftarrow \text{sign\_extend}(\text{tmp}[31:0])$

**Permission:**

M mode/S mode/U mode

**Exception:**

Illegal instruction.

**Instruction format:**

31	27	26	25	24	20	19	15	14	12	11	7	6	0
00100	10			rs2		rs1		001		rd			0001011

### 17.3.5 MULS: a multiply-subtract instruction

#### Syntax:

muls rd, rs1, rs2

#### Operation:

$rd \leftarrow rd - (rs1 * rs2)[63:0]$

#### Permission:

M mode/S mode/U mode

#### Exception:

Illegal instruction.

#### Instruction format:

31	27	26	25	24	20	19	15	14	12	11	7	6	0
00100	01			rs2		rs1		001		rd			0001011

### 17.3.6 MULSH: a multiply-subtract instruction that operates on the lower 16 bits

#### Syntax:

mulsh rd, rs1, rs2

#### Operation:

$tmp[31:0] \leftarrow rd[31:0] - (rs1[15:0] * rs2[15:0])$

$rd \leftarrow sign\_extend(tmp[31:0])$

#### Permission:

M mode/S mode/U mode

#### Exception:

Illegal instruction.

#### Instruction format:

31	27	26	25	24	20	19	15	14	12	11	7	6	0
00101	01			rs2		rs1		001		rd			0001011

### 17.3.7 MULSW: a multiply-subtract instruction that operates on the lower 32 bits

**Syntax:**

```
mulaw rd, rs1, rs2
```

**Operation:**

$$\text{tmp}[31:0] \leftarrow \text{rd}[31:0] - (\text{rs1}[31:0] * \text{rs2}[31:0])$$

$$\text{rd} \leftarrow \text{sign\_extend}(\text{tmp}[31:0])$$
**Permission:**

M mode/S mode/U mode

**Exception:**

Illegal instruction.

**Instruction format:**

31	27	26	25	24	20	19	15	14	12	11	7	6	0
00100	11			rs2		rs1		001		rd		0001011	

### 17.3.8 MVEQZ: an instruction that sends a message when the register is 0

**Syntax:**

```
mveqz rd, rs1, rs2
```

**Operation:** if (rs2 == 0)

$$\text{rd} \leftarrow \text{rs1}$$

else

$$\text{rd} \leftarrow \text{rd}$$
**Permission:**

M mode/S mode/U mode

**Exception:**

Illegal instruction.

**Instruction format:**

31	27	26	25	24	20	19	15	14	12	11	7	6	0
01000	00			rs2		rs1		001		rd		0001011	

### 17.3.9 MVNEZ: an instruction that sends a message when the register is not 0

**Syntax:**

```
mvnez rd, rs1, rs2
```

**Operation:**

```
if (rs2 != 0)
```

```
    rd ← rs1
```

```
else
```

```
    rd ← rd
```

**Permission:**

M mode/S mode/U mode

**Exception:**

Illegal instruction.

**Instruction format:**

31	27	26	25	24	20	19	15	14	12	11	7	6	0
01000	01	rs2			rs1			001	rd			0001011	

### 17.3.10 SRRI: an instruction that implements a cyclic right shift operation on a linked list

**Syntax:**

```
srri rd, rs1, imm6
```

**Operation:**

```
rd ← rs1 >>>> imm6
```

Shifts the original value of rs1 to the right, disconnects the last value on the list, and re-attaches the value to the start of the linked list.

**Permission:**

M mode/S mode/U mode

**Exception:**

Illegal instruction.

**Instruction format:**

31	26	25	20	19	15	14	12	11	7	6	0		
000100			imm6			rs1		001		rd		0001011	



### 17.3.11 SRRIW: an instruction that implements a cyclic right shift operation on a linked list of low 32 bits of registers.

#### Syntax:

`srriw rd, rs1, imm5`

#### Operation:

$rd \leftarrow \text{sign\_extend}(rs1[31:0] \gg \gg \gg \text{imm5})$

Shifts the original value of `rs1[31:0]` to the right, disconnects the last value on the list, and re-attaches the value to the start of the linked list.

#### Permission:

M mode/S mode/U mode

#### Exception:

Illegal instruction.

#### Instruction format:

31	25	24	20	19	15	14	12	11	7	6	0
0001010	imm5			rs1		001	rd		0001011		

## 17.4 Appendix B-4 Bitwise operation instructions

The bitwise operation instruction set extends bitwise operation instructions. Each instruction has 32 bits.

Arithmetic operation instructions in this instruction set are described in alphabetical order.

### 17.4.1 EXT: a signed extension instruction that extracts consecutive bits of a register

#### Syntax:

`ext rd, rs1, imm1,imm2`

#### Operation:

$rd \leftarrow \text{sign\_extend}(rs1[\text{imm1}:\text{imm2}])$

#### Permission:

M mode/S mode/U mode

#### Exception:

Illegal instruction.

#### Notes:

If imm1 is smaller than imm2, the action of this instruction is not predictable.

**Instruction format:**

31	26	25	20	19	15	14	12	11	7	6	0
imm1			imm2			rs1		010	rd		0001011

### 17.4.2 EXTU: a zero extension instruction that extracts consecutive bits of a register

**Syntax:**

extu rd, rs1, imm1,imm2

**Operation:**

$rd \leftarrow \text{zero\_extend}(rs1[imm1:imm2])$

**Permission:**

M mode/S mode/U mode

**Exception:**

Illegal instruction.

**Notes:**

If imm1 is smaller than imm2, the action of this instruction is not predictable.

**Instruction format:**

31	26	25	20	19	15	14	12	11	7	6	0
imm1			imm2			rs1		011	rd		0001011

### 17.4.3 FF0: an instruction that finds the first bit with the value of 0 in a register

**Syntax:**

ff0 rd, rs1

**Operation:**

Finds the first bit with the value of 0 from the highest bit of rs1 and writes the result back into the rd register. If the highest bit of rs1 is 0, the result 0 is returned. If all the bits in rs1 are 1, the result 64 is returned.

**Permission:**

M mode/S mode/U mode

**Exception:**

Illegal instruction.

Instruction format:

31	27	26	25	24	20	19	15	14	12	11	7	6	0
10000	10	00000	rs1	001	rd	0001011							

#### 17.4.4 FF1: an instruction that finds the bit with the value of 1

Syntax:

ff1 rd, rs1

Operation:

Finds the first bit with the value of 1 from the highest bit of rs1 and writes the index of this bit back into rd. If the highest bit of rs1 is 1, the result 0 is returned. If all the bits in rs1 are 1, the result 64 is returned.

Permission:

M mode/S mode/U mode

Exception:

Illegal instruction.

Instruction format:

31	27	26	25	24	20	19	15	14	12	11	7	6	0
10000	11	00000	rs1	001	rd	0001011							

#### 17.4.5 REV: an instruction that reverses the byte order in a word stored in the register

Syntax:

rev rd, rs1

Operation:

$rd[63:56] \leftarrow rs1[7:0]$

$rd[55:48] \leftarrow rs1[15:8]$

$rd[47:40] \leftarrow rs1[23:16]$

$rd[39:32] \leftarrow rs1[31:24]$

$rd[31:24] \leftarrow rs1[39:32]$

$rd[23:16] \leftarrow rs1[47:40]$

$rd[15:8] \leftarrow rs1[55:48]$

$rd[7:0] \leftarrow rs1[63:56]$

Permission:

M mode/S mode/U mode

**Exception:**

Illegal instruction.

**Instruction format:**

31	27	26	25	24	20	19	15	14	12	11	7	6	0
10000	01	00000	rs1	001	rd	0001011							

### 17.4.6 REVW: an instruction that reverses the byte order in a low 32-bit word

**Syntax:**

revw rd, rs1

**Operation:**

tmp[31:24]  $\leftarrow$  rs1[7:0]

tmp [23:16]  $\leftarrow$  rs1[15:8]

tmp [15:8]  $\leftarrow$  rs1[23:16]

tmp [7:0]  $\leftarrow$  rs1[31:24]

rd  $\leftarrow$  sign\_extend(tmp[31:0])

**Permission:**

M mode/S mode/U mode

**Exception:**

Illegal instruction.

**Instruction format:**

31	27	26	25	24	20	19	15	14	12	11	7	6	0
10010	00	00000	rs1	001	rd	0001011							

### 17.4.7 TST: an instruction that tests bits with the value of 0

**Syntax:**

tst rd, rs1, imm6

**Operation:**

if(rs1[imm6] == 1)

rd  $\leftarrow$  1

else

$$rd \leftarrow 0$$
**Permission:**

M mode/S mode/U mode

**Exception:**

Illegal instruction.

**Instruction format:**

31	26	25	20	19	15	14	12	11	7	6	0
100010			imm6		rs1		001		rd		0001011

**17.4.8 TSTNBZ: an instruction that tests bytes with the value of 0****Syntax:**

tstnbz rd, rs1

**Operation:**

$$rd[63:56] \leftarrow (rs1[63:56] == 0) ? 8'hff : 8'h0$$

$$rd[55:48] \leftarrow (rs1[55:48] == 0) ? 8'hff : 8'h0$$

$$rd[47:40] \leftarrow (rs1[47:40] == 0) ? 8'hff : 8'h0$$

$$rd[39:32] \leftarrow (rs1[39:32] == 0) ? 8'hff : 8'h0$$

$$rd[31:24] \leftarrow (rs1[31:24] == 0) ? 8'hff : 8'h0$$

$$rd[23:16] \leftarrow (rs1[23:16] == 0) ? 8'hff : 8'h0$$

$$rd[15:8] \leftarrow (rs1[15:8] == 0) ? 8'hff : 8'h0$$

$$rd[7:0] \leftarrow (rs1[7:0] == 0) ? 8'hff : 8'h0$$
**Permission:**

M mode/S mode/U mode

**Exception:**

Illegal instruction.

**Instruction format:**

31	27	26	25	24	20	19	15	14	12	11	7	6	0
10000		00		00000		rs1		001		rd		0001011	

## 17.5 Appendix B-5 Storage instructions

The storage instruction set extends storage instructions. Each instruction has 32 bits.

Arithmetic operation instructions in this instruction set are described in alphabetical order.

### 17.5.1 FLRD: a load doubleword instruction that shifts floating-point registers

#### Syntax:

`fldr rd, rs1, rs2, imm2`

#### Operation:

$rd \leftarrow \text{mem}[(rs1+rs2 \ll imm2)+7: (rs1+rs2 \ll imm2)]$

#### Permission:

M mode/S mode/U mode

#### Exception:

Unaligned access, access error, page error, or illegal instruction.

#### Notes:

If the value of `mxstatus.theadisaee` is 1' b0 or the value of `mstatus.fs` is 2' b00, executing this instruction causes an exception of illegal instruction.

#### Instruction format:

31	27	26	25	24	20	19	15	14	12	11	7	6	0
01100	imm2	rs2	rs1	110	rd	0001011							

### 17.5.2 FLRW: a load word instruction that shifts floating-point registers

#### Syntax:

`flrw rd, rs1, rs2, imm2`

#### Operation:

$rd \leftarrow \text{one\_extend}(\text{mem}[(rs1+rs2 \ll imm2)+3: (rs1+rs2 \ll imm2)])$

#### Permission:

M mode/S mode/U mode

#### Exception:

Unaligned access, access error, page error, or illegal instruction.

#### Notes:

If the value of `mxstatus.theadisaee` is 1' b0 or the value of `mstatus.fs` is 2' b00, executing this instruction causes an exception of illegal instruction.

**Instruction format:**

31	27	26	25	24	20	19	15	14	12	11	7	6	0
01000	imm2	rs2	rs1	110	rd	0001011							

### 17.5.3 FLURD: a load doubleword instruction that shifts low 32 bits of floating-point registers

**Syntax:**

`flurd rd, rs1, rs2, imm2`

**Operation:**

$rd \leftarrow \text{mem}[(rs1 + rs2[31:0] \ll imm2) + 7: (rs1 + rs2[31:0] \ll imm2)]$

**Permission:**

M mode/S mode/U mode

**Exception:**

Unaligned access, access error, page error, or illegal instruction.

**Notes:**

`rs2[31:0]` specifies an unsigned value. 0s are added to the high bits [63:32] for address calculation.

If the value of `mxstatus.theadisaee` is 1' b0 or the value of `mstatus.fs` is 2' b00, executing this instruction causes an exception of illegal instruction.

**Instruction format:**

31	27	26	25	24	20	19	15	14	12	11	7	6	0
01110	imm2	rs2	rs1	110	rd	0001011							

### 17.5.4 FLURW: a load word instruction that shifts low 32 bits of floating-point registers

**Syntax:**

`flurw rd, rs1, rs2, imm2`

**Operation:**

$rd \leftarrow \text{one\_extend}(\text{mem}[(rs1 + rs2[31:0] \ll imm2) + 3: (rs1 + rs2[31:0] \ll imm2)])$

**Permission:**

M mode/S mode/U mode

**Exception:**

Unaligned access, access error, page error, or illegal instruction.

**Notes:**

rs2[31:0] specifies an unsigned value. 0s are added to the high bits [63:32] for address calculation.

If the value of mxstatus.theadisaee is 1' b0 or the value of mstatus.fs is 2' b00, executing this instruction causes an exception of illegal instruction.

**Instruction format:**

31	27	26	25	24	20	19	15	14	12	11	7	6	0
01010	imm2	rs2	rs1	110	rd	0001011							

**17.5.5 FSRD: a store doubleword instruction that shifts floating-point registers****Syntax:**

```
fsrd rd, rs1, rs2, imm2
```

**Operation:**

$$\text{mem}[(\text{rs1} + \text{rs2} \ll \text{imm2}) + 7: (\text{rs1} + \text{rs2} \ll \text{imm2})] \leftarrow \text{rd}[63:0]$$
**Permission:**

M mode/S mode/U mode

**Exception:**

Unaligned access, access error, page error, or illegal instruction.

**Notes:**

If the value of mxstatus.theadisaee is 1' b0 or the value of mstatus.fs is 2' b00, executing this instruction causes an exception of illegal instruction.

**Instruction format:**

31	27	26	25	24	20	19	15	14	12	11	7	6	0
01100	imm2	rs2	rs1	111	rd	0001011							

**17.5.6 FSRW: a store word instruction that shifts floating-point registers.****Syntax:**

```
fsrw rd, rs1, rs2, imm2
```

**Operation:**

$$\text{mem}[(\text{rs1} + \text{rs2} \ll \text{imm2}) + 3: (\text{rs1} + \text{rs2} \ll \text{imm2})] \leftarrow \text{rd}[31:0]$$



**Permission:**

M mode/S mode/U mode

**Exception:**

Unaligned access, access error, page error, or illegal instruction.

**Notes:**

If the value of `mxstatus.theadisaee` is 1' b0 or the value of `mstatus.fs` is 2' b00, executing this instruction causes an exception of illegal instruction.

**Instruction format:**

31	27	26	25	24	20	19	15	14	12	11	7	6	0
01000	imm2	rs2	rs1	111	rd	0001011							

### 17.5.7 FSURD: a store doubleword instruction that shifts low 32 bits of floating-point registers

**Syntax:**

fsurd rd, rs1, rs2, imm2

**Operation:**

$$\text{mem}[(\text{rs1} + \text{rs2}[31:0] \ll \text{imm2}) + 7: (\text{rs1} + \text{rs2}[31:0] \ll \text{imm2})] \leftarrow \text{rd}[63:0]$$
**Permission:**

M mode/S mode/U mode

**Exception:**

Unaligned access, access error, page error, or illegal instruction.

**Notes:**

`rs2[31:0]` specifies an unsigned value. 0s are added to the high bits [63:32] for address calculation.

If the value of `mxstatus.theadisaee` is 1' b0 or the value of `mstatus.fs` is 2' b00, executing this instruction causes an exception of illegal instruction.

**Instruction format:**

31	27	26	25	24	20	19	15	14	12	11	7	6	0
01110	imm2	rs2	rs1	111	rd	0001011							

### 17.5.8 FSURW: a store word instruction that shifts low 32 bits of floating-point registers

**Syntax:**

fsurw rd, rs1, rs2, imm2

**Operation:**

$$\text{mem}[(\text{rs1} + \text{rs2}[31:0] \ll \text{imm2}) + 3: (\text{rs1} + \text{rs2}[31:0] \ll \text{imm2})] \leftarrow \text{rd}[31:0]$$

**Permission:**

M mode/S mode/U mode

**Exception:**

Unaligned access, access error, page error, or illegal instruction.

**Notes:**

rs2[31:0] specifies an unsigned value. 0s are added to the high bits [63:32] for address calculation.

If the value of mxstatus.theadisaee is 1' b0 or the value of mstatus.fs is 2' b00, executing this instruction causes an exception of illegal instruction.

**Instruction format:**

31	27	26	25	24	20	19	15	14	12	11	7	6	0
01010	imm2			rs2			rs1			111	rd		0001011

### 17.5.9 LBIA: a base-address auto-increment instruction that extends signed bits and loads bytes

**Syntax:**

lbia rd, (rs1), imm5,imm2

**Operation:**

$$\text{rd} \leftarrow \text{sign\_extend}(\text{mem}[\text{rs1}])$$

$$\text{rs1} \leftarrow \text{rs1} + \text{sign\_extend}(\text{imm5} \ll \text{imm2})$$

**Permission:**

M mode/S mode/U mode

**Exception:**

Unaligned access, access error, page error, or illegal instruction.

**Notes:**

The values of rd and rs1 must not be the same.

**Instruction format:**

31	27	26	25	24	20	19	15	14	12	11	7	6	0
00011	imm2			imm5			rs1			100	rd		0001011

### 17.5.10 LBIB: a load byte instruction that auto-increments the base address and extends signed bits

#### Syntax:

lbib rd, (rs1), imm5,imm2

#### Operation:

$rs1 \leftarrow rs1 + \text{sign\_extend}(imm5 \ll imm2)$

$rd \leftarrow \text{sign\_extend}(\text{mem}[rs1])$

#### Permission:

M mode/S mode/U mode

#### Exception:

Unaligned access, access error, page error, or illegal instruction.

#### Notes:

The values of rd and rs1 must not be the same.

#### Instruction format:

31	27	26	25	24	20	19	15	14	12	11	7	6	0
00001	imm2	imm5	rs1	100	rd	0001011							

### 17.5.11 LBUIA: a base-address auto-increment instruction that extends zero bits and loads bytes

#### Syntax:

lbuia rd, (rs1), imm5,imm2

#### Operation:

$rd \leftarrow \text{zero\_extend}(\text{mem}[rs1])$

$rs1 \leftarrow rs1 + \text{sign\_extend}(imm5 \ll imm2)$

#### Permission:

M mode/S mode/U mode

#### Exception:

Unaligned access, access error, page error, or illegal instruction.

#### Notes:

The values of rd and rs1 must not be the same.

**Instruction format:**

31	27	26	25	24	20	19	15	14	12	11	7	6	0
10011	imm2	imm5	rs1	100	rd	0001011							

### 17.5.12 LBUIB: a load byte instruction that auto-increments the base address and extends zero bits

**Syntax:**

lbuib rd, (rs1), imm5,imm2

**Operation:**

$rs1 \leftarrow rs1 + \text{sign\_extend}(\text{imm5} \ll \text{imm2})$

$rd \leftarrow \text{zero\_extend}(\text{mem}[rs1])$

**Permission:**

M mode/S mode/U mode

**Exception:**

Unaligned access, access error, page error, or illegal instruction.

**Notes:**

The values of rd and rs1 must not be the same.

**Instruction format:**

31	27	26	25	24	20	19	15	14	12	11	7	6	0
10001	imm2	imm5	rs1	100	rd	0001011							

### 17.5.13 LDD: an instruction that loads double registers

**Syntax:**

ldd rd1,rd2, (rs1),imm2

**Operation:**

$\text{address} \leftarrow rs1 + \text{zero\_extend}(\text{imm2} \ll 4)$

$rd1 \leftarrow \text{mem}[\text{address}+7:\text{address}]$

$rd2 \leftarrow \text{mem}[\text{address}+15:\text{address}+8]$

**Permission:**

M mode/S mode/U mode

**Exception:**

Unaligned access, access error, page error, or illegal instruction.

**Notes:**

The values of rd1, rd2, and rs1 must not be the same.

**Instruction format:**

31	27	26	25	24	20	19	15	14	12	11	7	6	0
11111	imm2	rd2	rs1	100	rd1	0001011							

### 17.5.14 LDIA: a base-address auto-increment instruction that loads doublewords and extends signed bits

**Syntax:**

ldia rd, (rs1), imm5,imm2

**Operation:**

$rd \leftarrow \text{sign\_extend}(\text{mem}[\text{rs1}+7:\text{rs1}])$

$\text{rs1} \leftarrow \text{rs1} + \text{sign\_extend}(\text{imm5} \ll \text{imm2})$

**Permission:**

M mode/S mode/U mode

**Exception:**

Unaligned access, access error, page error, or illegal instruction.

**Notes:**

The values of rd and rs1 must not be the same.

**Instruction format:**

31	27	26	25	24	20	19	15	14	12	11	7	6	0
01111	imm2	imm5	rs1	100	rd	0001011							

### 17.5.15 LDIB: a load doubleword instruction that auto-increments the base address and extends the signed bits

**Syntax:**

ldib rd, (rs1), imm5,imm2

**Operation:**

$$rs1 \leftarrow rs1 + \text{sign\_extend}(\text{imm5} \ll \text{imm2})$$

$$rd \leftarrow \text{sign\_extend}(\text{mem}[rs1+7:rs1])$$
**Permission:**

M mode/S mode/U mode

**Exception:**

Unaligned access, access error, page error, or illegal instruction.

**Notes:**

The values of rd and rs1 must not be the same.

**Instruction format:**

31	27	26	25	24	20	19	15	14	12	11	7	6	0
01101	imm2			imm5			rs1		100		rd		0001011

### 17.5.16 LHIA: a base-address auto-increment instruction that loads halfwords and extends signed bits

**Syntax:**

lhia rd, (rs1), imm5,imm2

**Operation:**

$$rd \leftarrow \text{sign\_extend}(\text{mem}[rs1+1:rs1])$$

$$rs1 \leftarrow rs1 + \text{sign\_extend}(\text{imm5} \ll \text{imm2})$$
**Permission:**

M mode/S mode/U mode

**Exception:**

Unaligned access, access error, page error, or illegal instruction.

**Notes:**

The values of rd and rs1 must not be the same.

**Instruction format:**

31	27	26	25	24	20	19	15	14	12	11	7	6	0
00111	imm2			imm5			rs1		100		rd		0001011

### 17.5.17 LHIB: a load halfword instruction that auto-increments the base address and extends signed bits

#### Syntax:

lhib rd, (rs1), imm5,imm2

#### Operation:

$$rs1 \leftarrow rs1 + \text{sign\_extend}(\text{imm5} \ll \text{imm2})$$

$$rd \leftarrow \text{sign\_extend}(\text{mem}[rs1+1:rs1])$$

#### Permission:

M mode/S mode/U mode

#### Exception:

Unaligned access, access error, page error, or illegal instruction.

#### Notes:

The values of rd and rs1 must not be the same.

#### Instruction format:

31	27	26	25	24	20	19	15	14	12	11	7	6	0
00101	imm2		imm5		rs1		100		rd		0001011		

### 17.5.18 LHUIA: a base-address auto-increment instruction that extends zero bits and loads halfwords

#### Syntax:

lhuiA rd, (rs1), imm5,imm2

#### Operation:

$$rd \leftarrow \text{zero\_extend}(\text{mem}[rs1+1:rs1])$$

$$rs1 \leftarrow rs1 + \text{sign\_extend}(\text{imm5} \ll \text{imm2})$$

#### Permission:

M mode/S mode/U mode

#### Exception:

Unaligned access, access error, page error, or illegal instruction.

#### Notes:

The values of rd and rs1 must not be the same.

**Instruction format:**

31	27	26	25	24	20	19	15	14	12	11	7	6	0
10111	imm2	imm5	rs1	100	rd	0001011							

### 17.5.19 LHUIB: a load halfword instruction that auto-increments the base address and extends zero bits

**Syntax:**

lhuib rd, (rs1), imm5,imm2

**Operation:**

$rs1 \leftarrow rs1 + \text{sign\_extend}(\text{imm5} \ll \text{imm2})$

$rd \leftarrow \text{zero\_extend}(\text{mem}[rs1+1:rs1])$

**Permission:**

M mode/S mode/U mode

**Exception:**

Unaligned access, access error, page error, or illegal instruction.

**Notes:**

The values of rd and rs1 must not be the same.

**Instruction format:**

31	27	26	25	24	20	19	15	14	12	11	7	6	0
10101	imm2	imm5	rs1	100	rd	0001011							

### 17.5.20 LRB: a load byte instruction that shifts registers and extends signed bits

**Syntax:**

lrb rd, rs1, rs2, imm2

**Operation:**

$rd \leftarrow \text{sign\_extend}(\text{mem}[(rs1+rs2 \ll \text{imm2})])$

**Permission:**

M mode/S mode/U mode

**Exception:**

Unaligned access, access error, page error, or illegal instruction.

**Instruction format:**



31	27	26	25	24	20	19	15	14	12	11	7	6	0
00000	imm2	rs2	rs1	100	rd	0001011							

### 17.5.21 LRBU: a load byte instruction that shifts registers and extends zero bits

#### Syntax:

lrbu rd, rs1, rs2, imm2

#### Operation:

$rd \leftarrow \text{zero\_extend}(\text{mem}[(rs1+rs2 \ll \text{imm2}]])$

#### Permission:

M mode/S mode/U mode

#### Exception:

Unaligned access, access error, page error, or illegal instruction.

#### Instruction format:

31	27	26	25	24	20	19	15	14	12	11	7	6	0
10000	imm2	rs2	rs1	100	rd	0001011							

### 17.5.22 LRD: a load doubleword instruction that shifts registers

#### Syntax:

lrd rd, rs1, rs2, imm2

#### Operation:

$rd \leftarrow \text{mem}[(rs1+rs2 \ll \text{imm2})+7: (rs1+rs2 \ll \text{imm2})]$

#### Permission:

M mode/S mode/U mode

#### Exception:

Unaligned access, access error, page error, or illegal instruction.

#### Instruction format:

31	27	26	25	24	20	19	15	14	12	11	7	6	0
01100	imm2	rs2	rs1	100	rd	0001011							

### 17.5.23 LRH: a load halfword instruction that shifts registers and extends signed bits

#### Syntax:

lrh rd, rs1, rs2, imm2

**Operation:**

$rd \leftarrow \text{sign\_extend}(\text{mem}[(rs1+rs2 \ll imm2)+1: (rs1+rs2 \ll imm2)])$

**Permission:**

M mode/S mode/U mode

**Exception:**

Unaligned access, access error, page error, or illegal instruction.

**Instruction format:**

31	27	26	25	24	20	19	15	14	12	11	7	6	0
00100	imm2	rs2	rs1	100	rd	0001011							

### 17.5.24 LRHU: a load halfword instruction that shifts registers and extends zero bits

**Syntax:**

lrhu rd, rs1, rs2, imm2

**Operation:**

$rd \leftarrow \text{zero\_extend}(\text{mem}[(rs1+rs2 \ll imm2)+1: (rs1+rs2 \ll imm2)])$

**Permission:**

M mode/S mode/U mode

**Exception:**

Unaligned access, access error, page error, or illegal instruction.

**Instruction format:**

31	27	26	25	24	20	19	15	14	12	11	7	6	0
10100	imm2	rs2	rs1	100	rd	0001011							

### 17.5.25 LRW: a load word instruction that shifts registers and extends signed bits

**Syntax:**

lrw rd, rs1, rs2, imm2

**Operation:**

$rd \leftarrow \text{sign\_extend}(\text{mem}[(rs1+rs2 \ll imm2)+3: (rs1+rs2 \ll imm2)])$

**Permission:**

M mode/S mode/U mode

**Exception:**

Unaligned access, access error, page error, or illegal instruction.

**Instruction format:**

31	27	26	25	24	20	19	15	14	12	11	7	6	0
01000	imm2	rs2	rs1	100	rd	0001011							

**17.5.26 LRWU: a load word instruction that shifts registers and extends zero bits****Syntax:**

lrwu rd, rs1, rs2, imm2

**Operation:**

$rd \leftarrow \text{zero\_extend}(\text{mem}[(rs1+rs2 \ll \text{imm2})+3: (rs1+rs2 \ll \text{imm2})])$

**Permission:**

M mode/S mode/U mode

**Exception:**

Unaligned access, access error, page error, or illegal instruction.

**Instruction format:**

31	27	26	25	24	20	19	15	14	12	11	7	6	0
11000	imm2	rs2	rs1	100	rd	0001011							

**17.5.27 LURB: a load byte instruction that shifts low 32 bits of registers and extends signed bits****Syntax:**

lurb rd, rs1, rs2, imm2

**Operation:**

$rd \leftarrow \text{sign\_extend}(\text{mem}[(rs1+rs2[31:0] \ll \text{imm2})])$

**Permission:**

M mode/S mode/U mode

**Exception:**

Unaligned access, access error, page error, or illegal instruction.

**Notes:**

rs2[31:0] specifies an unsigned value. 0s are added to the high bits [63:32] for address calculation.

**Instruction format:**

31	27	26	25	24	20	19	15	14	12	11	7	6	0
00010		imm2		rs2		rs1		100		rd		0001011	

### 17.5.28 LURBU: a load byte instruction that shifts low 32 bits of registers and extends zero bits

**Syntax:**

`lurbu rd, rs1, rs2, imm2`

**Operation:**

$rd \leftarrow \text{zero\_extend}(\text{mem}[(rs1+rs2[31:0] \ll imm2)])$

**Permission:**

M mode/S mode/U mode

**Exception:**

Unaligned access, access error, page error, or illegal instruction.

**Notes:**

`rs2[31:0]` specifies an unsigned value. 0s are added to the high bits [63:32] for address calculation.

**Instruction format:**

31	27	26	25	24	20	19	15	14	12	11	7	6	0
10010		imm2		rs2		rs1		100		rd		0001011	

### 17.5.29 LURD: a load doubleword instruction that shifts low 32 bits of registers

**Syntax:**

`lurd rd, rs1, rs2, imm2`

**Operation:**

$rd \leftarrow \text{mem}[(rs1+rs2[31:0] \ll imm2)+7: (rs1+rs2[31:0] \ll imm2)]$

**Permission:**

M mode/S mode/U mode

**Exception:**

Unaligned access, access error, page error, or illegal instruction.

**Notes:**

`rs2[31:0]` specifies an unsigned value. 0s are added to the high bits [63:32] for address calculation.

Instruction format:

31	27	26	25	24	20	19	15	14	12	11	7	6	0
01110	imm2	rs2	rs1	100	rd	0001011							

### 17.5.30 LURH: a load halfword instruction that shifts low 32 bits of registers and extends signed bits

Syntax:

`lurh rd, rs1, rs2, imm2`

Operation:

$rd \leftarrow \text{sign\_extend}(\text{mem}[(rs1+rs2[31:0] \ll imm2)+1: (rs1+rs2[31:0] \ll imm2)])$

Permission:

M mode/S mode/U mode

Exception:

Unaligned access, access error, page error, or illegal instruction.

Notes:

$rs2[31:0]$  specifies an unsigned value. 0s are added to the high bits [63:32] for address calculation.

Instruction format:

31	27	26	25	24	20	19	15	14	12	11	7	6	0
00110	imm2	rs2	rs1	100	rd	0001011							

### 17.5.31 LURHU: a load halfword instruction that shifts low 32 bits of registers and extends zero bits

Syntax:

`lurhu rd, rs1, rs2, imm2`

Operation:

$rd \leftarrow \text{zero\_extend}(\text{mem}[(rs1+rs2[31:0] \ll imm2)+1: (rs1+rs2[31:0] \ll imm2)])$

Permission:

M mode/S mode/U mode

**Exception:**

Unaligned access, access error, page error, or illegal instruction.

**Notes:**

rs2[31:0] specifies an unsigned value. 0s are added to the high bits [63:32] for address calculation.

**Instruction format:**

31	27	26	25	24	20	19	15	14	12	11	7	6	0
10110	imm2	rs2	rs1	100	rd	0001011							

### 17.5.32 LURW: a load word instruction that shifts low 32 bits of registers and extends signed bits

**Syntax:**

lurw rd, rs1, rs2, imm2

**Operation:**

$$rd \leftarrow \text{sign\_extend}(\text{mem}[(rs1 + rs2[31:0] \ll imm2) + 3: \\ (rs1 + rs2[31:0] \ll imm2)])$$
**Permission:**

M mode/S mode/U mode

**Exception:**

Unaligned access, access error, page error, or illegal instruction.

**Notes:**

rs2[31:0] specifies an unsigned value. 0s are added to the high bits [63:32] for address calculation.

**Instruction format:**

31	27	26	25	24	20	19	15	14	12	11	7	6	0
01010	imm2	rs2	rs1	100	rd	0001011							

### 17.5.33 LURWU: a load word instruction that shifts 32 bits of registers and extends zero bits

**Syntax:**

lwd rd1, rd2, (rs1),imm2

**Operation:**

```

address ← rs1 + zero_extend(imm2 << 3)

rd1 ← sign_extend(mem[address+3: address])

rd2 ← sign_extend(mem[address+7: address+4])

```

**Permission:**

M mode/S mode/U mode

**Exception:**

Unaligned access, access error, page error, or illegal instruction.

**Notes:**

The values of rd1, rd2, and rs1 must not be the same.

**Instruction format:**

31	27	26	25	24	20	19	15	14	12	11	7	6	0
11010	imm2	rs2	rs1	100	rd	0001011							

### 17.5.34 LWD: a load word instruction that loads double registers and extends signed bits

**Syntax:**

```
lwd rd, imm7(rs1)
```

**Operation:**

```

address ← rs1 + sign_extend(imm7)

rd ← sign_extend(mem[address+31: address])

rd+1 ← sign_extend(mem[address+63: address+32])

```

**Permission:**

M mode/S mode/U mode

**Exception:**

Unaligned access, access error, page error, or illegal instruction.

**Instruction format:**

31	27	26	25	24	20	19	15	14	12	11	7	6	0
11100	imm2	rd2	rs1	100	rd1	0001011							

### 17.5.35 LWIA: a base-address auto-increment instruction that extends signed bits and loads words

#### Syntax:

lwia rd, (rs1), imm5,imm2

#### Operation:

$rd \leftarrow \text{sign\_extend}(\text{mem}[\text{rs1}+3:\text{rs1}])$

$\text{rs1} \leftarrow \text{rs1} + \text{sign\_extend}(\text{imm5} \ll \text{imm2})$

#### Permission:

M mode/S mode/U mode

#### Exception:

Unaligned access, access error, page error, or illegal instruction.

#### Notes:

The values of rd and rs1 must not be the same.

#### Instruction format:

31	27	26	25	24	20	19	15	14	12	11	7	6	0
01011	imm2			imm5			rs1			100	rd		0001011

### 17.5.36 LWIB: a load word instruction that auto-increments the base address and extends signed bits

#### Syntax:

lwib rd, (rs1), imm5,imm2

#### Operation:

$\text{rs1} \leftarrow \text{rs1} + \text{sign\_extend}(\text{imm5} \ll \text{imm2})$

$rd \leftarrow \text{sign\_extend}(\text{mem}[\text{rs1}+3:\text{rs1}])$

#### Permission:

M mode/S mode/U mode

#### Exception:

Unaligned access, access error, page error, or illegal instruction.

#### Notes:

The values of rd and rs1 must not be the same.



Instruction format:

31	27	26	25	24	20	19	15	14	12	11	7	6	0
01001	imm2			imm5		rs1		100		rd			0001011

### 17.5.37 LWUD: a load word instruction that loads double registers and extends zero bits

Syntax:

lwud rd1,rd2, (rs1),imm2

Operation:

address  $\leftarrow$  rs1 + zero\_extend(imm2 << 3)

rd1  $\leftarrow$  zero\_extend(mem[address+3: address])

rd2  $\leftarrow$  zero\_extend(mem[address+7: address+4])

Permission:

M mode/S mode/U mode

Exception:

Unaligned access, access error, page error, or illegal instruction.

Notes:

The values of rd1, rd2, and rs1 must not be the same.

Instruction format:

31	27	26	25	24	20	19	15	14	12	11	7	6	0
11110	imm2			rd2		rs1		100		rd1			0001011

### 17.5.38 LWUIA: a base-address auto-increment instruction that extends zero bits and loads words

Syntax:

lwuia rd, (rs1), imm5,imm2

Operation:

rd  $\leftarrow$  zero\_extend(mem[rs1+3:rs1])

rs1  $\leftarrow$  rs1 + sign\_extend(imm5 << imm2)

Permission:

M mode/S mode/U mode

**Exception:**

Unaligned access, access error, page error, or illegal instruction.

**Notes:**

The values of rd and rs1 must not be the same.

**Instruction format:**

31	27	26	25	24	20	19	15	14	12	11	7	6	0
11011	imm2	imm5	rs1	100	rd	0001011							

### 17.5.39 LWUIB: a load word instruction that auto-increments the base address and extends zero bits

**Syntax:**

lwuib rd, (rs1), imm5,imm2

**Operation:**

$rs1 \leftarrow rs1 + \text{sign\_extend}(imm5 \ll imm2)$

$rd \leftarrow \text{zero\_extend}(\text{mem}[rs1+3:rs1])$

**Permission:**

M mode/S mode/U mode

**Exception:**

Unaligned access, access error, page error, or illegal instruction.

**Notes:**

The values of rd and rs1 must not be the same.

**Instruction format:**

31	27	26	25	24	20	19	15	14	12	11	7	6	0
11001	imm2	imm5	rs1	100	rd	0001011							

### 17.5.40 SBIA: a base-address auto-increment instruction that stores bytes

**Syntax:**

sbia rs2, (rs1), imm5,imm2

**Operation:**

$\text{mem}[rs1] \leftarrow rs2[7:0]$

$rs1 \leftarrow rs1 + \text{sign\_extend}(imm5 \ll imm2)$

**Permission:**

M mode/S mode/U mode

**Exception:**

Unaligned access, access error, page error, or illegal instruction.

**Instruction format:**

31	27	26	25	24	20	19	15	14	12	11	7	6	0
00011	imm2	imm5	rs1	101	rs2	0001011							

**17.5.41 SBIB: a store byte instruction that auto-increments the base address****Syntax:**

sbib rs2, (rs1), imm5,imm2

**Operation:** $rs1 \leftarrow rs1 + \text{sign\_extend}(imm5 \ll imm2)$  $mem[rs1] \leftarrow rs2[7:0]$ **Permission:**

M mode/S mode/U mode

**Exception:**

Unaligned access, access error, page error, or illegal instruction.

**Instruction format:**

31	27	26	25	24	20	19	15	14	12	11	7	6	0
00001	imm2	imm5	rs1	101	rs2	0001011							

**17.5.42 SDD: an instruction that stores double registers****Syntax:**

sdd rd1,rd2, (rs1),imm2

**Operation:** $address \leftarrow rs1 + \text{zero\_extend}(imm2 \ll 4)$  $mem[address+7:address] \leftarrow rd1$  $mem[address+15:address+8] \leftarrow rd2$ **Permission:**

M mode/S mode/U mode

**Exception:**

Unaligned access, access error, page error, or illegal instruction.

**Instruction format:**

31	27	26	25	24	20	19	15	14	12	11	7	6	0
11111	imm2	rd2	rs1	101	rd1	0001011							

**17.5.43 SDIA: a base-address auto-increment instruction that stores doublewords****Syntax:**

sdia rs2, (rs1), imm5,imm2

**Operation:**

$\text{mem}[\text{rs1}+7:\text{rs1}] \leftarrow \text{rs2}[63:0]$

$\text{rs1} \leftarrow \text{rs1} + \text{sign\_extend}(\text{imm5} \ll \text{imm2})$

**Permission:**

M mode/S mode/U mode

**Exception:**

Unaligned access, access error, page error, or illegal instruction.

**Instruction format:**

31	27	26	25	24	20	19	15	14	12	11	7	6	0
01111	imm2	imm5	rs1	101	rs2	0001011							

**17.5.44 SDIB: a store doubleword instruction that auto-increments the base address****Syntax:**

sdib rs2, (rs1), imm5,imm2

**Operation:**

$\text{rs1} \leftarrow \text{rs1} + \text{sign\_extend}(\text{imm5} \ll \text{imm2})$

$\text{mem}[\text{rs1}+7:\text{rs1}] \leftarrow \text{rs2}[63:0]$

**Permission:**

M mode/S mode/U mode

**Exception:**

Unaligned access, access error, page error, or illegal instruction.

**Instruction format:**

31	27	26	25	24	20	19	15	14	12	11	7	6	0
01101	imm2	imm5	rs1	101	rs2	0001011							

### 17.5.45 SHIA: a base-address auto-increment instruction that stores halfwords

#### Syntax:

shia rs2, (rs1), imm5,imm2

#### Operation:

$\text{mem}[\text{rs1}+1:\text{rs1}] \leftarrow \text{rs2}[15:0]$

$\text{rs1} \leftarrow \text{rs1} + \text{sign\_extend}(\text{imm5} \ll \text{imm2})$

#### Permission:

M mode/S mode/U mode

#### Exception:

Unaligned access, access error, page error, or illegal instruction.

#### Instruction format:

31	27	26	25	24	20	19	15	14	12	11	7	6	0
00111	imm2	imm5	rs1	101	rs2	0001011							

### 17.5.46 SHIB: a store halfword instruction that auto-increments the base address

#### Syntax:

shib rs2, (rs1), imm5,imm2

#### Operation:

$\text{rs1} \leftarrow \text{rs1} + \text{sign\_extend}(\text{imm5} \ll \text{imm2})$

$\text{mem}[\text{rs1}+1:\text{rs1}] \leftarrow \text{rs2}[15:0]$

#### Permission:

M mode/S mode/U mode

#### Exception:

Unaligned access, access error, page error, or illegal instruction.

#### Instruction format:

31	27	26	25	24	20	19	15	14	12	11	7	6	0
00101	imm2	imm5	rs1	101	rs2	0001011							

### 17.5.47 SRB: a store byte instruction that shifts registers

**Syntax:**

srb rd, rs1, rs2, imm2

**Operation:**

$\text{mem}[(\text{rs1} + \text{rs2} \ll \text{imm2})] \leftarrow \text{rd}[7:0]$

**Permission:**

M mode/S mode/U mode

**Exception:**

Unaligned access, access error, page error, or illegal instruction.

**Instruction format:**

31	27	26	25	24	20	19	15	14	12	11	7	6	0
00000	imm2			imm5			rs1		101		rd		0001011

### 17.5.48 SRD: a store doubleword instruction that shifts registers

**Syntax:**

srd rd, rs1, rs2, imm2

**Operation:**

$\text{mem}[(\text{rs1} + \text{rs2} \ll \text{imm2}) + 7: (\text{rs1} + \text{rs2} \ll \text{imm2})] \leftarrow \text{rd}[63:0]$

**Permission:**

M mode/S mode/U mode

**Exception:**

Unaligned access, access error, page error, or illegal instruction.

**Instruction format:**

31	27	26	25	24	20	19	15	14	12	11	7	6	0
01100	imm2			rs2			rs1		101		rd		0001011

### 17.5.49 SRH: a store halfword instruction that shifts registers

**Syntax:**

srh rd, rs1, rs2, imm2

**Operation:**

$$\text{mem}[(\text{rs1}+\text{rs2}<<\text{imm2})+1: (\text{rs1}+\text{rs2}<<\text{imm2})] \leftarrow \text{rd}[15:0]$$
**Permission:**

M mode/S mode/U mode

**Exception:**

Unaligned access, access error, page error, or illegal instruction.

**Instruction format:**

31	27	26	25	24	20	19	15	14	12	11	7	6	0
00100	imm2	rs2	rs1	101	rd	0001011							

**17.5.50 SRW: a store word instruction that shifts registers****Syntax:**

srw rd, rs1, rs2, imm2

**Operation:**

$$\text{mem}[(\text{rs1}+\text{rs2}<<\text{imm2})+3: (\text{rs1}+\text{rs2}<<\text{imm2})] \leftarrow \text{rd}[31:0]$$
**Permission:**

M mode/S mode/U mode

**Exception:**

Unaligned access, access error, page error, or illegal instruction.

**Instruction format:**

31	27	26	25	24	20	19	15	14	12	11	7	6	0
01000	imm2	rs2	rs1	101	rd	0001011							

**17.5.51 SURB: a store byte instruction that shifts low 32 bits of registers****Syntax:**

surb rd, rs1, rs2, imm2

**Operation:**

$$\text{mem}[(\text{rs1}+\text{rs2}[31:0]<<\text{imm2})] \leftarrow \text{rd}[7:0]$$
**Permission:**

M mode/S mode/U mode

**Exception:**

Unaligned access, access error, page error, or illegal instruction.

**Notes:**

rs2[31:0] specifies an unsigned value. 0s are added to the high bits [63:32] for address calculation.

**Instruction format:**

31	27	26	25	24	20	19	15	14	12	11	7	6	0
00010		imm2		rs2		rs1		101		rd		0001011	

**17.5.52 SURD: a store doubleword instruction that shifts low 32 bits of registers****Syntax:**

surd rd, rs1, rs2, imm2

**Operation:**

$\text{mem}[(\text{rs1} + \text{rs2}[31:0] \ll \text{imm2}) + 7: (\text{rs1} + \text{rs2}[31:0] \ll \text{imm2})] \leftarrow \text{rd}[63:0]$

**Permission:**

M mode/S mode/U mode

**Exception:**

Unaligned access, access error, page error, or illegal instruction.

**Notes:**

rs2[31:0] specifies an unsigned value. 0s are added to the high bits [63:32] for address calculation.

**Instruction format:**

31	27	26	25	24	20	19	15	14	12	11	7	6	0
01110		imm2		rs2		rs1		101		rd		0001011	

**17.5.53 SURH: a store halfword instruction that shifts low 32 bits of registers****Syntax:**

surh rd, rs1, rs2, imm2

**Operation:**

$\text{mem}[(\text{rs1} + \text{rs2}[31:0] \ll \text{imm2}) + 1: (\text{rs1} + \text{rs2}[31:0] \ll \text{imm2})] \leftarrow \text{rd}[15:0]$

**Permission:**

M mode/S mode/U mode

**Exception:**

Unaligned access, access error, page error, or illegal instruction.

**Notes:**



rs2[31:0] specifies an unsigned value. 0s are added to the high bits [63:32] for address calculation.

**Instruction format:**

31	27	26	25	24	20	19	15	14	12	11	7	6	0
00110	imm2			rs2			rs1		101		rd		0001011

### 17.5.54 SURW: a store word instruction that shifts low 32 bits of registers

**Syntax:**

surw rd, rs1, rs2, imm2

**Operation:**

$\text{mem}[(\text{rs1} + \text{rs2}[31:0] \ll \text{imm2}) + 3: (\text{rs1} + \text{rs2}[31:0] \ll \text{imm2})] \leftarrow \text{rd}[31:0]$

**Permission:**

M mode/S mode/U mode

**Exception:**

Unaligned access, access error, page error, or illegal instruction.

**Notes:**

rs2[31:0] specifies an unsigned value. 0s are added to the high bits [63:32] for address calculation.

**Instruction format:**

31	27	26	25	24	20	19	15	14	12	11	7	6	0
01010	imm2			rs2			rs1		101		rd		0001011

### 17.5.55 SWIA: a base-address auto-increment instruction that stores words

**Syntax:**

swia rs2, (rs1), imm5, imm2

**Operation:**

$\text{mem}[\text{rs1} + 3: \text{rs1}] \leftarrow \text{rs2}[31:0]$

$\text{rs1} \leftarrow \text{rs1} + \text{sign\_extend}(\text{imm5} \ll \text{imm2})$

**Permission:**

M mode/S mode/U mode

**Exception:**

Unaligned access, access error, page error, or illegal instruction.

**Instruction format:**

31	27	26	25	24	20	19	15	14	12	11	7	6	0
01011	imm2	imm5	rs1	101	rs2	0001011							

### 17.5.56 SWIB: a store word instruction that auto-increments the base address

#### Syntax:

swib rs2, (rs1), imm5,imm2

#### Operation:

$rs1 \leftarrow rs1 + \text{sign\_extend}(imm5 \ll imm2)$

$\text{mem}[rs1+3:rs1] \leftarrow rs2[31:0]$

#### Permission:

M mode/S mode/U mode

#### Exception:

Unaligned access, access error, page error, or illegal instruction.

#### Instruction format:

31	27	26	25	24	20	19	15	14	12	11	7	6	0
01001	imm2	imm5	rs1	101	rs2	0001011							

### 17.5.57 SWD: an instruction that stores the low 32 bits of double registers

#### Syntax:

swd rd1,rd2,(rs1),imm2

#### Operation:

$\text{address} \leftarrow rs1 + \text{zero\_extend}(imm2 \ll 3)$

$\text{mem}[\text{address}+3:\text{address}] \leftarrow rd1[31:0]$

$\text{mem}[\text{address}+7:\text{address}+4] \leftarrow rd2[31:0]$

#### Permission:

M mode/S mode/U mode

#### Exception:

Unaligned access, access error, page error, or illegal instruction.

#### Instruction format:

31	27	26	25	24	20	19	15	14	12	11	7	6	0
11100	imm2	rd2	rs1	101	rd1	0001011							

## 17.6 Appendix B-6 Half-precision floating-point instructions

You can use instructions in this instruction set to process floating-point half-precision data. Each instruction has 32 bits. Instructions in this instruction set are described in alphabetical order.

### 17.6.1 FADD.H: a half-precision floating-point add instruction

**Syntax:**

fadd.h fd, fs1, fs2, rm

**Operation:**

$fd \leftarrow fs1 + fs2$

**Permission:**

M mode/S mode/U mode

**Exception:**

Illegal instruction.

**Affected flag bits:**

Floating-point status bits NV, OF, and NX

**Notes:**

RM determines the round-off mode:

- 3' b000: Rounds off to the nearest even number. The corresponding assembler instruction is fadd.h fd, fs1, fs2, rne.
- 3' b001: Rounds off to zero. The corresponding assembler instruction is fadd.h fd, fs1, fs2, rtz.
- 3' b010: Rounds off to negative infinity. The corresponding assembler instruction is fadd.h fd, fs1, fs2, rdn.
- 3' b011: Rounds off to positive infinity. The corresponding assembler instruction is fadd.h fd, fs1, fs2, rup.
- 3' b100: Rounds off to the nearest large value. The corresponding assembler instruction is fadd.h fd, fs1, fs2, rmm.
- 3' b101: This code is reserved and not used.
- 3' b110: This code is reserved and not used.
- 3' b111: Dynamically rounds off based on the rm bit of the fcsr register. The corresponding assembler instruction is fadd.h fd, fs1, fs2.

**Instruction format:**

31	25	24	20	19	15	14	12	11	7	6	0
0000010		fs2		fs1		rm		fd		1010011	

### 17.6.2 FCLASS.H: a half-precision floating-point classification instruction

#### Syntax:

fclass.h rd, fs1

#### Operation:

if ( fs1 = -inf)

rd  $\leftarrow$  64' h1

if ( fs1 = -norm)

rd  $\leftarrow$  64' h2

if ( fs1 = -subnorm)

rd  $\leftarrow$  64' h4

if ( fs1 = -zero)

rd  $\leftarrow$  64' h8

if ( fs1 = +zero)

rd  $\leftarrow$  64' h10

if ( fs1 = +subnorm)

rd  $\leftarrow$  64' h20

if ( fs1 = +norm)

rd  $\leftarrow$  64' h40

if ( fs1 = +inf)

rd  $\leftarrow$  64' h80

if ( fs1 = sNaN)

rd  $\leftarrow$  64' h100

if ( fs1 = qNaN)

rd  $\leftarrow$  64' h200

#### Permission:

M mode/S mode/U mode

**Exception:**

Illegal instruction.

**Affected flag bits:**

None

**Instruction format:**

31	25	24	20	19	15	14	12	11	7	6	0
1110010	00000	fs1	001	rd	1010011						

### 17.6.3 FCVT.H.L: an instruction that converts a signed long integer into a half-precision floating-point number

**Syntax:**

fcvt.h.l fd, rs1, rm

**Operation:**

$fd \leftarrow \text{signed\_long\_convert\_to\_half}(rs1)$

**Permission:**

M mode/S mode/U mode

**Exception:**

Illegal instruction.

**Affected flag bits:**

Floating-point status bits NX and OF

**Notes:**

RM determines the round-off mode:

- 3' b000: Rounds off to the nearest even number. The corresponding assembler instruction is fcvt.h.l fd, rs1, rne.
- 3' b001: Rounds off to zero. The corresponding assembler instruction is fcvt.h.l fd, rs1, rtz.
- 3' b010: Rounds off to negative infinity. The corresponding assembler instruction is fcvt.h.l fd, rs1, fdn.
- 3' b011: Rounds off to positive infinity. The corresponding assembler instruction is fcvt.h.l fd, rs1, rup.
- 3' b100: Rounds off to the nearest large value. The corresponding assembler instruction is fcvt.h.l fd, rs1, rmm.
- 3' b101: This code is reserved and not used.

- 3' b110: This code is reserved and not used.
- 3' b111: Dynamically rounds off based on the rm bit of the fcsr register. The corresponding assembler instruction is `fcvt.h.l fd, rs1`.

**Instruction format:**

31	25	24	20	19	15	14	12	11	7	6	0
1101010	00010	rs1	rm	fd	1010011						

### 17.6.4 FCVT.H.LU: an instruction that converts an unsigned long integer into a half-precision floating-point number

**Syntax:**

`fcvt.h.lu fd, rs1, rm`

**Operation:**

$fd \leftarrow \text{unsigned\_long\_convert\_to\_half\_fp}(rs1)$

**Permission:**

M mode/S mode/U mode

**Exception:**

Illegal instruction.

**Affected flag bits:**

Floating-point status bits NX and OF

**Notes:**

RM determines the round-off mode:

- 3' b000: Rounds off to the nearest even number. The corresponding assembler instruction is `fcvt.h.lu fd, rs1, rne`.
- 3' b001: Rounds off to zero. The corresponding assembler instruction is `fcvt.h.lu fd, rs1, rtz`.
- 3' b010: Rounds off to negative infinity. The corresponding assembler instruction is `fcvt.h.lu fd, rs1, fdn`.
- 3' b011: Rounds off to positive infinity. The corresponding assembler instruction is `fcvt.h.lu fd, rs1, rup`.
- 3' b100: Rounds off to the nearest large value. The corresponding assembler instruction is `fcvt.h.lu fd, rs1, rmm`.
- 3' b101: This code is reserved and not used.
- 3' b110: This code is reserved and not used.

- 3' b111: Dynamically rounds off based on the rm bit of the fcsr register. The corresponding assembler instruction is `fcvt.h.lu fd, rs1`.

**Instruction format:**

31	25	24	20	19	15	14	12	11	7	6	0
1101010	00011	rs1	rm	fd	1010011						

### 17.6.5 FCVT.H.S: an instruction that converts single precision floating-point data to half-precision floating-point data

**Syntax:**

`fcvt.h.s fd, fs1, rm`

**Operation:**

$fd \leftarrow \text{single\_convert\_to\_half}(fs1)$

**Permission:**

M mode/S mode/U mode

**Exception:**

Illegal instruction.

**Affected flag bits:**

Floating-point status bits NV, OF, UF, and NX

**Notes:**

RM determines the round-off mode:

- 3' b000: Rounds off to the nearest even number. The corresponding assembler instruction is `fcvt.h.s fd, fs1, rne`.
- 3' b001: Rounds off to zero. The corresponding assembler instruction is `fcvt.h.s fd, fs1, rtz`.
- 3' b010: Rounds off to negative infinity. The corresponding assembler instruction is `fcvt.h.s fd, fs1, fdn`.
- 3' b011: Rounds off to positive infinity. The corresponding assembler instruction is `fcvt.h.s fd, fs1, rup`.
- 3' b100: Rounds off to the nearest large value. The corresponding assembler instruction is `fcvt.h.s fd, fs1, rmm`.
- 3' b101: This code is reserved and not used.
- 3' b110: This code is reserved and not used.

- 3' b111: Dynamically rounds off based on the rm bit of the fcsr register. The corresponding assembler instruction is fcvt.h.s fd, fs1.

**Instruction format:**

31	25	24	20	19	15	14	12	11	7	6	0
0100010	00000	fs1	rm	fd	1010011						

### 17.6.6 FCVT.H.W: an instruction that converts a signed integer into a half-precision floating-point number

**Syntax:**

fcvt.h.w fd, rs1, rm

**Operation:**

$fd \leftarrow \text{signed\_int\_convert\_to\_half}(rs1)$

**Permission:**

M mode/S mode/U mode

**Exception:**

Illegal instruction.

**Affected flag bits:**

Floating-point status bits NX and OF

**Notes:**

RM determines the round-off mode:

- 3' b000: Rounds off to the nearest even number. The corresponding assembler instruction is fcvt.h.w fd, rs1, rne.
- 3' b001: Rounds off to zero. The corresponding assembler instruction is fcvt.h.w fd, rs1, rtz.
- 3' b010: Rounds off to negative infinity. The corresponding assembler instruction is fcvt.h.w fd, rs1, fdn.
- 3' b011: Rounds off to positive infinity. The corresponding assembler instruction is fcvt.h.w fd, rs1, rup.
- 3' b100: Rounds off to the nearest large value. The corresponding assembler instruction is fcvt.h.w fd, rs1, rmm.
- 3' b101: This code is reserved and not used.
- 3' b110: This code is reserved and not used.



- 3' b111: Dynamically rounds off based on the rm bit of the fcsr register. The corresponding assembler instruction is `fcvt.h.w fd, rs1`.

**Instruction format:**

31	25	24	20	19	15	14	12	11	7	6	0
1101010	00000	rs1	rm	fd	1010011						

### 17.6.7 FCVT.H.WU: an instruction that converts an unsigned integer into a half-precision floating-point number

**Syntax:**

`fcvt.h.wu fd, rs1, rm`

**Operation:**

$fd \leftarrow \text{unsigned\_int\_convert\_to\_half\_fp}(rs1)$

**Permission:**

M mode/S mode/U mode

**Exception:**

Illegal instruction.

**Affected flag bits:**

Floating-point status bits NX and OF

**Notes:**

RM determines the round-off mode:

- 3' b000: Rounds off to the nearest even number. The corresponding assembler instruction is `fcvt.h.wu fd, rs1, rne`.
- 3' b001: Rounds off to zero. The corresponding assembler instruction is `fcvt.h.wu fd, rs1, rtz`.
- 3' b010: Rounds off to negative infinity. The corresponding assembler instruction is `fcvt.h.wu fd, rs1, fdn`.
- 3' b011: Rounds off to positive infinity. The corresponding assembler instruction is `fcvt.h.wu fd, rs1, rup`.
- 3' b100: Rounds off to the nearest large value. The corresponding assembler instruction is `fcvt.h.wu fd, rs1, rmm`.
- 3' b101: This code is reserved and not used.
- 3' b110: This code is reserved and not used.

- 3' b111: Dynamically rounds off based on the rm bit of the fcsr register. The corresponding assembler instruction is `fcvt.h.wu fd, rs1`.

**Instruction format:**

31	25	24	20	19	15	14	12	11	7	6	0
1101010	00001	rs1	rm	fd	1010011						

### 17.6.8 FCVT.L.H: an instruction that converts a half-precision floating-point number to a signed long integer

**Syntax:**

`fcvt.l.h rd, fs1, rm`

**Operation:**

$rd \leftarrow \text{half\_convert\_to\_signed\_long}(fs1)$

**Permission:**

M mode/S mode/U mode

**Exception:**

Illegal instruction.

**Affected flag bits:**

Floating-point status bits NV and NX

**Notes:**

RM determines the round-off mode:

- 3' b000: Rounds off to the nearest even number. The corresponding assembler instruction is `fcvt.l.h rd, fs1, rne`.
- 3' b001: Rounds off to zero. The corresponding assembler instruction is `fcvt.l.h rd, fs1, rtz`.
- 3' b010: Rounds off to negative infinity. The corresponding assembler instruction is `fcvt.l.h rd, fs1, rdn`.
- 3' b011: Rounds off to positive infinity. The corresponding assembler instruction is `fcvt.l.h rd, fs1, rup`.
- 3' b100: Rounds off to the nearest large value. The corresponding assembler instruction is `fcvt.l.h rd, fs1, rmm`.
- 3' b101: This code is reserved and not used.
- 3' b110: This code is reserved and not used.

- 3' b111: Dynamically rounds off based on the rm bit of the fcsr register. The corresponding assembler instruction is `fcvt.l.h rd, fs1`.

**Instruction format:**

31	25	24	20	19	15	14	12	11	7	6	0
1100010		00010		fs1		rm		rd		1010011	

### 17.6.9 FCVT.LU.H: an instruction that converts a half-precision floating-point number to an unsigned long integer

**Syntax:**

`fcvt.lu.h rd, fs1, rm`

**Operation:**

$rd \leftarrow \text{half\_convert\_to\_unsigned\_long}(fs1)$

**Permission:**

M mode/S mode/U mode

**Exception:**

Illegal instruction.

**Affected flag bits:**

Floating-point status bits NV and NX

**Notes:**

RM determines the round-off mode:

- 3' b000: Rounds off to the nearest even number. The corresponding assembler instruction is `fcvt.lu.h rd, fs1, rne`.
- 3' b001: Rounds off to zero. The corresponding assembler instruction is `fcvt.lu.h rd, fs1, rtz`.
- 3' b010: Rounds off to negative infinity. The corresponding assembler instruction is `fcvt.lu.h rd, fs1, rdn`.
- 3' b011: Rounds off to positive infinity. The corresponding assembler instruction is `fcvt.lu.h rd, fs1, rup`.
- 3' b100: Rounds off to the nearest large value. The corresponding assembler instruction is `fcvt.lu.h rd, fs1, rmm`.
- 3' b101: This code is reserved and not used.
- 3' b110: This code is reserved and not used.

- 3' b111: Dynamically rounds off based on the rm bit of the fcsr register. The corresponding assembler instruction is `fcvt.lu.h rd, fs1`.

**Instruction format:**

31	25	24	20	19	15	14	12	11	7	6	0
1100010		00011		fs1		rm		rd		1010011	

### 17.6.10 FCVT.S.H: an instruction that converts half-precision floating-point data to single precision floating-point data

**Syntax:**

`fcvt.s.h fd, fs1`

**Operation:**

$fd \leftarrow \text{half\_convert\_to\_single}(fs1)$

**Permission:**

M mode/S mode/U mode

**Exception:**

Illegal instruction.

**Affected flag bits:**

None

**Instruction format:**

31	25	24	20	19	15	14	12	11	7	6	0
0100000		00010		fs1		000		fd		1010011	

### 17.6.11 FCVT.W.H: an instruction that converts a half-precision floating-point number to a signed integer

**Syntax:**

`fcvt.w.h rd, fs1, rm`

**Operation:**

$tmp \leftarrow \text{half\_convert\_to\_signed\_int}(fs1)$

$rd \leftarrow \text{sign\_extend}(tmp)$

**Permission:**

M mode/S mode/U mode

**Exception:**

Illegal instruction.

**Affected flag bits:**

Floating-point status bits NV and NX

**Notes:**

RM determines the round-off mode:

- 3' b000: Rounds off to the nearest even number. The corresponding assembler instruction is `fcvt.w.h rd, fs1, rne`.
- 3' b001: Rounds off to zero. The corresponding assembler instruction is `fcvt.w.h rd, fs1, rtz`.
- 3' b010: Rounds off to negative infinity. The corresponding assembler instruction is `fcvt.w.h rd, fs1, rdn`.
- 3' b011: Rounds off to positive infinity. The corresponding assembler instruction is `fcvt.w.h rd, fs1, rup`.
- 3' b100: Rounds off to the nearest large value. The corresponding assembler instruction is `fcvt.w.h rd, fs1, rmm`.
- 3' b101: This code is reserved and not used.
- 3' b110: This code is reserved and not used.
- 3' b111: Dynamically rounds off based on the `rm` bit of the `fcsr` register. The corresponding assembler instruction is `fcvt.w.h rd, fs1`.

**Instruction format:**

31	25	24	20	19	15	14	12	11	7	6	0
1100010				00000				fs1		rm	
								rd		1010011	

### 17.6.12 FCVT.WU.H: an instruction that converts a half-precision floating-point number to an unsigned integer

**Syntax:**

```
fcvt.wu.h rd, fs1, rm
```

**Operation:**

```
tmp ← half_convert_to_unsigned_int(fs1)
rd ← sign_extend(tmp)
```

**Permission:**

M mode/S mode/U mode

**Exception:**

Illegal instruction.

**Affected flag bits:**

Floating-point status bits NV and NX

**Notes:**

RM determines the round-off mode:

- 3' b000: Rounds off to the nearest even number. The corresponding assembler instruction is `fcvt.wu.h rd, fs1, rne`.
- 3' b001: Rounds off to zero. The corresponding assembler instruction is `fcvt.wu.h rd, fs1, rtz`.
- 3' b010: Rounds off to negative infinity. The corresponding assembler instruction is `fcvt.wu.h rd, fs1, rdn`.
- 3' b011: Rounds off to positive infinity. The corresponding assembler instruction is `fcvt.wu.h rd, fs1, rup`.
- 3' b100: Rounds off to the nearest large value. The corresponding assembler instruction is `fcvt.wu.h rd, fs1, rmm`.
- 3' b101: This code is reserved and not used.
- 3' b110: This code is reserved and not used.
- 3' b111: Dynamically rounds off based on the `rm` bit of the `fcsr` register. The corresponding assembler instruction is `fcvt.wu.h rd, fs1`.

**Instruction format:**

31	25	24	20	19	15	14	12	11	7	6	0
1100010			00001		fs1		rm	rd		1010011	

**17.6.13 FDIV.H: a half-precision floating-point division instruction****Syntax:**

`fddiv.h fd, fs1, fs2, rm`

**Operation:**

$fd \leftarrow fs1 / fs2$

**Permission:**

M mode/S mode/U mode

**Exception:**

Illegal instruction.

**Affected flag bits:**

Floating-point status bits NV, DZ, OF, UF, and NX

**Notes:**

RM determines the round-off mode:

- 3' b000: Rounds off to the nearest even number. The corresponding assembler instruction is `fdiv.h fs1, fs2, rne`.
- 3' b001: Rounds off to zero. The corresponding assembler instruction is `fdiv.h fd fs1, fs2, rtz`.
- 3' b010: Rounds off to negative infinity. The corresponding assembler instruction is `fdiv.h fd, fs1, fs2, rdn`.
- 3' b011: Rounds off to positive infinity. The corresponding assembler instruction is `fdiv.h fd, fs1, fs2, rup`.
- 3' b100: Rounds off to the nearest large value. The corresponding assembler instruction is `fdiv.h fd, fs1, fs2, rmm`.
- 3' b101: This code is reserved and not used.
- 3' b110: This code is reserved and not used.
- 3' b111: Dynamically rounds off based on the `rm` bit of the `fcsr` register. The corresponding assembler instruction is `fdiv.h fd, fs1, fs2`.

**Instruction format:**

31	25	24	20	19	15	14	12	11	7	6	0
0001110		fs2		fs1		rm		fd		1010011	

**17.6.14 FEQ.H: an equal instruction that compares two half-precision numbers****Syntax:**

`feq.h rd, fs1, fs2`

**Operation:**

`if(fs1 == fs2)`

`rd ← 1`

`else`

`rd ← 0`

**Permission:**

M mode/S mode/U mode

**Exception:**

Illegal instruction.

**Affected flag bits:**

Floating-point status bit NV

**Instruction format:**

31	25	24	20	19	15	14	12	11	7	6	0
1010010		fs2		fs1		010		rd		1010011	

### 17.6.15 FLE.H: a less than or equal to instruction that compares two half-precision floating-point numbers

**Syntax:**

fle.h rd, fs1, fs2

**Operation:**

if(fs1 <= fs2)

rd ← 1

else

rd ← 0

**Permission:**

M mode/S mode/U mode

**Exception:**

Illegal instruction.

**Affected flag bits:**

Floating-point status bit NV

**Instruction format:**

31	25	24	20	19	15	14	12	11	7	6	0
1010010		fs2		fs1		000		rd		1010011	

### 17.6.16 FLH: an instruction that loads half-precision floating-point data

**Syntax:**

flh fd, imm12(rs1)

**Operation:**



$$\text{address} \leftarrow \text{rs1} + \text{sign\_extend}(\text{imm12})$$

$$\text{fd}[15:0] \leftarrow \text{mem}[(\text{address}+1):\text{address}]$$

$$\text{fd}[63:16] \leftarrow 48' \text{ hfffffffffff}$$
**Permission:**

M mode/S mode/U mode

**Exception:**

Unaligned access, access error, page error, or illegal instruction.

**Affected flag bits:**

None

**Instruction format:**

31	20	19	15	14	12	11	7	6	0
imm12[11:0]				rs1	001	rd	0000111		

### 17.6.17 FLT.H: a less than instruction that compares two half-precision floating-point numbers

**Syntax:**

flt.h rd, fs1, fs2

**Operation:**

if(fs1 &lt; fs2)

rd ← 1

else

rd ← 0

**Permission:**

M mode/S mode/U mode

**Exception:**

Illegal instruction.

**Affected flag bits:**

Floating-point status bit NV

**Instruction format:**

31	25	24	20	19	15	14	12	11	7	6	0
1010010			fs2		fs1		001	rd	1010011		

### 17.6.18 FMADD.H: a half-precision floating-point multiply-add instruction

**Syntax:**

```
fmadd.h fd, fs1, fs2, fs3, rm
```

**Operation:**

$$fd \leftarrow fs1 * fs2 + fs3$$
**Permission:**

M mode/S mode/U mode

**Exception:**

Illegal instruction.

**Affected flag bits:**

Floating-point status bits NV, OF, UF, and IX

**Notes:**

RM determines the round-off mode:

- 3' b000: Rounds off to the nearest even number. The corresponding assembler instruction is `fmadd.h fd, fs1, fs2, fs3, rne`.
- 3' b001: Rounds off to zero. The corresponding assembler instruction is `fmadd.h fd, fs1, fs2, fs3, rtz`.
- 3' b010: Rounds off to negative infinity. The corresponding assembler instruction is `fmadd.h fd, fs1, fs2, fs3, rdn`.
- 3' b011: Rounds off to positive infinity. The corresponding assembler instruction is `fmadd.h fd, fs1, fs2, fs3, rup`.
- 3' b100: Rounds off to the nearest large value. The corresponding assembler instruction is `fmadd.h fd, fs1, fs2, fs3, rmm`.
- 3' b101: This code is reserved and not used.
- 3' b110: This code is reserved and not used.
- 3' b111: Dynamically rounds off based on the `rm` bit of the `fcsr` register. The corresponding assembler instruction is `fmadd.h fd, fs1, fs2, fs3`.

**Instruction format:**

31	27	26	25	24	20	19	15	14	12	11	7	6	0
fs3				10	fs2				fs1		rm	fd	1000011

### 17.6.19 FMAX.H: a half-precision floating-point maximum instruction

**Syntax:**

```
fmax.h fd, fs1, fs2
```

**Operation:**

```
if(fs1 >= fs2)
```

```
    fd ← fs1
```

```
else
```

```
    fd ← fs2
```

**Permission:**

M mode/S mode/U mode

**Exception:**

Illegal instruction.

**Affected flag bits:**

Floating-point status bit NV

**Instruction format:**

31	25	24	20	19	15	14	12	11	7	6	0
0010110	fs2			fs1			001	fd			1010011

### 17.6.20 FMIN.H: a half-precision floating-point minimum instruction

**Syntax:**

```
fmin.h fd, fs1, fs2
```

**Operation:**

```
if(fs1 >= fs2)
```

```
    fd ← fs2
```

```
else
```

```
    fd ← fs1
```

**Permission:**

M mode/S mode/U mode

**Exception:**

Illegal instruction.

**Affected flag bits:**

Floating-point status bit NV

**Instruction format:**

31	25	24	20	19	15	14	12	11	7	6	0
0010110		fs2		fs1		000		fd		1010011	

### 17.6.21 FMSUB.H: a half-precision floating-point multiply-subtract instruction

**Syntax:**

fmsub.h fd, fs1, fs2, fs3, rm

**Operation:**

$fd \leftarrow fs1 * fs2 - fs3$

**Permission:**

M mode/S mode/U mode

**Exception:**

Illegal instruction.

**Affected flag bits:**

Floating-point status bits NV, OF, UF, and IX

**Notes:**

RM determines the round-off mode:

- 3' b000: Rounds off to the nearest even number. The corresponding assembler instruction is fmsub.h fd, fs1, fs2, fs3, rne.
- 3' b001: Rounds off to zero. The corresponding assembler instruction is fmsub.h fd, fs1, fs2, fs3, rtz.
- 3' b010: Rounds off to negative infinity. The corresponding assembler instruction is fmsub.h fd, fs1, fs2, fs3, rdn.
- 3' b011: Rounds off to positive infinity. The corresponding assembler instruction is fmsub.h fd, fs1, fs2, fs3, rup.
- 3' b100: Rounds off to the nearest large value. The corresponding assembler instruction is fmsub.h fd, fs1, fs2, fs3, rmm.
- 3' b101: This code is reserved and not used.
- 3' b110: This code is reserved and not used.
- 3' b111: Dynamically rounds off based on the rm bit of the fcsr register. The corresponding assembler instruction is fmsub.h fd, fs1, fs2, fs3.

**Instruction format:**

31	27	26	25	24	20	19	15	14	12	11	7	6	0	
fs3				10	fs2			fs1		rm	fd			1000111

### 17.6.22 FMUL.H: a half-precision floating-point multiply instruction

**Syntax:**

fmul.h fd, fs1, fs2, rm

**Operation:**

$fd \leftarrow fs1 * fs2$

**Permission:**

M mode/S mode/U mode

**Exception:**

Illegal instruction.

**Affected flag bits:**

Floating-point status bits NV, OF, UF, and NX

**Notes:**

RM determines the round-off mode:

- 3' b000: Rounds off to the nearest even number. The corresponding assembler instruction is `fmul.h fd, fs1, fs2, rne`.
- 3' b001: Rounds off to zero. The corresponding assembler instruction is `fmul.h fd, fs1, fs2, rtz`.
- 3' b010: Rounds off to negative infinity. The corresponding assembler instruction is `fmul.h fd, fs1, fs2, rdn`.
- 3' b011: Rounds off to positive infinity. The corresponding assembler instruction is `fmul.h fd, fs1, fs2, rup`.
- 3' b100: Rounds off to the nearest large value. The corresponding assembler instruction is `fmul.h fd, fs1, fs2, rmm`.
- 3' b101: This code is reserved and not used.
- 3' b110: This code is reserved and not used.
- 3' b111: Dynamically rounds off based on the rm bit of the fcsr register. The corresponding assembler instruction is `fmul.h fs1, fs2`.

**Instruction format:**

31	25	24	20	19	15	14	12	11	7	6	0
0001010		fs2		fs1		rm		fd		1010011	

### 17.6.23 FMV.H.X: a half-precision floating-point write transmit instruction

**Syntax:**

fmv.h.x fd, rs1

**Operation:**

fd[15:0]  $\leftarrow$  rs1[15:0]

fd[63:16]  $\leftarrow$  48' hffffffff

**Permission:**

M mode/S mode/U mode

**Exception:**

Illegal instruction.

**Affected flag bits:**

None

**Instruction format:**

31	25	24	20	19	15	14	12	11	7	6	0
1111010		00000		rs1		000		fd		1010011	

### 17.6.24 FMV.X.H: a transmission instruction that reads half-precision floating-point registers

**Syntax:**

fmv.x.h rd, fs1

**Operation:**

tmp[15:0]  $\leftarrow$  fs1[15:0]

rd  $\leftarrow$  sign\_extend(tmp[15:0])

**Permission:**

M mode/S mode/U mode

**Exception:**

Illegal instruction.

Affected flag bits:

None

Instruction format:

31	25	24	20	19	15	14	12	11	7	6	0
1110010	00000	fs1	000	rd	1010011						

### 17.6.25 FNMADD.H: a half-precision floating-point negate-(multiply-add) instruction

Syntax:

fnmadd.h fd, fs1, fs2, fs3, rm

Operation:

$fd \leftarrow -(fs1 * fs2 + fs3)$

Permission:

M mode/S mode/U mode

Exception:

Illegal instruction.

Affected flag bits:

Floating-point status bits NV, OF, UF, and IX

Notes:

RM determines the round-off mode:

- 3' b000: Rounds off to the nearest even number. The corresponding assembler instruction is fnmadd.h fd,fs1, fs2, fs3, rne.
- 3' b001: Rounds off to zero. The corresponding assembler instruction is fnmadd.h fd,fs1, fs2, fs3, rtz.
- 3' b010: Rounds off to negative infinity. The corresponding assembler instruction is fnmadd.h fd,fs1, fs2, fs3, rdn.
- 3' b011: Rounds off to positive infinity. The corresponding assembler instruction is fnmadd.h fd,fs1, fs2, fs3, rup.
- 3' b100: Rounds off to the nearest large value. The corresponding assembler instruction is fnmadd.h fd,fs1, fs2, fs3, rmm.
- 3' b101: This code is reserved and not used.
- 3' b110: This code is reserved and not used.
- 3' b111: Dynamically rounds off based on the rm bit of the fcsr register. The corresponding assembler instruction is fnmadd.h fd,fs1, fs2, fs3.

Instruction format:

31	27	26	25	24	20	19	15	14	12	11	7	6	0
fs3	10	fs2	fs1	rm	fd	1001111							

### 17.6.26 FNMSUB.H: a half-precision floating-point negate-(multiply-subtract) instruction

Syntax:

fnmsub.h fd, fs1, fs2, fs3, rm

Operation:

$fd \leftarrow -(fs1 * fs2 - fs3)$

Permission:

M mode/S mode/U mode

Exception:

Illegal instruction.

Affected flag bits:

Floating-point status bits NV, OF, UF, and IX

Notes:

RM determines the round-off mode:

- 3' b000: Rounds off to the nearest even number. The corresponding assembler instruction is fnmsub.h fd,fs1, fs2, fs3, rne.
- 3' b001: Rounds off to zero. The corresponding assembler instruction is fnmsub.h fd,fs1, fs2, fs3, rtz.
- 3' b010: Rounds off to negative infinity. The corresponding assembler instruction is fnmsub.h fd,fs1, fs2, fs3, rdn.
- 3' b011: Rounds off to positive infinity. The corresponding assembler instruction is fnmsub.h fd,fs1, fs2, fs3, rup.
- 3' b100: Rounds off to the nearest large value. The corresponding assembler instruction is fnmsub.h fd,fs1, fs2, fs3, rmm.
- 3' b101: This code is reserved and not used.
- 3' b110: This code is reserved and not used.
- 3' b111: Dynamically rounds off based on the rm bit of the fcsr register. The corresponding assembler instruction is fnmsub.h fd,fs1, fs2, fs3.

Instruction format:



31	27	26	25	24	20	19	15	14	12	11	7	6	0					
fs3				10	fs2				fs1				rm	fd	1001011			

### 17.6.27 FSGNJ.H: a half-precision floating-point sign-injection instruction

**Syntax:**

fsgnj.h fd, fs1, fs2

**Operation:**

$fd[14:0] \leftarrow fs1[14:0]$

$fd[15] \leftarrow fs2[15]$

$fd[63:16] \leftarrow 48' \text{ hffffffffffff}$

**Permission:**

M mode/S mode/U mode

**Exception:**

Illegal instruction.

**Affected flag bits:**

None

**Instruction format:**

31	25	24	20	19	15	14	12	11	7	6	0	
0010010			fs2		fs1		000		fd		1010011	

### 17.6.28 FSGNJN.H: a half-precision floating-point sign-injection negate instruction

**Syntax:**

fsgnjn.h fd, fs1, fs2

**Operation:**

$fd[14:0] \leftarrow fs1[14:0]$

$fd[15] \leftarrow ! fs2[15]$

$fd[63:16] \leftarrow 48' \text{ hffffffffffff}$

**Permission:**

M mode/S mode/U mode

**Exception:**

Illegal instruction.

**Affected flag bits:**

None

**Instruction format:**

31	25	24	20	19	15	14	12	11	7	6	0
0010010		fs2		fs1		001		fd		1010011	

**17.6.29 FSGNJX.H: a half-precision floating-point sign-injection XOR instruction****Syntax:**

fsgnjx.h fd, fs1, fs2

**Operation:**

$$fd[14:0] \leftarrow fs1[14:0]$$

$$fd[15] \leftarrow fs1[15] \wedge fs2[15]$$

$$fd[63:16] \leftarrow 48'h\text{ffffffff}$$
**Permission:**

M mode/S mode/U mode

**Exception:**

Illegal instruction.

**Affected flag bits:**

None

**Instruction format:**

31	25	24	20	19	15	14	12	11	7	6	0
0010010		fs2		fs1		010		fd		1010011	

**17.6.30 FSH: an instruction that stores half-precision floating point numbers****Syntax:**

fsh fs2, imm12(fs1)

**Operation:**

$$\text{address} \leftarrow fs1 + \text{sign\_extend}(\text{imm12})$$

$$\text{mem}[(\text{address}+1):\text{address}] \leftarrow fs2[15:0]$$

**Permission:**

M mode/S mode/U mode

**Exception:**

Unaligned access, access error, page error, or illegal instruction.

**Instruction format:**

31	25	24	20	19	15	14	12	11	7	6	0
imm12[11:5]				fs2		fs1		001	imm12[4:0]		0100111

### 17.6.31 FSQRT.H: an instruction that calculates the square root of the half-precision floating-point number

**Syntax:**

fsqrt.h fd, fs1, rm

**Operation:** $fd \leftarrow \text{sqrt}(fs1)$ **Permission:**

M mode/S mode/U mode

**Exception:**

Illegal instruction.

**Affected flag bits:**

Floating-point status bits NV and NX

**Notes:**

RM determines the round-off mode:

- 3' b000: Rounds off to the nearest even number. The corresponding assembler instruction is fsqrt.h fd, fs1,rne.
- 3' b001: Rounds off to zero. The corresponding assembler instruction is fsqrt.h fd, fs1,rtz.
- 3' b010: Rounds off to negative infinity. The corresponding assembler instruction is fsqrt.h fd, fs1,rdn.
- 3' b011: Rounds off to positive infinity. The corresponding assembler instruction is fsqrt.h fd, fs1,rup.
- 3' b100: Rounds off to the nearest large value. The corresponding assembler instruction is fsqrt.h fd, fs1,rmm.
- 3' b101: This code is reserved and not used.
- 3' b110: This code is reserved and not used.

- 3' b111: Dynamically rounds off based on the rm bit of the fcsr register. The corresponding assembler instruction is `fsqrt.h fd, fs1`.

**Instruction format:**

31	25	24	20	19	15	14	12	11	7	6	0
0101110	00000	fs1	rm	fd	1010011						

### 17.6.32 FSUB.H: a half-precision floating-point subtract instruction

**Syntax:**

`fsub.h fd, fs1, fs2, rm`

**Operation:**

$fd \leftarrow fs1 - fs2$

**Permission:**

M mode/S mode/U mode

**Exception:**

Illegal instruction.

**Affected flag bits:**

Floating-point status bits NV, OF, and NX

**Notes:**

RM determines the round-off mode:

- 3' b000: Rounds off to the nearest even number. The corresponding assembler instruction is `fsub.h fd, fs1,fs2,rne`.
- 3' b001: Rounds off to zero. The corresponding assembler instruction is `fsub.h fd, fs1,fs2,rtz`.
- 3' b010: Rounds off to negative infinity. The corresponding assembler instruction is `fsub.h fd, fs1,fs2,rdn`.
- 3' b011: Rounds off to positive infinity. The corresponding assembler instruction is `fsub.h fd, fs1,fs2,rup`.
- 3' b100: Rounds off to the nearest large value. The corresponding assembler instruction is `fsub.h fd, fs1,fs2,rmm`.
- 3' b101: This code is reserved and not used.
- 3' b110: This code is reserved and not used.
- 3' b111: Dynamically rounds off based on the rm bit of the fcsr register. The corresponding assembler instruction is `fsub.h fd, fs1,fs2`.

Instruction format:

31	25	24	20	19	15	14	12	11	7	6	0
0000110			fs2		fs1		rm		fd		1010011

---

## Appendix C Control Registers

---

This section describes the M-mode control registers, S-mode control registers, and U-mode control registers.

### 18.1 Appendix C-1 M-mode control registers

M-mode control registers are classified by function into M-mode information registers, M-mode exception configuration registers, M-mode exception handling registers, M-mode memory protection registers, M-mode counter registers, and M-mode counter configuration registers.

#### 18.1.1 M-mode information register group

##### 18.1.1.1 Machine vendor ID register (mvendorid)

The mvendorid register stores the vendor IDs of T-Head Semiconductor Co., Ltd. It is not defined and the values are all zero.

This register is 64 bits wide and is read-only in M-mode. Accesses in non-M-mode and writes in M-mode will cause an illegal instruction exception.

##### 18.1.1.2 Machine architecture ID register (marchid)

The marchid register stores the architecture IDs of CPU cores. It stores internal IDs of T-Head Semiconductor Co., Ltd. and its reset value is subject to the product.

This register is 64 bits wide and is read-only in M-mode. Accesses in non-M-mode and writes in M-mode will cause an illegal instruction exception.

### 18.1.1.3 Machine implementation ID register (mimpid)

The mimpid register stores hardware implementation IDs of CPU cores. This register is not implemented by C910/C920, and its read access value is 0.

This register is 64 bits wide and is read-only in M-mode. Accesses in non-M-mode and writes in M-mode will cause an illegal instruction exception.

### 18.1.1.4 Machine hart ID register (mhartid)

The mhartid register stores hart IDs of CPU cores.

This register is 64 bits wide and is read-only in M-mode. Accesses in non-M-mode and writes in M-mode will cause an illegal instruction exception.

## 18.1.2 M-mode exception configuration register group

### 18.1.2.1 Machine status register (mstatus)

The mstatus register stores status and control information of the CPU in M-mode, including the global interrupt enable bit, exception preserve interrupt enable bit, and exception preserve privilege mode bit.

This register is 64 bits wide and is readable and writable in M-mode. Accesses in non-M-mode will cause an illegal instruction exception.

	63	62																													36	35	34	33	32
	SD	0																												SXL		UXL			
Reset	0	0																												2		2			

	31																		25	24	23	22	21	20	19	18	17	16	15	14	13	12	11	10	9	8	7	6	5	4	3	2	1	0
	0																			VS	TSR	TW					XS	FS	MPP	0	SPP		0		0	MIE	0	SIE	0					
																				TVM		SUM												MPIE		SPIE								
																				MXR		MPRV																						
Reset	0																			0	0	0	0	0	0	0	0	0	0	3	0	1	0	0	0	0	0	0	0	0	0	0		

Machine status register (mstatus)

**SIE: supervisor interrupt enable bit**

- When SIE is 0, S-mode interrupts are invalid.
- When SIE is 1, S-mode interrupts are valid.

This bit is reset to 0 when the CPU is downgraded to the S-mode response interrupt, and is set to the value of SPIE when the CPU exits the interrupt service program.

**MIE: machine interrupt enable bit**

- When MIE is 0, M-mode interrupts are invalid.
- When MIE is 1, M-mode interrupts are valid.

This bit is reset to 0 when the response is interrupted in M-mode on the CPU, and is set to the value of MPIE when the CPU exits the interrupt service program.

**SPIE: supervisor preserve interrupt enable bit**

This bit stores the value of the SIE bit before the response is interrupted in S-mode on the CPU.

This bit will be reset to 0, and set to 1 when the CPU exits the interrupt service program.

**MPIE: machine preserve interrupt enable bit**

This bit stores the value of the MIE bit before the response is interrupted in M-mode on the CPU.

This bit will be reset to 0, and set to 1 when the CPU exits the interrupt exception service program.

**SPP: supervisor preserve privilege bit**

This bit stores the privilege status before the CPU accesses the exception service program in S-mode.

- When SPP is 2' b00, the CPU is in U-mode before accessing the exception service program.
- When SPP is 2' b01, the CPU is in S-mode before accessing the exception service program.

This bit will be reset to 2' b01.

**MPP: machine preserve privilege bit**

This bit stores the privilege status before the CPU accesses the exception service program in M-mode.

- When MPP is 2' b00, the CPU is in U-mode before entering the exception service program.
- When MPP is 2' b01, the CPU is in S-mode before accessing the exception service program.
- When MPP is 2' b11, the CPU is in M-mode before entering the exception service program.

This bit will be reset to 2' b11.

**FS: floating-point status bit**

This bit determines whether to store floating-point registers during context switching.

- When FS is 2' b00, the floating-point unit is in the Off state and exceptions will occur for accesses to related floating-point registers.
- When FS is 2' b01, the floating-point unit is in the Initial state.



- When FS is 2' b10, the floating-point unit is in the Clean state.
- When FS is 2' b11, the floating-point unit is in the Dirty state, which means the floating-point and control registers have been modified.

**XS: extension status bit**

C910/C920 has no extension units, and therefore this bit is fixed to 0.

**MPRV: modify privilege mode**

- When MPRV is 1, load and store requests are executed based on the privilege mode in MPP.
- When MPRV is 0, load and store requests are executed based on the current privilege mode of the CPU.

**SUM: allow S-mode accesses to U-mode virtual memory spaces**

- When SUM is 1, load, storage, and value-taking requests can be initiated in S-mode to access U-mode virtual memory spaces.
- When SUM is 0, load, storage, and value-taking requests cannot be initiated in S-mode to access virtual memory spaces marked as U-mode.

**MXR: allow accesses of load requests to memory spaces marked as executable**

- When MXR is 1, accesses of load requests are allowed to virtual memory spaces marked as executable or readable.
- When MXR is 0, accesses of load requests are allowed only to virtual memory spaces marked as readable.

**TVM: trap virtual memory**

- When TVM is 1, an illegal instruction exception occurs for reads and writes to the satp control register and for the execution of the sfence instruction in S-mode.
- When TVM is 0, reads and writes to the satp control register and the execution of the sfence instruction are allowed in S-mode.

**TW: timeout wait**

- When TW is 1, an illegal instruction exception occurs if the WFI instruction is executed in S-mode.
- When TW is 0, the WFI instruction can be executed in S-mode.

**TSR: trap sret**

- When TSR is 1, an illegal instruction exception occurs if the sret instruction is executed in S-mode.

When TSR is 0, the sret instruction can be executed in S-mode.

**VS: vector status bit**

This bit determines whether to store vector registers during context switching.

- When VS is 2' b00, the vector unit is in the Off state and exceptions will occur for accesses to related vector registers.
- When VS is 2' b01, the vector unit is in the Initial state.
- When VS is 2' b10, the vector unit is in the Clean state.
- When VS is 2' b11, the vector unit is in the Dirty state, which means the vector registers and vector control registers have been modified.

The VS bit is valid only when the vector execution unit is configured, and is always 0 if the vector execution unit is not configured.

#### **UXL: register width**

This bit is read-only and always 2, which means the register is 64 bits wide in U-mode.

#### **SXL: register width**

This bit is read-only and always 2, which means the register is 64 bits wide in S-mode.

#### **SD: dirty state sum bit of the floating-point, vector, and extension units**

- When SD is 1, the floating-point unit, vector unit, or extension unit is in the Dirty state.
- When SD is 0, none of the floating-point, vector, and extension units is in the Dirty state.

### **18.1.2.2 M-mode instruction set architecture register (misa)**

The misa register stores the features of the instruction set architecture supported by the CPU.

This register is 64 bits wide and is readable and writable in M-mode. Accesses in non-M-mode will cause an illegal instruction exception.

C910/C920 supports the RV64GC instruction set architecture, and the reset value of the MISA register is 0x80000000094112d. For more information about the assignment rules, see the official document of RISC-V riscv-privileged.

C910/C920 does not support the dynamic configuration of the MISA register. Writes to this register do not take effect.


### **18.1.2.3 M-mode exception downgrade control register (medeleg)**

The medeleg register can downgrade exceptions that occur in S-mode and U-mode to S-mode responses. The lower 16 bits of the medeleg register are in one-to-one correspondence to exception vector tables. Exceptions to be downgraded to S-mode responses can be selected.

This register is 64 bits wide and is readable and writable in M-mode. Accesses in non-M-mode will cause an illegal instruction exception.

#### 18.1.2.4 M-mode interrupt downgrade control register (mideleg)

The mideleg register can downgrade S-mode interrupts to S-mode responses.




M-mode interrupt downgrade control register (mideleg)

This register is 64 bits wide and is readable and writable in M-mode. Accesses in non-M-mode will cause an illegal instruction exception.

#### 18.1.2.5 M-mode interrupt-enable register (mie)

The mie register enables and masks different types of interrupts. This register is 64 bits wide and is readable and writable in M-mode. Accesses in non-M-mode will cause an illegal instruction exception.



M-mode interrupt-enable register (mie)

**SSIE: S-mode software interrupt enable bit**

- When SSIE is 0, S-mode software external interrupts are invalid.
- When SSIE is 1, S-mode software external interrupts are valid.

**MSIE:** M-mode software interrupt enable bit

- When MSIE is 0, M-mode software interrupts are invalid.
- When MSIE is 1, M-mode software interrupts are valid.

**STIE: S-mode timer interrupt enable bit**

- When STIE is 0, S-mode timer interrupts are invalid.
- When STIE is 1, S-mode timer interrupts are valid.

**MTIE: M-mode timer interrupt enable bit**

- When MTIE is 0, M-mode timer interrupts are invalid.
- When MTIE is 1, M-mode timer interrupts are valid.

**SEIE: S-mode external interrupt enable bit**

- When SEIE is 0, S-mode external interrupts are invalid.
- When SEIE is 1, S-mode external interrupts are valid.

**MEIE: M-mode external interrupt enable bit**

- When MEIE is 0, M-mode external interrupts are invalid.
- When MEIE is 1, M-mode external interrupts are valid.


**MOIE: M-mode overflow interrupt enable bit**

- When MOIE is 0, M-mode counter overflow interrupts are invalid.
- When MOIE is 1, M-mode counter overflow interrupts are valid.

**18.1.2.6 M-mode trap vector base address register (mtvec)**

The mtvec register stores the entry address of the exception service program.

This register is 64 bits wide and is readable and writable in M-mode. Accesses in non-M-mode will cause an illegal instruction exception.



M-mode trap vector base address register (mtvec)

**BASE: vector base address bit**

The BASE bit indicates the upper 62 bits of the entry address of the exception service program. Combining this base address with 2' b00 obtains the entry address of the exception service program.

This bit will be reset to 0.

#### MODE: vector entry mode bit

- When MODE[1:0] is 2' b00, the base address is used as the entry address for both exceptions and interrupts.
- When MODE[1:0] is 2' b01, the base address is used as the entry address for exceptions, and  $\text{BASE} + 4 * \text{cause}$  is used as the entry address for interrupts.

#### 18.1.2.7 M-mode counter access enable register (mcounteren)

The mcounteren register determines whether U-mode counters can be accessed in S-mode.

For more information, see ref:*performance\_test*.

### 18.1.3 M-mode exception handling register group

#### 18.1.3.1 M-mode scratch register (mscratch)

The mscratch register is used by the CPU to back up temporary data in the exception service program. It is usually used to store the entry pointer to the local context space in M-mode.

This register is 64 bits wide and is readable and writable in M-mode. Accesses in non-M-mode will cause an illegal instruction exception.

#### 18.1.3.2 M-mode exception program counter register (mepc)


The mepc register stores the program counter value (PC value) when the CPU exits from the exception service program. C910/C920 supports 16 bits wide instructions. The MEPC value is aligned with 16 bits and the lowest bit is 0.

This register is 64 bits wide and is readable and writable in M-mode. Accesses in non-M-mode will cause an illegal instruction exception.

#### 18.1.3.3 M-mode cause register (mcause)

The mcause register stores the vector numbers of events that trigger exceptions. The vector numbers are used to handle corresponding events in the exception service program.

This register is 64 bits wide and is readable and writable in M-mode. Accesses in non-M-mode will cause an illegal instruction exception.



M-mode cause register (mcause)

### Interrupt: interrupt bit

- When the Interrupt bit is 0, the corresponding exception is not triggered by an interrupt. The exception code is parsed as an exception.
- When the Interrupt bit is 1, the corresponding exception is triggered by an interrupt. The exception code is parsed as an interrupt.

### Exception Code: exception vector number

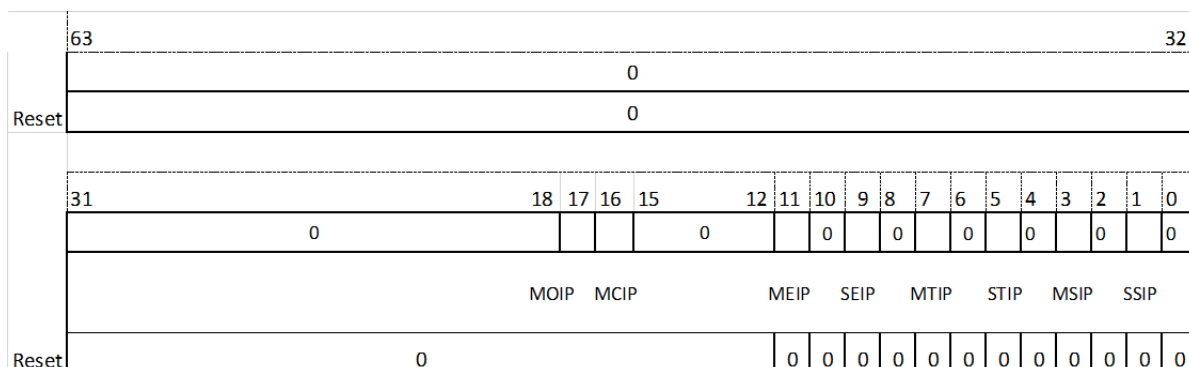
When the CPU encounters an exception, the Exception Code bit will be updated to the value of the exception source.

#### 18.1.3.4 M-mode interrupt-pending register (mip)

The mip register stores information about pending interrupts. When the CPU cannot immediately respond to an interrupt, the corresponding bit in the mip register will be set.

Writing the msip and ssip registers in the CLINT interrupt controller can trigger corresponding interrupts. After the interrupts become valid, the msip bit and ssip bit can be queried based on the corresponding bits in the mip register.

This register is 64 bits wide and is readable and writable in M-mode. Accesses in non-M-mode will cause an illegal instruction exception.



M-mode interrupt-pending register (mip)

### SSIP: supervisor software interrupt pending bit

- When SSIP is 0, there is no pending S-mode software interrupt on the CPU.
- When SSIP is 1, there are pending S-mode software interrupts on the CPU.

The SSIP bit is readable and writable in M-mode. After it is delegated to S-mode, it is readable and writable in S-mode. Otherwise, it is read-only in S-mode.

#### **MSIP: M-mode software interrupt pending bit**

- When MSIP is 0, there is no pending M-mode software interrupt on the CPU.
- When MSIP is 1, there are pending M-mode software interrupts on the CPU.

This bit is read-only.

#### **STIP: S-mode timer interrupt pending bit**

- When STIP is 0, there is no pending S-mode timer interrupt on the CPU.
- When STIP is 1, there are pending S-mode timer interrupts on the CPU.

#### **MTIP: M-mode timer interrupt pending bit**

- When MTIP is 0, there is no pending M-mode timer interrupt on the CPU.
- When MTIP is 1, there are pending M-mode timer interrupts on the CPU.

#### **SEIP: S-mode external interrupt pending bit**

- When SEIP is 0, there is no pending S-mode external interrupt on the CPU.
- When SEIP is 1, there are pending S-mode external interrupts on the CPU.

#### **MEIP: machine external interrupt pending bit**

- When MEIP is 0, there is no pending M-mode external interrupt on the CPU.
- When MEIP is 1, there are pending M-mode external interrupts on the CPU.

#### **MOIP: M-mode overflow interrupt pending bit**

- When MOIP is 0, there is no pending M-mode counter overflow interrupt on the CPU.
- When MOIP is 1, there are pending M-mode counter overflow interrupts on the CPU.

### **18.1.4 M-mode memory protection registers**

M-mode memory protection registers are related to the settings of the memory protection unit.

#### **18.1.4.1 Physical memory protection configuration register (pmpcfg)**

The pmpcfg register configures access permissions and address matching mode for the physical memory.

This register is 64 bits wide and is readable and writable in M-mode. Accesses in non-M-mode will cause an illegal instruction exception.

For more information, see ref:*physical\_mem\_pmpcfg*.

#### 18.1.4.2 Physical memory address register (pmpaddr)

The pmpaddr register configures the address range of each entry of the physical memory.

This register is 64 bits wide and is readable and writable in M-mode. Accesses in non-M-mode will cause an illegal instruction exception.

For more information, see ref:*physical\_mem\_pmpaddr*.

### 18.1.5 M-mode counter registers

M-mode counter registers belong to the PMU and collect software and hardware information during a program operation for software development personnel to optimize programs.

#### 18.1.5.1 M-mode cycle counter (mcycle)

The mcycle counter stores the cycles executed by the CPU. When the CPU is in the execution state (non-low power state), the mcycle register increases the count upon each execution cycle.

The mcycle counter is 64 bits wide and will be reset to 0.

For more information, see ref:*performance\_test\_cont*.

#### 18.1.5.2 M-mode instructions-retired counter (minstret)

The minstret counter stores the number of retired instructions of the CPU. The minstret register increases the count when each instruction retires.

The minstret counter is 64 bits wide and will be reset to 0.

For more information, see ref:*performance\_test\_cont*.

#### 18.1.5.3 M-mode event counter (mhpmcountern)

The mhpmcountern counter counts events.

The mhpmcountern counter is 64 bits wide and will be reset to 0.

For more information, see ref:*performance\_test\_cont*.

### 18.1.6 M-mode counter configuration registers

The M-mode counter configuration register (mhpmeventn) selects events for M-mode event counters.



### 18.1.6.1 M-mode event selector (mhpmeventn)

The events mhpmevent3-31 selected by the mhpmeventn register for M-mode event counters mhpmcouter3-31 are in one-to-one correspondence. In C910/C920, event counters can count only specified events. Therefore, only specified values can be written to mhpmevent3-31.

The mhpmeventn counter is 64 bits wide and will be reset to 0.

For more information, see *ref:performance\_test\_mhpmevent*.

## 18.1.7 M-mode CPU control and status extension registers

C910/C920 extends some registers for the CPU and status, including the M-mode extension status register (mxstatus) and M-mode hardware control register (mhcr), M-mode hardware operation register (mcor), M-mode L2 Cache control register (mccr2), M-mode implicit operation register (mhint), M-mode reset vector base address register (mrvbr), S-mode counter write enable register (mcounterwen), M-mode event interrupt enable register (mcounterinten), and M-mode event overflow mark register (mcounterof).

### 18.1.7.1 M-mode extension status register (mxstatus)

The mxstatus register stores the current privilege mode of the CPU and the enable bit of the extension functions of C910/C920.

This register is 64 bits wide and is readable and writable in M-mode. Accesses in non-M-mode will cause an illegal instruction exception.

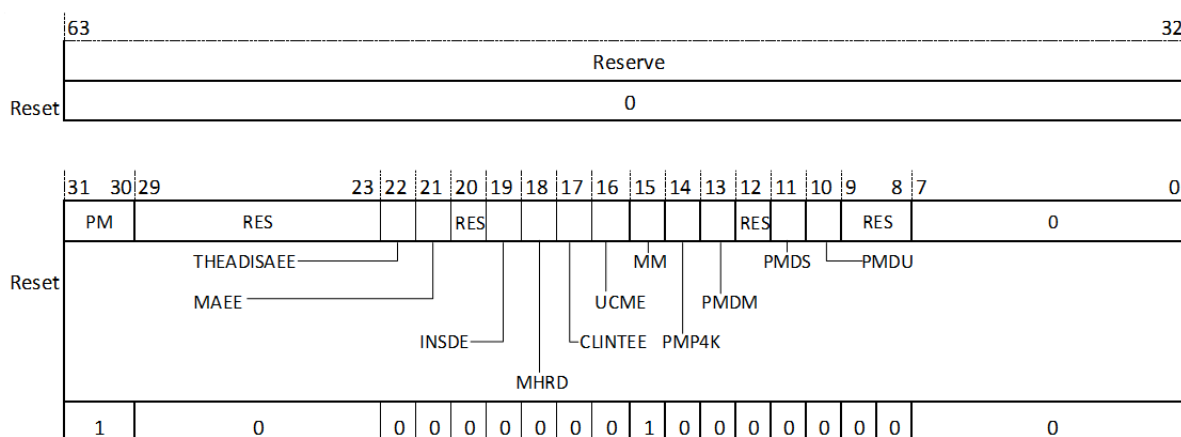


Fig. 18.1: M-mode extension status register (mxstatus)

### PMDU: U-mode performance monitoring count enable bit

When PMDU is 0, performance counters are allowed to count in U-mode.

When PMDU is 1, performance counters are not allowed to count in U-mode.

**PMDS: S-mode performance monitoring count enable bit**

When PMDS is 0, performance counters are allowed to count in S-mode.

When PMDS is 1, performance counters are not allowed to count in S-mode.

**PMDM: M-mode performance monitoring count enable bit**

When PMDM is 0, performance counters are allowed to count in M-mode.

When PMDM is 1, performance counters are not allowed to count in M-mode.

**PMP4K: PMP minimum granularity control bit**

The minimum PMP granularity supported by C910/C920 is 4 KB, which is not affected by this bit.

**MM: misaligned access enable bit**

When MM is 0, misaligned accesses are not supported and cause misaligned exceptions.

When MM is 1, misaligned accesses are supported and processed by hardware. (The default value of this bit is 1 in C910/C920.)

**UCME: execute extended cache instructions in U-mode**

When UCME is 0, extended cache instructions cannot be executed in U-mode. Otherwise, instruction exceptions may occur.

When UCME is 1, extended cache instructions can be executed in U-mode.

**CLINTEE: Clint timer/software interrupt supervisor extension enable bit**

When CLINTEE is 0, supervisor software interrupts and timer interrupts initiated by Clint are not responded to.

When CLINTEE is 1, supervisor software interrupts and timer interrupts initiated by Clint can be responded to.

**MHRD: disable hardware writeback**

When MHRD is 0, hardware writeback is performed if the TLB is missing.

When MHRD is 1, hardware writeback is not performed after the TLB is missing.

**INSDE: disable Icache snoop D-Cache**

When INSDE is 0, D-Cache is snooped after I-Cache is missing.

When INSDE is 1, D-Cache is not snooped after I-Cache is missing.

**MAEE: extend MMU address attribute**

When MAEE is 0, the MMU address attribute is not extended.

When MAEE is 1, the address attribute is extended in the PTE of the MMU. Users can configure the address attribute of pages.

**THEADISAE: enables extended instruction sets**

When THEADISAEE is 0, using C910/C920 extended instruction sets causes instruction exceptions.

When THEADISAEE is 1, C910/C920 extended instruction sets can be used.

**PM: privilege mode**

When PM is 2' b00, the CPU is running in U-mode.

When PM is 2' b01, the CPU is running in S-mode.

When PM is 2' b11, the CPU is running in M-mode. (The PM bit will be reset to M-mode.)

**18.1.7.2 M-mode hardware configuration register (mhcr)**

The mhcr register configures the performance and functions of the CPU.

This register is 64 bits wide and is readable and writable in M-mode. Accesses in non-M-mode will cause an illegal instruction exception.

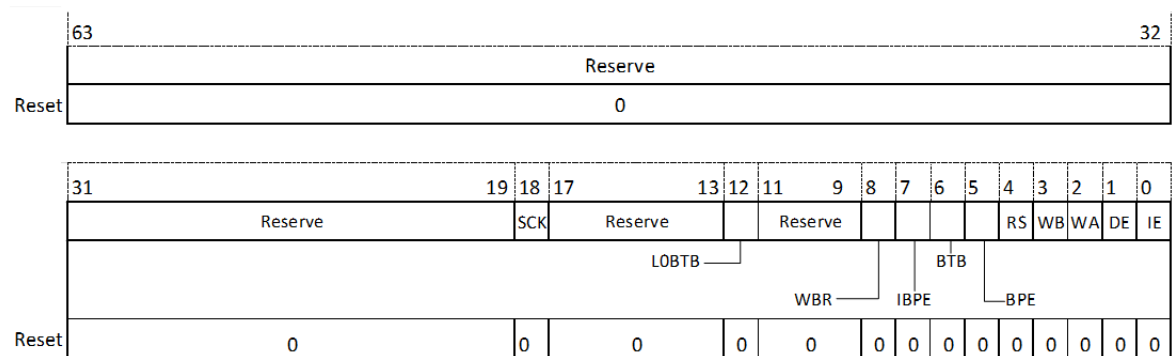


Fig. 18.2: M-mode hardware configuration register (mhcr)

**IE: I-Cache enable bit**

When IE is 0, I-Cache is disabled.

When IE is 1, I-Cache is enabled.

**DE: D-Cache enable bit**

When DE is 0, D-Cache is disabled.

When DE is 1, D-Cache is enabled.

**WA: cache write allocate set bit**

When WA is 0, the data cache is in write non-allocate mode.

When WA is 1, the data cache is in write allocate mode.

**WB: cache writeback set bit**

When WB is 0, the data cache is in write-through mode.

When WB is 1, the data cache is in writeback mode.

C910/C920 supports only the writeback mode. Therefore, the WB bit is fixed to 1.

**RS: address return stack set bit**

When RS is 0, the return stack is disabled.

When RS is 1, the return stack is enabled.

**BPE: branch prediction enable bit**

When BPE is 0, branch prediction is disabled.

When BPE is 1, branch prediction is enabled.

**BTB: branch target prediction enable bit**

When BTB is 0, branch target prediction is disabled.

When BTB is 1, branch target prediction is enabled.

**IBPE: indirect branch prediction enable bit**

When IBPE is 0, indirect branch prediction is disabled.

When IBPE is 1, indirect branch prediction is enabled.

**WBR: write burst transmission enable bit**

When WBR is 0, write burst transmission is not supported.

When WBR is 1, write burst transmission is supported.

The WBR bit is fixed to 1 by default in C910/C920, and cannot be modified.

**L0BTB: level-1 branch target prediction enable bit**

When L0BTB is 0, level-1 branch target prediction is disabled.

When L0BTB is 1, level-1 branch target prediction is enabled.

**SCK: ratio of system clock to CPU clock**

This bit indicates the ratio of the system clock to the CPU clock. The calculation format is SCK+1. There are corresponding pins on the CPU. The SCK bit is configured during a reset and cannot be modified later.

**18.1.7.3 M-mode hardware operation register (mcor)**

The mcor register performs related operations on caches and branch predictors.

This register is 64 bits wide and is readable and writable in M-mode. Accesses in non-M-mode will cause an illegal instruction exception.

**CACHE\_SEL: cache select bit**

When CACHE\_SEL is 2'b01, the instruction cache is selected.

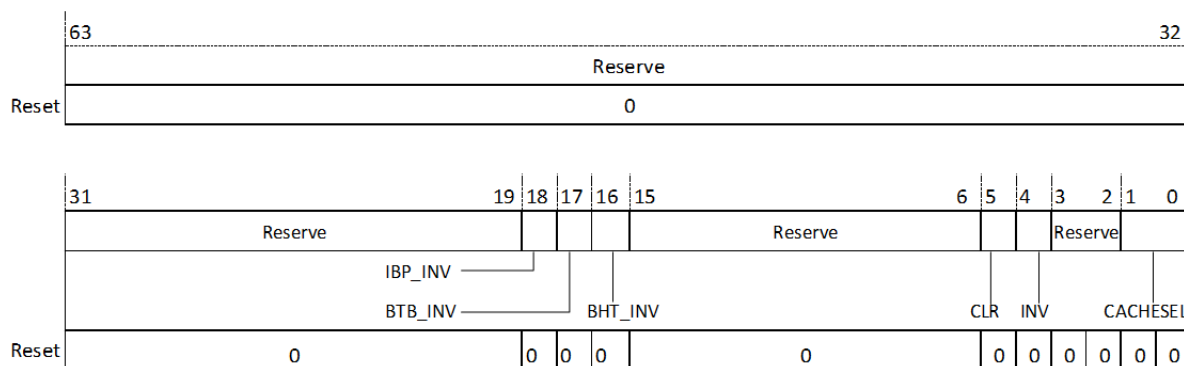


Fig. 18.3: M-mode hardware operation register (mcor)

When `CACHE_SEL` is 2' b10, the data cache is selected.

When `CACHE_SEL` is 2' b11, the instruction and data caches are selected.

#### INV: cache invalidate bit

When `INV` is 0, caches are not invalidated.

When `INV` is 1, caches are invalidated.

#### CLR: dirty entry clear bit

When `CLR` is 0, dirty entries in caches are not written out of the chip.

When `CLR` is 1, dirty entries in caches are written out of the chip.

#### BHT\_INV: BHT invalidate bit

When `BHT_INV` is 0, data in branch history tables (BHTs) is not invalidated.

When `BHT_INV` is 1, data in BHTs is invalidated.

#### BTB\_INV: BTB invalidate bit

When `BTB_INV` is 0, data in branch target buffers (BTBs) is not invalidated.

When `BTB_INV` is 1, data in BTBs is invalidated.

#### IBP\_INV: IBP invalidate bit

When `IBP_INV` is 0, indirect branch prediction (IBP) data is not invalidated.

When `IBP_INV` is 1, IBP data is invalidated.

All the preceding invalidate and dirty entry clear bits are set to 1 when corresponding operations are in progress and reset to 0 when the operations are completed.

#### 18.1.7.4 M-mode L2 Cache control register (mccr2)

The mccr2 register configures the access delays of memories in the shared L2 Cache, including L2 Cache enable/disable, instruction prefetch, and TLB prefetch enable.

This register is 64 bits wide and is readable and writable in M-mode. Accesses in non-M-mode will cause an illegal instruction exception.

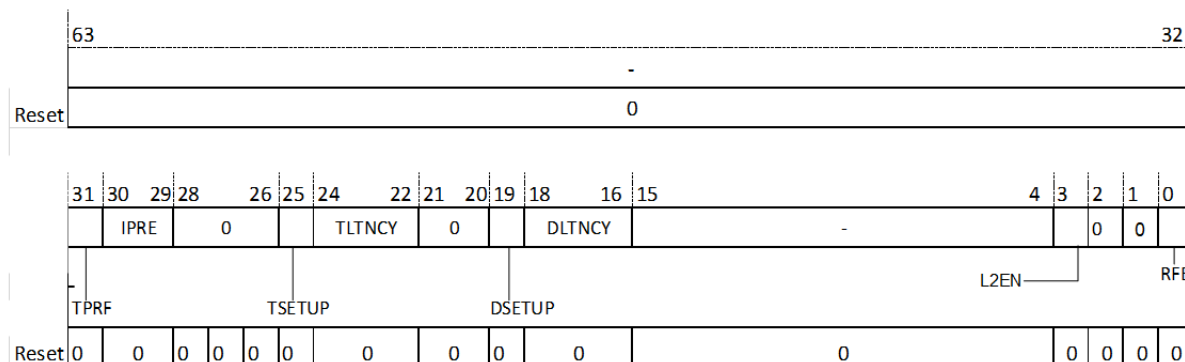


Fig. 18.4: M-mode L2 Cache control register (mccr2)

##### RFE: read allocation enable bit

When RFE is 0, if accessed data is missing in the L2 Cache, the data is not written back to the L2 Cache.

When RFE is 1, if accessed data is missing in the L2 Cache, the data is written back to the L2 Cache.

##### L2EN: L2 Cache enable bit

When L2EN is 0, the L2 Cache is disabled.

When L2EN is 1, the L2 Cache is enabled. (This bit is fixed to 1 in C910/C920.)

##### DLTNCY: data RAM access cycle configure bit for the L2 Cache

When DLTNCY is 0, the data RAM access cycle is 1.

When DLTNCY is 1, the data RAM access cycle is 2.

When DLTNCY is 2, the data RAM access cycle is 3.

When DLTNCY is 3, the data RAM access cycle is 4.

When DLTNCY is 4, the data RAM access cycle is 5.

When DLTNCY is 5, the data RAM access cycle is 6.

When DLTNCY is 6, the data RAM access cycle is 7.

When DLTNCY is 7, the data RAM access cycle is 8.

##### DSETUP: data RAM setup configure bit for the L2 Cache

When DSETUP is 0, the data RAM does not require an additional setup cycle.

When DSETUP is 1, the data RAM requires an additional setup cycle.

**TLTNCY: tag RAM access cycle configure bit for the L2 Cache**

When TLTNCY is 0, the tag RAM access cycle is 1.

When TLTNCY is 1, the tag RAM access cycle is 2.

When TLTNCY is 2, the tag RAM access cycle is 3.

When TLTNCY is 3, the tag RAM access cycle is 4.

When TLTNCY is 4, the tag RAM access cycle is 5.

**TSETUP: tag RAM setup configure bit for the L2 Cache**

When TSETUP is 0, the tag RAM does not require an additional setup cycle.

When TSETUP is 1, the tag RAM requires an additional setup cycle.

**IPRF: instruction prefetch capability of the L2 Cache**

This bit indicates the number of prefetched cache lines when desired data of a value-taking request is missing in the L2 Cache.

When IPRF is 0, instruction prefetch is disabled for the L2 Cache.

When IPRF is 1, one cache line is prefetched.

When IPRF is 2, two cache lines are prefetched.

When IPRF is 3, three cache lines are prefetched.

**TPRF: TLB prefetch enable bit for the L2 Cache**

When TPRF is 0, TLB prefetch is disabled for the L2 Cache.

When TPRF is 1, TLB prefetch is enabled for the L2 Cache.

### 18.1.7.5 M-mode implicit operation register (mhint)

The mhint register controls the enable/disable of multiple functions of caches.

This register is 64 bits wide and is readable and writable in M-mode. Accesses in non-M-mode will cause an illegal instruction exception.

**DPLD: D-Cache prefetch enable bit**

When DPLD is 0, D-Cache prefetch is disabled.

When DPLD is 1, D-Cache prefetch is enabled.

**AMR: write allocate policy automatic adjustment enable bit for the L1 Cache**

When AMR is 0, the write allocate policy is subject to the page attribute WA of the access address.

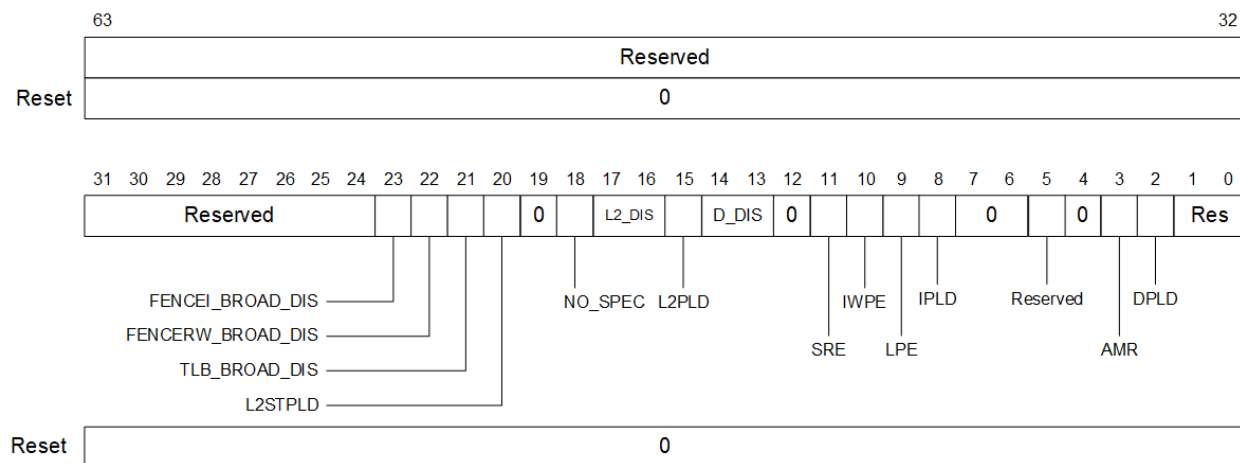


Fig. 18.5: M-mode implicit operation register (mhint)

When AMR is 1, if multiple cache lines are stored continuously, desired data of subsequent storage operations of continuous addresses is no longer written to the L1 Cache.

#### IPLD: I-Cache prefetch enable bit

When IPLD is 0, I-Cache prefetch is disabled.

When IPLD is 1, I-Cache prefetch is enabled.

#### LPE: loop acceleration enable bit

When LPE is 0, loop acceleration is disabled.

When LPE is 1, loop acceleration is enabled.

#### IWPE: I-Cache way prediction enable bit

When IWPE is 0, way prediction is disabled for I-Cache.

When IWPE is 1, way prediction is enabled for I-Cache.

#### SRE: single retirement enable bit

When SRE is 0, single retirement mode is disabled.

When SRE is 1, single retirement mode is enabled.

#### D\_DIS: number of prefetched cache lines in D-Cache

When D\_DIS is 0, two cache lines are prefetched.

When D\_DIS is 1, four cache lines are prefetched.

When D\_DIS is 2, eight cache lines are prefetched.

When D\_DIS is 3, 16 cache lines are prefetched.

The default value is 0.



**L2PLD: the L2 Cache prefetch enable bit**

When L2PLD is 0, L2 Cache prefetch is disabled.

When L2PLD is 1, L2 Cache prefetch is enabled.

**L2\_DIS: number of prefetched cache lines in the L2 Cache**

When L2\_DIS is 0, eight cache lines are prefetched.

When L2\_DIS is 1, 16 cache lines are prefetched.

When L2\_DIS is 2, 32 cache lines are prefetched.

When L2\_DIS is 3, 64 cache lines are prefetched.

The L2 Cache prefetch is based on the L1 Cache prefetch.

**NO\_SPEC: spec fail prediction enable bit**

When NO\_SPEC is 0, spec fail prediction is disabled.

When NO\_SPEC is 1, spec fail prediction is enabled.

**L2STPLD: store prefetch enable bit for the L2 Cache**

When L2STPLD is 0, store prefetch is disabled for the L2 Cache.

When L2STPLD is 1, store prefetch is enabled for the L2 Cache.

**TLB\_BROAD\_DIS: the TLB fence operation broadcast disable bit**

When TLB\_BROAD\_DIS is 0, sfence.vma instruction operations are broadcast to other cores.

When TLB\_BROAD\_DIS is 1, sfence.vma instruction operations are not broadcast.

This bit is invalid when there is only one core.

**FENCERW\_BROAD\_DIS: fence operation broadcast disable bit**

When FENCERW\_BROAD\_DIS is 0, fence instruction operations are broadcast to other cores.

When FENCERW\_BROAD\_DIS is 1, fence instruction operations are not broadcast.

This bit is invalid when there is only one core.

**FENCEI\_BROAD\_DIS: fence.i operation broadcast disable bit**

When FENCEI\_BROAD\_DIS is 0, fence.i instruction operations are broadcast to other cores.

When FENCEI\_BROAD\_DIS is 1, fence.i instruction operations are not broadcast.

This bit is invalid when there is only one core.

**18.1.7.6 M-mode reset vector base address register (mrabr)**

The mrabr register stores base addresses of reset exception vectors. Each C910/C920 core has an independent mrabr register.

This register is 64 bits wide and is read-only in M-mode. Accesses in non-M-mode will cause an illegal instruction exception.

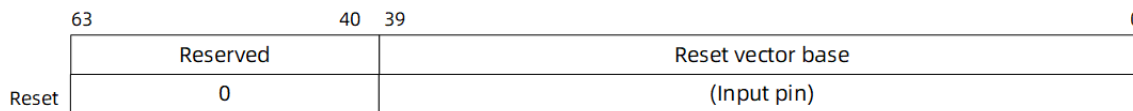


Fig. 18.6: M-mode reset vector base address register (mrivr)

## Reset vector base: reset base address

It controls the reset base address of a core.

#### 18.1.7.7 S-mode counter write enable register (mcounterwen)

The mcounterwen register determines whether S-mode event counters can be written in S-mode.

This register is 64 bits wide and is readable and writable in M-mode. Accesses in non-M-mode will cause an illegal instruction exception.

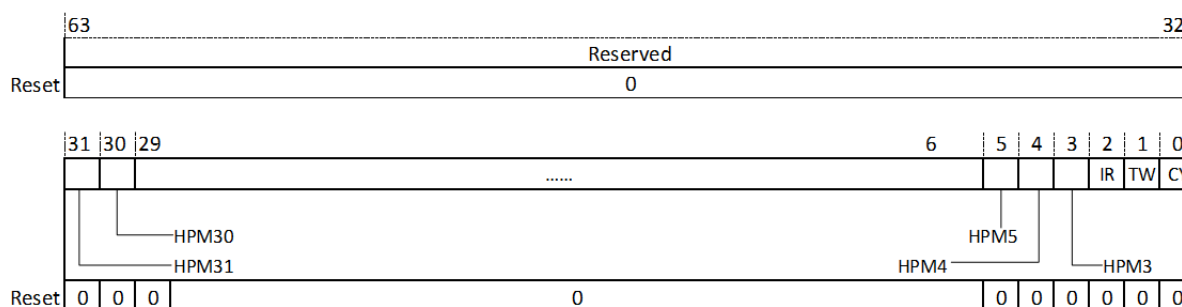


Fig. 18.7: S-mode counter write enable register (mcounterwen)

When `mcounterwen.bit[n]` is 1, writes to the corresponding `shpmcounter` are allowed in S-mode.

When `mcounterwen.bit[n]` is 0, writes to the corresponding `shpmcounter` are not allowed in S-mode, and cause instruction exceptions.

#### 18.1.7.8 M-mode event interrupt enable register (mcounterinten)

The mcounterinten register enables the triggering of interrupts when event counters overflow.

This register is 64 bits wide and is readable and writable in M-mode. Accesses in non-M-mode will cause an illegal instruction exception.

When `mcounterinten.bit[n]` is 1, an interrupt is triggered when the corresponding `mhpmcounter` overflows.

When `mcounterinten.bit[n]` is 0, an interrupt is not triggered when the corresponding `mhpmcounter` overflows.

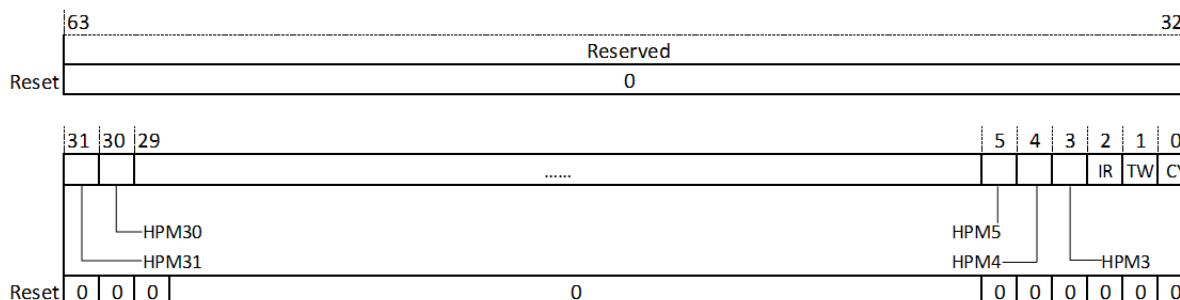


Fig. 18.8: M-mode event interrupt enable register (mcounterinten)

#### 18.1.7.9 M-mode event overflow mark register (mcounteren)

The mcounteren register marks whether event counters overflow.

This register is 64 bits wide and is readable and writable in M-mode. Accesses in non-M-mode will cause an illegal instruction exception.

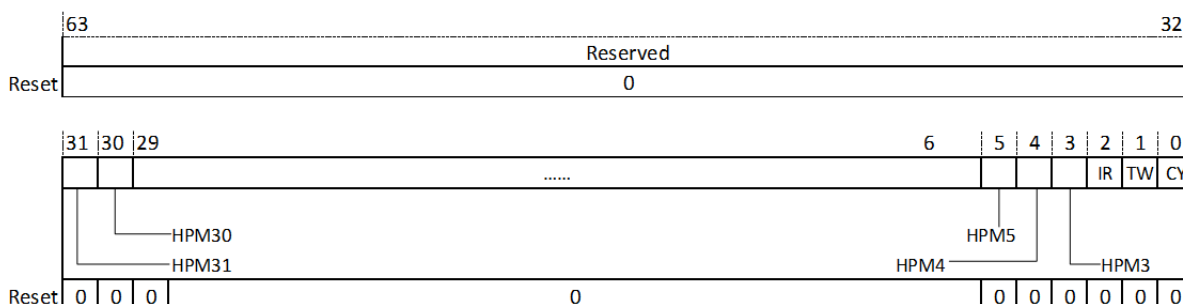


Fig. 18.9: M-mode event overflow mark register (mcounteren)

When `mcounterof.bit[n]` is 1, the corresponding `mhpmcounter` overflows.

When `mcounterof.bit[n]` is 0, the corresponding `mhpcounter` does not overflow.


### 18.1.8 M-mode cache access extension registers

M-mode cache access extension registers directly read content in the L1 Cache and the L2 Cache for cache debugging.

#### 18.1.8.1 M-mode cache instruction register (mcins)

The mcins register initiates read requests to the L1 or L2 Cache.

This register is 64 bits wide and is readable and writable in M-mode. Accesses in non-M-mode will cause an illegal instruction exception.



M-mode cache instruction register (mcins)


**R: cache read access**

- When R is 0, cache read requests are not initiated.
- When R is 1, cache read requests are initiated.

18.1.8.2 M-mode cache access index register (mcindex)

The mcindex register configures the location of a cache accessed by read requests.

This register is 64 bits wide and is readable and writable in M-mode. Accesses in non-M-mode will cause an illegal instruction exception.



M-mode cache access index register (mcindex)

**RID: RAM flag bit**

This bit specifies the accessed RAM.

- When RID is 0, I-Cache tag RAM is accessed.
- When RID is 1, I-Cache data RAM is accessed.
- When RID is 2, D-Cache tag RAM is accessed.
- When RID is 3, D-Cache data RAM is accessed.
- When RID is 4, L2 Cache tag RAM is accessed.
- When RID is 5, L2 Cache data RAM is accessed.
- When RID is 12, D-Cache LD tag RAM is accessed.

**WAY: cache way information**

This bit specifies the RAM access way.


**INDEX: cache index**

This bit specifies the RAM access index.

**18.1.8.3 M-mode cache data register (mcdata0/1)**

The mcdata0/1 register records data read from the L1 or L2 Cache.

This register is 64 bits wide and is readable and writable in M-mode. Accesses in non-M-mode will cause an illegal instruction exception.



M-mode cache access data register (mcdata)

Table 18.1: Correspondence between cache data content and RAM types

RAM type	CDATA content
ICACHE TAG	CDATA0[39:12]: TAG CDATA0[0]: VALID
ICACHE DATA	CDATA0–CDATA1: 128bit DATA
DCACHE TAG	CDATA0[39:12]: TAG CDATA0[2]: DIRTY CDATA0[1]: SHARED CDATA0[0]: VALID
DCACHE DATA	CDATA0–CDATA1: 128bit DATA
L2CACHE TAG	CDATA0[39:12]: TAG CDATA[1]: DIRTY CDATA0[0]: VALID
L2CACHE DATA	CDATA0–CDATA1: 128bit DATA

**18.1.9 M-mode CPU model registers**

**18.1.9.1 M-mode CPU model register (mcpuid)**

The mcpuid registers store CPU models. The reset value is subject to the product.

### 18.1.9.2 On-chip bus base address register (mapbaddr)

The mapbaddr register stores the base addresses of on-chip registers (CLINT and PLIC) of the CPU. The value of this register is subject to hardware integration.

## 18.1.10 Multi-core extension registers

### 18.1.10.1 Snoop listening enable register (msmpr)

The msmp register controls whether cores can process listening requests. The listening request processing capability is configured for each core separately. The consistency bus of the L2 subsystem controls the sending of listening requests based on the listening status of each core. This register is readable and writable in M-mode.

The msmp register is 64 bits wide. Only bit 0 is defined and the other bits are reserved.

#### Bit 0: SMPEN: core listening enable bit

- When SMPEN is 1' b0, the corresponding core cannot process listening requests, and the L2 subsystem masks the listening requests bound for the core. (This is the reset value.)
- When SMPEN is 1' b1, the corresponding core can process listening requests, and the L2 subsystem sends the listening requests bound for the core.

Before a CPU core is powered off, its SMPEN bit must be set to 0 to disable listening. After a core is powered on, its SMPEN bit must be set to 1 before D-Cache and MMU are enabled. The SMPEN bit must be set to 1 when a core runs properly (including WFI mode). Otherwise, unexpected results may be caused.

## 18.2 Appendix C-2 S-mode control registers

S-mode control registers are classified by function into S-mode exception configuration registers, S-mode exception handling registers, and S-mode address translation registers.

### 18.2.1 S-mode exception configuration registers

When exceptions and interrupts are downgraded to S-mode responses, exceptions must be configured through the S-mode exception configuration registers, like in M-mode.

#### 18.2.1.1 S-mode status register (sstatus)

The sstatus register stores status and control information of the CPU in S-mode, including the global interrupt enable bit, exception preserve interrupt enable bit, and exception preserve privilege mode bit. The sstatus register is a partial mapping of the mstatus register.

This register is 64 bits wide and is readable and writable in S-mode. Accesses in U-mode will cause an illegal instruction exception.

	63	62																															36	35	34	33	32
	SD	0																																	0	UXL	
Reset	0	0																																	2	2	

	31											25	24	23	22	21	20	19	18	17	16	15	14	13	12	11	10	9	8	7	6	5	4	3	2	1	0
	0										VS		0	0	0			0	XS	FS	0		0	SPP	0	0		0	0	0	SIE	0					
											MXR					SUM										SPIE											
Reset	0										0	0	0	0	0	0	0	0	0	0	0	3	0	1	0	0	0	0	0	0	0	0	0	0			

S-mode status register (sstatus)

For more information, see ref:[appendix\\_c12\\_mstatus](#).

### 18.2.1.2 S-mode interrupt-enable register (sie)

The sie register controls the enable and mask of different types of interrupts, and is a partial mapping of the mie register. This register is 64 bits wide and readable in S-mode. The write permission in S-mode is determined by the mideleg register of the corresponding bit. Accesses in U-mode will cause an illegal instruction exception.

	63																															32
	0																															
Reset	0																															

	31											18	17	16	15											10	9	8	7	6	5	4	3	2	1	0				
	0												0	0											0	0	0		0	0	0		0	0	0		0			
											MOIE															SEIE					STIE					SSIE				
Reset	0											0	0	0										0	0	0	0	0	0	0	0	0	0	0	0	0	0			


S-mode interrupt-enable register (sie)

For more information, see ref:[appendix\\_c12\\_mie](#).

### 18.2.1.3 S-mode trap vector base address register (stvec)

The stvec register stores the entry address of the exception service program.

This register is 64 bits wide and is readable and writable in S-mode. Accesses in U-mode will cause an illegal instruction exception.



S-mode trap vector base address register (stvec)

For more information, see [ref:appendix\\_c12\\_mtvec](#).

#### 18.2.1.4 S-mode counter access enable register (scounteren)

The scounteren register determines whether U-mode counters can be accessed in U-mode.

For more information, see [ref:performance\\_test\\_scounteren](#).

### 18.2.2 S-mode exception handling registers

#### 18.2.2.1 S-mode scratch register (sscratch)

The sscratch register is used by the CPU to back up temporary data in the exception service program. It is usually used to store the entry pointer to the local context space in S-mode.

This register is 64 bits wide and is readable and writable in S-mode. Accesses in U-mode will cause an illegal instruction exception.

#### 18.2.2.2 S-mode exception program counter register (sepc)

The sepc register stores the program counter value (PC value) when the CPU exits from the exception service program. C910/C920 supports 16 bits wide instructions. The values of sepc are aligned with 16 bits and the lowest bit is 0.

This register is 64 bits wide and is readable and writable in S-mode. Accesses in U-mode will cause an illegal instruction exception.

#### 18.2.2.3 S-mode cause register (scause)

The scause register stores the vector numbers of events that trigger exceptions. The vector numbers are used to handle corresponding events in the exception service program.


This register is 64 bits wide and is readable and writable in S-mode. Accesses in U-mode will cause an illegal instruction exception.



### 18.2.2.4 S-mode interrupt-pending register (sip)

The sip register stores information about pending interrupts. When the CPU cannot immediately respond to an interrupt, the corresponding bit in the sip register will be set.

This register is 64 bits wide and readable in S-mode. The write permission is determined by the mideleg register of the corresponding bit. Accesses in U-mode will cause an illegal instruction exception.



S-mode interrupt-pending register (sip)

## 18.2.3 S-mode address translation registers

Virtual memory spaces need to be accessed in S-mode. The S-mode address translation register (satp) controls MMU mode switching, hardware writeback base address, and process ID.

### 18.2.3.1 S-mode address translation register (satp)

The S-mode address translation register (satp) controls MMU mode switching, hardware writeback base address, and process ID.

This register is 64 bits wide and is readable and writable in S-mode. Accesses in U-mode will cause an illegal instruction exception.

For more information, see [ref:virtual\\_mem\\_manage\\_satp](#).

## 18.2.4 S-mode CPU control and status extension registers

### 18.2.4.1 S-mode extension status register (sxstatus)

The sxstatus register is the mapping of the mxstatus register. For more information, see [ref:appendix\\_c17\\_mxstatus](#).

This register is 64 bits wide and is readable in S-mode. Only the MM bit is writable. Accesses in U-mode will cause an illegal instruction exception.

#### 18.2.4.2 S-mode hardware control register (shcr)

The shcr register is the mapping of the mhcr register. For more information, see [ref:appendix\\_c17\\_mhcr](#).

This register is 64 bits wide and is readable in S-mode. Accesses in U-mode will cause an illegal instruction exception.

#### 18.2.4.3 S-mode event overflow interrupt enable register (scounterinten)

The scounterinten register is the mapping of the mcounterinten register. For more information, see [ref:appendix\\_c17\\_mcounterinten](#).

This register is 64 bits wide and is readable in S-mode. Accesses in U-mode will cause an illegal instruction exception.

When `mcounterwen.bit[n]` is 1, `scounterinten.bit[n]` determines whether to generate an interrupt when the corresponding shpmcounter overflows.

#### 18.2.4.4 S-mode event overflow mark register (scounterof)

The scounterof register is the mapping of the mcounterof register. For more information, see [ref:appendix\\_c17\\_mcounterof](#).

This register is 64 bits wide and is readable in S-mode. Accesses in U-mode will cause an illegal instruction exception.

When `mcounterwen.bit[n]` is 1, `scounterof.bit[n]` indicates whether the corresponding shpmcounter overflows.

#### 18.2.4.5 S-mode cycle counter (scycle)

The scycle counter stores the cycles executed by the CPU. When the CPU is in the execution state (non-low power state), the scycle register increases the count upon each execution cycle.

The mcycle counter is 64 bits wide and will be reset to 0.

For more information, see [ref:performance\\_test\\_cont](#).

#### 18.2.4.6 S-mode instructions-retired counter (sinstret)

The sinstret counter stores the number of retired instructions of the CPU. The sinstret register increases the count when each instruction retires.

The sinstret counter is 64 bits wide and will be reset to 0.

For more information, see [ref:performance\\_test\\_cont](#).

#### 18.2.4.7 S-mode event counter (shpmcountern)

The shpmcountern counter is the mapping of the mhpcountern counter.

For more information, see [ref:performance\\_test\\_cont](#).

### 18.2.5 S-mode MMU extension register

C910/C920 extends MMU related registers to implement software writeback. Software can directly write and read the TLB.

#### 18.2.5.1 S-mode MMU control register (smcir)

This register is 64 bits wide and is readable in S-mode. Accesses in U-mode will cause an illegal instruction exception.

For more information, see [ref:virtual\\_mem\\_manage\\_smcir](#).

#### 18.2.5.2 S-mode MMU control register (smir)

This register is 64 bits wide and is readable in S-mode. Accesses in U-mode will cause an illegal instruction exception.

For more information, see [ref:virtual\\_mem\\_manage\\_smir](#).

#### 18.2.5.3 S-mode MMU control register (smeh)

This register is 64 bits wide and is readable in S-mode. Accesses in U-mode will cause an illegal instruction exception.

For more information, see [ref:virtual\\_mem\\_manage\\_sme](#).

#### 18.2.5.4 S-mode MMU control register (smel)

This register is 64 bits wide and is readable in S-mode. Accesses in U-mode will cause an illegal instruction exception.

For more information, see [ref:virtual\\_mem\\_manage\\_smel](#).

## 18.3 Appendix C-3 U-mode control registers

U-mode control registers are classified by function into floating-point registers, counter registers, and vector control registers.

18.3.1 U-mode floating-point control registers

18.3.1.1 Floating-point accrued exceptions register (fflags)

The fflags register is the mapping of accrued exceptions of the fcsr register. For more information, see [ref:appendix\\_c31\\_fcsr](#).

18.3.1.2 Floating-point dynamic rounding mode register (frm)

The frm register is the mapping of the rounding mode of the fcsr register. For more information, see [ref:appendix\\_c31\\_fcsr](#).

18.3.1.3 Floating-point control and status register (fcsr)

The fcsr register records floating-point accrued exceptions and the rounding mode.

This register is 64 bits wide and readable and writable in any mode.

	63																					32	
	-																						
Reset	0																						
	31											11	10	9	8	7	5	4					0
	-										Vxrm			Rounding Mode(frm)	Accrued Exceptions(fflags)								
															NV	DZ	OF	UF	NX				
																Vxsat							
Reset	0										0		0	0	0	0	0	0					

Floating-point control and status register (fcsr)

**NX: imprecise exception**

- When NX is 0, no imprecise exception occurs.

When NX is 1, imprecise exceptions occur.

**UF: underflow exception**

- When UF is 0, no underflow exception occurs.
- When UF is 1, underflow exceptions occur.

**OF: overflow exception**

- When OF is 0, no overflow exception occurs.
- When OF is 1, overflow exceptions occur.

**DZ: division by zero exception**

- When DZ is 0, no division by zero exception occurs.
- When DZ is 1, division by zero exceptions occur.

**NV: illegal instruction operand exception**

- When NV is 0, no exception of illegal instruction operands occurs.
- When NV is 1, exceptions of illegal instruction operands occur.

**RM: rounding mode**

- When RM is 0, the RNE rounding mode takes effect, and values are rounded off to the nearest even number.
- When RM is 1, the RTZ rounding mode takes effect, and values are rounded off to zero.
- When RM is 2, the RDN rounding mode takes effect, and values are rounded off to negative infinity.
- When RM is 3, the RUP rounding mode takes effect, and values are rounded off to positive infinity.
- When RM is 4, the RMM rounding mode takes effect, and values are rounded off to the nearest number.

**VXSAT: vector overflow flag bit**

This register is the mapping of the corresponding bit.

**VXRM: vector rounding mode bit**

This register is the mapping of the corresponding bit.

## 18.3.2 U-mode counter/timer registers

### 18.3.2.1 User cycle register (cycle)

The cycle register stores the cycles executed by the CPU. When the CPU is in the execution state (non-low power state), the cycle register increases the count upon each execution cycle.

The mcycle counter is 64 bits wide and will be reset to 0.

For more information, see [ref:performance\\_test\\_cont](#).

### 18.3.2.2 U-mode timer register (time)

The time register is the read-only mapping of the mtime register.

For more information, see [ref:performance\\_test\\_cont](#).

18.3.2.3 User instructions-retired counter (instret)

The instret counter stores the number of retired instructions of the CPU. The instret register increases the count when each instruction retires.

The sinstret counter is 64 bits wide and will be reset to 0.

For more information, see ref:performance\_test\_cont.

18.3.2.4 User event counter (hpmcountern)


The hpmcountern counter is the mapping of the mhpcountern counter.

For more information, see ref:performance\_test\_cont.

18.3.3 U-mode floating-point extension control registers

18.3.3.1 U-mode floating-point extension control register (focr)

The focr register controls the floating-point extension function and floating-point exception accrue bit.



Floating-point extension control register (focr)

**NX: imprecise exception**

It is the mapping of the corresponding bit of the focr register.

**UF: underflow exception**

It is the mapping of the corresponding bit of the focr register.

**OF: overflow exception**

It is the mapping of the corresponding bit of the focr register.

**DZ: division by zero exception**

It is the mapping of the corresponding bit of the focr register.

**NV: illegal instruction operand exception**

It is the mapping of the corresponding bit of the fcsr register.

**FE: floating-point exception accrue bit**

This bit is set to 1 when any floating-point exception occurs.

**DQNaN: output QNaN mode bit**

When DQNaN is 0, the output QNaN value is the default value.

When DQNaN is 1, the output QNaN value is consistent with the IEEE754 standard.

**RM: rounding mode**

It is the mapping of the corresponding bit of the fcsr register.